



CCIE Routing and Switching v5.0 Configuration Practice Labs

Martin J. Duggan

Cisco Press

About This eBook

ePUB is an open, industry-standard format for eBooks. However, support of ePUB and its many features varies across reading devices and applications. Use your device or app settings to customize the presentation to your liking. Settings that you can customize often include font, font size, single or double column, landscape or portrait mode, and figures that you can click or tap to enlarge. For additional information about the settings and features on your reading device or app, visit the device manufacturer's Web site.

Many titles include programming code or configuration examples. To optimize the presentation of these elements, view the eBook in single-column, landscape mode and adjust the font size to the smallest setting. In addition to presenting code and configurations in the reflowable text format, we have included images of the code that mimic the presentation found in the print book; therefore, where the reflowable format may compromise the presentation of the code listing, you will see a "Click here to view code image" link. Click the link to view the print-fidelity code image. To return to the previous page viewed, click the Back button on your device or app.

Cisco CCIE Routing and Switching v5.0

Configuration Practice Labs

Third Edition

Martin James Duggan

Cisco Press
800 East 96th Street
Indianapolis, IN 46240

Cisco CCIE Routing and Switching v5.0 Configuration Practice Labs, Third Edition

Martin James Duggan

Copyright© 2014 Pearson Education, Inc.

Published by:

Cisco Press
800 East 96th Street
Indianapolis, IN 46240 USA

All rights reserved. No part of this book may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage and retrieval system, without written permission from the publisher, except for the inclusion of brief quotations in a review.

Printed in the United States of America

First Printing February, 2014

ISBN-13: 978-0-13-378631-6

ISBN-10: 0-13-378631-5

Warning and Disclaimer

This book is designed to provide information about the Cisco Certified Internetwork Expert (CCIE) Routing and Switching (R&S) Lab 5.0 Exam. Every effort has been made to make this book as complete and as accurate as possible, but no warranty or fitness is implied.

The information is provided on an “as is” basis. The authors, Cisco Press, and Cisco Systems, Inc. shall have neither liability nor responsibility to any person or entity with respect to any loss or damages arising from the information contained in this book or from the use of the discs or programs that may accompany it.

The opinions expressed in this book belong to the author and are not necessarily those of Cisco Systems, Inc.

Trademark Acknowledgments

All terms mentioned in this book that are known to be trademarks or service marks have been appropriately capitalized. Cisco Press or Cisco Systems, Inc., cannot attest to the accuracy of this information. Use of a term in this book should not be regarded as affecting the validity of any trademark or service mark.

Special Sales

For information about buying this title in bulk quantities, or for special sales opportunities (which may include electronic versions; custom cover designs; and content particular to your business, training goals, marketing focus, or branding interests), please contact our corporate sales department at corpsales@pearsoned.com or (800) 382-3419.

For government sales inquiries, please contact governmentsales@pearsoned.com.

For questions about sales outside the U.S., please contact international@pearsoned.com.

Feedback Information

At Cisco Press, our goal is to create in-depth technical books of the highest quality and value. Each book is crafted with care and precision, undergoing rigorous development that involves the unique expertise of members from the professional technical community.

Readers' feedback is a natural continuation of this process. If you have any comments regarding how we could improve the quality of this book, or otherwise alter it to better suit your needs, you can contact us through email at feedback@ciscopress.com. Please make sure to include the book title and ISBN in your message.

We greatly appreciate your assistance.

Publisher: Paul Boger

Associate Publisher: Dave Dushtimer

Business Operation Manager, Cisco Press: Jan Cornelissen

Executive Editor: Brett Bartow

Managing Editor: Sandra Schroeder

Senior Development Editor: Christopher Cleveland

Senior Project Editor: Tonya Simpson

Copy Editor: Keith Cline

Technical Editor: Neil Shaw

Editorial Assistant: Vanessa Evans

Composition: Bronkella Publishing

Proofreader: Sarah Kearns



Americas Headquarters

Cisco Systems, Inc.
San Jose, CA

Asia Pacific Headquarters

Cisco Systems (USA) Pte. Ltd.
Singapore

Europe Headquarters

Cisco Systems International BV

Amsterdam, The Netherlands

Cisco has more than 200 offices worldwide. Addresses, phone numbers, and fax numbers are listed on the Cisco Website at www.cisco.com/go/offices.

CCDE, CCENT, Cisco Eos, Cisco HealthPresence, the Cisco logo, Cisco Lumin, Cisco Nexus, Cisco StadiumVision, Cisco Telepresence, Cisco WebEx, DCE, and Welcome to the Human Network are trademarks; Changing the Way We Work, Live, Play, and Learn and Cisco Store are service marks; and Access Registrar, Aironet, AsyncOS, Bringing the Meeting To You, Catalyst, CCDA, CCDP, CCIE, CCIP, CCNA, CCNP, CCSP, CCVP, Cisco, the Cisco Certified Internetwork Expert logo, Cisco IOS, Cisco Press, Cisco Systems, Cisco Systems Capital, the Cisco Systems logo, Cisco Unity, Collaboration Without Limitation, EtherFast, EtherSwitch, Event Center, Fast Step, Follow Me Browsing, FormShare, GigaDrive, HomeLink, Internet Quotient, IOS, Phone, iQuick Study, IronPort, the IronPort logo, LightStream, Linksys, MediaTone, MeetingPlace, MeetingPlace Chime Sound, MGX, Networkers, Networking Academy, Network Registrar, PCNow, PIX, PowerPanels, ProConnect, ScriptShare, SenderBase, SMARTnet, Spectrum Expert, StackWise, The Fastest Way to Increase Your Internet Quotient, TransPath, WebEx, and the WebEx logo are registered trademarks of Cisco Systems, Inc. and/or its affiliates in the United States and certain other countries.

All other trademarks mentioned in this document or website are the property of their respective owners. The use of the word partner does not imply a partnership relationship between Cisco and any other company. (0812R)

About the Author

Martin James Duggan, CCIE No. 7942, is a network architect for AT&T. He designs network solutions for customers globally and specializes in data center networking and QoS. Martin mentors colleagues through their Cisco qualifications and holds regular internal training classes. Previous to this, Martin was a network architect for IBM, performing IP network designs and global network reviews. Martin has been in the industry for 23 years, focusing on Cisco solutions for the past 14 years. Martin is the author of the Cisco Press *CCIE Routing and Switching v4.0 Troubleshooting Practice Labs*, First Edition.

About the Technical Reviewer

Neil Shaw, CCIE No. 38866, is a network consultant for Alternative Networks plc. Prior to this, Neil worked for both IBM and AT&T, with more than 16 years of IT industry experience. Specializing in data center and enterprise networking for the past 8 years, Neil's experience includes design, implementation, and troubleshooting of IP networks with key skills in security, routing, and switching.

Dedications

I dedicate this publication to my children. Anna and James, you two are such fun to be with, and I simply cannot get enough of your time, whether we are chasing each other down a snow-covered mountain or just playing a game of Boggle at home; hanging out with you makes my life complete. Jake, welcome to the family, your thirst for milk is only comparable with Anna's and James's thirst for Apple products. I can't wait to hear your first words and to teach you how to ride a bike. All I ask when you win the Tour de France is an autograph for your old man.

Lotte, you brought me happiness, and now you have brought me Jake; I am so lucky to have you in my life. Mum and Dad, I want to thank you for all the future babysitting; Jake will be in very capable hands when he is with you.

Acknowledgments

I want to thank Brett Bartow for once again providing me with this enviable opportunity to write for Cisco Press.

Bruno, it has been great working with you again. Cisco is very lucky to have you; your knowledge on this subject is second to none. Next time we meet, the sushi is definitely on me.

To Mike Jones, my Strava cycling buddy, I blame this book if you beat my time up Bolswood Lane!

To Mike Randall, who has the largest white board in Salt Lake City, thanks for the inspiration and opportunities you provided me. You had better start getting some miles in on that Di2 of yours, as I will be heading out and expecting a fast pace when I get there.

To Neil Shaw, who reviewed my work, thanks for your attention to detail. It's easy to see how you got your number!

Contents at a Glance

[Introduction](#)

[Chapter 1 Practice Lab 1](#)

[Chapter 2 Practice Lab 2](#)

[Chapter 3 Practice Lab 3](#)

[Chapter 4 Summary](#)

Appendix A Practice Lab 1 Switch and Router Initial and Final Configurations (online only)

Appendix B Practice Lab 2 Switch and Router Initial and Final Configurations (online only)

Appendix C Practice Lab 3 Switch and Router Initial and Final Configurations (online only)

[Where are the Companion Content Files?](#)

Contents

[Introduction](#)

[Practice Lab 1](#)

[Equipment List](#)

[Setting Up the Lab](#)

[Lab Topology](#)

[Switch Instructions](#)

[Serial Link](#)

[IP Address Instructions](#)

[Pre-Lab Tasks](#)

[General Guidelines](#)

[Practice Lab One](#)

[Section 1: LAN Switching \(25 Points\)](#)

[Section 2: IPv4 IGP Protocols \(24 Points\)](#)

[Section 2.1: OSPF](#)

[Section 2.2: EIGRP](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP \(14 Points\)](#)

[Section 4: IPv6 \(15 Points\)](#)

[Section 4.1: EIGRPv6](#)

[Section 4.2: OSPFv3](#)

[Section 4.3: Redistribution](#)

[Section 5: QoS \(8 Points\)](#)

[Section 6: Security \(6 Points\)](#)

[Section 7: Multicast \(4 Points\)](#)

[IP Services \(4 Points\)](#)

[“Ask the Proctor”](#)

[Section 1: LAN Switching](#)

[Section 2: IPv4 IGP Protocols](#)

[Section 2.1: OSPF](#)

[Section 2.2: EIGRP](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP](#)

[Section 4: IPv6](#)

[Section 4.3: Redistribution](#)

[Section 5: QoS](#)

[Section 6: Security](#)

[Section 7: Multicast](#)

[Section 8: IP Services](#)

[Lab Debrief](#)

[Section 1: LAN Switching \(25 Points\)](#)

[Section 2: IPv4 IGP Protocols \(24 Points\)](#)

[Section 2.1: OSPF](#)

[Section 2.2: EIGRP](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP \(14 Points\)](#)

[Section 4: IPv6 \(15 Points\)](#)

[Section 4.1: EIGRPv6](#)

[Section 4.2: OSPFv3](#)

[Section 4.3: Redistribution](#)

[Section 5: QoS \(8 Points\)](#)

[Section 6: Security \(6 Points\)](#)

[Section 7: Multicast \(4 Points\)](#)

[IP Services \(4 Points\)](#)

[Lab Wrap-Up](#)

[**Practice Lab 2**](#)

[Equipment List](#)

[Setting Up the Lab](#)

[Lab Topology](#)

[Switch Instructions](#)

[IP Address Instructions](#)

[Pre-Lab Tasks](#)

[General Guidelines](#)

[Practice Lab Two](#)

[Section 1: LAN Switching \(22 Points\)](#)

[Section 2: IPv4 IGP Protocols \(26 Points\)](#)

[Section 2.1: EIGRP](#)

[Section 2.2: OSPF](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP \(15 Points\)](#)

[Section 4: IPv6 \(12 Points\)](#)

[Section 4.1: EIGRPv6](#)

[Section 4.2: OSPFv3](#)

[Section 4.3: Redistribution](#)

[Section 5: QoS \(6 Points\)](#)

[Section 6: Multicast \(9 Points\)](#)

[Section 7: Security \(10 Points\)](#)

[“Ask the Proctor”](#)

[Section 1: LAN Switching](#)

[Section 2: IPv4 IGP Protocols](#)

[Section 2.1: EIGRP](#)

[Section 2.2: OSPF](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP](#)

[Section 4: IPv6](#)

[Section 5: QoS](#)

[Section 6: Multicast](#)

[Section 7: Security](#)

[Practice Lab Debrief](#)

[Section 1: LAN Switching \(22 Points\)](#)

[Section 2: IPv4 IGP Protocols \(26 Points\)](#)

[Section 2.1: EIGRP](#)

[Section 2.2: OSPF](#)

[Section 2.3: Redistribution](#)

[Section 3: BGP \(15 Points\)](#)

[Section 4: IPv6 \(12 Points\)](#)

[Section 4.1: EIGRPv6](#)

[Section 4.2: OSPFv3](#)

[Section 4.3: Redistribution](#)

[Section 5: QoS \(6 Points\)](#)

[Section 6: Multicast \(9 Points\)](#)

[Section 7: Security \(10 Points\)](#)

[Lab Wrap-Up](#)

[Practice Lab 3](#)

[Equipment List](#)

[Setting Up the Lab](#)

[Lab Topology](#)

[Switch Instructions](#)

[IP Address Instructions](#)

[Pre-Lab Tasks](#)

[General Guidelines](#)

[Practice Lab Three](#)

[Section 1: LAN Switching \(4 Points\)](#)

[Section 2: MPLS and OSPF \(27 Points\)](#)

[Section 3: BGP \(5 Points\)](#)

[Section 4: EIGRP and MP-BGP \(3 Points\)](#)

[Section 5: OSPF and MP-BGP \(6 Points\)](#)

[Section 6: MPLS \(7 Points\)](#)

[Section 7: VPLS Simulation \(10 Points\)](#)

[Section 8: Multicast \(10 Points\)](#)

[Section 9: IPv6 \(6 Points\)](#)

[Section 10: QoS \(7 Points\)](#)

[Section 11: Security \(15 Points\)](#)

[Practice Lab 3: “Ask the Proctor”](#)

[Section 1: LAN Switching and Frame Relay](#)

[Section 2: MPLS and OSPF](#)

[Section 3: BGP](#)

[Section 4: EIGRP and MP-BGP](#)

[Section 5: OSPF and MP-BGP](#)

[Section 6: MPLS](#)

[Section 7: VPLS Simulation](#)

[Section 8: Multicast](#)

[Section 9: IPv6](#)

[Section 10: QoS](#)

[Section 11: Security](#)

[Practice Lab 3 Debrief](#)

[Section 1: LAN Switching \(4 Points\)](#)

[Section 2: MPLS and OSPF \(27 Points\)](#)

[Section 3: BGP \(5 Points\)](#)

[Section 4: EIGRP and MP-BGP \(3 Points\)](#)

[Section 5: OSPF and MP-BGP \(6 Points\)](#)

[Section 6: MPLS \(7 Points\)](#)

[Section 7: VPLS Simulation \(10 Points\)](#)

[Section 8: Multicast \(10 Points\)](#)

[Section 9: IPv6 \(6 Points\)](#)

[Section 10: QoS \(7 Points\)](#)

[Section 11: Security \(15 Points\)](#)

[Lab 3 Wrap-Up](#)

[Summary](#)

[Are You Ready?](#)

[Further Reading](#)

[Help and Advice](#)

[How Can I Schedule My CCIE Lab Exam?](#)

[The Day Before](#)

[The Day of the Exam](#)

[Pass or Fail, What Next?](#)

[Where are the Companion Content Files?](#)

Command Syntax Conventions

The conventions used to present command syntax in this book are the same conventions used in the IOS Command Reference. The Command Reference describes these conventions as follows:

- **Boldface** indicates commands and keywords that are entered literally as shown. In actual configuration examples and output (not general command syntax), boldface indicates commands that are manually input by the user (such as a **show** command).
- *Italic* indicates arguments for which you supply actual values.
- Vertical bars (|) separate alternative, mutually exclusive elements.
- Square brackets ([]) indicate an optional element.
- Braces ({ }) indicate a required choice.
- Braces within brackets ([{ }]) indicate a required choice within an optional element.

Introduction

For more than 10 years, the CCIE program has identified networking professionals with the highest level of expertise. Less than 3 percent of all Cisco certified professionals actually achieve CCIE status. The majority of candidates who take the exam fail at the first attempt because they are not fully prepared; they generally find that their study plan did not match what was expected of them in the exam. This practice exam has been designed to take you as close as possible to actually taking the actual lab exam. It can show whether you are ready to schedule your lab or if you need to reevaluate your study plan.

Exam Overview

The CCIE qualification consists of two separate exams, a 2-hour written exam and an 8-hour hands-on lab exam that includes troubleshooting, diagnostic, and configuration sections. Written exams are computer-based, multiple-choice exams lasting 2 hours and available at hundreds of authorized testing centers worldwide. The written exam tests your theoretical knowledge to ensure that you are ready to take the lab exam; therefore, you are eligible to schedule the lab exam only after you have passed the written exam. Having purchased this publication, you have most likely already passed the written exam and are ready to practice for the lab exam. The lab exam is a 5.5-hour, hands-on exam in which you configure a series of complex scenarios in strict accordance with the questions; it is tough, but achievable. Troubleshooting is now included for 2 hours, and you also face a series of further diagnostic questions for a 30-minutes period of the exam. You can find current lab blueprint content information at the following URL:

https://learningnetwork.cisco.com/community/certifications/ccie_routing_switching/lab_exam_v5

The diagnostic section that has been added to the lab exam focuses on the skills required to properly diagnose network issues. The time for this new section is fixed at 30 minutes. The sequence of delivery of the lab sections is fixed as follows: Troubleshooting, Diagnostics, and Configuration. Although the time constraint is an essential component of the Troubleshooting section, the system enables some flexibility by granting candidates the option of adding up to 30 minutes to complete the section. The system doesn't permit toggling between modules, however. To maintain the total exam time of 8 hours, the optional extra time used in the Troubleshooting section is automatically deducted from the time credit allotted for the Configuration section. On the other hand, if the time spent in the Troubleshooting section is less than 2 hours, the Configuration section is credited by the time gained.

Scoring Point System

In the actual exam, a higher number of available points for certain questions would generally indicate that the required solution would take more time to achieve or that there would be multiple lines of configuration involved. If you find yourself running short on time, try to get the smaller tasks completed, and then return to the more complex questions.

Study Roadmap

Taking the lab exam is all about experience; you cannot expect to take it and pass after just completing your written exam relying on your theoretical knowledge. You must spend countless hours of rack time configuring features and learning how protocols interact with one another. To be confident enough to schedule your lab exam, review the following outlined points.

Assessing Your Strengths

Using the content blueprint, determine your experience and knowledge in the major topic areas. For areas of strength, practicing for speed should be your focus. For weak areas, you might need training or book study in addition to practice.

Study Materials

Choose lab materials that provide configuration examples and take a hands-on approach. Look for materials approved or provided by Cisco and its Learning Partners.

Hands-On Practice

Build and practice your lab scenarios on a per-topic basis. Go beyond the basics and practice additional features. Learn the **show** and **debug** commands along with each topic. If a protocol has multiple ways of configuring a feature, practice all of them.

Cisco Documentation

Familiarize yourself with content from the following URL, but note that access to this URL likely will be restricted in the real exam:

<http://www.cisco.com/cisco/web/psa/configure.html>

Home Labs

Although acquiring a personal home lab is ideal, it can be costly to gather all the equipment you will need.

Cisco 360 Program

The Cisco 360 Learning Program encompasses six stages of activity to support successful learning for students:

- 1. Assessment:** Students take a diagnostic preassessment lab to benchmark their knowledge of various networking topics.
- 2. Planning:** Based on the preassessment, students create a learning plan that uses a mix of learning components to focus their study.
- 3. Learning:** Students learn by participating in lessons and lectures, reading materials, and working with peers and instructors.
- 4. Practice:** Students use the practice exercises to apply learning on actual network equipment.
- 5. Mastery:** Students measure their understanding by completing assessments of knowledge and skill for various approaches to solving network problems.

6. Review: Students review their work with a mentor or instructor and tune their skills with tips and best practices.

You can find detailed information about the 360 program at the following URL:

https://learningnetwork.cisco.com/community/learning_center/cisco_360/360-rs

Equipment List and IOS Requirements

The lab exam tests any feature that can be configured on the equipment and the IOS versions indicated here:

- 3925 series routers - IOS 15.3(T) – Advanced Enterprise Services
- Catalyst 3560X series switches running IOS Version 15.0S – Advanced IP Services

Chapter Overview

Each chapter consists of questions related to the v5.0 Blueprint within the following sections:

Chapter Topic Sections	v5.0 Topic Section
LAN Switching	Layer 2 Technologies
IGPs	Layer 3 Technologies
Redistribution	
BGP	
Multicast	
IPv6	
MPLS, Tunneling, DMVPN	VPN Technologies
QoS	Infrastructure Services
Security	Infrastructure Security

[Chapters 1](#) and [2](#) follow a similar format but without questions about Multiprotocol Label Switching (MPLS), whereas [Chapter 3](#) emphasizes virtual private networks (VPN) and questions specifically related to MPLS.

Each chapter begins with an overview of the required equipment and topology and provides general guidelines to follow while taking the practice exam. The actual practice lab exam is then presented, which should be undertaken in an advised 8-hour window. An “Ask the Proctor” section is included at the end of the question section to provide some clues as to how to answer the question (which you might find helpful if you cannot immediately solve a particular question). A full lab debrief is then provided that explains the optimal solution for each question, along with the associated **show** commands to provide full verification of the working solution. Each chapter then closes with a brief summary offering tips and advice.

The appendixes provide initial and final configurations for each router and switch per chapter to enable you to preconfigure and check your final working configurations at the beginning and end of each practice lab exam.

The appendixes are provided to you online from the book’s registered product page in .txt format, to

make it easier for you to copy/paste them into your own lab gear. To access these files, see the “[Where Are the Companion Files?](#)” page at the end of this eBook document.

Practice Lab 1

The CCIE exam commences with 2 hours of troubleshooting followed by 5 1/2 hours of configuration and a final 30 minutes of additional questions. This lab consists of 100 points and has been timed to last for 8 hours of configuration and self-troubleshooting, so aim to complete the lab within this period. Then either score yourself at this point or continue until you believe you have met all the objectives. You will now be guided through the equipment requirements and pre-lab tasks in preparation for taking this practice lab.

If you do not own six routers and four switches, consider using the equipment available and additional lab exercises and training facilities available within the CCIE R&S 360 program. You can find detailed information on the 360 program and CCIE R&S exam on the following URLs, respectively:

https://learningnetwork.cisco.com/community/learning_center/cisco_360/360-rs

https://learningnetwork.cisco.com/community/certifications/ccie_routing_switching

Equipment List

You need the following hardware and software components to begin this practice lab:

- Six routers loaded with Cisco IOS Software Release 15.3T Advanced Enterprise image and the minimum interface configuration, as documented in [Table 1-1](#)
- Four 3560X switches with IOS 15.0S IP Services

Router	Model	Ethernet I/F	Serial I/F
R1	3925	2	—
R2	3925	2	1
R3	3925	2	—
R4	3925	2	—
R5	3925	2	1
R6	3925	2	—

Table 1-1 Hardware Required per Router

Setting Up the Lab

You can use any combination of routers as long as you fulfill the requirements within the topology diagram, as shown in [Figure 1-1](#). However, you should use the same model of routers because this can make life easier if you load configurations directly from those supplied with your own devices. If your router interface speeds do not match those used in this lab, consider reconfiguring the bandwidth statement accordingly to provide symmetry with the routing protocol metrics.

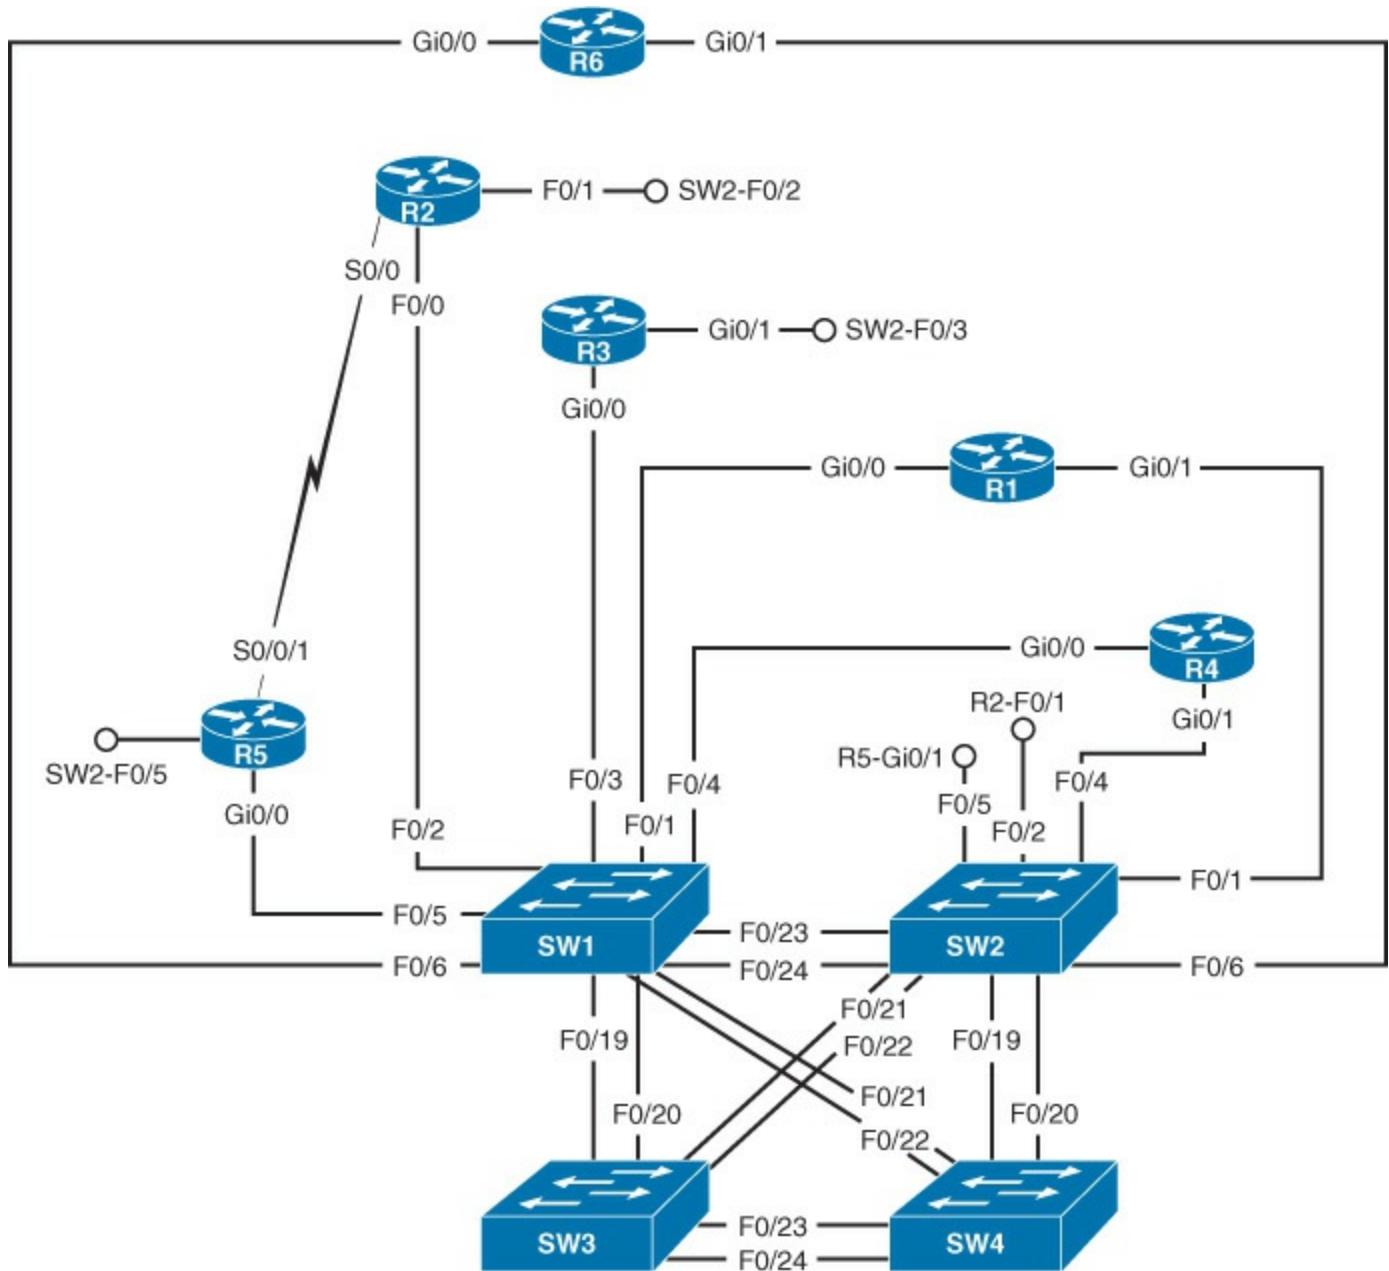


Figure 1-1 Lab Topology

Note

The CCIE Assessor topology version B is used for this lab. Additional interfaces available on the Assessor that are not required for this lab were omitted from [Figure 1-1](#). If you are not using the CCIE Assessor, use [Figure 1-1](#) and [Figure 1-4](#) to determine how many interfaces you need to complete your own topology.

Note

Notice in the initial configurations supplied that some interfaces will not have IP address preconfigured. This is because you either will not be using that interface or you need to configure this interface from default within the exercise. The initial configurations supplied should be used to preconfigure your routers and switch before the lab starts.

If your routers have different interface speeds than those used within this book, adjust the bandwidth statements on the relevant interfaces to keep all interface speeds in line. This can ensure that you do not get unwanted behavior due to differing IGP metrics.

Lab Topology

This practice lab uses the topology outlined in [Figure 1-1](#), which you must re-create with your own equipment or by simply using the CCIE Assessor.

Switch Instructions

Configure VLAN assignments from the configurations supplied or from [Table 1-2](#), with the exception of Switch2 Fa0/4 (which will be configured during the lab).

VLAN	Switch 1	Switch 2	Switch 3	Switch 4
34	Fa0/3, Fa0/4	—	—	—
45	Fa0/5	See questions	—	—
46	Fa0/6	See questions	—	—
100	—	Fa0/1	—	—
132	Fa0/1, Fa0/2	Fa0/3	—	—
200	—	Fa0/2	—	—
300	I/F VLAN 300	Fa0/5, Fa0/6, I/F VLAN 300	I/F VLAN 300	I/F VLAN 300

Table 1-2 *VLAN Assignment*

Note

Switch 2 will be configured during the actual lab questions for VLAN 45 and 46 interface Fa0/4.

Connect your switches with RJ-45 Ethernet cross-over cables, as shown in [Figure 1-2](#).

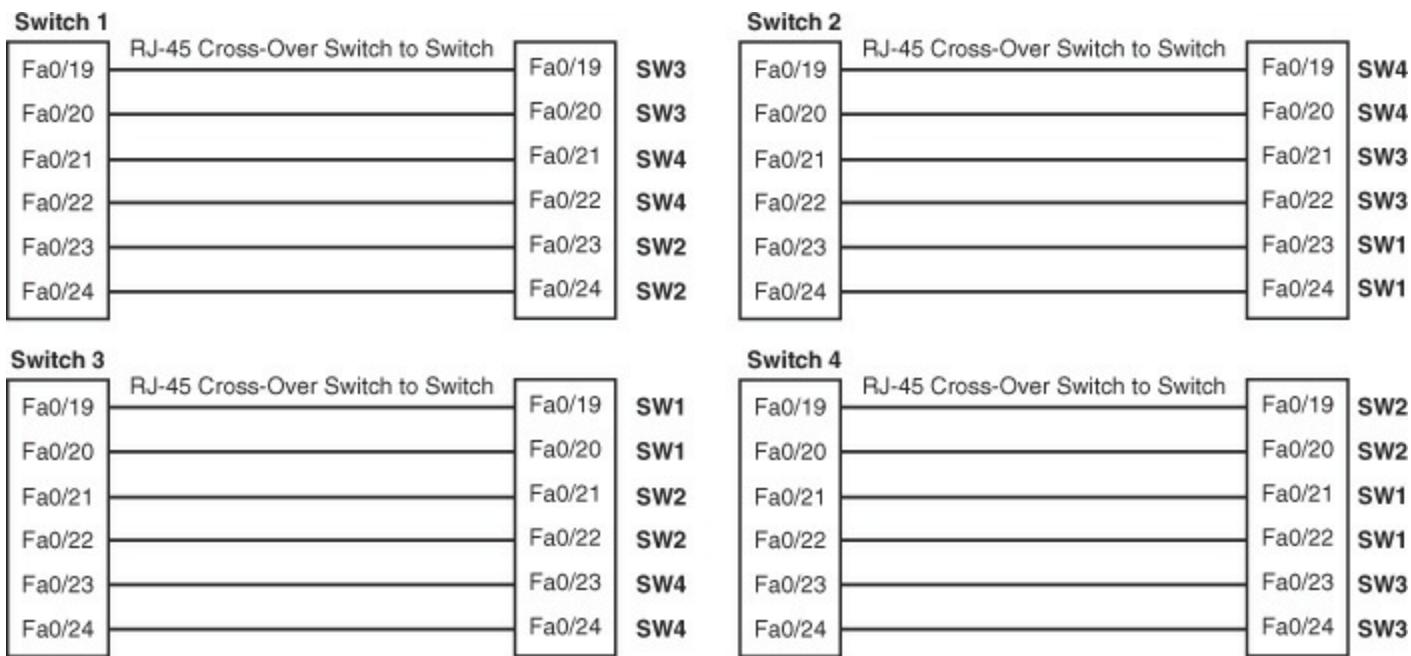


Figure 1-2 Switch Cabling

Serial Link

A preconfigured PPP back-to-back serial link exists between R2 and R5, and R2 has been configured to provide the clocking for the connection in the initial configuration files. Therefore, R2 should have the DCE serial cable and R5 the DTE serial cable for the back-to-back connectivity.

IP Address Instructions

In the real CCIE lab, the majority of your IP addresses will be preconfigured. For this exercise, however, you are required to configure your IP addresses, as shown in [Figure 1-3](#), or load the initial router configurations supplied. If you are manually configuring your equipment, ensure that you include the following loopback addresses:

R1 Lo0 120.100.1.1/24

R2 Lo0 120.100.2.1/24

R3 Lo0 120.100.3.1/24

R4 Lo0 120.100.4.1/24

R5 Lo0 120.100.5.1/24

R6 Lo0 120.100.6.1/24

SW1 Lo0 120.100.7.1/24

SW2 Lo0 120.100.8.1/24

SW3 Lo0 120.100.9.1/24

SW4 Lo0 120.100.10.1/24

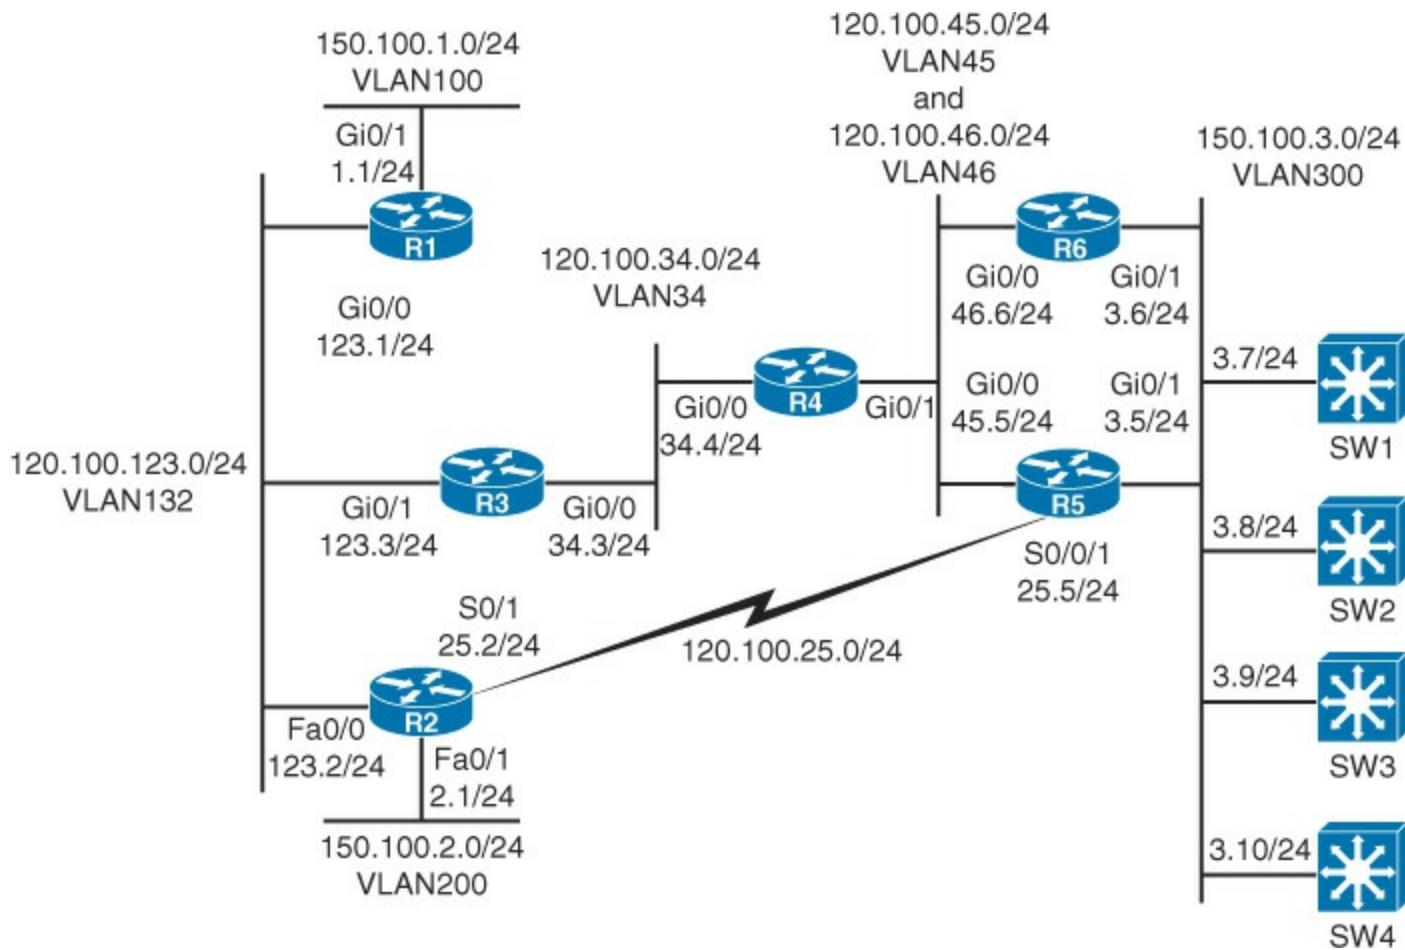


Figure 1-3 IP Addressing Diagram

Pre-Lab Tasks

- Build the lab topology as per [Figure 1-1](#) and [Figure 1-2](#).
- Configure the IP addresses on each router, as shown in [Figure 1-3](#), and add the loopback addresses. Alternatively, you can load the initial configuration files supplied if your router is compatible with those used to create this exercise. R1 requires a secondary IP address on its Gigabit Ethernet 0/1 interface for this lab; you can find details on the accompanying initial configuration for R1.

General Guidelines

- Read the whole lab before you start.
- Do not configure any static/default routes unless otherwise specified.
- Ensure full IP visibility between routers for ping testing/Telnet access to your devices (except for the switch loopback addresses, which will not be visible to the majority of your network because of the configuration tasks).
- If you find yourself running out of time, choose questions that you are confident you can answer; failing this, choose questions with a higher point rating to maximize your potential score.
- Get into a comfortable and quiet environment where you can focus for the next 8 hours.
- Take a 30-minute break midway through the exercise.
- Have available a Cisco documentation CD-ROM or access online the latest documentation from <http://www.cisco.com/cisco/web/psa/configure.html>. Note that access to this URL is

likely to be restricted within the real exam.

Note

Access only this URL, not the whole Cisco.com website (because if you are permitted to use documentation during your CCIE lab exam, it will be restricted). To save time during your lab exam, consider opening several windows with the pages you are likely to look at.

Practice Lab One

You will now answer questions in relation to the network topology, as shown in [Figure 1-4](#).

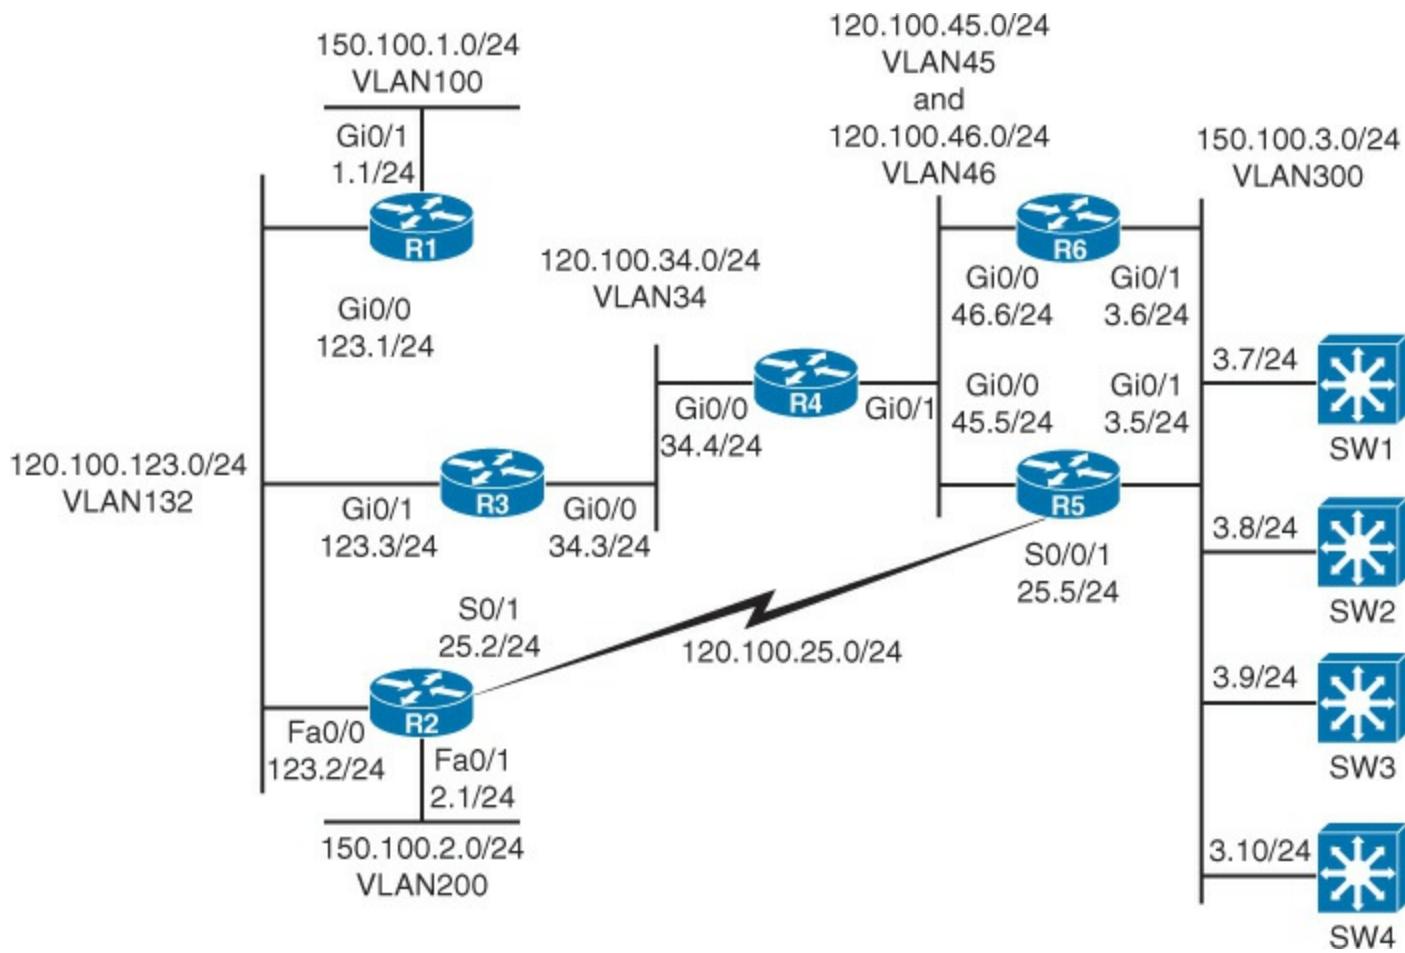


Figure 1-4 Network Topology for Practice Lab One

Section 1: LAN Switching (25 Points)

- Configure your switches as a collapsed backbone network with Switches 1 and 2 performing core and distribution functionality and Switches 3 and 4 as access switches in your topology. Switches 3 and 4 should connect only to the core switches. (2 points)
- Switch 1 and 2 should run spanning tree in 802.1w mode; Switches 3 and 4 should operate in their default spanning-tree mode. (2 points)
- Configure Switch 1 to be the root bridge and Switch 2 the secondary root bridge for VLANs 1 and 300. Ensure that Switches 3 and 4 can never become root bridges for any VLANs for which Switch 1 and Switch 2 are root bridges, by configuring only Switches 1 and 2. (2 points)
- Make sure that you fully use the available bandwidth between switches by grouping together

your interswitch links as trunks. Ensure that only dot1q and EtherChannel are supported. (3 points)

- Ensure that traffic is distributed on individual Ethernet trunks between switches based on the destination MAC address of individual flows. (2 points)
- Ensure that user interfaces, should they toggle excessively, are shut down dynamically by all switches; if they remain stable for 35 seconds, they should be reenabled. Configure Fast Ethernet Port 0/10 on each switch so that if multicast traffic is received on this port the port is automatically disabled. (2 points)
- Fast Ethernet Ports 0/11–17 will be used for future connectivity on each switch. Configure these ports as access ports for VLAN 300, which should begin forwarding traffic immediately upon connection. Devices connected to these ports will dynamically receive IP addresses from a DHCP server due to be connected to Port 0/18 on SW1. For security purposes, this is the only port on the network from which DHCP addresses should be allocated. Ensure that the switches intercept the DHCP requests and add the ingress port and VLAN and switch MAC address before sending onward to the DHCP server. Limit DHCP requests to 600 packets per minute per user port. (6 points)
- For additional security, ensure that the user ports on Switches 1–4 and 11–17 can communicate only with the network with IP addresses gained from the DHCP feature configured previously. Use a dynamic feature to ensure that the only information forwarded upon connection is DHCP request packets, and then, for additional security, any traffic that matches the DHCP IP information received from the DHCP binding. (3 points)
- R5 and R6 have been preconfigured with IP addresses on their Ethernet interfaces. Configure R4 and its associated switch port accordingly without using secondary addressing to communicate with R5 and R6. Configure R4 with an IP address of 120.100.45.4/24 to communicate with R5, and configure R4 with an IP address of 120.100.46.4/24 to communicate with R6. Configure R4 Gi0/1 and Switch 2 FE0/4 only. (3 points)

Section 2: IPv4 IGP Protocols (24 Points)

Section 2.1: OSPF

Refer to [Figure 1-5](#).

- Use a process ID of 1; all OSPF configuration where possible should not be configured under the process ID. The loopback interfaces of Routers R1, R2, and R3 should be configured to be in Area 0. R4 should be in Area 34 and R5 in Area 5. (2 points)
- No loopback networks should be advertised as host routes. (1 point)
- Ensure that R1 does not advertise the preconfigured secondary address under interface Gigabit 0/1 of 120.100.100.1/24 to the OSPF network. Do not use any filtering techniques to achieve this. (2 points)
- R5 should use the serial link within Area 5 for its primary communication to the OSPF network. If this network should fail either at Layer 1 or Layer 2, R5 should form a neighbor relationship with R4 under Area 5 to maintain connectivity. Your solution should be dynamic, ensuring that while the Area 5 serial link is operational there is no neighbor relationship between R4 and R5; however, the Ethernet interfaces of R4 and R5 must remain up. To confirm the operational status

of the serial network, ensure that the serial interface of R5 is reachable by configuration of R5. You are permitted to define neighbor statements between R5 and R4. (4 points)

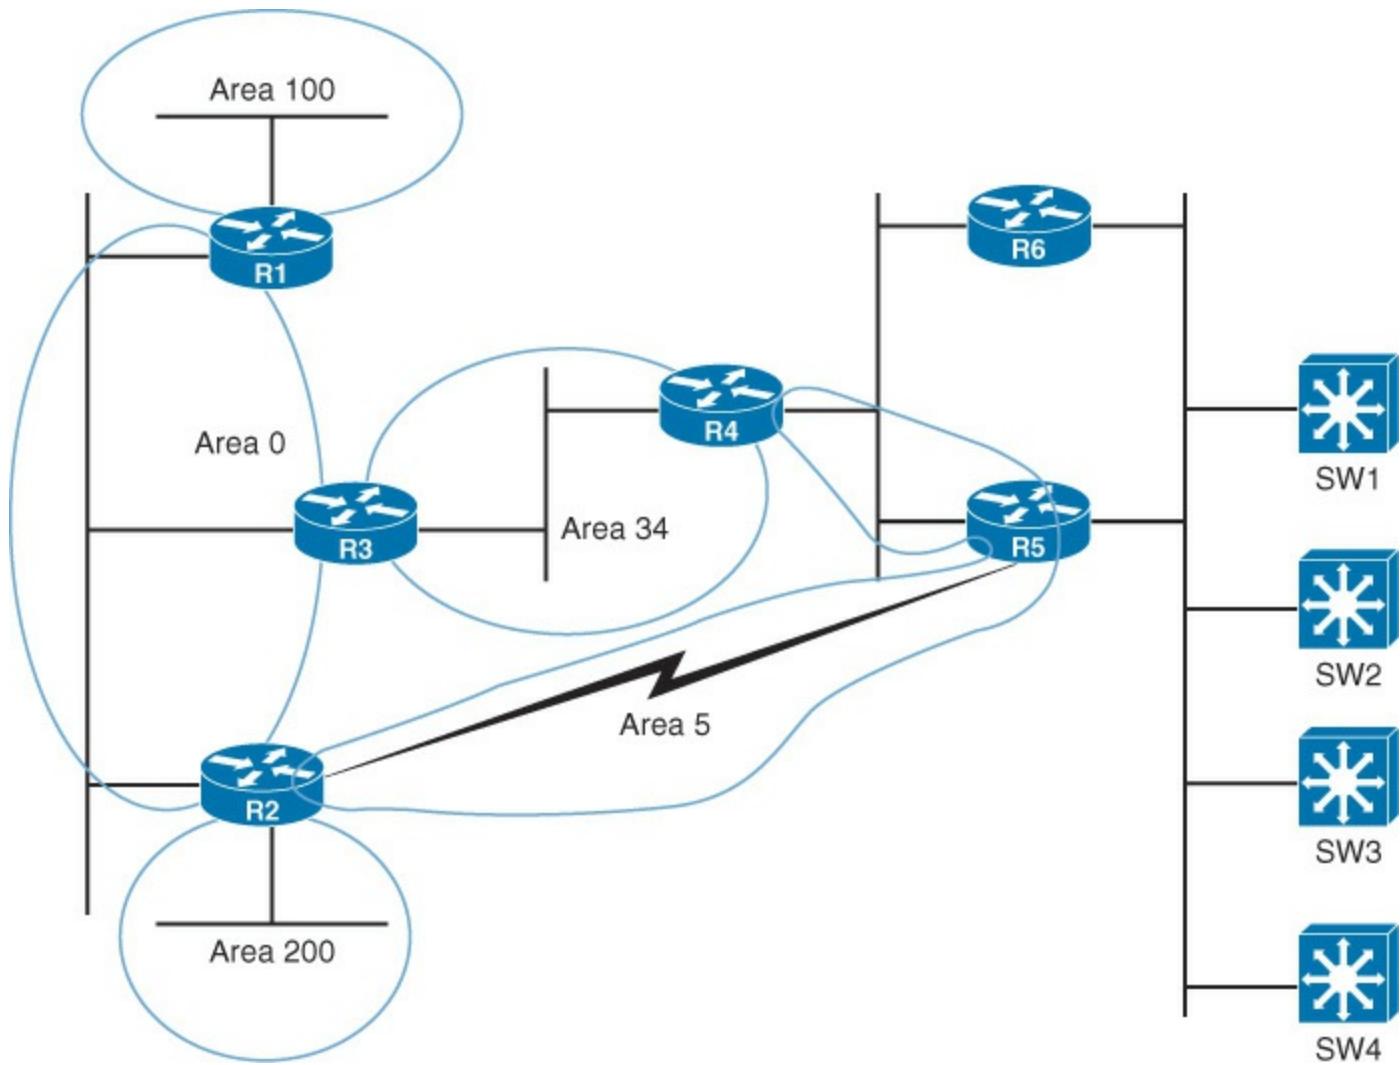


Figure 1-5 OSPF Topology

Section 2.2: EIGRP

Refer to [Figure 1-6](#).

- Configure EIGRP with an instance name of CCIE where possible using an autonomous system number of 1. The loopback interfaces of all routers and switches should be advertised within EIGRP. (2 points)
- Ensure that R4 does not install any of the EIGRP loopback routes from any of the switches into its routing table; these routes should also not be present in the OSPF network post redistribution. Do not use any route-filtering ACLs, prefix lists, or admin distance manipulation to achieve this, and perform configuration only on R4. (4 points)
- R4 will have dual equal-cost routes to VLAN 300 (network 150.100.3.0) from R5 and R6. Ensure that R4 sends traffic to this destination network to R5 instead of load sharing. If the route from R5 becomes unavailable, traffic should be sent to R6. You cannot policy route, alter the bandwidth or delay statements on R4's interfaces, or use an offset list. Perform your configuration on R4 only. Your solution should be applied to all routes received from R5 and R6, as opposed to solely the route to network VLAN 300. (4 points)

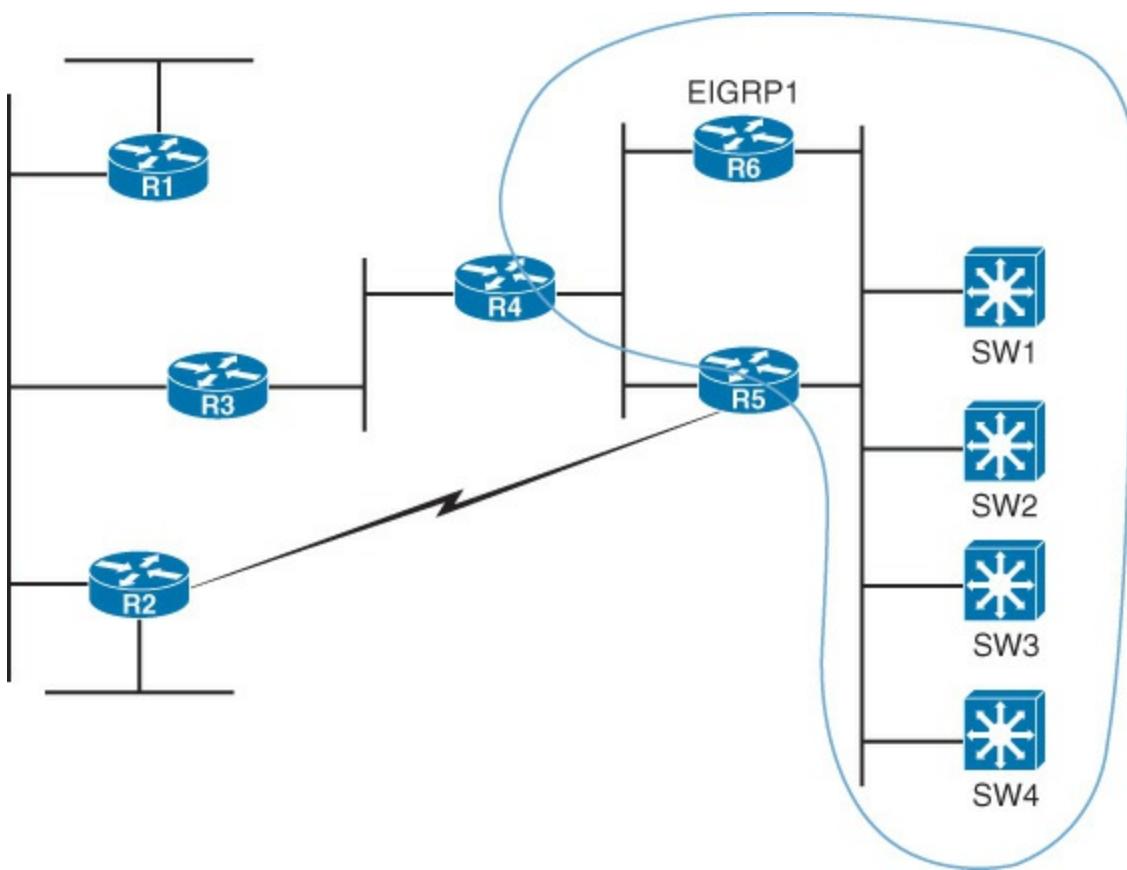


Figure 1-6 EIGRP Topology

Section 2.3: Redistribution

- Perform mutual redistribution of IGPs on R4. All routes should be accessible except for the switch loopback networks (because these should not be visible via R4 from an earlier question). EIGRP routes redistributed within the OSPF network should remain with a fixed cost of 5000 throughout the network. (3 points)
- Configure R4 to redistribute only up to five EIGRP routes and generate a system warning when the fourth route is redistributed. Do not use any access lists in your solution. (2 points)

Section 3: BGP (14 Points)

Refer to [Figure 1-7](#).

- Configure iBGP peering as follows: R1-R3, R2-R3, R6-R5, SW1-R6, and SW1-R5. Use minimal configuration and use loopback interfaces for your peering. Configure eBGP peering as follows: R3-R4, R4-R6, R4-R5, and R5-R2. Use minimal configuration and use loopback interfaces for your peering with the exception of R4 to R5. (2 points)
- Use the autonomous system numbers supplied in [Figure 1-7](#). (2 points)

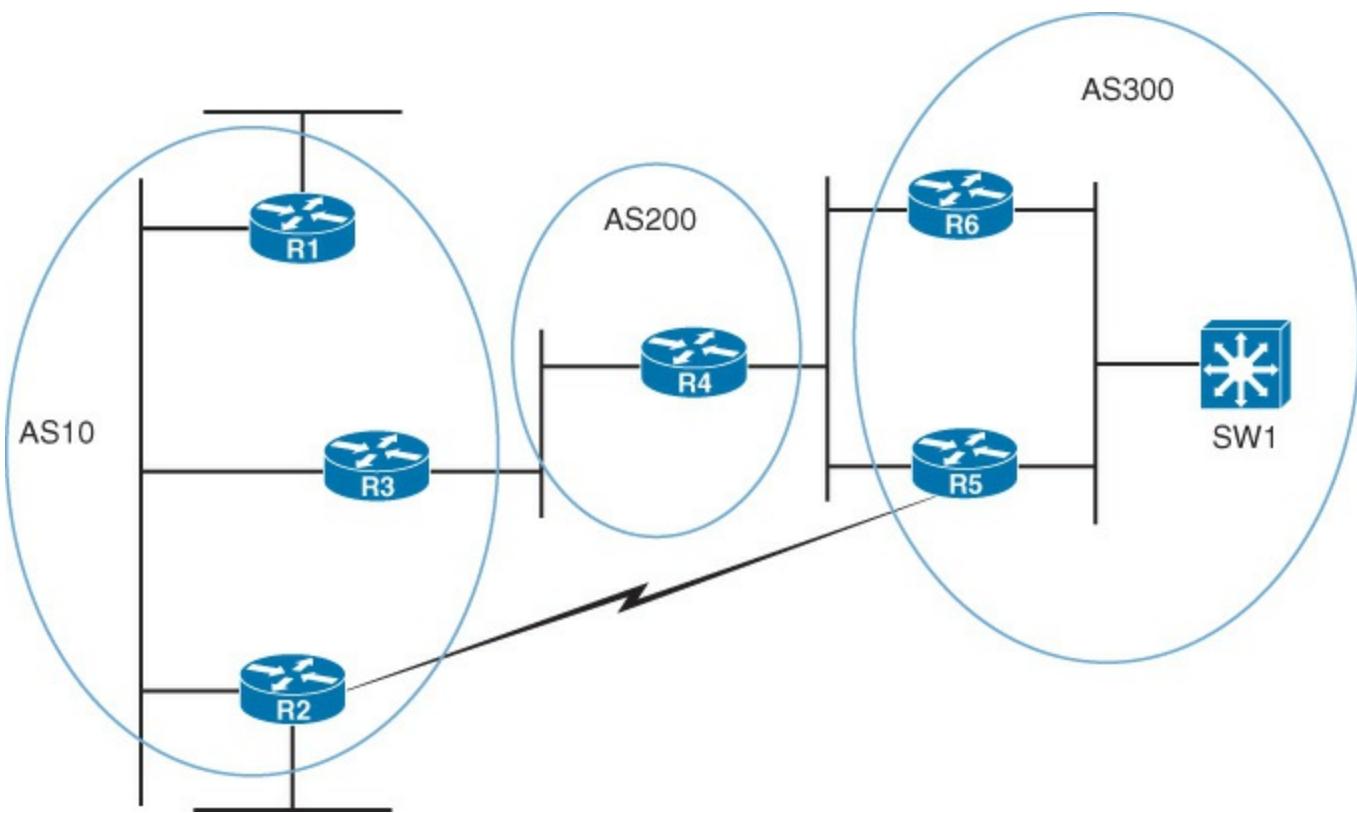


Figure 1-7 BGP Topology

- AS200 is to be used as a backup transit network for traffic between AS10 and AS300; therefore, if the serial network between R5 and R2 fails, ensure that the peering between R2 and R5 is not maintained via the Ethernet network. Do not use any ACL type restrictions or change the existing peering. (2 points)
- Configure a new loopback interface 2 on R2 of 130.100.200.1/24, and advertise this into BGP using the network command. Configure R2 in such a way that if the serial link between R2 and R5 fails, AS300 no longer receives this route. Do not use any route filtering between neighbors to achieve this. (3 points)
- Configure HSRP between R5 and R6 on VLAN300 with R5 active for .1/24. If the network 130.100.200.0/24 is no longer visible to AS300, R6 should dynamically become the HSRP active. Configure R5 to achieve this solution. (4 points)
- Configure two new loopback interfaces on R1 and R2 of 126.1.1.1/24 and 130.1.1.1/24, respectively, and advertise these into BGP using the network command. R3 should be configured to enable only BGP routes originated from R1 up to network 128.0.0.0 and from above network 128.0.0.0 originated from R2. Use only a single ACL on R3 as part of your solution. (3 points)

Section 4: IPv6 (15 Points)

Refer to [Figure 1-8](#).

- Configure IPv6 addresses on your network as follows:
 - 2007:C15:C0:10::1/64 – R1 Gi0/1
 - 2007:C15:C0:11::1/64 – R1 Gi0/0
 - 2007:C15:C0:11::2/64 – R2 FE0/0
 - 2007:C15:C0:11::3/64 – R3 Gi0/1

2007:C15:C0:12::2/64 - R2 FE0/1
 2007:C15:C0:14::2/64 – R2 S0/1
 2007:C15:C0:14::5/64 – R5 S0/0/1
 2007:C15:C0:15::3/64 – R3 Gi0/0
 2007:C15:C0:15::4/64 – R4 Gi0/0
 2007:C15:C0:16::5/64 – R5 Gi0/1
 2007:C15:C0:16::6/64 – R6 Gi0/1

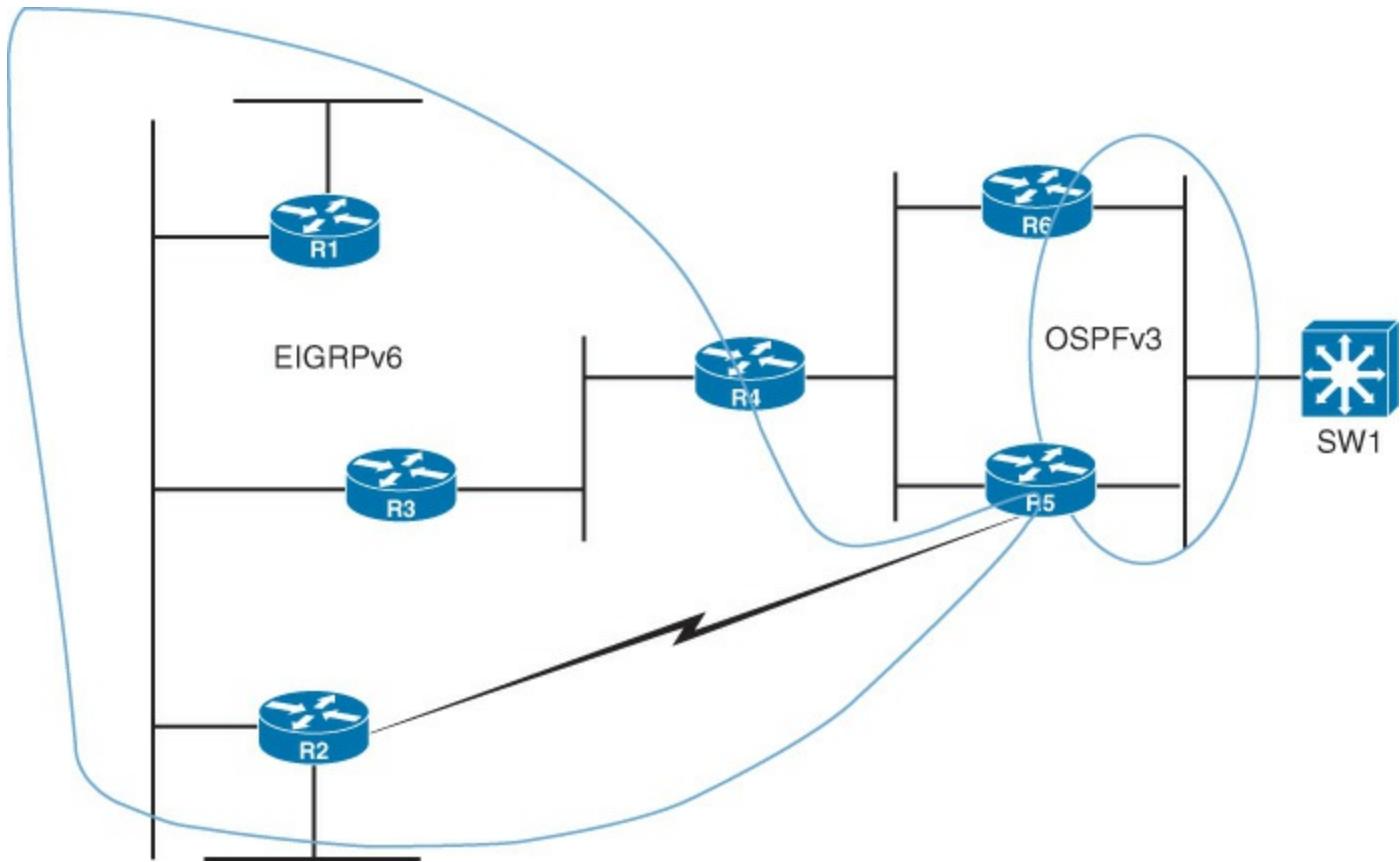


Figure 1-8 IPv6 Topology

Section 4.1: EIGRPv6

Configure EIGRPv6 under the instance of CCIE with a primary autonomous system of 1. R1 must not form any neighbor relationship with R2 on VLAN 132 (without the use of any ACL, static neighbor relationship, or multicast blocking feature). R1 must dynamically learn a default route over EIGRPv6 via R3 on VLAN 132 by which to communicate with the IPv6 network. (4 points)

Section 4.2: OSPFv3

- Configure OSPFv3 with a process ID of 1, with all OSPF interfaces assigned to Area 0. (2 points)
- The IPv6 network is deemed to be stable; therefore, reduce the number of LSAs flooded within the OSPF domain. (2 points)

Section 4.3: Redistribution

- Redistribute EIGRPv6 routes into the OSPFv3 demand (one way). EIGRPv6 routes should have a fixed cost of 5000 associated with them within the OSPF network. (1 point)

- Ensure that the OSPF3 network is reachable from the EIGRPv6 network by a single route of 2007::/16, which should be seen within the EIGRPv6 domain. Configure R5 only to achieve this. The OSPF domain should continue to receive specific EIGRPv6 subnets. (2 points)
- Ensure that if the serial link fails between the OSPF and EIGRPv6 domain, routing is still possible between R5 and R4 over VLAN 45. Do not enable EIGRPv6 on the VLAN 45 interfaces of R4 and R5. Configure R4 and R5 to achieve this, which should be considered as an alternative path only if a failure occurs. (3 points)
- Ensure that the summary route configured previously is not seen back on the routing table of R5; configure only R5 to achieve this. (1 point)

Section 5: QoS (8 Points)

- You are required to configure QoS on Switch 1 according to the Cisco QoS baseline model. Create a Modular QoS configuration for all user ports (Fast Ethernet 1–24) that facilitates the following requirements (3 points):
 1. All ports should trust the DSCP values received from their connecting devices.
 2. Packets received from the user ports with DSCP values of 48, 46, 34, 32, 24, 28, 16, and 10 should be re-marked to DSCP 8 (PHB CS1) in the event of traffic flowing above 5 Mbps on a per-port basis. This traffic could be a combination of any of the preceding DSCP values with any source/destination combination. Ensure a minimum burst value is configured above the 5 Mbps.
- Switch 1 will be connected to a new trusted domain in the future using interface Gigabit 0/1. A DSCP value received locally on SW1 of AF43 should be mapped to AF42 when destined for the new domain. (2 points)
- Configure Cisco Modular QoS as follows on R2 for the following traffic types based on their associated per-hop behavior into classes. Incorporate these into an overall policy that should be applied to the T1 interface S0/1. Allow each class the effective bandwidth as detailed, entered as a percentage. (2 points)

Class	PHB	Assigned Speed
Routing	CS6	46 Kbps
VoIP	EF	247 Kbps
Interactive Video	AF41	247 Kbps
Mission Critical Data	AF31	247 Kbps
Call-Signaling	CS3	46 Kbps
Transactional Data	AF21	216 Kbps
Network-Mgmt	CS2	46 Kbps
Bulk Data	Af11	46 Kbps
Scavenger	CS1	15 Kbps
Default	0	386 Kbps

- Configure R2 so that traffic can be monitored on the serial network with a view to a dynamic policy being generated in the future that trusts the DSCP value of traffic identified on this media. (1 point)

Section 6: Security (6 Points)

- Configure R3 to identify and discard the following custom virus; the virus is characterized by the ASCII characters Hastings_Beer within the payload and uses UDP Ports 11664 to 11666. The ID of the virus begins on the third character of the payload. The virus originated on VLAN 34. (2 points)
- An infected host is on VLAN 200 of 150.100.2.100; ensure that only within BGP AS10, traffic destined for this host is directed to null0 of each local router. You cannot use any ACLs to block traffic to this host specifically, but you can use a static route pointing to null0 for traffic destined to 192.0.2.0 /24 on routers within AS10. R2 can have an additional static route pointing to null0. Use a BGP feature on R2 to ensure that traffic to this source is blocked. Prevent unnecessary replies when traffic is passed to the null0 interface for users residing on VLAN 100. (3 points)
- To protect the control plane on router R6, configure CoPP so that IP packets with a TTL of 0 or 1 are dropped rather than processed, with a resulting ICMP redirect sent to the originator. (1 point)

Section 7: Multicast (4 Points)

- Configure routers R1, R2, R3, and R4 for IPv4 Multicast; configure R3 to send multicast advertisements of its own time by use of NTP sourced from interface Gig 0/0. Configure PIM sparse mode on all required interfaces. R3 should also be used to advertise its own gigabit interface IP address as an RP. R3 should also advertise the IP address you are using for the NTP advertisements that will be 224.0.1.1. Do not use the command ntp server in any configurations. Routers R1, R2, and R4 should all show a clock synchronized to that of R3. (4 points)

IP Services (4 Points)

- Configure the following commands on router R1:
aaa new-model
logging buffered
logging 120.100.99.1
- Configure a policy on router R1 so that if a user tries to remove AAA services or disable logging via the CLI that a syslog message of UNAUTHORIZED-COMMAND-ENTERED is generated. The policy should ensure that neither command is executed and should consist of a single-line command for the CLI pattern detection. The policy and CLI should run asynchronously. The policy should also generate an email from the router to a mail server residing on IP address 120.100.99.2 (to security@lab-exam.net from eem@lab-exam.net subject “User-Issue” with the message body consisting of details of who was logged on the time either of the commands were entered). (4 points)

“Ask the Proctor”

Note

This section should be used only if you require clues to complete the questions. In the actual CCIE lab, the proctor will not enter into any discussions about the questions or answers. He or she will be present to ensure that you do not have problems with the lab environment and to maintain the timing element of the exam.

Section 1: LAN Switching

- Q.** Do you want me to configure the collapsed backbone network by manipulating spanning tree to ensure that Switch 1 and Switch 2 are the cores for each VLAN in use?
- A.** You are requested to configure root bridges in a later question.
- Q.** All the switches are already connected, so I can't change this unless I shut down some of the connections between switches. Is this acceptable?
- A.** Yes.
- Q.** If I explicitly configure Switches 1 and 2 as root bridges, surely this will never enable Switches 3 and 4 to become root bridges.
- A.** No, it won't. If a superior BPDU is received on ports connecting to Switches 3 and 4 from Switches 1 and 2, Switches 3 and 4 could become root bridges; use a feature that effectively ignores a superior BPDU if received.
- Q.** Do you want me to disable spanning tree down to Switches 3 and 4? Is this acceptable?
- A.** No, spanning tree must remain in operation.
- Q.** Can I configure a MAC address type access list to block all multicast at Layer 2?
- A.** No, this wouldn't disable the port if multicast traffic was present on it; look for a dynamic solution that does not require an ACL.
- Q.** Can I configure the switchport block multicast command?
- A.** No, this would block the traffic but wouldn't disable the port.
- Q.** Would you like me to VLAN load balance to utilize bandwidth?
- A.** No, the question directs you how to use the trunks.
- Q.** Would you like me to configure Switch 1 to allocate DHCP addresses?
- A.** No, the question relates to a fictitious DHCP server that would be connected to Fa0/18 on Switch 1.
- Q.** Can I manipulate a helper-address function to answer the DHCP question by using ACLs?
- A.** No, use a recognized DHCP security-related solution.
- Q.** Can I configure port security to bind my MAC addresses?
- A.** No, use a feature that complements your DHCP solution.

Q. Can I just configure R4 to trunk to Switch 2 and have a subinterface in both VLAN 45 and VLAN 46?

A. Yes.

Q. I've configured my trunk on Switch 2 to R4 and I can't ping between R4 and R5; similarly, I can't ping between R4 and R6. Is there anything else I need to do?

A. Remember that the switches are in VTP transparent mode; you might want to check that Switch 2 has the required VLANs configured to enable propagation within your switched network.

Section 2: IPv4 IGP Protocols

Section 2.1: OSPF

Q. I am used to configuring OSPF under the process. Surely this is the only place I can configure the parameters.

A. There have been recent advances in OSPF enabling you to configure it purely under specific areas of the router, rather like with IPv6. Take a look at the commands available to you under the interfaces.

Q. My neighbor relationship is down over the serial network. I notice I have different OSPF network types preconfigured. Can I change these?

A. No, use an alternative method of bringing the interface parameters back into line.

Q. My secondary address is advertised automatically under OSPF. Can I use a distribute list or prefix type list to block it?

A. No, use an OSPF feature to disable the advertisement of this secondary address.

Q. I've attempted to form a neighbor relationship with R4 from R5 using a backup interface. Is this okay?

A. No, the question states that your solution should cater for either Layer 1 or Layer 2 failures and that the Ethernet should remain up. Backup interfaces would be fine for a Layer 1 failure but not for a Layer 2 type issue if you had problems with PPP that caused neighbor failures over the serial network. This feature would also ensure that the Ethernet network would be down until the backup interface is activated.

Q. How about an OSPF demand circuit between R4 and R5?

A. No, this would involve a neighbor relationship being maintained. You need to allow the neighbor relationship to be formed only if a failure condition occurs.

Q. Can I use BFD between R4 and R5?

A. No, this might aid in failure detection, but it does not meet the objectives of the question.

Q. To confirm the operation status of R5's serial interface, can I just ping it?

A. You can use ICMP, but you need to ensure that your solution is dynamic.

Q. If I use IP SLA to automatically ping R5 to check the status, is this okay?

A. Yes.

Q. Okay, I have IP SLA running, but I'm stuck. Is this anything to do with tracking the response to the ping?

A. Yes.

Q. How about if I use policy routing with the next hop based on the tracking status?

A. This is fine; just remember that this traffic will be based locally on the router when applying any policies.

Q. I've worked out how to do this and managed to get a neighbor up when the serial network fails, but my OSPF connectivity is still not perfect through the Ethernet. Is this normal?

A. Not if you have configured correctly. Take a look at your topology and areas. Something might have changed when R5 connects over the Ethernet.

Section 2.2: EIGRP

Q. I can't configure my switches with an EIGRP instance name. Is the legacy method with just an autonomous system acceptable for the switches?

A. Yes, this is fine and in accordance with the question.

Q. If I advertise my loopbacks into EIGRP, won't that mean that R4 and R5 will have their loopbacks advertised by both OSPF and EIGRP?

A. Yes, this is fine and in accordance with the question.

Q. To stop R4 from receiving the switch loopbacks, can I stop advertising them from the switches?

A. No, you should use a feature on R4 to block them.

Q. Can I use a neighbor prefix list to block the loopbacks?

A. No, you cannot use any type of ACLs or prefix lists.

Q. I've noticed when I look at the specific loopback routes that they have a hop count associated with them. It's unusual to associate hop counts with EIGRP, but can I block routes based on their hop count?

A. Yes.

Q. If I can't change the bandwidth and delay on R4, can I use a route map to manipulate the EIGRP K values associated on a per-neighbor basis?

A. Yes.

Section 2.3: Redistribution

Q. Do you require a distribute list to block the switch loopbacks from entering the OSPF domain?

A. No, you should have blocked these from entering your IP routing table within R4 previously, so additional blocking would not be required.

Q. I have only one redistribution point, and there is no benefit in creating filtering to protect against potential routing loops between protocols. Is this acceptable?

A. Yes, in this scenario, this would be superfluous.

Q. Can I use a route map to enable five specific EIGRP routes to be redistributed into OSPF?

A. No, the question doesn't guide you to redistribute specific routes. Use a more general method of allowing a specific number of routes.

Section 3: BGP

Q. Is it okay to disable autosynchronization in BGP?

A. You need to determine whether you need this feature on or off. Remember that you should have synchronization on only when you are fully redistributing between BGP and your IGP.

Q. Do you want me to configure eBGP multihop but limit it to a value of 2 on R3 for a TTL security check?

A. There is a specific security configuration feature within BGP to perform the TTL check.

Q. If I use the TTL security hops with a value of 2, is this all you are looking for?

A. You must ensure that your peering still works effectively between R3 and R4 when you have configured this feature.

Q. I find that when the serial network fails, my neighbor relationship is still maintained between R2 and R5. This is because the loopback routes are still available over the alternative path through the network. Can I block my loopbacks or policy route at some point to effectively break the peering?

A. You do need to effectively break the peering, but there is a much simpler method of achieving this that still maintains unaltered communication between R2 and R5. Think about what you need to configure when you have EBGP peers.

Q. I might have been a little generous with my original multihop value between R2 and R5. If I reduce this to a TTL of 2, I can break the peering. Is this okay?

A. Yes.

Q. I think I can stop the loopback on R2 being advertised by using the community value of **no-export**, but if I enable this to R2, it wouldn't make it to R5 even when the serial network is working?

A. Correct, it wouldn't be advertised to R5 AS300 from R2. Just think about whether R2 is the best place to send the community to originally.

Q. For the HSRP question, is this some form of conditional advertising?

A. No, the clue is in the question; just find a way of tracking the BGP route and manipulate the HSRP process.

Q. If I enable IP SLA to track a route in the routing table, can I use this to control HSRP?

A. Yes.

Q. You haven't told me what address I should use for HSRP. Is it okay to use the first address in the subnet?

A. Yes.

Q. I have configured my two new loopbacks. Can I use two route maps inbound from R1 and R2 both pointing to different ACLs so that each route map calls only one ACL?

A. No, you still have two ACLs.

Q. Can I set community values on the routes and match on these using a single ACL?

A. No, you are instructed to use an ACL; your solution would require additional configuration.

Q. Can I use a prefix list to achieve this?

A. No, you are instructed to use an ACL.

Q. So, I need an ACL with a mask suitable for both ranges?

A. Not necessarily. You would need to match only one requirement on the permit functionality; the other could be met by deny.

Section 4: IPv6

Q. Should I use the eui-64 address format when configuring my addresses?

A. No, if these were required, the question would have instructed you to use them.

Q. Can I form an EIGRPv6 neighbor relationship between R1 and R3 and also R3 and R2?

A. Yes.

Q. Can I use different autonomous systems and then redistribute at R3?

A. Yes.

Q. You are not requesting mutual redistribution between EIGRPv6 and OSPFv3. How will my EIGRPv6 domain communicate with the OSPFv3 domain?

A. This issue is addressed in the following task.

Q. If I can't use EIGRPv6 directly on VLAN 45 between R4 and R5, can I configure OSPFv3 on VLAN 45?

A. No, find a way to still run EIGRPv6 between routers without enabling it on the physical interfaces.

Q. Can I tunnel between R4 and R5?

A. Yes.

Section 4.3: Redistribution

Q. I have redistributed EIGRPv6 into OPSFv3 on R5, which is the only suitable location, and noticed that in my OSPFv3 domain I do not see the IPv6 network configured on the serial network between R2 and R5. Is this okay?

A. No, this network should be advertised to the OPSFv3 domain. Use a feature within the OPSFv3 process as you would to overcome this if this were IPv4 redistribution.

Q. Can I redistribute a static IPv6 route on R5 into RIPng for 2007::/16?

A. No, static routes are permitted unless specified. What would you do if this were IPv4?

Q. If I can't enable EIGRPv6 on VLAN 45 between R4 and R5, can I enable OSPFv3?

A. No, this would also require you to perform redistribution at this point?

Q. How about tunneling again and enabling EIGRPv6 over the tunnel. Is this okay?

A. Yes.

Q. I have created my tunnel and found that this is now the primary route rather than an alternative path. Can I perform some kind of backup interface to make this come up only if a failure occurs on the serial link?

A. No, you haven't been given sufficient information to make this judgment. This approach would also break your IPv4 network. Think why the Ethernet path is preferred and manipulate it.

Q. Can I use a prefix list to block the summary and permit all other IPv6 routes?

A. Yes, this is fine.

Section 5: QoS

Q. Can I just trust DSCP on my physical ports?

A. No, this should be completed as part of your policy.

Q. Shall I rate limit my ports to 5M on a per-port basis?

A. No, this should be completed as part of your policy.

Q. You haven't indicated what the minimum burst size should be; is this correct?

A. Yes, just use the available limits within the command options.

Q. I believe I can use a DSCP mutation map to convert the DSCP values for the future, but the command won't take the values AF43 and AF42.

A. No, it won't because these are Assured Forwarding values. You need to convert these to DSCP values. Search your documentation CD or available Cisco.com pages.

Q. I am trying to assign bandwidth within my class with the speeds supplied, but I can see only a percentage option. Is this correct?

A. Yes, you must do some math. You are supplied with the information you require and just need to remember how fast a T1 line is.

Section 6: Security

Q. Can I use a route map and ACLs to identify the traffic by port number?

A. No, this would identify the UDLD traffic but not the virus payload as per the question. Investigate the options open to you with NBAR.

Q. Can I policy route traffic destined to the infected host to null0?

A. No, you must use a BGP-related feature.

Q. A static route for 192.0.2.0/24 won't have any bearing on traffic destined to the infected host.

Why is this relevant?

- A. Think about the way BGP works. It's the only routing protocol where you don't need to be directly connected to form a neighbor relationship; therefore, you transport next-hop information with your updates.
- Q. I have configured CoPP on R6 and seem to have lost all my routes. Is this expected behavior? Do you want me to fix this as part of the CoPP question?
- A. If you have lost your routes, think about why this has happened. Yes, provide a fix; otherwise, you would lose points in other sections.

Section 7: Multicast

- Q. If I can't configure ntp server on R1, R2, and R4, there won't be a way I can get these routers to peer with R3. Is this correct?
- A. Yes, you don't need to specifically peer with R3 as the server. However, you should aim to receive the NTP stream that R3 should be configured to multicast.
- Q. Do you want me to create and announce the group 224.0.1.1 on R3?
- A. Yes.

Section 8: IP Services

- Q. Based on the email address, I guess this is an EEM question?
- A. Correct.
- Q. Do you need me to set up a route to 120.100.99.0/24?
- A. No.
- Q. I can't get both commands onto a single CLI pattern event. Is it okay to configure two?
- A. No, you are directed to configure a single CLI pattern event command that will pick up either command.

Lab Debrief

This section analyzes each question, showing you what was required and how to achieve the desired results. You should use this section to produce an overall score for this practice lab.

Section 1: LAN Switching (25 Points)

- Configure your switches as a collapsed backbone network with Switches 1 and 2 performing core and distribution functionality and Switches 3 and 4 as access switches in your topology. Switches 3 and 4 should connect to only the core switches. (2 points)

This is a simple start to the exercise. The switches are fully meshed to begin with. To create a collapsed backbone topology, the core switches should be connected together, and each access switch should be dual-homed to the core switches. The only switches that should not connect directly to each other are the access switches (SW3 and SW4). By shutting down the interfaces between SW3 and SW4, you create the required topology. If you have configured this correctly, as shown in [Example 1-1](#), you have scored 2 points. Even though the resulting topology is not looped at this stage, you can

verify route bridge assignment by using the **show spanning tree root** command.

Example 1-1 SW3 and SW4 Configuration

[Click here to view code image](#)

```
SW3(config)# interface range fastethernet 0/23-24
SW3(config-if-range)# shut

SW4(config)# interface range fastethernet 0/23-24
SW4(config-if-range)# shut
```

- Switch 1 and 2 should run spanning tree in 802.1w mode. Switches 3 and 4 should operate in their default spanning-tree mode. (2 points)

802.1w is Rapid Spanning Tree, which is backward compatible with the switches' default (PVST). So, if you configure Switches 1 and 2 into Rapid Spanning Tree mode, spanning tree can still operate effectively with Switches 3 and 4. If you have configured this correctly, as shown in [Example 1-2](#), you have earned another 2 points.

Example 1-2 SW1 and SW2 Configuration

[Click here to view code image](#)

```
SW1(config)# spanning-tree mode rapid-pvst
SW2(config)# spanning-tree mode rapid-pvst
```

- Configure Switch 1 to be the root bridge and Switch 2 the secondary root bridge for VLANs 1 and 300. Ensure that Switches 3 and 4 can never become root bridges for any VLANs for which Switch 1 and Switch 2 are root bridges by configuring only Switches 1 and 2. (2 points)

This is a straightforward question for the core switches. The root bridge prioritization root guard is configured on the ports that connect Switches 1 and 2 to Switches 3 and 4. This ensures that if a superior BPDU is received on these ports, it is ignored. If you have configured this correctly, as shown in [Example 1-3](#), you have 2 points.

Example 1-3 SW1 and SW2 Root Bridge Configuration

[Click here to view code image](#)

```
SW1(config)# spanning-tree vlan 1 root primary
SW1(config)# spanning-tree vlan 300 root primary
SW1(config-if)# interface fastethernet 0/19
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/20
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/21
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/22
SW1(config-if)# spanning-tree guard root
```

```
SW2(config)# spanning-tree vlan 1 root secondary
SW2(config)# spanning-tree vlan 300 root secondary
SW2(config-if)# interface fastethernet 0/19
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/20
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/21
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/22
SW2(config-if)# spanning-tree guard root
```

-
- Make sure that you fully use the available bandwidth between switches by grouping your interswitch links as trunks. Ensure that only dot1q and EtherChannel are supported. (3 points)

This is another straightforward question for all switches to create EtherChannels between devices. Using the command **channel-group n mode on** under the physical interfaces ensures that only EtherChannel is supported, as opposed to Port Aggregation Protocol (PAgP) or Link Aggregation Control Protocol (LACP), and dot1q is the trunking protocol. For Layer 2 EtherChannels, you do not have to create a port-channel interface first by using the **interface port-channel** configuration command before assigning a physical port to a channel group. You can use the channel-group interface configuration command that automatically creates the port-channel interface, although a manual port channel configuration has been shown here for clarity. Remember that now that you have EtherChannels between switches, you will need to configure root guard on these interfaces to ensure that Switches 3 and 4 cannot become root bridges. This is over and above the physical interface configuration completed previously. If you have configured this correctly, as shown in [Example 1-4](#), you have scored 3 points.

Example 1-4 Switch 1, 2, 3, and 4 EtherChannel Configuration

[Click here to view code image](#)

```
SW1(config-if)# interface range fastethernet0/19-20
SW1(config-if)# channel-group 1 mode on
SW1(config-if)# interface range fastethernet0/21-22
SW1(config-if)# channel-group 2 mode on
SW1(config-if)# interface range fastethernet0/23-24
SW1(config-if)# channel-group 3 mode on
SW1(config-if)# interface Port-channel1
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface Port-channel2
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface Port-channel3
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
```

```
SW2(config-if)# interface range fastethernet0/19-20
```

```

SW2(config-if)# channel-group 1 mode on
SW2(config-if)# interface range fastethernet0/21-22
SW2(config-if)# channel-group 2 mode on
SW2(config-if)# interface range fastethernet0/23-24
SW2(config-if)# channel-group 3 mode on
SW2(config-if)# interface Port-channel1
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk
SW2(config-if)# interface Port-channel2
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk
SW2(config-if)# interface Port-channel3
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk

```

```

SW3(config-if)# interface range fastethernet0/19-20
SW3(config-if)# channel-group 1 mode on
SW3(config-if)# interface range fastethernet0/21-22
SW3(config-if)# channel-group 2 mode on
SW3(config-if)# interface Port-channel1
SW3(config-if)# switchport trunk encapsulation dot1q
SW3(config-if)# switchport mode trunk
SW3(config-if)# interface Port-channel2
SW3(config-if)# switchport trunk encapsulation dot1q
SW3(config-if)# switchport mode trunk

```

```

SW4(config-if)# interface range fastethernet0/19-20
SW4(config-if)# channel-group 1 mode on
SW4(config-if)# interface range fastethernet0/21-22
SW4(config-if)# channel-group 2 mode on
SW4(config-if)# interface Port-channel1
SW4(config-if)# switchport trunk encapsulation dot1q
SW4(config-if)# switchport mode trunk
SW4(config-if)# interface Port-channel2
SW4(config-if)# switchport trunk encapsulation dot1q
SW4(config-if)# switchport mode trunk

```

SW1# show interfaces port-channel 1 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW1# show interfaces port-channel 2 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW1# show interfaces port-channel 3 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po3		connected	trunk	a-full	a-100	

SW1# show etherchannel summary

Number of channel-groups in use: 3
Number of aggregators: 3

Group	Port-channel	Protocol	Ports
-------	--------------	----------	-------

1	Po1 (SU)	-	Fa0/19 (P)	Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P)	Fa0/22 (P)
3	Po3 (SU)	-	Fa0/23 (P)	Fa0/24 (P)

SW2# **show interfaces port-channel 1 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW2# **show interfaces port-channel 2 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW2# **show interfaces port-channel 3 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po3		connected	trunk	a-full	a-100	

SW2# **show etherchannel summary**

Number of channel-groups in use: 3

Number of aggregators: 3

Group Port-channel Protocol Ports

1	Po1 (SU)	-	Fa0/19 (P)	Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P)	Fa0/22 (P)
3	Po3 (SU)	-	Fa0/23 (P)	Fa0/24 (P)

SW3# **show interface port-channel 1 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW3# **show interface port-channel 2 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW3# **show etherchannel summary**

Number of channel-groups in use: 2

Number of aggregators: 2

Group Port-channel Protocol Ports

1	Po1 (SU)	-	Fa0/19 (P)	Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P)	Fa0/22 (P)

SW4# **show interface port-channel 1 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW4# **show interface port-channel 2 status**

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW4# **show etherchannel summary**

Number of channel-groups in use: 2

Number of aggregators: 2

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)

- Ensure that traffic is distributed on individual Ethernet trunks between switches based on the destination MAC address of individual flows. (2 points)

A common problem with EtherChannels is traffic not being distributed equally among the physical interfaces. Configuring channel load balancing based on the destination MAC address of an individual flow is just one method available to distribute traffic. If you have configured this correctly, as shown in [Example 1-5](#), you have scored 2 points.

Example 1-5 Switch 1, 2, 3, and 4 EtherChannel Load-Balancing Configuration

[Click here to view code image](#)

```
SW1(config)# port-channel load-balance dst-mac
SW2(config)# port-channel load-balance dst-mac
SW3(config)# port-channel load-balance dst-mac
SW4(config)# port-channel load-balance dst-mac

SW1# show etherchannel load-balance
EtherChannel Load-Balancing Operational State (dst-mac):
Non-IP: Destination MAC address
  IPv4: Destination MAC address
  IPv6: Destination IP address
```

- Ensure that user interfaces, if they toggle excessively, are shut down dynamically by all switches; if they remain stable for 35 seconds, they should be reenabled. Configure Fast Ethernet Port 0/10 on each switch so that if multicast traffic is received on this port, the port is automatically disabled. (3 points)

Interfaces that flap can cause problems in a network. Toggling would usually indicate a problem such as a faulty connecting network interface card (NIC) or faulty cable. Placing the ports into error disable is a way to stabilize the environment. To disable a port when multicast traffic is present, you need to configure storm control with the multicast option set to 0. If you have configured this correctly, as shown in [Example 1-6](#), you have scored 3 points.

Example 1-6 Switch 1, 2, 3, and 4 Configuration

[Click here to view code image](#)

```
SW1(config)# errdisable recovery cause link-flap
SW1(config)# errdisable recovery interval 35
SW1(config)# interface fastethernet 0/10
SW1(config-if)# storm-control multicast level 0
SW1(config-if)# storm-control action shutdown
```

```

SW2(config)# errdisable recovery cause link-flap
SW2(config)# errdisable recovery interval 35
SW2(config)# interface fastethernet 0/10
SW2(config-if)# storm-control multicast level 0
SW2(config-if)# storm-control action shutdown

SW3(config)# errdisable recovery cause link-flap
SW3(config)# errdisable recovery interval 35
SW3(config)# interface fastethernet 0/10
SW3(config-if)# storm-control multicast level 0
SW3(config-if)# storm-control action shutdown

SW4(config)# errdisable recovery cause link-flap
SW4(config)# errdisable recovery interval 35
SW3(config)# interface fastethernet 0/10
SW3(config-if)# storm-control multicast level 0
SW3(config-if)# storm-control action shutdown

```

- Fast Ethernet ports 0/11–17 will be used for future connectivity on each switch. Configure these ports as access ports for VLAN 300, which should begin forwarding traffic immediately upon connection. Devices connected to these ports will dynamically receive IP addresses from a DHCP server due to be connected to port 0/18 on SW1. For security purposes, this is the only port on the network from which DHCP addresses should be allocated. Ensure that the switches intercept the DHCP requests and add the ingress port and VLAN and switch MAC address before sending forward to the DHCP server. Limit DHCP requests to 600 packets per minute per user port. (6 points)

This is a Dynamic Host Control Protocol (DHCP) snooping question. This is a useful security feature that protects the network from rogue DHCP servers. When the DHCP option-82 feature is enabled on the switch with the command **ip dhcp snooping information option**, a subscriber is identified by the switch port through which it connects to the network and by its MAC address. DHCP snooping also facilitates a rate-limiting feature for DHCP requests to prevent a DHCP denial of service by excessive false requests from a host, which would have the “gobbler effect” of requesting numerous leases from the same port. The question includes a couple of points that could easily be overlooked if you are suffering from exam pressure, namely that the ports are required to be configured with **switchport host** (or by configuring portfast) to set the port mode to access and to forward immediately. The rate limiting is configured in packets per second, not per minute as implied, so you need to pay attention to detail. If you have configured this correctly, as shown in [Example 1-7](#), you have scored 6 points.

Example 1-7 Switch 1, 2, 3, and 4 DHCP Snooping Configuration

[Click here to view code image](#)

```

SW1(config)# ip dhcp snooping
SW1(config)# ip dhcp snooping vlan 300
SW1(config)# ip dhcp snooping information option
SW1(config)# int fastethernet 0/18
SW1(config-if)# ip dhcp snooping trust
SW1(config)# interface range fastethernet 0/11-17

```

```

SW1(config-if-range)# ip dhcp snooping limit rate 10
SW1(config)# interface range fastethernet 0/11-18
SW1(config-if-range)# switchport host
SW1(config-if-range)# switchport access vlan 300

SW2(config)# ip dhcp snooping
SW2(config)# ip dhcp snooping vlan 300
SW2(config)# ip dhcp snooping information option
SW2(config)# interface range fastethernet 0/11-17
SW2(config-if-range)# ip dhcp snooping limit rate 10
SW2(config-if-range)# switchport host
SW2(config-if-range)# switchport access vlan 300

SW3(config)# ip dhcp snooping
SW3(config)# ip dhcp snooping vlan 300
SW3(config)# ip dhcp snooping information option
SW3(config)# interface range fastethernet 0/11-17
SW3(config-if-range)# ip dhcp snooping limit rate 10
SW3(config-if-range)# switchport host
SW3(config-if-range)# switchport access vlan 300

SW4(config)# ip dhcp snooping
SW4(config)# ip dhcp snooping vlan 300
SW4(config)# ip dhcp snooping information option
SW4(config)# interface range fastethernet 0/11-17
SW4(config-if-range)# ip dhcp snooping limit rate 10
SW4(config-if-range)# switchport host
SW4(config-if-range)# switchport access vlan 300

SW1# sh ip dhcp snooping
Switch DHCP snooping is enabled
DHCP snooping is configured on following VLANs:
300
Insertion of option 82 is enabled
  circuit-id format: vlan-mod-port
  remote-id format: MAC
Option 82 on untrusted port is not allowed
Verification of hwaddr field is enabled
Interface          Trusted      Rate limit (pps)
-----            -----
fastethernet0/11    no           10
fastethernet0/12    no           10
fastethernet0/13    no           10
fastethernet0/14    no           10
fastethernet0/15    no           10
fastethernet0/16    no           10
fastethernet0/17    no           10
fastethernet0/18    yes          unlimited

```

- For additional security ensure that the user ports on Switches 1–4 and 11–17 can communicate only with the network with IP addresses gained from the DHCP feature configured previously. Use a dynamic feature to ensure that the only information forwarded upon connection is DHCP request packets and then, for additional security, any traffic that matches the DHCP IP information received from the DHCP binding. (3 points)

A complementary feature to DHCP snooping is IP Source Guard. This feature binds the information received from the DHCP address offered and effectively builds a dynamic VACL on a per-port basis

to enable only source traffic matched from the DHCP offer to ingress the switch port for additional security. If you have configured this correctly, as shown in [Example 1-8](#), you have scored 3 points.

Example 1-8 Switch 1, 2, 3, and 4 IP Source Guard Configuration

[Click here to view code image](#)

```
SW1(config)# interface range fast 0/11-17
SW1(config-if-range)# ip verify source

SW2(config)# interface range fast 0/11-17
SW2(config-if-range)# ip verify source

SW3(config)# interface range fast 0/11-17
SW3(config-if-range)# ip verify source

SW4(config)# interface range fast 0/11-17
SW4(config-if-range)# ip verify source
```

- R5 and R6 have been preconfigured with IP addresses on their Ethernet interfaces. Configure R4 and its associated switch port accordingly without using secondary addressing to communicate with R5 and R6. Configure R4 with an IP address of 120.100.45.4/24 to communicate with R5, and configure R4 with an IP address of 120.100.46.4/24 to communicate with R6. Configure R4 Gi0/1 and Switch 2 FE0/4 only. (3 points)

This is just a simple trunking question on Switch 2 to R4 to enable R4 to connect to VLAN 45 and VLAN 46. One point to remember is that Switch 2 does not have VLAN 45 and VLAN 46 configured locally within the default configuration, so you will need to create the VLANs locally before configuring the trunk. If you have configured this correctly, as shown in [Example 1-9](#), you have scored 3 points.

Example 1-9 Switch 2 and R4 Trunking Configuration

[Click here to view code image](#)

```
R4(config)# interface GigabitEthernet0/1.45
R4(config-if)# encapsulation dot1Q 45
R4(config-if)# ip address 120.100.45.4 255.255.255.0
R4(config-if)# interface GigabitEthernet0/1.46
R4(config-if)# encapsulation dot1Q 46
R4(config-if)# ip address 120.100.46.4 255.255.255.0

SW2(config)# vlan 45-46
SW2(config)# interface fastethernet0/4
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport trunk allowed vlan 45,46
SW2(config-if)# switchport mode trunk
```

Section 2: IPv4 IGP Protocols (24 Points)

Section 2.1: OSPF

- Use a process ID of 1; all OSPF configuration where possible should not be configured under the process ID. The loopback interfaces of routers R1, R2, and R3 should be configured to be in Area 0. R4 should be in Area 34 and R5 in Area 5. (2 points)

Recent advances in OSPF have enabled configuration of the network area directly under the interface as opposed to within the OSPF process. [Example 1-10](#) details the Open Shortest Path First (OSPF) configuration.

Example 1-10 OSPF Configuration

[Click here to view code image](#)

```
R1(config)# interface GigabitEthernet 0/0
R1(config-if)# ip ospf 1 area 100
R1(config)# interface GigabitEthernet 0/1
R1(config-if)# ip ospf 1 area 0
R1(config-if)# interface Loopback 0
R1(config-if)# ip ospf 1 area 0

R2(config)# interface Loopback 0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# interface fastethernet 0/0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# interface Serial 0/1
R2(config-if)# ip ospf 1 area 5
R2(config-if)# interface fastethernet 0/1
R2(config-if)# ip ospf 1 area 200

R3(config)# interface loopback 0
R3(config-if)# ip ospf 1 area 0
R3(config-if)# interface GigabitEthernet 0/1
R3(config-if)# ip ospf 1 area 0
R3(config-if)# interface GigabitEthernet 0/0
R3(config-if)# ip ospf 1 area 34

R4(config)# interface Loopback 0
R4(config-if)# ip ospf 1 area 34
R4(config-if)# interface GigabitEthernet 0/0
R4(config-if)# ip ospf 1 area 34
R4(config-if)# interface GigabitEthernet 0/1.45
R4(config-if)# ip ospf 1 area 5

R5(config)# interface Loopback 0
R5(config-if)# ip ospf 1 area 5
R5(config-if)# interface GigabitEthernet 0/0
R5(config-if)# ip ospf 1 area 5
R5(config-if)# interface Serial 0/0/1
R5(config-if)# ip ospf 1 area 5
```

If you have configured OSPF correctly, as shown in [Example 1-10](#), you have scored 2 points.

- No loopback networks should be advertised as host routes. (1 point)

Loopback interfaces within OSPF are by default advertised as host routes. To manipulate this behavior, you need to override the network type that the IOS associates with the loopback interface. [Example 1-11](#) shows the host routes learned on R2. If you have configured this correctly, as shown in [Example 1-11](#), you have scored 1 point.

Example 1-11 OSPF Loopback Interface Host Routes and Configuration

[Click here to view code image](#)

```
R2# sh ip route | include /32
O      120.100.5.1/32 [110/65] via 120.100.25.5, 00:04:34, Serial0/1
O IA    120.100.4.1/32 [110/66] via 120.100.123.3, 00:00:42, Serial0/0
O      120.100.1.1/32 [110/129] via 120.100.123.3, 00:01:00, Serial0/0
O      120.100.3.1/32 [110/65] via 120.100.123.3, 00:01:00, Serial0/0

R2# sh ip route | include /32
C      120.100.25.5/32 is directly connected, Serial0/1
O      120.100.5.1/32 [110/65] via 120.100.25.5, 00:39:59, Serial0/1
O IA    120.100.4.1/32 [110/3] via 120.100.123.3, 00:47:32, GigabitEthernet0/1
O      120.100.1.1/32 [110/2] via 120.100.123.1, 00:50:56, GigabitEthernet0/1
O      120.100.3.1/32 [110/2] via 120.100.123.3, 00:49:20, GigabitEthernet0/1

R1# conf t
R1(config)# int Loopback 0
R1(config-if)# ip ospf network point-to-point

R2# conf t
R2(config)# interface Loopback 0
R2(config-if)# ip ospf network point-to-point

R3# conf t
R3(config)# int Loopback 0
R3(config-if)# ip ospf network point-to-point

R4# conf t
R4(config)# int Loopback 0
R4(config-if)# ip ospf network point-to-point

R5# conf t
R4(config)# int Loopback 0
R4(config-if)# ip ospf network point-to-point

R2# sh ip route ospf 1 | include /24
O IA    120.100.4.0/24 [110/3] via 120.100.123.3, 01:42:09, fastethernet0/0
O      120.100.5.0/24 [110/65] via 120.100.25.5, 00:17:09, Serial0/1
O      120.100.1.0/24 [110/2] via 120.100.123.1, 01:43:00, fastethernet0/0
O      120.100.3.0/24 [110/2] via 120.100.123.3, 01:42:26, fastethernet0/0
O      120.100.45.0/24 [110/65] via 120.100.25.5, 02:52:46, Serial0/1
O IA    120.100.34.0/24 [110/2] via 120.100.123.3, 01:43:00, fastethernet0/0
O IA    120.100.100.0/24 [110/2] via 120.100.123.1, 00:00:04, fastethernet0/0
```

- Ensure that R1 does not advertise the preconfigured secondary address under interface Gigabit 0/1 of 120.100.100.1/24 to the OSPF network. Do not use any filtering techniques to achieve

this. (2 points)

The associated behavior with configuring OSPF directly under the interface is that it will by default advertise any secondary addresses assigned to the interface. R1 has a preconfigured secondary address on interface Gigabit 0/1 that is therefore advertised. Because you cannot filter this advertisement, you need to inform OSPF not to include the secondary addresses under the interface command. If you have configured this correctly, as shown in [Example 1-12](#), you have scored 2 points.

Example 1-12 OSPF Secondary Address Advertisement and Configuration

[Click here to view code image](#)

```
R1# show ip ospf int GigabitEthernet 0/1
GigabitEthernet0/1 is up, line protocol is up
  Internet Address 150.100.1.1/24, Area 100
  Process ID 1, Router ID 120.100.1.1, Network Type BROADCAST, Cost: 1
  Enabled by interface config, including secondary ip addresses
  Transmit Delay is 1 sec, State DR, Priority 1
  Designated Router (ID) 120.100.1.1, Interface address 150.100.1.1
  No backup designated router on this network
  Timer intervals configured, Hello 10, Dead 40, Retransmit 5
    oob-resync timeout 40
    Hello due in 00:00:00
  Supports Link-local Signaling (LLS)
  Cisco NSF helper support enabled
  IETF NSF helper support enabled
  Index 1/1, flood queue length 0
  Next 0x0(0)/0x0(0)
  Last flood scan length is 0, maximum is 0
  Last flood scan time is 0 msec, maximum is 0 msec
  Neighbor Count is 0, Adjacent neighbor count is 0
  Suppress hello for 0 neighbor(s)

R1(config)# interface GigabitEthernet 0/1
R1(config-if)# ip ospf 1 area 100 secondaries none

R2# sh ip route 120.100.100.0
% Subnet not in table
```

- R5 should use the serial link within Area 5 for its primary communication to the OSPF network. If this network should fail either at Layer 1 or Layer 2, R5 should form a neighbor relationship with R4 under Area 5 to maintain connectivity. Your solution should be dynamic, ensuring that while the Area 5 serial link is operational, there is no neighbor relationship between R4 and R5; however, the Ethernet interfaces of R4 and R5 must remain up. To confirm the operational status of the serial network, ensure that the serial interface of R5 is reachable by configuration of R5. You are permitted to define **neighbor** statements between R5 and R4. (4 points)

This is a complex scenario that can consume your time, but all the clues are in the question, so some lateral thinking is required. You can rule out a backup interface solution because the Ethernet needs to remain up, and the solution must cater for Layer 1 and Layer 2 rather than purely Layer 1. Similarly, a demand scenario is also out because this would involve a neighbor relationship being formed. You are also requested to confirm operational status of the serial interface on R5 with your overall solution being dynamic. This would take a great deal of effort and trial and error, but you will find

that you can use the IP SLA feature to monitor the IP address of the serial interface on R5 by R5 itself. If this responds to the automatic polling with Internet Control Message Protocol (ICMP), you know the serial link is up at Layers 1 and 2. If the polling fails, you know the interface is down. IP SLA can then be used to inform the router, and a forwarding decision can be manipulated; this feature is known as policy-based routing (PBR) support with multiple tracking options. This gives PBR access to all the objects that are available through the tracking process.

The tracking process provides the ability to track individual objects, such as ICMP ping reachability, and inform the required PBR process when an object state changes. In summary, if the object status changes, R5 can simply manipulate the way it sends traffic by policy routing. The traffic it manipulates needs to be OSPF that should be directed to R4 to form the adjacency over the Ethernet network (VLAN 45). So, when the R5 serial link is up and running, we just need to break the adjacency between R5 and R4. When the serial link fails, we need to allow the adjacency between R5 and R4 to form. The first step in this solution is to configure the IP SLA object tracking on R5. This configuration is detailed in [Example 1-13](#).

Example 1-13 R5 IP SLA Configuration and Status

[Click here to view code image](#)

```
R5(config)# ip sla 1
R5(config-ip-sla)# icmp-echo 120.100.25.5
R5(config-ip-sla-echo)# exit
R5(config)# ip sla schedule 1 life forever start-time now
R5(config)# track 1 rtr 1 reachability

R5# show ip sla statistics

Round Trip Time (RTT) for      Index 1
    Latest RTT: 4 milliseconds
Latest operation start time: *21:17:10.683 UTC Mon Aug 05 2013Latest operation
return code: OK
Number of successes: 2
Number of failures: 0
Operation time to live: Forever
```

Note

OSPF should have already been configured between R4 and R5 within your original peering configuration. The neighbor adjacency takes a while waiting for the dead time to expire (120 seconds after changing of the OSPF network type).

OSPF needs to be configured between R4 and R5 with manual neighbor statements as directed in the question, which ensures the routers unicast traffic to each other. To do this, you must change the network type to **non-broadcast**. The unicast traffic between neighbors can be identified by an ACL that the PBR process can match. Then, instead of allowing normal traffic flow between R5 and R4 to form the neighbor relationship, the next hop can be modified, and because the OSPF TTL is set to 1 by default, the traffic will effectively be dropped by the next hop and the OSPF between R5 and R4 will never establish. Similarly, when the object tracking fails, the PBR process will be overridden

and traffic can flow as normal. This will then allow R5 and R4 to form an OSPF adjacency. So, if you use the PBR command **set ip next-hop verify-availability 120.100.25.2 10 track 1**, R5 can forward normal OSPF traffic to 120.100.25.2 (R2 serial to effectively discard the traffic) if the tracked object (1) is up. If the object status changes to down, the PBR process is informed, and the OSPF traffic to 120.100.25.2 would follow the usual next hop. R5 must be configured to locally policy route traffic because normal PBR behavior is for traffic manipulation for traffic that flows through the router rather than traffic generated by the router itself. [Example 1-14](#) shows the required OSPF configuration on R4 and R5, the PBR on R5, a debug of R2 sending TTL expired to R5 after the OSPF traffic is sent to R2 instead of R5, and the resulting neighbor partial adjacency that is formed between R4 and R5.

Example 1-14 R4 and R5 OSPF and PBR Configuration

[Click here to view code image](#)

```
R4(config)# interface GigabitEthernet0/1.45
R4(config-if)# ip ospf network non-broadcast
R4(config-if)# router ospf 1
R4(config-router)# neighbor 120.100.45.5

R5(config)# interface GigabitEthernet0/0
R5(config-if)# ip ospf network non-broadcast
R5(config-if)# router ospf 1
R5(config-router)# neighbor 120.100.45.4
R5(config-router)# exit
R5(config)# access-list 100 permit ospf host 120.100.45.5 host 120.100.45.4
R5(config)# route-map TEST permit 10
R5(config-route-map)# match ip address 100
R5(config-route-map)# set ip next-hop verify-availability 120.100.25.2 10 track 1
R5(config-route-map)# interface GigabitEthernet0/0
R5(config-if)# ip policy route-map TEST
R5(config-if)# exit
R5(config)# ip local policy route-map TEST

R2# debug ip icmp
ICMP packet debugging is on
R2#
*Feb 26 22:17:12.847: ICMP: time exceeded (time to live) sent to 120.100.45.5 (dest was 120.100.45.4)
R2#

R5# show ip ospf neigh

Neighbor ID      Pri   State            Dead Time     Address          Interface
120.100.2.1      0     FULL/ -          00:00:37      120.100.25.2   Serial0/0/1
120.100.4.1      1     INIT/DROTHER  00:01:45      120.100.45.4   GigabitEthernet0/0
```

[Example 1-15](#) shows the OSPF adjacency formed when the serial link between R2 and R5 is shut down on R5. The PBR is overridden and normal routing occurs because the next hop is not verified by the object tracking. Your routing table needs to be an exact replica as that shown in [Example 1-15](#). You must remember that when an OSPF adjacency forms between R5 and R2, you are joining Area 5

into Area 34, and a virtual link between R3 and R4 is required to extend area 0. If you had not configured a virtual link, it would have been an easy mistake that would take your points away. (This was a difficult question, but a good one to practice with and examine how features operate and interact with each other; you might have been scratching your head or cursing me, but I'd be surprised if you didn't learn something new from this question.)

If you configured this correctly, including the virtual link, you have scored 4 points (definitely a question worth leaving to the end of your exam when you might have some time left over to experiment).

Example 1-15 R3 and R4 OSPF Virtual Link Configuration and R5 Test

[Click here to view code image](#)

```
R3(config)# router ospf 1
R3(config-router)# area 34 virtual-link 120.100.4.1

R4(config)# router ospf 1
R4(config-router)# area 34 virtual-link 120.100.3.1

R5(config)# interface s0/0/1
R5(config-if)# shut
R5(config-if)#
*Jan  2 21:58:16.811: %OSPF-5-ADJCHG: Process 1, Nbr 120.100.2.1 on Serial0/0/1 from
    FULL to DOWN, Neighbor Down: Interface down or detached
*Jan  2 21:58:18.807: %LINK-5-CHANGED: Interface Serial0/0/1, changed state to
    administratively down
*Jan  2 21:58:19.807: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial0/0/1,
    changed state to down
R5(config-if)# do show ip ospf neigh

Neighbor ID      Pri      State            Dead Time           Address          Interface
N/A              0        ATTEMPT/DROTHER 00:00:33    120.100.45.4 GigabitEthernet0/0
R5(config-if)#
*Jan  2 21:59:43.547: %OSPF-5-ADJCHG: Process 1, Nbr 0.0.0.0 on GigabitEthernet0/0
    from ATTEMPT to DOWN, Neighbor Down: Dead timer expired
R5(config-if)#
*Jan  2 22:00:08.135: %OSPF-5-ADJCHG: Process 1, Nbr 120.100.4.1 on
    GigabitEthernet0/0 from LOADING to FULL, Loading Done
R5(config-if)#

R5# sh ip route ospf
  150.100.0.0/24 is subnetted, 2 subnets
O IA    150.100.2.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
     120.0.0.0/24 is subnetted, 9 subnets
O IA    120.100.4.0 [110/2] via 120.100.45.4, 00:04:49, GigabitEthernet0/0
O IA    120.100.1.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.2.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.3.0 [110/3] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.34.0 [110/2] via 120.100.45.4, 00:04:49, GigabitEthernet0/0
O IA    120.100.123.0 [110/3] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.100.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
```

Section 2.2: EIGRP

- Configure EIGRP with an instance name of CCIE where possible using an autonomous system number of 1. The loopback interfaces of all routers and switches should be advertised within EIGRP. (2 points)

This is not a difficult question by any means, just one that has a magnitude of configuration and sets up your Enhanced Interior Gateway Routing Protocol (EIGRP) network using the named instance and address family IPv4 for the following questions. You need to remember to include your preconfigured loopback interfaces and enable routing on the Layer 3 switches. Use the **show ip eigrp neighbor** command to verify your peering before moving on to the next question. If you have configured this correctly, as shown in [Example 1-16](#), you have scored 2 points.

Example 1-16 EIGRP Configuration

[Click here to view code image](#)

```
R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# network 120.100.4.0 0.0.0.255
R4(config-router-af)# network 120.100.45.0 0.0.0.255
R4(config-router-af)# network 120.100.46.0 0.0.0.255
```

```
R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# network 120.100.5.0 0.0.0.255
R5(config-router-af)# network 120.100.45.0 0.0.0.255
R5(config-router-af)# network 120.100.3.0 0.0.0.255
```

```
R6(config)# router eigrp CCIE
R6(config-router)# address-family ipv4 unicast autonomous-system 1
R6(config-router-af)# network 120.100.6.0 0.0.0.255
R6(config-router-af)# network 120.100.46.0 0.0.0.255
R6(config-router-af)# network 120.100.3.0 0.0.0.255
```

```
SW1(config)# ip routing
SW1(config)# exit
SW1# sh run | beg eigrp
router eigrp 1
  network 120.100.7.1 0.0.0.0
  network 150.100.3.7 0.0.0.0
  no auto-summary
```

```
SW2(config)# ip routing
SW2(config)# exit
SW2# sh run | beg eigrp
router eigrp 1
  network 120.100.8.1 0.0.0.0
  network 150.100.3.8 0.0.0.0
  no auto-summary
```

```

SW3(config)# ip routing
SW3(config)# exit
SW3# sh run | beg eigrp
router eigrp 1
  network 120.100.9.1 0.0.0.0
  network 150.100.3.9 0.0.0.0
  no auto-summary

SW4(config)# ip routing
SW4(config)# exit
SW4# sh run | beg eigrp
router eigrp 1
  network 120.100.10.1 0.0.0.0
  network 150.100.3.10 0.0.0.0
  no auto-summary

```

- Ensure that R4 does not install any of the EIGRP loopback routes from any of the switches into its routing table; these routes should also not be present in the OSPF network post redistribution. Do not use any route-filtering ACLs, prefix lists, or admin distance manipulation to achieve this, and perform configuration only on R4. (4 points)

A distribute or prefix list would have been the obvious choice here, but this is not permitted. Upon close inspection of the loopback routes within [Example 1-17](#), you will notice that the routes have a hop count of 2 associated with them. Hop count isn't something you would naturally assimilate with EIGRP, but you can configure the process to ignore routes received with a hop count larger than a configured threshold with the command **metric maximum-hops**. By configuring the maximum hop count of 1 on R4, you can simply stop the loopback routes from entering the process. If you have configured this correctly, as shown in [Example 1-17](#), you have scored 4 points.

Example 1-17 EIGRP Maximum-Hops Configuration

[Click here to view code image](#)

```

R4# show ip route eigrp
  150.100.0.0/24 is subnetted, 3 subnets
D      150.100.3.0
        [90/30720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/30720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
  120.0.0.0/8 is variably subnetted, 16 subnets, 2 masks
D      120.100.8.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.9.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.10.0/24
        [90/158720] via 120.100.46.6, 00:01:07, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:01:07, GigabitEthernet0/1.45
D      120.100.5.0/24
        [90/156160] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.6.0/24
        [90/156160] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
D      120.100.7.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45

```

```
R4# show ip route 120.100.8.0
Routing entry for 120.100.8.0/24
Known via "eigrp 1", distance 90, metric 158720, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:15 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:15 ago, via GigabitEthernet0/1.46
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 2
```

```
R4# show ip route 120.100.9.0
Routing entry for 120.100.9.0/24
Known via "eigrp 1", distance 90, metric 158720, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:25 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:25 ago, via GigabitEthernet0/1.46
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 2
```

```
R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topology)# metric maximum-hops 1
R4(config-router-af-topology)# do show ip route eigrp
    150.100.0.0/24 is subnetted, 3 subnets
D      150.100.3.0
        [90/30720] via 120.100.46.6, 00:00:04, GigabitEthernet0/1.46
        [90/30720] via 120.100.45.5, 00:00:04, GigabitEthernet0/1.45
    120.0.0.0/8 is variably subnetted, 13 subnets, 2 masks
D      120.100.5.0/24
        [90/156160] via 120.100.45.5, 00:00:04, GigabitEthernet0/1.45
D      120.100.6.0/24
        [90/156160] via 120.100.46.6, 00:00:04, GigabitEthernet0/1.46
```

- R4 will have dual equal-cost routes to VLAN 300 (network 150.100.3.0) from R5 and R6. Ensure that R4 sends traffic to this destination network to R5 instead of load sharing. Should the route from R5 become unavailable, traffic should be sent to R6. You may not policy route, alter the **bandwidth** or **delay** statements on R4's interfaces or use an offset list. Perform your configuration on R4 only. Your solution should be applied to all routes received from R5 and R6 as opposed to solely the route to network VLAN 300. (4 points)

To receive identical routes your topology must have identical interface types or **bandwidth** statements used on R4, R5, and R6. [Example 1-18](#) shows the VLAN 300 route (150.100.3.0/24) received on R4 from both R5 and R6 with a metric of 30720. If you want to manipulate this route, the usual best practice method is to modify the bandwidth or delay on one of the Ethernet interfaces, but this is not permitted. In fact, you are left with only one method that can be applied on R4, which will

influence all routes from R5 and R6, as opposed to just this individual route. A route map is required to override the EIGRP-assigned metrics assigned to routes on one interface by manipulating the bandwidth assigned to Gigabit 1/0.45. Gigabit 1/0.46 will, by default, have a lower bandwidth assigned to routes received from it from the **permit 20** statement in the route map. The route map is applied inbound to the process as a distribute list. [Example 1-18](#) also shows that when the interface Gigabit 0/0 is shut down on R5 that the route for VLAN 300 is still received from R6 (R4's feasible successor), so the route is still available but with a different metric. If you have configured this correctly, as shown in [Example 1-18](#), you have scored 4 points. (You could have also manipulated the delay within the route map or created a statement for each individual interface as opposed to just Gigabit 1/0.45.)

Example 1-18 EIGRP Metric Manipulation Configuration

[Click here to view code image](#)

```
R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 30720, type internal
  Redistributing via ospf 1, eigrp 1
  Advertised by ospf 1 metric 5000 subnets
  Last update from 120.100.45.5 on GigabitEthernet0/1.45, 00:25:40 ago
  Routing Descriptor Blocks:
    * 120.100.46.6, from 120.100.46.6, 00:25:40 ago, via GigabitEthernet0/1.46
      Route metric is 30720, traffic share count is 1
      Total delay is 200 microseconds, minimum bandwidth is 100000 Kbit
      Reliability 254/255, minimum MTU 1500 bytes
      Loading 1/255, Hops 1
    120.100.45.5, from 120.100.45.5, 00:25:40 ago, via GigabitEthernet0/1.45
      Route metric is 30720, traffic share count is 1
      Total delay is 200 microseconds, minimum bandwidth is 100000 Kbit
      Reliability 252/255, minimum MTU 1500 bytes
      Loading 1/255, Hops 1
```

```
R4(config)# route-map CHANGEMETRIC permit 10
R4(config-route-map)# match interface gigabitEthernet 0/1.45
R4(config-route-map)# set metric 2000 10 255 1 1500
R4(config-route-map)# route-map CHANGEMETRIC permit 20
R4(config-route-map)# set metric 1000 10 255 1 1500
R4(config-route-map)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topolgy)# distribute-list route-map CHANGEMETRIC in
R4(config-router-af-topolgy)# ^Z
R4# clear ip route *
R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 1282560, type internal
  Redistributing via ospf 1, eigrp 1
  Advertised by ospf 1 metric 5000 subnets
  Last update from 120.100.45.5 on GigabitEthernet0/1.45, 00:03:10 ago
  Routing Descriptor Blocks:
    * 120.100.45.5, from 120.100.45.5, 00:03:10 ago, via GigabitEthernet0/1.45
      Route metric is 1282560, traffic share count is 1
      Total delay is 100 microseconds, minimum bandwidth is 2000 Kbit
```

```
Reliability 255/255, minimum MTU 1500 bytes
Loading 1/255, Hops 1
```

```
R5(config)# int gig0/0
R5(config-if)# shutdown

R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 2562560, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:10 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:10 ago, via GigabitEthernet0/1.46
    Route metric is 2562560, traffic share count is 1
    Total delay is 100 microseconds, minimum bandwidth is 1000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

Section 2.3: Redistribution

- Perform mutual redistribution of IGPs on R4. All routes should be accessible except for the switch loopback networks (because these should not be visible via R4 from an earlier question). EIGRP routes redistributed within the OSPF network should remain with a fixed cost of 5000 throughout the network. (3 points)

A simple redistribution question for the warm-up lab, you have only a single redistribution point (R4), so have no concerns when using protocols such as EIGRP and OSPF, with their inherent protection against routing loops. The fixed cost of 5000 is achieved by advertising redistributed routes into OSPF using a metric type of 2, which is the default, so no specific configuration is required for this. The only points you need to consider when redistributing into OSPF are to use the **subnets** command to ensure classless redistribution and to use default metrics in each protocol. If you have configured this correctly, as shown in [Example 1-19](#), you have scored 3 points.

Example 1-19 R4 Redistribution Configuration and Verification

[Click here to view code image](#)

```
R4(config-route-map)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topology)# redistribute ospf 1
R4(config-router-af-topology)# default-metric 10000 100 255 1 1500
R4(config-router-af-topology)# router ospf 1
R4(config-router)# redistribute eigrp 1 subnets
R4(config-router)# default-metric 5000

R1# show ip route ospf | include E2
O E2      150.100.3.0 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet0/0
O E2      120.100.6.0/24 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet 0/0
O E2      120.100.46.0/24 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet 0/0

SW1# show ip route eigrp | include EX
D EX      150.100.2.0 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
```

```
D EX 150.100.1.0 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.25.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.1.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.2.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.3.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.34.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.123.3/32 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX 120.100.123.0/24 [170/284416] via 150.100.3.6, 00:01:44, Vlan300
```

- Configure R4 to only redistribute up to five EIGRP routes, and generate a system warning when the fourth route is redistributed. Do not use any access lists in your solution. (2 points)

You can limit the number of prefixes redistributed into OSPF and generate a warning when the number of prefixes reaches a defined maximum by use of the **redistribute maximum-prefix** command. To generate the warning on the fourth route, you must configure a percentage threshold (80 percent). If you have configured this correctly, as shown in [Example 1-20](#), you have scored 2 points.

Example 1-20 R4 Prefix Configuration

[Click here to view code image](#)

```
R4(config)# router ospf 1
R4(config-router)# redistribute maximum-prefix 5 80
```

Section 3: BGP (14 Points)

- Configure iBGP peering as follows: R1-R3, R2-R3, R6-R5, SW1-R6, and SW1-R5. Use minimal configuration and use loopback interfaces for your peering. Configure eBGP peering as follows: R3-R4, R4-R6, R4-R5, and R5-R2. Use minimal configuration and use loopback interfaces for your peering with the exception of R4 to R5. (2 points) Use the autonomous system numbers supplied in [Figure 1-7](#). For your eBGP peering on R3, use the TTL security feature, which will not permit a session from R4 to become established if R4 is more than 2 hops away. This feature must be configured only on R3 and not on R4. (2 points)

You can get some easy peering points to begin with, but you'll have to do a lot of typing to earn them. You must remember to use peer groups to minimize configuration where possible, namely on R3, R6, and Switch 1, and follow the peering instructions closely because these are relevant for the following questions. You should have noticed that R3 was required to be a route reflector for iBGP peers R1 and R2 in AS10 and that **no synchronization** is required because the underlying IGP is not redistributed into BGP. Remember to verify your peering with the **show ip bgp neighbor** command. The peering becomes complicated when the TTL security feature is enabled by use of the command **neighbor 120.100.4.1 ttl-security hops 2** on R3. This command is a neat feature that will not permit the peering session if the received neighbor TTL value is less than 253 in this case, which would suggest that the incoming session could be some form of remote attack with spoofed source IP address of the original neighbor. Because you are not permitted to configure the same feature on R4, the peering will break, of course, even if you have configured the eBGP multihop feature on R4 with a value of 2. (Of course, this will simply increment the TTL value from a default value of 0.)

[Example 1-21](#) shows a debug on R3 for the eBGP peering. The field highlighted is the Time To Live (TTL) hex value displayed from the hidden command (dump) when performing the debug. You need to

get the hex value to FD (253 decimal) by configuring the multihop value to 255 on R4, to show R3 that the R4 can only be a maximum of two hops away. If you have configured this correctly, as shown in [Example 1-21](#), you have scored 2 points.

Example 1-21 BGP Peering Configuration

[Click here to view code image](#)

```
R1# sh run | begin bgp
router bgp 10
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.3.1 update-source Loopback0
no auto-summary

R2# sh run | begin bgp
router bgp 10
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.5.1 remote-as 300
neighbor 120.100.5.1 ebgp-multihop 2
neighbor 120.100.5.1 update-source Loopback0
no auto-summary

R3# sh run | begin bgp
router bgp 10
no synchronization
neighbor IBGP peer-group
neighbor IBGP remote-as 10
neighbor IBGP update-source Loopback0
neighbor IBGP route-reflector-client
neighbor 120.100.1.1 peer-group IBGP
neighbor 120.100.2.1 peer-group IBGP
neighbor 120.100.4.1 remote-as 200
neighbor 120.100.4.1 ttl-security hops 2
neighbor 120.100.4.1 update-source Loopback0
no auto-summary

R4# sh run | begin bgp
router bgp 200
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.3.1 ebgp-multihop 2
neighbor 120.100.3.1 update-source Loopback0
neighbor 120.100.6.1 remote-as 300
neighbor 120.100.6.1 ebgp-multihop 2
neighbor 120.100.6.1 update-source Loopback0
neighbor 120.100.45.5 remote-as 300
no auto-summary

R3(config)# access-list 100 permit ip host 120.100.4.1 host 120.100.3.1
R3(config)# exit
R3# debug ip packet 100 detail dump
IP packet debugging is on (detailed) (dump) for access list 100

R3# TCP src=42692, dst=179, seq=2600279946, ack=0, win=163
84 SYN
0F400C00: C204 07400000 B..@..
```

```
0F400C10: C20211E0 00100800 45C0002C 6A870000 B..`....E@.,j...
0F400C20: 0106467E 01010101 03030303 A6C400B3 ..F~.....&D.3
0F400C30: 9AFD1F8A 00000000 60024000 F1BB0000 .}.....`.@.q;..
0F400C40: 02040218 ....
```

! The TTL from R4 is decremented to 01 Hex = 01 decimal as R4 has ebgp-multihop 2 configured and the BGP session will not be established as R3 has the TTL security check enabled, from R3's perspective R4 could be 254 hops away!
! Configure R4 so the TTL value will read 253 decimal (FD hex) by configuring an ebgp multihop value of 255 (this value will decrement down to 253 when it is processed by R3).

```
R4(config)# router bgp 200
R4(config)# neighbor 120.100.3.1 ebgp-multihop 255
```

```
R3# TCP src=44109, dst=179, seq=3925370469, ack=3209854606
, win=16263 ACK
```

```
0F7CBB60: C204 07400000 B..@..
0F7CBB70: C20211E0 00100800 45C00028 8C9A0000 B..`....E@..(....
0F7CBB80: FD06286E 01010101 03030303 AC4D00B3 }.(n.....,M.3
0F7CBB90: E4028565 BF527E8E 50103F87 13FC0000 d..e?R~.P.?..|..
0F7CBBA0:
```

! Now a hex value of FD (253 Decimal) can be seen at R3 from R4, this shows that R4 can not be further than 2 hops away from R3 and the security check passes and BGP is established.

```
R3# sh ip bgp neighbor | include hops | TTL
External BGP neighbor may be up to 2 hops away.
Connection is ECN Disabled, Minimum incoming TTL 253, Outgoing TTL 255
```

```
R5# sh run | begin bgp
router bgp 300
no synchronization
neighbor 120.100.2.1 remote-as 10
neighbor 120.100.2.1 ebgp-multihop 2
neighbor 120.100.2.1 update-source Loopback0
neighbor 120.100.6.1 remote-as 300
neighbor 120.100.6.1 update-source Loopback0
neighbor 120.100.45.6 remote-as 200
neighbor 120.100.7.1 remote-as 300
neighbor 120.100.7.1 update-source Loopback0
no auto-summary
```

```
R6# sh run | beg bgp
router bgp 300
no synchronization
neighbor IBGP peer-group
neighbor IBGP remote-as 300
neighbor IBGP update-source Loopback0
neighbor 120.100.4.1 remote-as 200
neighbor 120.100.4.1 ebgp-multihop 2
neighbor 120.100.4.1 update-source Loopback0
neighbor 120.100.5.1 peer-group IBGP
neighbor 120.100.7.1 peer-group IBGP
no auto-summary
```

```
SW1# sh run | begin bgp
router bgp 300
no synchronization
```

```
neighbor IBGP peer-group
neighbor IBGP remote-as 300
neighbor IBGP update-source Loopback0
neighbor 120.100.5.1 peer-group IBGP
neighbor 120.100.6.1 peer-group IBGP
no auto-summary
```

- AS200 is to be used as a backup transit network for traffic between AS10 and AS300; therefore, if the serial network between R5 and R2 fails, ensure that the peering between R2 and R5 is not maintained via the Ethernet network. Do not use any ACL type restrictions or change the existing peering. (2 points)

As R2 and R5 peer to each other using their loopback interfaces, the peering is maintained if the serial network between R2 and R5 fails. [Example 1-22](#) shows the path taken between R5 and R2 when the serial interface is shut down on R5. To break the peering without using ACLs, you just need to ensure that the **ebgp-multihop** count used in the original peering is set at 2 and no greater. [Example 1-22](#) also shows the ICMP debug with the TTL expiration messages, which indicate the peering will have failed, even though there is IP connectivity between loopbacks. If your **ebgp-multihop** count is set at 2 between R2 and R5, you have scored 2 points.

Example 1-22 eBGP TTL Expiration

[Click here to view code image](#)

```
R5(config)# int s0/0/1
R5(config-if)# shut

R5# trace 120.100.2.1

Type escape sequence to abort.
Tracing the route to 120.100.2.1

1 120.100.45.4 0 msec 0 msec 0 msec
2 120.100.34.3 0 msec 4 msec 0 msec
3 120.100.123.2 4 msec * 4 msec

R5# debug ip icmp
ICMP packet debugging is on
R5#
*Jan 17 21:32:32.455: ICMP: time exceeded rcvd from 120.100.34.3
R5#
*Jan 17 21:32:34.179: ICMP: time exceeded rcvd from 120.100.34.3
R5#

R2# debug ip icmp
ICMP packet debugging is on
R2#
Jan 17 21:26:11.310: ICMP: time exceeded rcvd from 120.100.34.4
R2#
Jan 17 21:26:13.306: ICMP: time exceeded rcvd from 120.100.34.4
```

- Configure a new loopback interface 2 on R2 of 130.100.200.1/24, and advertise this into BGP using the **network** command. Configure R2 in such a way that if the serial connection between R2 and R5 fails, AS300 no longer receives this route. Do not use any route filtering between

neighbors to achieve this. (3 points)

If the peering between R2 and R5 fails, the new network route will flow from AS10 to AS300 via AS200 instead of flowing directly from AS10 to AS300. Therefore, a simple use of communities can be used to ensure that the route is not exported to AS200. You simply need to apply a **no-export** value to the route as it is advertised on R2 toward R3; this way the route is not advertised to AS200 if a failure occurs. Under normal conditions, AS200 would still see the route from AS300. If you have configured this correctly, as shown in [Example 1-23](#), you have scored 3 points.

Example 1-23 Route Advertisement and no-export Configuration on R2

[Click here to view code image](#)

```
R5# sh ip bgp
Origin codes: i - IGP, e - EGP, ? - incomplete

Network           Next Hop            Metric LocPrf Weight Path
*>i130.100.200.0/24 120.100.4.1          0      100      0 200 10 i

R2(config)# interface Loopback2
R2(config-if)# ip address 130.100.200.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 130.100.200.0 mask 255.255.255.0
R2(config-router)# neighbor 120.100.3.1 route-map NO-EXPORT out
R2(config-router)# neighbor 120.100.3.1 send-community
R2(config-router)# exit
R2(config)# access-list 5 permit 130.100.200.0
R2(config)# route-map NO-EXPORT permit 10
R2(config-route-map)# match ip address 5
R2(config-route-map)# set community no-export
R2(config-route-map)# route-map NO-EXPORT permit 20

R3# sh ip bgp 130.100.200.1
BGP routing table entry for 130.100.200.0/24, version 4
Paths: (1 available, best #1, table Default-IP-Routing-Table, not advertised to
EBGP peer)
    Advertised to update-groups:
        2
    Local, (Received from a RR-client)
        120.100.2.1 (metric 65) from 120.100.2.1 (130.100.200.1)
            Origin IGP, metric 0, localpref 100, valid, internal, best
            Community: no-export

R5# conf t
Enter configuration commands, one per line. End with CNTL/Z.
R5(config)# int s0/0/1
R5(config-if)# shut
R5(config-if)# ^Z
R5# show ip bgp

R5#
```

- Configure HSRP between R5 and R6 on VLAN 300 with R5 active for .1/24. If the network 130.100.200.0/24 is no longer visible to AS300, R6 should dynamically become the HSRP active. Configure R5 to achieve this solution. (4 points)

The clue is in the question; all you need to do is track the specific route with the IP SLA object tracking feature and inform the Hot-Standby Router Protocol (HSRP) process whether the Border Gateway Protocol (BGP) route is withdrawn. You might feel that this is not strictly a BGP question, but because the IOS section has been removed from the exam, it is possible that topics and features such as this will crop up within other sections, so it's best to be aware of as many features as possible.

Because the question does not specifically instruct you to configure an exact IP address for your HSRP, you are free to use an unallocated IP address. R5 should be the HSRP active under normal conditions, so this should be configured with the **preempt** command to reinstate control when the route becomes visible once again post withdrawal. Similarly, R6 also requires **preempt** to take control when the priority of R5 decrements. R5 hasn't been configured with a priority in this example because it uses the default value of 100. [Example 1-24](#) shows the configuration and testing steps involved to withdraw the route by shutting down the serial interface on R5 and toggling the HSRP functionality between R5 and R6. If you have configured this correctly, as shown in [Example 1-24](#), you have scored 4 points.

Example 1-24 IP SLA Tracking and HSRP Configuration on R5 and R6

[Click here to view code image](#)

```
R5(config)# track 2 ip route 130.100.200.0 255.255.255.0 reachability
R5(config-track)# interface GigabitEthernet0/1
R5(config-if)# standby 1 ip 150.100.3.1
R5(config-if)# standby 1 preempt
R5(config-if)# standby 1 track 2 decrement 20

R6(config)# interface GigabitEthernet0/1
R6(config-if)# standby 1 ip 150.100.3.1
R6(config-if)# standby 1 priority 90
R6(config-if)# standby 1 preempt

R5# sh standby gigabitEthernet 0/1
GigabitEthernet0/1 - Group 1
  State is Active
    23 state changes, last state change 00:20:11
    Virtual IP address is 150.100.3.1
    Active virtual MAC address is 0000.0c07.ac01
    Local virtual MAC address is 0000.0c07.ac01 (v1 default)
    Hello time 3 sec, hold time 10 sec
      Next hello sent in 0.460 secs
    Preemption enabled
    Active router is local
    Standby router is 150.100.3.6, priority 90 (expires in 8.472 sec)
    Priority 100 (default 100)
      Track object 2 state Up decrement 20
    IP redundancy name is "hsrp-Gi0/1-1" (default)
R5#
R5# conf t
R5(config)# int s0/0/1
```

```

R5(config-if)# shut
R5(config-if)#

R5%#BGP-3-NOTIFICATION: sent to neighbor 120.100.2.1 4/0 (hold time expired) 0 bytes
R5%#HSRP-6-STATECHANGE: GigabitEthernet0/1 Grp 1 state Active -> Speak
R5%#HSRP-6-STATECHANGE: GigabitEthernet0/1 Grp 1 state Speak -> Standby
R5# sh standby gigabitEthernet 0/1
GigabitEthernet0/1 - Group 1
  State is Standby
    25 state changes, last state change 00:00:10
  Virtual IP address is 150.100.3.1
  Active virtual MAC address is 0000.0c07.ac01
    Local virtual MAC address is 0000.0c07.ac01 (v1 default)
  Hello time 3 sec, hold time 10 sec
    Next hello sent in 1.880 secs
  Preemption enabled
  Active router is 150.100.3.6, priority 90 (expires in 8.980 sec)
  Standby router is local
  Priority 80 (default 100)
    Track object 2 state Down decrement 20
  IP redundancy name is "hsrp-Gi0/1-1" (default)

```

- Configure two new loopback interfaces on R1 and R2 of 126.1.1.1/24 and 130.1.1.1/24, respectively, and advertise these into BGP using the **network** command. R3 should be configured to enable only BGP routes originated from R1 up to network 128.0.0.0 and from above network 128.0.0.0 originated from R2. Use only a single ACL on R3 as part of your solution. (3 points)

This is quite an intricate question because you are permitted to use only a single access control list (ACL) to filter the routes on R3. The way to do this is to use an ACL that matches networks up to 128.0.0.0 and permits this through one route map while denying through a separate route map. The route maps should be applied on a per-neighbor basis and both call up the same single ACL. [Example 1-26](#) shows the configuration for the new loopbacks on R1 and R2 and the filtering on R3. Further testing is detailed in [Example 1-26](#) to substantiate the filtering process on R3. If you have configured this correctly, as shown in [Example 1-25](#), you have scored 3 points.

Example 1-25 Route Map Filtering on R3

[Click here to view code image](#)

```

R1(config)# interface Loopback1
R1(config-if)# ip address 126.1.1.1 255.255.255.0
R1(config-if)# router bgp 10
R1(config-router)# network 126.1.1.0 mask 255.255.255.0

R2(config)# interface Loopback1
R2(config-if)# ip address 130.1.1.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 130.1.1.0 mask 255.255.255.0

R3(config)# access-list 1 permit 0.0.0.0 127.255.255.255
R3(config)# route-map UPTO128 permit 10
R3(config-route-map)# match ip add 1

```

```

R3(config)# route-map ABOVE128 deny 10
R3(config-route-map)# match ip add 1
R3(config-route-map)# route-map ABOVE128 permit 20

R3(config)# router bgp 10
R3(config-router)# neighbor 120.100.1.1 route-map UPTO128 in
R3(config-router)# neighbor 120.100.2.1 route-map ABOVE128 in

R3# sh ip bgp
BGP table version is 8, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*->i126.1.1.0/24    120.100.1.1        0       100     0 i
*->i130.1.1.0/24    120.100.2.1        0       100     0 i
*->i130.100.200.0/24 120.100.2.1        0       100     0 i
R3#

```

Further testing of the filtering requires additional interfaces to be configured and advertised on R1 and R2. [Example 1-26](#) shows an interface higher than 128.0.0.0 advertised on R1 and one lower advertised on R2; R3 simply blocks these from entering BGP.

Note

This additional testing configuration is not present on the supplied, final configuration.

Example 1-26 Route Map Filtering Verification

[Click here to view code image](#)

```

R1(config)# interface Loopback3
R1(config-if)# ip address 132.1.1.1 255.255.255.0
R1(config-if)# router bgp 10
R1(config-router)# network 132.1.1.0 mask 255.255.255.0
R1(config-router)# ^Z
R1# sh ip bgp neighbors 120.100.3.1 advertised
BGP table version is 7, local router ID is 126.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*> 126.1.1.0/24    0.0.0.0             0       32768   i
*> 132.1.1.0/24    0.0.0.0             0       32768   i

```

Total number of prefixes 2

```

R3# sh ip bgp
BGP table version is 4, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

```

Network	Next Hop	Metric	LocPrf	Weight	Path
---------	----------	--------	--------	--------	------

```

*>i126.1.1.0/24          120.100.1.1          0    100      0 i
*>i130.1.1.0/24          120.100.2.1          0    100      0 i
*>i130.100.200.0/24 120.100.2.1          0    100      0 i

R2# conf t
R2(config)# int Loopback3
R2(config-if)# ip add 100.1.1.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 100.1.1.0 mask 255.255.255.0
R2(config-router)# ^Z
R2# sh ip bgp neighbor 120.100.3.1 advertised
BGP table version is 5, local router ID is 130.100.200.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop          Metric LocPrf Weight Path
*> 100.1.1.0/24      0.0.0.0            0        32768 i
*> 130.1.1.0/24      0.0.0.0            0        32768 i
*> 130.100.200.0/24 0.0.0.0            0        32768 i

Total number of prefixes 3

R3# sh ip bgp
BGP table version is 4, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop          Metric LocPrf Weight Path
*>i126.1.1.0/24     120.100.1.1          0    100      0 i
*>i130.1.1.0/24     120.100.2.1          0    100      0 i
*>i130.100.200.0/24 120.100.2.1          0    100      0 i

```

Section 4: IPv6 (15 Points)

The prerequisite to the questions is configuration of the IPv6 addresses. You should test your IPv6 connectivity to ensure that you are ready to progress to the routing questions. [Example 1-27](#) shows the required IPv6 configuration to progress to the routing questions. Consider using the **show ipv6 interfaces brief** command for a quick check of your interface configuration.

Example 1-27 IPv6 Testing and Initial Configuration

[Click here to view code image](#)

```

R1(config)# ipv6 unicast-routing
R1(config)# interface gigabitEthernet 0/1
R1(config-if)# ipv6 address 2007:C15:C0:10::1/64
R1(config-if)# gigabitEthernet 0/0
R1(config-if)# ipv6 address 2007:C15:C0:11::1/64

```

```

R2(config)# ipv6 unicast-routing
R2(config)# interface fastethernet 0/1
R2(config-if)# ipv6 address 2007:C15:C0:12::2/64
R2(config-if)# interface fastethernet 0/0

```

```
R2(config-if)# ipv6 address 2007:C15:C0:11::2/64
R2(config-if)# interface serial 0/1
R2(config-if)# ipv6 address 2007:C15:C0:14::2/64
```

```
R3(config)# ipv6 unicast-routing
R3(config)# interface gigabitEthernet 0/0
R3(config-if)# ipv6 address 2007:C15:C0:15::3/64
R3(config-if)# gigabitEthernet 0/1
R3(config-if)# ipv6 address 2007:C15:C0:11::3/64
```

```
R4(config)# ipv6 unicast-routing
R4(config)# interface gigabitEthernet 0/0
R4(config-if)# ipv6 address 2007:C15:C0:15::4/64
```

```
R5(config)# ipv6 unicast-routing
R5(config)# interface gigabitEthernet 0/1
R5(config)# ipv6 address 2007:C15:C0:16::5/64
R5(config-if)# interface Serial0/0/1
R5(config-if)# ipv6 address 2007:C15:C0:14::5/64
```

```
R6(config)# ipv6 unicast-routing
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 address 2007:C15:C0:16::6/64
```

Section 4.1: EIGRPv6

- Configure EIGRPv6 under the instance of CCIE with a primary autonomous system of 1. R1 must not form any neighbor relationship with R2 on VLAN 132 (without the use of any ACL, static neighbor relationships, or multicast blocking feature). R1 must dynamically learn a default route over EIGRPv6 via R3 on VLAN 132 in which to communicate with the IPv6 network. (4 points)

EIGRP configuration is required under an instance of CCIE under the address family IPv6. You could usually stop routers on the same subnet forming a neighbor relationship by creating some static mapping or block the multicast and so on, but the question does not permit you to do this. The clue is in the question stating use a primary autonomous system, which suggests that you can use an additional autonomous system to provide connectivity between R1 and R3, completely bypassing R2. R3 can simply send a default route within the autonomous system to which R1 belongs on VLAN 132, which R2 will have no visibility of. Bear in mind that a named instance within EIGRP can run only one autonomous system, so an additional named instance could be created on R3 to communicate with R1, but the question dictates that the instance is effectively limited to that of CCIE. This leaves you no other option but to enable the secondary autonomous system on R3 under the physical interface. To ensure full visibility from R1 to R2, however, you are required to redistribute EIGRPv6 autonomous systems on R3. Because R1 will receive a default route, you do not require mutual redistribution on R3. Although you could simply perform a one-way redistribution within the protocol, it is better practice to call a route map and just reference the IPv6 network on R1 for redistribution. If you have configured this correctly, as shown in [Example 1-28](#), you have scored 4 points.

Example 1-28 EIGRPv6 Configuration and Testing

[Click here to view code image](#)

```
R1(config)# router eigrp CCIE
R1(config-router)# address-family ipv6 unicast autonomous-system 2
R1(config-router-af)# af-interface GigabitEthernet0/0
R1(config-router-af-interface)# no shutdown
R1(config-router-af-interface)# af-interface GigabitEthernet0/1
R1(config-router-af-interface)# no shutdown

R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv6 unicast autonomous-system 1
R2(config-router-af)# af-interface fastethernet0/1
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface fastethernet0/0
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface Serial0/1
R2(config-router-af-interface)# no shutdown
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 1
R3(config-router-af)# af-interface GigabitEthernet0/0
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface GigabitEthernet0/1
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# exit
R3(config-router-af)# exit
R3(config-router)# exit
R3(config)# interface GigabitEthernet0/1
R3(config-if)# ipv6 eigrp 2
R3(config-if)# ipv6 summary-address eigrp 2 ::/0
R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 1
R3(config-router)# topology base
R3(config-router-topology)# redistribute eigrp 2 route-map EIGRPv6-2-1
R3(config-router-topology)# exit
R3(config-router-af)# exit
R3(config-router)# ipv6 router eigrp 2
R3(config-rtr)# no shut
R3(config-rtr)# exit
R3(config)# route-map EIGRPv6-2-1 permit 10
R3(config-route-map)# match ipv6 address EIGRPv6-2
R3(config-route-map)# route-map EIGRPv6-2-1 deny 20
R3(config-route-map)# exit
R3(config)# ipv6 access-list EIGRPv6-2
R3(config-ipv6-acl)# permit ipv6 2007:C15:C0:10::/64 any

R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv6 unicast autonomous-system 1
R4(config-router-af)# af-interface GigabitEthernet0/0
R4(config-router-af-interface)# no shutdown
```

```

R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
R5(config-router-af)# af-interface Serial0/0/1
R5(config-router-af-interface)# no shutdown

R1# sh ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  ::/0 [90/30720]
  via FE80::216:47FF:FEBB:1E12, GigabitEthernet0/0

R2# sh ipv6 route
EX  2007:C15:C0:10::/64 [170/33280]
  via FE80::216:47FF:FEBB:1E12, fastethernet0/0
D  2007:C15:C0:15::/64 [90/30720]
  via FE80::216:47FF:FEBB:1E12, fastethernet0/0

R3# sh ipv6 route eigrp
D  ::/0 [5/28160]
  via ::, Null0
D  2007:C15:C0:10::/64 [90/30720]
  via FE80::214:69FF:FE61:5EF0, GigabitEthernet0/1
D  2007:C15:C0:12::/64 [90/30720]
  via FE80::215:C6FF:FEF2:ABF1, GigabitEthernet0/1
D  2007:C15:C0:14::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABF1, GigabitEthernet0/1

R4# sh ipv6 route eigrp
EX  2007:C15:C0:10::/64 [170/33280]
  via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D  2007:C15:C0:11::/64 [90/30720]
  via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D  2007:C15:C0:12::/64 [90/33280]
  via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D  2007:C15:C0:14::/64 [90/2174976]
  via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0

R5# sh ipv6 route eigrp
EX  2007:C15:C0:10::/64 [170/2177536]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D  2007:C15:C0:11::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D  2007:C15:C0:12::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D  2007:C15:C0:15::/64 [90/2174976]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1

```

Section 4.2: OSPFv3

- Configure OSPFv3 with a process ID of 1, with all OSPF interfaces assigned to Area 0. (2 points).

This is a clear-cut OSPFv3 configuration. If you have configured this correctly, as shown in [Example 1-29](#), you have scored 2 points.

Example 1-29 R5 and R6 OSPFv3 Configuration

[Click here to view code image](#)

```
R5(config)# interface gigabitEthernet 0/1
R5(config-if)# ipv6 ospf 1 area 0
```

```
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 ospf 1 area 0
```

```
R5# show ipv6 ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
120.100.6.1	1	FULL/DR	00:00:30	3	GigabitEthernet0/1

```
R6# show ipv6 ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
120.100.5.1	1	FULL/BDR	00:00:39	3	GigabitEthernet0/1

- The IPv6 network is deemed to be stable; therefore, reduce the number of LSAs flooded within the OSPF domain. (2 points)

To suppress the unnecessary flooding of link-state advertisements in stable topologies, the **ipv6 ospf flood-reduction** command is required under interface configuration mode. If you have configured this correctly, as shown in [Example 1-30](#), you have scored 2 points.

Example 1-30 R5 and R6 Flood-Reduction Configuration

[Click here to view code image](#)

```
R5(config)# interface gigabitEthernet 0/1
R5(config-if)# ipv6 ospf flood-reduction
```

```
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 ospf flood-reduction
```

Section 4.3: Redistribution

- Redistribute EIGRPv6 routes into the OSPFv3 demand (one way); EIGRPv6 routes should have a fixed cost of 5000 associated with them within the OSPF network. (1 point)

As per vanilla OSPF, the default behavior for OSPFv3 is for redistributed routes to be advertised with a fixed cost as type 2 external routes, so a simple redistribution configuration with a default metric of 5000 on R5 is required. [Example 1-31](#) shows the required configuration and routing table on R6 for the redistributed EIGRPv6 routes. Pay attention to ensure that you have full route visibility, because the serial network on R5 (2007:C15:C0:14::/64) will not be present within the OSPFv3 domain unless R5 specifically redistributes its own connected interfaces. If you have configured this

correctly, as shown in [Example 1-31](#), you have scored 1 point.

Example 1-31 R5 OSPFv3 Redistribution Configuration

[Click here to view code image](#)

```
R5(config)# ipv6 router ospf 1
R5(config-router)# redistribute eigrp 1 metric 5000

R6# sh ipv6 route ospf
IPv6 Routing Table - 10 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OE2 2007:C15:C0:10::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2 2007:C15:C0:11::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2 2007:C15:C0:12::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2 2007:C15:C0:13::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2 2007:C15:C0:15::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1

R5(config)# ipv6 router ospf 1
R5(config-rtr)# redistribute eigrp 1 metric 5000 include-connected
R6# show ipv6 route 2007:C15:C0:14::
IPv6 Routing Table - 10 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OE2 2007:C15:C0:14::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
```

- Ensure that the OSPF3 network is reachable from the EIGRPv6 network by a single route of 2007::/16, which should be seen within the EIGRPv6 domain. Configure R5 only to achieve this. The OSPF domain should continue to receive specific EIGRPv6 subnets. (2 points)

Because you are not mutually redistributing protocols, you are required to configure an IPv6 summary route into the EIGRPv6 domain on R5 to provide full connectivity from the EIGRPv6 domain into OSPFv3. If you have configured this correctly, as shown in [Example 1-32](#), you have scored 2 points.

Example 1-32 R5 EIGRPv6 Summary Configuration and Connectivity Testing

[Click here to view code image](#)

```
R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
```

```
R5(config-router-af)# af-interface Serial0/0/1
R5(config-router-af-interface)# summary-address 2007::/16
```

```
R3# sh ipv6 route | include /16
D 2007::/16 [90/2684416]
```

```
R3# ping ipv6 2007:C15:C0:16::5
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:16::5, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/12/16 ms
R3# ping ipv6 2007:C15:C0:16::6
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:16::6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms
```

-
- Ensure that if the serial link fails between the OSPF and EIGRPv6 domain that routing is still possible between R5 and R4 over VLAN 45. Do not enable EIGRPv6 on the VLAN 45 interfaces of R4 and R5; instead, configure R4 and R5 to achieve this, and this should be considered as an alternative path only if a failure occurs. (3 points)

R4 and R5 both belong to the EIGRPv6 domain. If you cannot enable EIGRPv6 on the VLAN 45 interfaces, all you can do is create a tunnel between the devices. You might have considered enabling OSPFv3 between routers, but you have not been given sufficient information to do this, and it would then create additional problems in terms of redistribution points. [Example 1-34](#) shows the required configuration to tunnel IPv6 through IPv4 on R4 and R5. R5 is still required to advertise the summary route to the EIGRPv6 network through the tunnel for reachability of the OSPFv3 network. If you have configured this correctly, as shown in [Example 1-33](#), you have scored 3 points.

Example 1-33 R4 and R5 Tunnel Configuration and Verification

[Click here to view code image](#)

```
R4(config)# interface Tunnel0
R4(config-if)# ipv6 address 2007:C15:C0:17::4/64
R4(config-if)# tunnel source GigabitEthernet0/1.45
R4(config-if)# tunnel destination 120.100.45.5
R4(config-if)# tunnel mode ipv6ip
R4(config-if)# router eigrp CCIE
R4(config-router)# address-family ipv6 unicast autonomous-system 1
R4(config-router-af)# af-interface Tunnel0
R4(config-router-af-interface)# no shutdown

R5(config)# interface Tunnel0
R5(config-if)# ipv6 address 2007:C15:C0:17::5/64
R5(config-if)# ipv6 eigrp 1
R5(config-if)# tunnel source GigabitEthernet0/0
R5(config-if)# tunnel destination 120.100.45.4
R5(config-if)# tunnel mode ipv6ip
```

```

R5(config-if)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
R5(config-router-af)# af-interface Tunnel0
R5(config-router-af-interface)# no shutdown
R5(config-router-af-interface)# summary-address 2007::/16

R5# sh ipv6 route eigrp
D 2007::/16 [5/2169856]
  via ::, Null0
EX 2007:C15:C0:10::/64 [170/2177536]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:11::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:12::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:15::/64 [90/2174976]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0

R5(config)# int s0/1/0
R5(config-if)# shut
R5(config-if)# do sh ipv6 route eigrp
D 2007::/16 [5/297244416]
  via ::, Null0
EX 2007:C15:C0:10::/64 [170/297252096]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:11::/64 [90/297249536]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:12::/64 [90/297252096]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:15::/64 [90/297246976]
  via FE80::7864:2D04, Tunnel0

```

Section 5: QoS (8 Points)

- You are required to configure QoS on Switch 1 according to the Cisco QoS baseline model. Create a Modular QoS configuration for all user ports (Fast Ethernet 1–24) that facilitates the following requirements (3 points):
 1. All ports should trust the DSCP values received from their connecting devices.
 2. Packets received from the user ports with DSCP values of 48, 46, 34, 32, 24, 28, 16, and 10 should be re-marked to DSCP 8 (PHB CS1) if traffic flowing occurs above 5 Mbps on a per-port basis. This traffic could be a combination of any of the preceding DSCP values with any source/destination combination. Ensure a minimum burst value is configured above the 5 Mbps.

It is acknowledged within the industry that a user port rarely generates more than 5 Mbps of traffic on a standard Fast Ethernet connection. If traffic rates increase above this threshold, it could be indicative of a denial-of-service (DoS) or Worm attack. One way to mitigate an attack is to create a Scavenger class that simply re-marks traffic DSCP values when the threshold has been exceeded. This will not block traffic but will ensure that mission-critical traffic remains unaffected from an attack by trusting the DSCP value for known traffic and re-marking unknown application traffic down to CS1.

To answer the question, you are required to create a Modular QoS policy that trusts the incoming differentiated services code point (DSCP) value received from the host within the policy rather than

by configuring the trust value on a per-interface basis and by policing traffic at a rate of 5 Mbps. When the minimum burst rate is exceeded, the DSCP values will be remapped according to the **police d-dscp** map to Scavenger class CS1 (DSCP8). Note that all DSCP baseline values are being remapped with the exception of DSCP26, which is generally reserved for mission-critical data. This approach enables traffic associated with this value to remain unchanged even when traffic rates exceed 5 Mbps. This approach also assumes that the virus does not itself re-mark traffic to this value to increase its chances of causing damage. However, the exclusion of DSCP26 is not relevant to the configuration and methodology you use to answer the question. The question requires you to configure a standard IP ACL that permits any traffic. For traffic matching this classification, the DSCP value in the incoming packet is trusted. If the matched traffic exceeds an average traffic rate of 5 Mbps and a normal burst size of 8000 bytes, its DSCP is marked down according to the policed DSCP map values and transmitted. If you have configured this correctly, as shown in [Example 1-34](#), you have scored 3 points.

Example 1-34 Switch 1 QoS Configuration and Verification

[Click here to view code image](#)

```
SW1(config)# mls qos
SW1(config)# mls qos map policed-dscp 48 46 34 32 24 28 16 10 to 8
SW1(config)# access-list 1 permit any
SW1(config)# class-map POLICE
SW1(config-cmap)# match access-group 1
SW1(config-cmap)# exit
SW1(config)# policy-map RE-MARK
SW1(config-pmap)# class POLICE
SW1(config-pmap-c)# trust dscp
SW1(config-pmap-c)# police 5000000 8000 exceed-action policed-dscp-transmit
SW1(config-pmap-c)# exit
SW1(config-pmap)# exit
SW1(config)# interface range fastethernet 0/1-24
SW1(config-if-range)# service-policy input RE-MARK

SW1# show policy-map RE-MARK
Policy Map RE-MARK
  Class POLICE
    police 5000000 8000 exceed-action policed-dscp-transmit
    trust dscp
```

- Switch 1 will be connected to a new trusted domain in the future using interface Gigabit 0/1. A DSCP value received locally on SW1 of AF43 should be mapped to AF42 when destined for the new domain. (2 points)

This requires a DSCP mutation map to convert DSCP values between environments. If you did not realize that AF43 is DSCP38 and AF42 is DSCP36, you would struggle to answer this question, but a search of your documentation CD should have assisted you. For the mutation map to function correctly, you need to explicitly trust DSCP values received on the interface on which you are configuring the map. If you have configured this correctly, as shown in [Example 1-35](#), you have scored 2 points.

Example 1-35 Switch 1 DSCP-mutation Map Configuration

[Click here to view code image](#)

```
SW1(config)# mls qos map dscp-mutation AF43-TO-AF42 38 to 36
SW1(config)# interface Gig0/1
SW1(config-if)# mls qos trust dscp
SW1(config-if)# mls qos dscp-mutation AF43-TO-AF42
```

- Configure Cisco Modular QoS as follows on R2 for the following traffic types based on their associated per-hop behavior into classes. Incorporate these into an overall policy that should be applied to the T1 interface S0/1. Allow each class the effective bandwidth as detailed, entered as a percentage. (2 points)

Class	PHB	Assigned Speed
Routing	CS6	46 Kbps
VoIP	EF	247 Kbps
Interactive Video	AF41	247 Kbps
Mission Critical Data	AF31	247 Kbps
Call-Signaling	CS3	46 Kbps
Transactional Data	AF21	216 Kbps
Network-Mgmt	CS2	46 Kbps
Bulk Data	Af11	46 Kbps
Scavenger	CS1	15 Kbps
Default	0	386 Kbps

You have 2 points available here, so you know it's either going to be complex or involve a great deal of configuration. This one is a bit of both, so there is a risk of configuration errors for those points to slip away. There is also some math involved because the **policy-map** command requires a percentage value of bandwidth as opposed to actual speed. Because you are using a T1 interface, you know that the maximum available bandwidth is 1544 Kbps, so the values required are as follows:

$$1\% = 15 \text{ Kbps}, 3\% = 46 \text{ Kbps}, 14\% = 216 \text{ Kbps}, 16\% = 247 \text{ Kbps}, 25\% = 386 \text{ Kbps}$$

A class map to match all values for the provided classes is required that is then associated with the policy map. The overall policy is then applied to the outgoing interface Serial0/1, and a nice little gotcha is that you must configure the interface with the command **max-reserved-bandwidth 100**; otherwise, the full bandwidth is not made available for the policy. Usually you would assign voice traffic into a real-time queue (low-latency queuing [LLQ]), but the question doesn't dictate this, so effectively all traffic types are being assigned with different proportions of class-based weighted fair queuing (CBWFQ). If you have configured this correctly, as shown in [Example 1-36](#), you have scored 2 points.

Example 1-36 Switch1 Modular QoS Configuration

[Click here to view code image](#)

```
R2# sh run class-map
!
class-map match-all VOIP
  match ip dscp ef
class-map match-all BULK-DATA
  match ip dscp af11
class-map match-all NET-MAN
  match ip dscp cs2
class-map match-all VIDEO
  match ip dscp af41
class-map match-all ROUTING
  match ip dscp cs6
class-map match-all SCAVENGER
  match ip dscp cs1
class-map match-all TRANS-DATA
  match ip dscp af21
class-map match-all MISSION-CRIT
  match ip dscp af31
class-map match-all CALL-SIG
  match ip dscp cs3
!
end

R2# sh run policy-map
!
policy-map QOS
  class VOIP
    bandwidth percent 16
  class VIDEO
    bandwidth percent 16
  class BULK-DATA
    bandwidth percent 3
    random-detect
  class TRANS-DATA
    bandwidth percent 14
  class NET-MAN
    bandwidth percent 3
  class ROUTING
    bandwidth percent 3
  class SCAVENGER
    bandwidth percent 1
  class MISSION-CRIT
    bandwidth percent 16

  class CALL-SIG
    bandwidth percent 3
  class class-default
    bandwidth percent 25
!
end

R2# sh run int s0/1 | begin max-reserved-bandwidth 100
max-reserved-bandwidth 100
service-policy output QOS
end

R2# show policy-map QOS
Policy Map QOS
```

```

Class VOIP
  Bandwidth 16 (%) Max Threshold 64 (packets)
Class VIDEO
  Bandwidth 16 (%) Max Threshold 64 (packets)
Class BULK-DATA
  Bandwidth 3 (%)
    exponential weight 9
    class      min-threshold      max-threshold      mark-probability
    -----
    0          -                  -                  1/10
    1          -                  -                  1/10
    2          -                  -                  1/10
    3          -                  -                  1/10
    4          -                  -                  1/10
    5          -                  -                  1/10
    6          -                  -                  1/10
    7          -                  -                  1/10
    rsvp       -                  -                  1/10

Class TRANS-DATA
  Bandwidth 14 (%) Max Threshold 64 (packets)
Class NET-MAN
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class ROUTING
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class SCAVENGER
  Bandwidth 1 (%) Max Threshold 64 (packets)
Class MISSION-CRIT
  Bandwidth 16 (%)
    exponential weight 9
    class      min-threshold      max-threshold      mark-probability
    -----
    0          -                  -                  1/10
    1          -                  -                  1/10
    2          -                  -                  1/10
    3          -                  -                  1/10
    4          -                  -                  1/10
    5          -                  -                  1/10
    6          -                  -                  1/10
    7          -                  -                  1/10
    rsvp       -                  -                  1/10

Class CALL-SIG
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class class-default
  Bandwidth 25 (%)
    exponential weight 9
    class      min-threshold      max-threshold      mark-probability
    -----
    0          -                  -                  1/10
    1          -                  -                  1/10
    2          -                  -                  1/10
    3          -                  -                  1/10
    4          -                  -                  1/10
    5          -                  -                  1/10
    6          -                  -                  1/10
    7          -                  -                  1/10

```

- Configure R2 so that traffic can be monitored on the serial network with a view to a dynamic policy being generated in the future that trusts the DSCP value of traffic identified on this media. (1 point)

This is a simple question that requires the command **auto discovery qos trust** be configured under the serial interface of R2. This command uses NBAR to inspect the application traffic that flows through the router with a view of generating a QoS policy based on the traffic flow profile. The keyword **trust** in the command ensures that the DSCP value of the traffic monitored on the network is trusted. If you have configured this correctly, you have scored 1 point.

Section 6: Security (6 Points)

- Configure R3 to identify and discard the following custom virus. The virus is characterized by the ASCII characters Hastings_Beer within the payload and uses UDP ports 11664 to 11666. The ID of the virus begins on the third character of the payload. The virus originated on VLAN 34. (4 points)

This fictitious virus requires the use of Network-Based Application Recognition (NBAR) with Packet Description Language Module (PDLM) to inspect a packet payload to identify the virus based on the information supplied within the question. Because the virus is located within the third ASCII character, you need to inform the custom NBAR list to ignore the first two characters, which ensures that it will begin to check the third packet. If you have configured this correctly, as shown in [Example 1-37](#), you have scored 3 points. You can use the **show policy-map** command to verify your configuration.

Example 1-37 R3 NBAR Configuration

[Click here to view code image](#)

```
R3(config)# ip nbar custom Hastings_Beer 2 ascii Hastings_Beer udp range 11664 11666
R3(config)# class-map match-all VIRUS
R3(config-cmap)# match protocol Hastings_Beer
R3(config-cmap)# policy-map BLOCK-VIRUS
R3(config-pmap)# class VIRUS
R3(config-pmap-c)# drop
R3(config-pmap-c)# interface gigabit0/0
R3(config-if)# Service-policy input BLOCK-VIRUS
```

- There is an infected host on VLAN 200 of 150.100.2.100. Ensure that only traffic destined for this host is directed to null0 of each local router. You may not use any ACLs to block traffic to this host specifically, but you may use a static route pointing to null0 for traffic destined to 192.0.2.0 /24 on routers within AS10. R2 may have an additional static route pointing to null0. Use a BGP feature on R2 to ensure traffic to this source is blocked. Prevent unnecessary replies when traffic is passed to the null0 interface for users residing on VLAN 100. (4 points)

This question is representative of black-hole routing. This is an effective method of discarding packets being sent to a known destination. This approach to discarding traffic is efficient because it

enables the edge routers to route traffic rather than use ACLs, and it can be deployed dynamically by making use of the next-hop field within BGP updates. You are permitted to create a static route on routers R1, R2, and R3 in AS10 for network 192.0.2.0/24 to null0 and one additional route on R2. This route would need to be directing traffic to the infected host to null0, to update routers R1 and R3. R2 simply advertises the host route for the infected host to AS10 and sets the next hop for this to 192.0.2.1. Routers R1 and R3 then direct traffic to null0 when traffic is destined to the infected host. To ensure that the solution is used only in AS10, you must set the community to **no-export** for the specific static route and tag the route with a value of 10 to identify it. You must therefore send the community values to neighbor R3 on R2, but this should have completed previously for an earlier BGP question. Use of the **no icmp unreachable** command on R1's Gigabit Ethernet interface prevents unnecessary replies when traffic is passed to the null0 interface. If you have configured this correctly, as shown in [Example 1-38](#), you have scored 3 points.

Example 1-38 BGP Black-Hole Routing Configuration and Verification

[Click here to view code image](#)

```
R2(config)# ip route 192.0.2.1 255.255.255.255 null0
R2(config)# ip route 150.100.2.100 255.255.255.255 Null0 Tag 10
R2(config)# router bgp 10
R2(config-router)# redistribute static route-map BLACKHOLE
R2(config-router)# route-map BLACKHOLE permit 10
R2(config-route-map)# match tag 10
R2(config-route-map)# set ip next-hop 192.0.2.1
R2(config-route-map)# set community no-export
R2(config-route-map)# exit

R2(config)# do show ip bgp neigh 120.100.3.1 advertised
BGP table version is 6, local router ID is 130.100.200.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
Network          Next Hop            Metric LocPrf Weight Path
*> 130.1.1.0/24    0.0.0.0                  0        32768 i
*> 130.100.200.0/24 0.0.0.0                 0        32768 i
*> 150.100.2.100/32 192.0.2.1                0        32768 i
Total number of prefixes 3

R2# show ip route 150.100.2.100
Routing entry for 150.100.2.100/32
Known via "static", distance 1, metric 0 (connected)
Tag 10
  Redistributing via bgp 10
  Advertised by bgp 10 route-map BLACKHOLE
  Routing Descriptor Blocks:
    * directly connected, via Null0
  Route metric is 0, traffic share count is 1
  Route tag 10

R3(config)# ip route 192.0.2.1 255.255.255.255 null0
R3(config)# do show ip bgp
BGP table version is 14, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
```

Origin codes: i - IGP, e - EGP, ? - incomplete	Network	Next Hop	Metric	LocPrf	Weight	Path
*>i126.1.1.0/24		120.100.1.1	0	100	0	i
*>i130.1.1.0/24		120.100.2.1	0	100	0	i
*>i130.100.200.0/24		120.100.2.1	0	100	0	i
* i150.100.2.100/32		192.0.2.1	0	100	0	i

```
R1(config)# ip route 192.0.2.1 255.255.255.255 null0
R1(config)# interface Gigabit0/1
R1(config-if)# no icmp unreachable
R1(config-if)# do show ip bgp
BGP table version is 8, local router ID is 126.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
Network          Next Hop           Metric LocPrf Weight Path
*> 126.1.1.0/24    0.0.0.0          0        32768  i
*>i130.1.1.0/24   120.100.2.1      0        100    0 i
*>i130.100.200.0/24 120.100.2.1      0        100    0 i
* i150.100.2.100/32 192.0.2.1       0        100    0 i
```

```
R1# show ip route 150.100.2.100
Routing entry for 150.100.2.100/32
Known via "bgp 10", distance 200, metric 0, type internal
Last update from 192.0.2.1 00:00:02 ago
Routing Descriptor Blocks:
* 192.0.2.1, from 120.100.3.1, 00:00:02 ago
Route metric is 0, traffic share count is 1
AS Hops 0
R1# show ip route 192.0.2.1
Routing entry for 192.0.2.1/32
Known via "static", distance 1, metric 0 (connected)
Routing Descriptor Blocks:
* directly connected, via Null0
Route metric is 0, traffic share count is 1
```

- To protect the control plane on router R6, configure CoPP so that IP packets with a TTL of 0 or 1 are dropped rather than processed, with a resulting ICMP redirect sent to the source. (1 point)

Cisco IOS Software sends all packets with a TTL of 0 or 1 to the process level to be processed. The device must then send an ICMP TTL expire message to the source. By filtering packets that have a TTL of 0 and 1, you can reduce the load on the process level. The control plane policing simply blocks packets with a TTL value of 0 and 1 as directed, but this will break your EIGRP and BGP peering. So, you must specifically permit these packets within your ACL; otherwise, you would have just lost valuable points. If you found yourself running short on time and couldn't justify further time to investigate how to maintain your routing peering, remember that this is a 1-point question, worth leaving and coming back to, if possible. If you have configured this correctly, as shown in [Example 1-39](#), you have scored 1 point.

Example 1-39 CoPP Configuration

[Click here to view code image](#)

```
R6(config)# ip access-list extended TTL
```

```
R6(config-ext-nacl)# deny eigrp any any
R6(config-ext-nacl)# deny tcp any any eq bgp
R6(config-ext-nacl)# deny tcp any eq bgp any
R6(config-ext-nacl)# permit ip any any ttl eq 0 1
R6(config-ext-nacl)# class-map DROP-TTL-0/1
R6(config-cmap)# match access-group name TTL
R6(config-cmap)# policy-map CoPP-TTL
R6(config-pmap)# class DROP-TTL-0/1
R6(config-pmap-c)# drop
R6(config-pmap-c)# control-plane
R6(config-cp)# service-policy input CoPP-TTL
```

Section 7: Multicast (4 Points)

- Configure routers R1, R2, R3, and R4 for IPv4 multicast. Configure R3 to send multicast advertisements of its own time by use of NTP sourced from interface Gig 0/0. Configure PIM sparse mode on all required interfaces. R3 should also be used to advertise its own gigabit interface IP address as an RP. R3 should also advertise the IP address you are using for the NTP advertisements, which will be 224.0.1.1. Do not use the command **ntp server** in any configurations. Routers R1, R2, and R4 should all show a clock synchronized to that of R3. (4 points)

Network Time Protocol (NTP) can be multicast on the reserved group IP address of 224.0.1.1 rather than the more familiar broadcast or unicast scenarios. The question requires you to configure R3 to become the NTP master and announce the group address to the NTP clients. You are not permitted to use the command **ntp server**, and so you must configure the clients with the command **ntp multicast client**. They will then have the capability to join the NTP group by use of Protocol Independent Multicast (PIM). It is good practice to TTL scope your multicast announcements so that they do not propagate past the domain you require. If you have not taken this into consideration in your solution, you would not be deducted points, but be aware of the facility in case you face a question that specifies this. If you have configured this correctly, as shown in [Example 1-40](#), you have scored 4 points.

Example 1-40 NTP Multicast Configuration and Verification

[Click here to view code image](#)

```
R3(config)# ip multicast-routing
R3(config)# ntp master
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-mode
R3(config-if)# ntp multicast ttl 2
R3(config-if)# GigabitEthernet0/1

R3(config-if)# ip pim sparse-mode
R3(config-if)# ip pim send-rp-announce GigabitEthernet0/0 scope 2 group-list 4
R3(config)# ip pim send-rp-discovery GigabitEthernet0/0 scope 2
R3(config)# access-list 4 permit 224.0.1.1

R3# show ntp status
Clock is synchronized, stratum 8, reference is 127.127.7.1
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
```

reference time is C98F1E61.2AE19310 (21:17:21.167 UTC Tue Feb 27 2007)
clock offset is 0.0000 msec, root delay is 0.00 msec
root dispersion is 0.02 msec, peer dispersion is 0.02 msec

R1(config)# **ip multicast-routing**

R1(config-if)# **interface**

GigabitEthernet0/0

R1(config-if)# **ip pim sparse-mode**

R1(config-if)# **ntp multicast client**

R1# **show ntp status**

Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1E79.9FB2321D (21:17:45.623 UTC Tue Feb 27 2007)
clock offset is 0.0157 msec, root delay is 3.88 msec
root dispersion is 0.06 msec, peer dispersion is 0.02 msec

R1(config-if)#[/b]

R1# **show ip igmp group**

IGMP Connected Group Membership

Group Address	Interface	Uptime	Expires	Last Reporter
224.0.1.1	Serial0/0/0	00:40:12	00:02:50	120.100.123.1
224.0.1.39	Serial0/0/0	00:07:21	00:02:51	120.100.123.3
224.0.1.40	Serial0/0/0	00:40:13	00:02:52	120.100.123.1

CHANGE interface

R2(config)# **ip multicast-routing**

R2(config-if)# **interface fastethernet 0/0**

R2(config-if)# **ip pim sparse-mode**

R2(config-if)# **ntp multicast client**

R2# **show ntp status**

Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1E73.83B73E68 (21:17:39.514 UTC Tue Feb 27 2007)
clock offset is 0.0182 msec, root delay is 4.14 msec
root dispersion is 15875.06 msec, peer dispersion is 15875.02 msec

R2# **show ip igmp group**

IGMP Connected Group Membership

Group Address	Interface	Uptime	Expires	Last Reporter
224.0.1.1	Serial0/0	00:41:08	00:02:59	120.100.123.2
224.0.1.39	Serial0/0	00:08:12	00:02:57	120.100.123.3
224.0.1.40	Serial0/0	00:41:09	00:01:59	120.100.123.2

Change IF

R4(config)# **ip multicast-routing**

R4(config-if)# **interface GigabitEthernet0/0**

R4(config-if)# **ip pim sparse-mode**

R4(config-if)# **ntp multicast client**

R4# **show ntp status**

Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1EF1.2B7DB1F2 (21:19:45.169 UTC Tue Feb 27 2007)
clock offset is -0.6937 msec, root delay is 1.37 msec
root dispersion is 7877.08 msec, peer dispersion is 7876.34 msec

R4# **show ip igmp group**

IGMP Connected Group Membership

Group Address	Interface	Uptime	Expires	Last Reporter
224.0.1.1	GigabitEthernet0/0	00:41:29	00:02:42	120.100.34.4
224.0.1.39	GigabitEthernet0/0	00:08:35	00:02:42	120.100.34.3

IP Services (4 Points)

- Configure the following commands on router R1:

aaa new-model

logging buffered

logging 120.100.99.1

Configure a policy on router R1 so that if a user tries to remove AAA services or disable logging via the CLI that a syslog message of UNAUTHORIZED-COMMAND-ENTERED is generated. The policy should ensure that neither command is executed and should consist of a single-line command for the CLI pattern detection. The policy and CLI should run asynchronously. The policy should also generate an email from the router to a mail server residing on IP address 120.100.99.2 (to security@lab-exam.net from eem@lab-exam.net, with the subject “User-Issue,” with the message body consisting of details of who was logged on the time either of the commands were entered). (2 points)

This is an intricate Embedded Events Manager (EEM) question. You are required to configure an EEM applet with a CLI pattern event on a single line to match on either of the commands (**no aaa xxx** and **no logging xxx**). This is achieved by a pattern of “**^no (aaa|logging).***”. The following **sync no skip yes** parameters simply state that the policy and CLI should run asynchronously and that the command entered should not be executed as directed. When the commands are matched via the CLI pattern, the policy requires the syslog message to be generated, a CLI command action to run **show users**, and a final action to send an email with the details of the previous **show** command (which is achieved by the command “**\$_cli_result**”). [Example 1-41](#) details the required configuration and resulting execution of the EEM when the commands **no aaa new-model** and **no logging buffered** are entered and not executed on the router. If you have configured this correctly, as shown in [Example 1-41](#), you have scored 4 points.

Example 1-41 R1 EEM Configuration and Verification Testing

[Click here to view code image](#)

```
R1(config)# aaa new-model
R1(config)# logging buffered
R1(config)# logging 120.100.99.1
R1(config)#
R1(config)# event manager applet CCIE-QUESTION
R1(config-applet)# event cli pattern "^(no (aaa|logging)).*" sync no skip yes
R1(config-applet)# action 1.0 syslog msg "UNAUTHORIZED-COMMAND-ENTERED"
R1(config-applet)# action 2.0 cli command "show user"
R1(config-applet)# action 3.0 mail server "120.100.99.2" to "security@lab-exam.net"
   from "eem@lab-exam.net" subject "User-Issue" body "$_cli_result"

R1(config-applet)# no aaa new-model
%HA_EM-6-LOG: CCIE-QUESTION: UNAUTHORISED-COMMAND-ENTERED
%HA_EM-3-FMPD_SMTP_CONNECT: Unable to connect to SMTP server: 120.100.99.2
%HA_EM-3-FMPD_ERROR: Error executing applet CCIE-QUESTION statement 3.0
R1(config)# no logging buffered
```

```
%HA_EM-6-LOG: CCIE-QUESTION: UNAUTHORISED-COMMAND-ENTERED  
%HA_EM-3-FMPD_SMTP_CONNECT: Unable to connect to SMTP server: 120.100.99.2  
%HA_EM-3-FMPD_ERROR: Error executing applet CCIE-QUESTION statement 3.0  
R1(config)# do show run | include aaa new-model  
aaa new-model  
R1(config)# do show run | include logging buffered  
logging buffered 4096 debugging
```

Lab Wrap-Up

So, how did it go? Did you run out of time? Did you manage to finish but miss what was actually required? If you scored over 80, well done. If you accomplished this within 8 hours or less, you will be prepared for any scenario that you are likely to face during the 5.5 hours of the Configuration section of the actual exam. Remember that the Troubleshooting section on the v5.0 exam is a separate section from the Configuration section and has a different scenario; you will have 2 hours to complete the Troubleshooting section.

This lab was designed to ensure that you troubleshoot your own work as you progress through the questions. What sets the CCIE exam apart within the industry is the complexity of the questions to test you further than you thought possible. The exam is not trying to trick you, but it will ensure that you have the ability to think laterally—an ability that will ensure that you exceed in your networking career and one that sets CCIEs apart. Spend the time to go back over the questions and practice with the configurations using debug and **show** commands to fully absorb any new areas you might have come across.

Did you anticipate and factor into your configuration items such as the maximum reserved bandwidth within QoS? If you did, congratulations, because this would have saved you time and secured you points. It also shows that you fully understand the protocols involved and adapt at testing your configurations. How can you ensure that you have the ability to spot any underlying issues related to a question? Well, it's all mileage; you'll get out of your study what you put into it.

Practice Lab 2

Equipment List

Practice Lab 2 follows an identical format to Lab 1 with timings and also consists of 100 points. You need the following hardware and software components to begin this practice lab.

- Six routers loaded with Cisco IOS Software Release 15.3T Advanced Enterprise image and the minimum interface configuration, as documented in [Table 2-1](#)

Router	Model	Ethernet I/F
R1	3925	2
R2	3925	2
R3	3925	2
R4	3925	2
R5	3925	2
R6	3925	2

Table 2-1 Hardware Required per Router

Note

Notice in the initial configurations supplied that some interfaces will not have IP addresses preconfigured. This is because you will either not be using that interface or you must configure it from default within the exercise. The initial configurations supplied should be used to preconfigure your routers and switch before the lab starts.

If your routers have different interface speeds than those used in this book, adjust the **bandwidth** statements on the relevant interfaces to keep all interface speeds in line. This will ensure that you do not get unwanted behavior because of differing IGP metrics.

- Four 3560X switches with IOS 15.0S IP Services

Setting Up the Lab

Use any combination of routers as long as you fulfill the requirements within the topology diagram, as shown in [Figure 2-1](#). However, you should use the same model of routers because this can make life easier if you load configurations directly from the supplied configurations into your own devices. If your router interface speeds do not match those used in this lab, consider reconfiguring the **bandwidth** statement accordingly to provide symmetry with the routing protocol metrics.

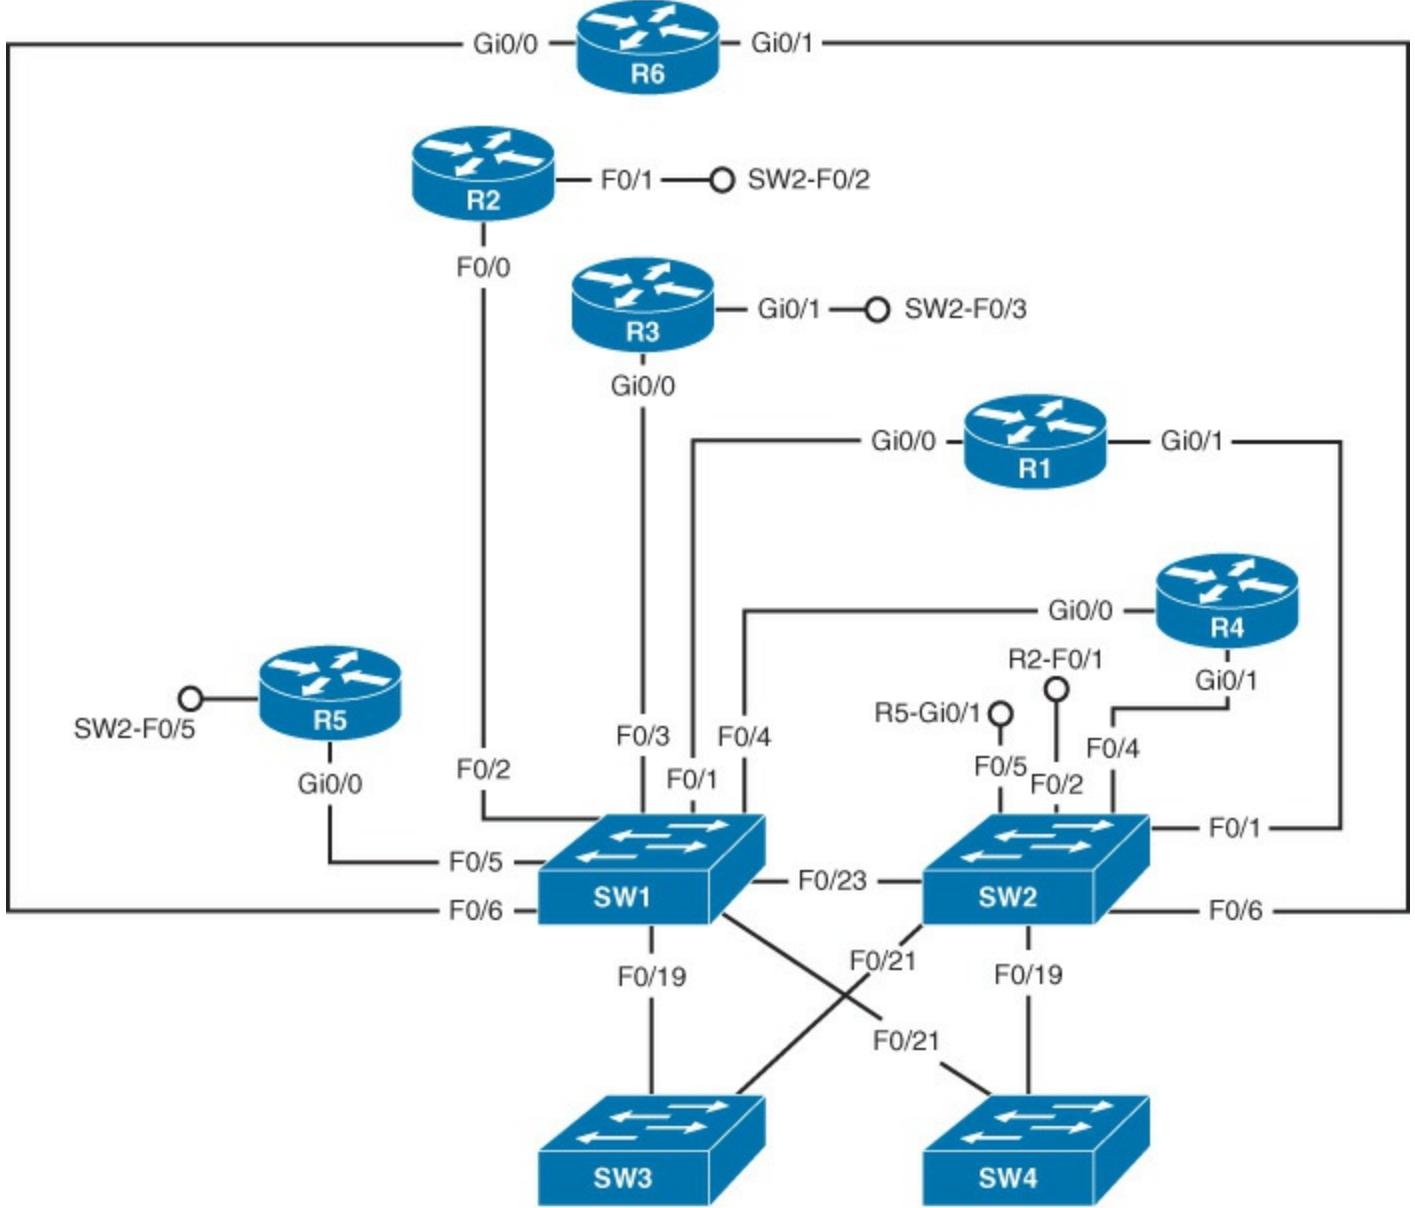


Figure 2-1 Practice Lab 2 Network Topology

Lab Topology

This practice lab uses the topology as outlined in [Figure 2-1](#), which you will need to re-create with your own equipment or by using lab equipment on the CCIE R&S 360 program.

Switch Instructions

Configure VLAN assignments from the configurations supplied or from [Table 2-2](#).

VLAN	Switch 1	Switch 2	Switch 3	Switch 4
34	Fa0/3, Fa0/4, Fa0/5	—	—	—
46	Fa0/6	Fa0/4	—	—
53	VLAN 53	Fa0/5	VLAN 53	—
63	—	Fa0/6	VLAN 63	VLAN 63
100	—	Fa0/1	—	—
132	Fa0/1, Fa0/2	Fa0/3	—	—
200	—	Fa0/2	—	—

Table 2-2 VLAN Assignment

Connect your switches with RJ-45 Ethernet cross-over cables, as shown in [Figure 2-2](#).

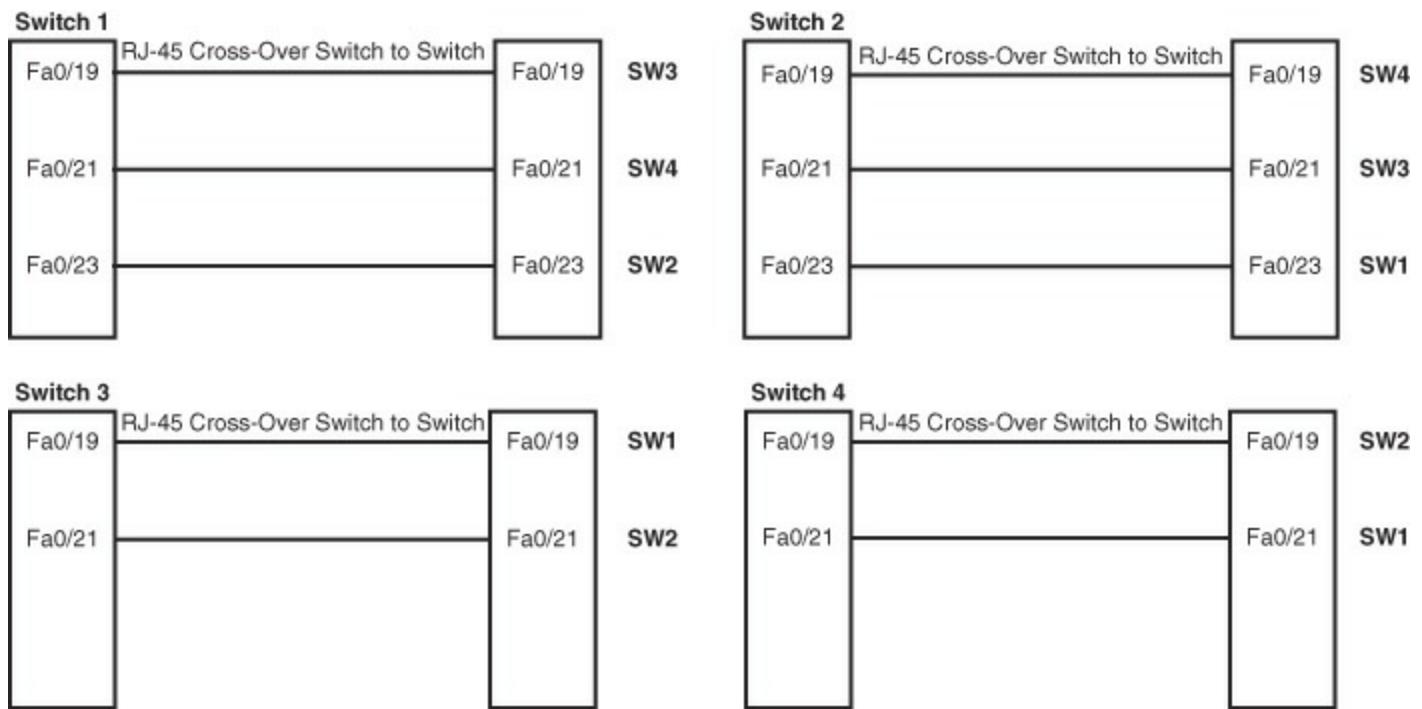


Figure 2-2 Switch-to-Switch Connectivity

IP Address Instructions

You will find in the actual CCIE lab that the majority of your IP addresses will be preconfigured. For this exercise, you are required to configure your IP addresses, as shown in [Figure 2-3](#), or load the initial router configurations supplied. If you are manually configuring your equipment, ensure that you include the following loopback addresses. (R1 and R3 use the same IP address for Loopback 255.)

R1 Lo0 120.100.1.1/24

Lo255 200.200.200.200/24

R2 Lo0 120.100.2.1/24

R3 Lo0 120.100.3.1/24

Lo255 200.200.200.200/24

R4 Lo0 120.100.4.1/24

R5 Lo0 120.100.5.1/24
R6 Lo0 120.100.6.1/24
SW1 Lo0 120.100.7.1/24
SW2 Lo0 120.100.8.1/24
SW3 Lo0 120.100.9.1/24
SW4 Lo0 120.100.10.1/24

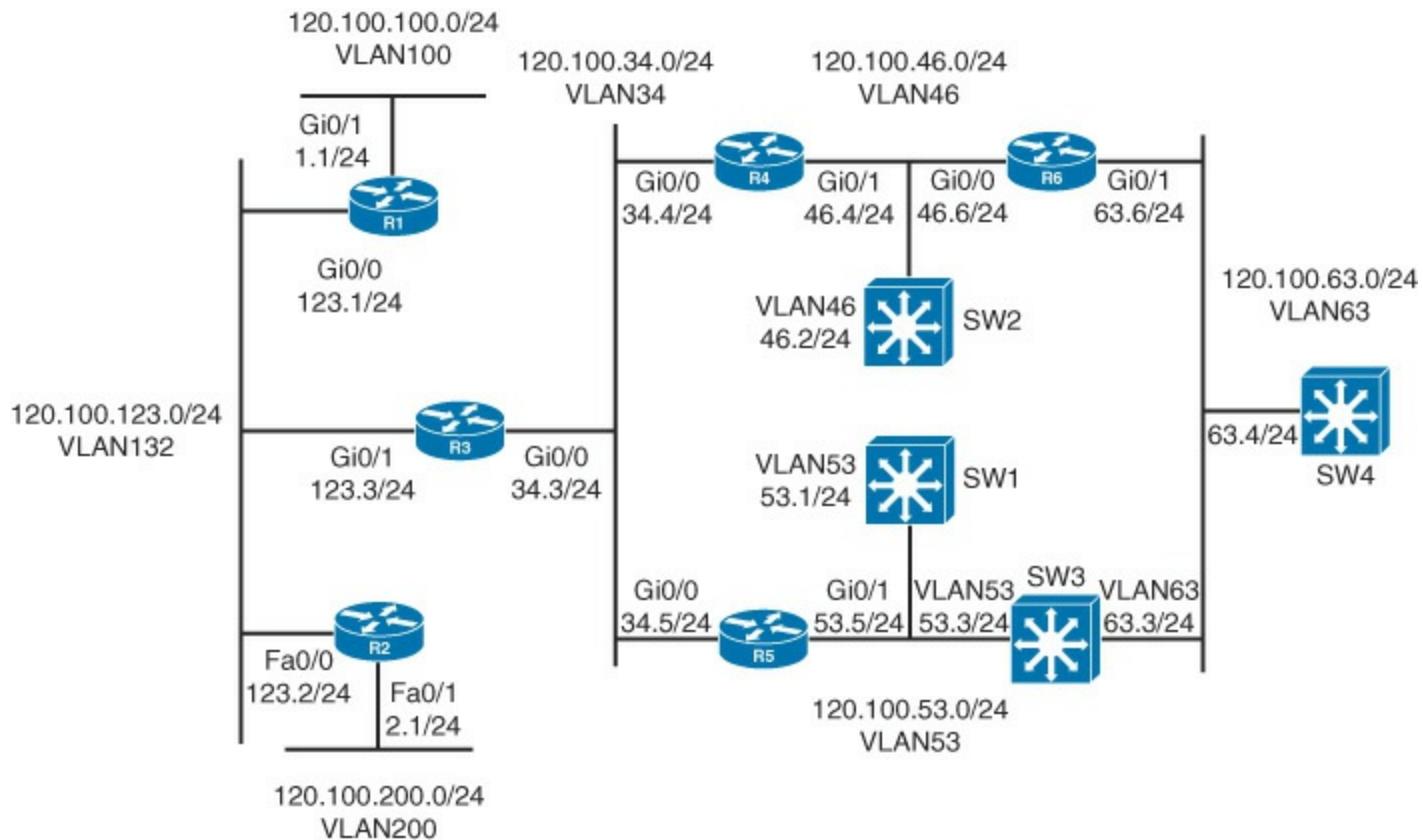


Figure 2-3 IP Addressing Diagram

Pre-Lab Tasks

- Build the lab topology per [Figure 2-1](#) and [Figure 2-2](#).
- Configure the IP addresses on each router as shown in [Figure 2-3](#) and add the loopback addresses. Alternatively, you can load the initial configuration files supplied if your router is compatible with those used to create this exercise.

General Guidelines

- Read the whole lab before you start.
- Do not configure any static/default routes unless otherwise specified.
- Ensure full IP visibility between routers for ping testing/Telnet access to your devices.
- If you run out of time, choose questions that you are confident you can answer, or choose questions with a higher point rating to maximize your potential score.
- Get into a comfortable and quiet environment where you can focus for the next 8 hours.
- Take a 30-minute break midway through the exercise.
- Have available a Cisco documentation CD-ROM or access online the latest documentation

from the following URL:

<http://www.cisco.com/cisco/web/psa/configure.html>

Note

Access only these URLs, not the whole [Cisco.com](#) website (because if you are permitted to use documentation during your CCIE lab exam, it will be restricted). To save time during your lab, consider opening several windows with the pages you are likely to look at.

Practice Lab Two

You will now be answering questions in relation to the network topology, as shown in [Figure 2-4](#).

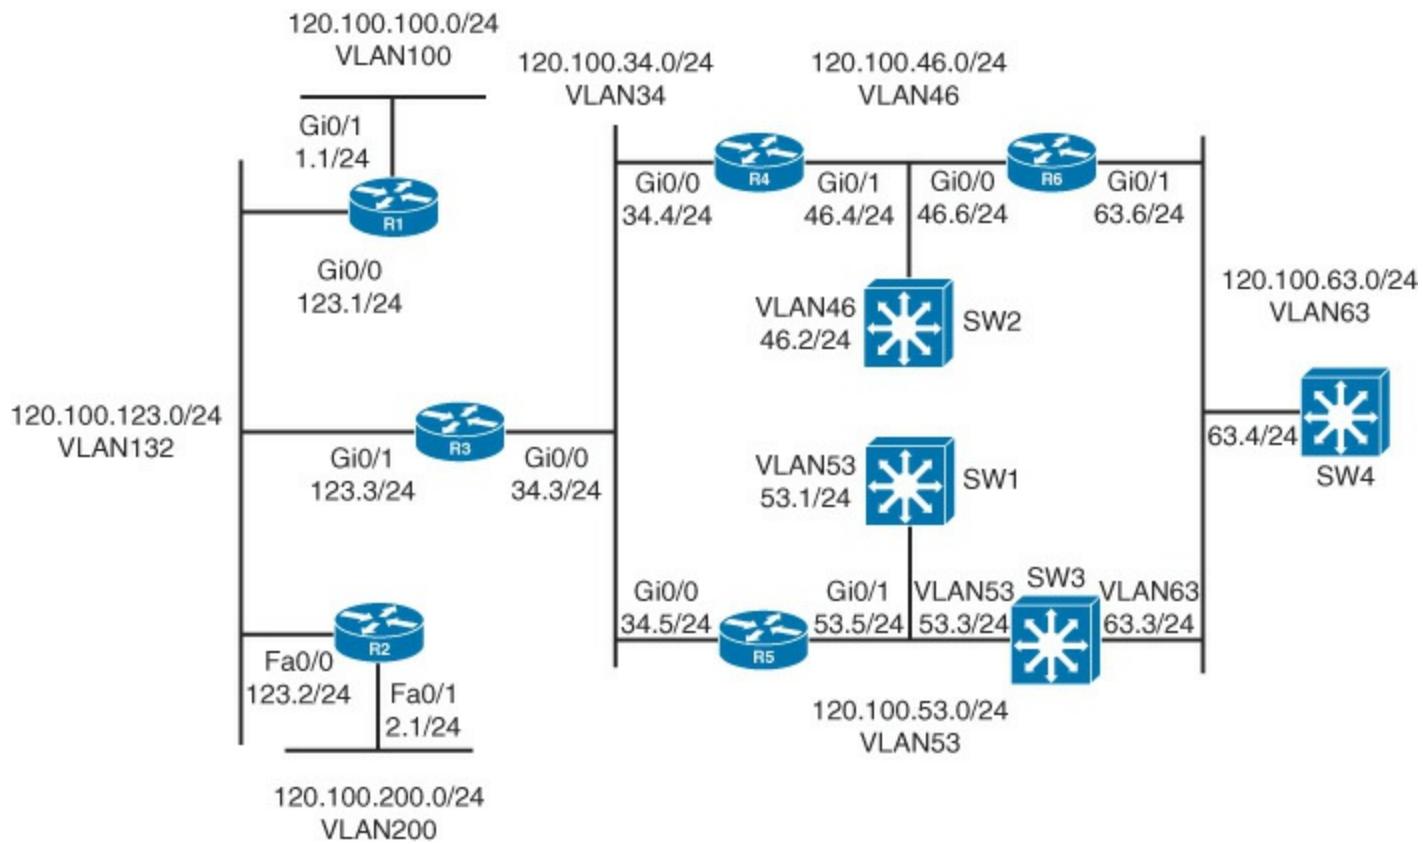


Figure 2-4 Lab Topology Diagram

Section 1: LAN Switching (22 Points)

- Configure your switched network to use 802.1w spanning tree. Switch 1 should be the root bridge for VLANs 34, 46, 53, 63, 100, 132, and 200, with Switch 2 being the secondary root bridge for all listed VLANs. (3 points)
- Switch 3 should use its interface directly connecting to Switch 2 (Fast Ethernet 0/21) for traffic directed toward even-numbered VLANs (34, 46, 100, 132, 200) and the interface directly connecting to Switch 1 (Fast Ethernet 0/19) for odd-numbered VLANs (53, 63). (3 points)
- Switch 4 should use its interface directly connecting to Switch 2 (Fast Ethernet 0/19) for traffic destined toward even-numbered VLANs (34, 46, 100, 132, 200) and the interface directly connected to Switch 1 (Fast Ethernet 0/21) for odd-numbered VLANs (53, 63). (3 points)
- Ensure a cable fault between Switches 1 and 2 could not result in one-way traffic between the

two switches, resulting in spanning-tree issues. (2 points)

- Configure Switch 1 and Switch 2 to enable connectivity of two further switches in the future to be connected to ports Fast Ethernet 0/18 on each switch. The new switches should be able to tunnel their own configured VLANs through a new VLAN (30) between Switch 1 and Switch 2. There is no requirement to configure a root bridge or VLAN load balancing for the new VLAN between Switch 1 and Switch 2. (4 points)
- Configure your switched network to monitor the VLAN 200 interface associated with R2 (Switch 2 Fast Ethernet 0/1), and send only traffic destined to R2 on this switch port across your network to Switch 3 port Fast Ethernet 0/17; use a new VLAN (20) to assist in this configuration. There is no requirement to configure a root bridge or VLAN load balancing for the new VLAN. (3 points)
- Configure the interface on Switch 2 that connects to R5 VLAN 53 (Fast Ethernet 0/5) in such a way that if all the trunks on Switch 2 connecting to Switch 1, Switch 3, and Switch 4 should fail, this Ethernet port transitions into error-disable state. (3 points)
- Configure interfaces Fast Ethernet 0/9 and 0/10 on Switch 1 so that even if they are configured to belong to the same VLAN, they will not be able to forward unicast, broadcast, or multicast traffic to one another. Do not use any form of ACL or configure the ports to belong to a PVLAN. (1 point)

Section 2: IPv4 IGP Protocols (26 Points)

Section 2.1: EIGRP

- Configure EIGRP per [Figure 2-5](#) using an instance name of CCIE and autonomous system of 1; each EIGRP router should have its Loopback 0 interface configured and advertised within EIGRP. (2 points)

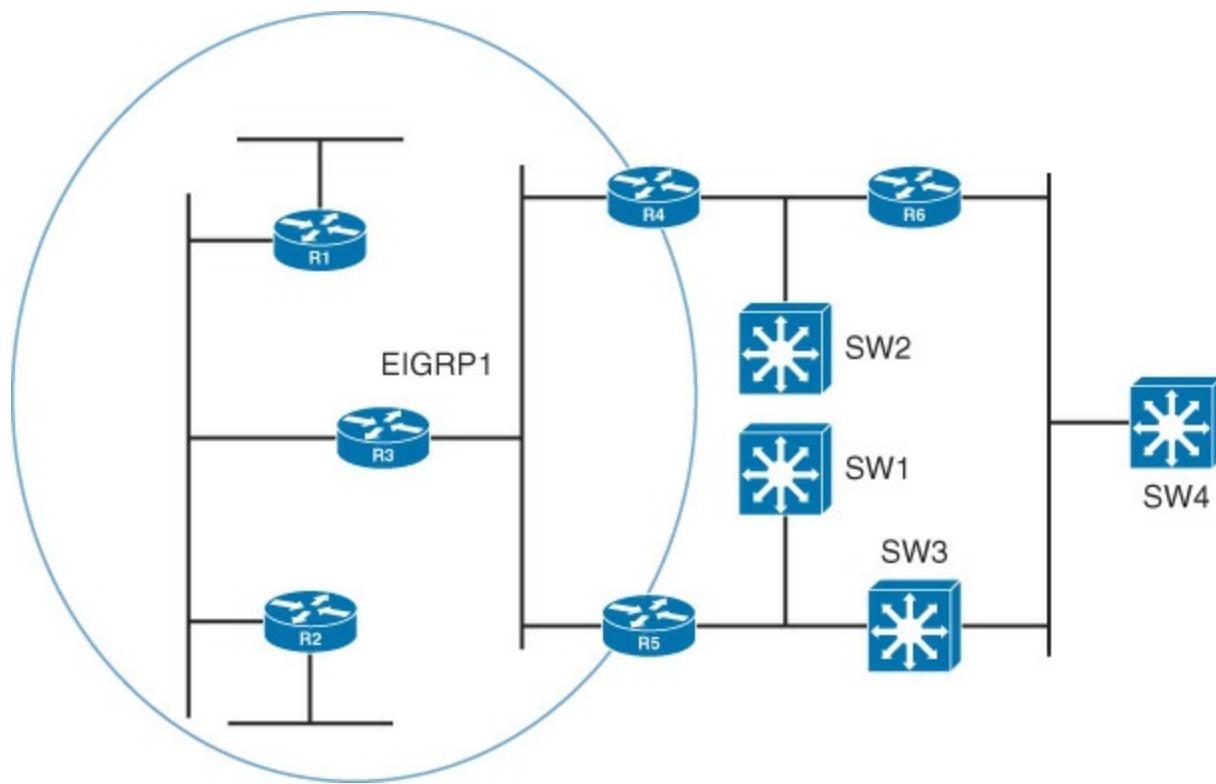


Figure 2-5 EIGRP Topology

- Configure R1 to advertise a summary route of 120.100.0.0/16 outbound on its VLAN 132 interface. R3 should see the original VLAN 100 and Loopback 0 individual routes in addition to the summary route. You may use only one summary route in your configuration; do not apply the **summary** command directly to the interface. (3 points)
- Ensure that the length of time that EIGRP considers neighbors to be valid without receiving a hello packet on the VLAN 132 network between R1, R2, and R3 is 200 seconds; do not change the hello-interval parameter. (2 points)
- Configure new loopback interfaces on R1 and R2 using a loopback interface 2 with an identical IP address of 150.101.1.1/24 on both routers; advertise this network into EIGRP on each router. Ensure that R3 prefers the route from R1 by manipulating the delay associated with this route. Do not manually adjust the delay associated with the interface by use of the **delay** command. You are only permitted to configure R2 to influence the delay. (3 points)
- Configure EIGRP with a new instance name of CCIE2 between R2 and R3 over VLAN 132 with an autonomous system of 2 and 256-bit encryption with a password of lake2aho3, any additional connections to AS2 should be encrypted using the same password without further configuration on R2 and R3. Configure a new loopback interface on R2 (Loopback 3) with an IP address of 150.101.2.1/24, and advertise this and only this network to R3 from R2. (2 points)

Section 2.2: OSPF

- Configure OSPF per [Figure 2-6](#) using a process ID of 1. All OSPF configuration, where possible, should not be configured under the process ID. Each OSPF router should also have its Loopback 0 interface configured and advertised within OSPF as follows: (2 points)

R4 Loopback 0 – Area 0
 R5 Loopback 0 – Area 0
 R6 Loopback 0 – Area 1
 SW1 Loopback 0 – Area 2
 SW2 Loopback 0 – Area 1
 SW3 Loopback 0 – Area 2
 SW4 Loopback 0 – Area 3

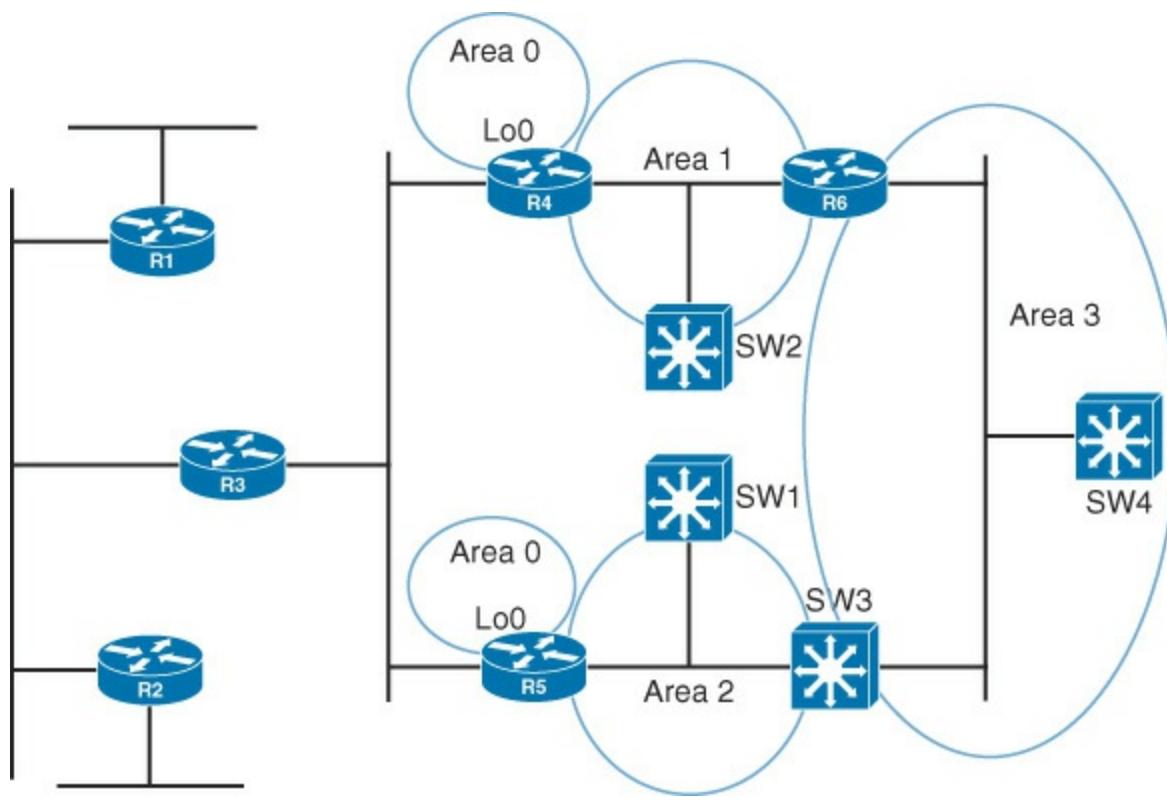


Figure 2-6 OSPF Topology

- Area 0 is partitioned between R4 and R5. Ensure that your network can accommodate this issue. You are not permitted to form any Area 0 neighbor relationship directly between R4 and R5 to join Area 0. (4 points)

Section 2.3: Redistribution

- Perform a one-way redistribution of EIGRP AS2 into EIGRP AS1 on R3 using the following default metric: 1544 20000 255 1 1500. Ensure that R1 shows a next hop for the AS2 advertised route of 150.101.2.0/24 of R2 and perform configuration only on R3 for this task. (3 points)
- Perform mutual redistribution of EIGRP AS1 and OSPF on R4 and R5. Use a metric of 5000 for redistributed routes into OSPF that should appear as external type 2 routes and the following K values for OSPF routes redistributed into EIGRP: 1544 20000 255 1 1500. (2 points)
- R3 will have equal cost external EIGRP routes to the redistributed OSPF subnet 120.100.63.0/24 (VLAN 63). Configure only R3 to ensure that R3 routes via a next hop of R5 (120.100.34.5) for this destination subnet. If this route fails, the route advertised from R4 (120.100.34.4) should be used dynamically. (3 points)

Section 3: BGP (15 Points)

- Configure BGP peering per [Figure 2-7](#) as follows: iBGP R1-R3, R2-R3, R4-R6, R4-SW2. R5-SW1 R5-SW3. eBGP R3-R4, R3-R5, SW4-SW3. R6-SW4. Use loopback interfaces to peer on all routers with the exception of peering between R3-R4 and R3-R5. Do not use the command **ebgp-multihop** within your configurations. (3 points)

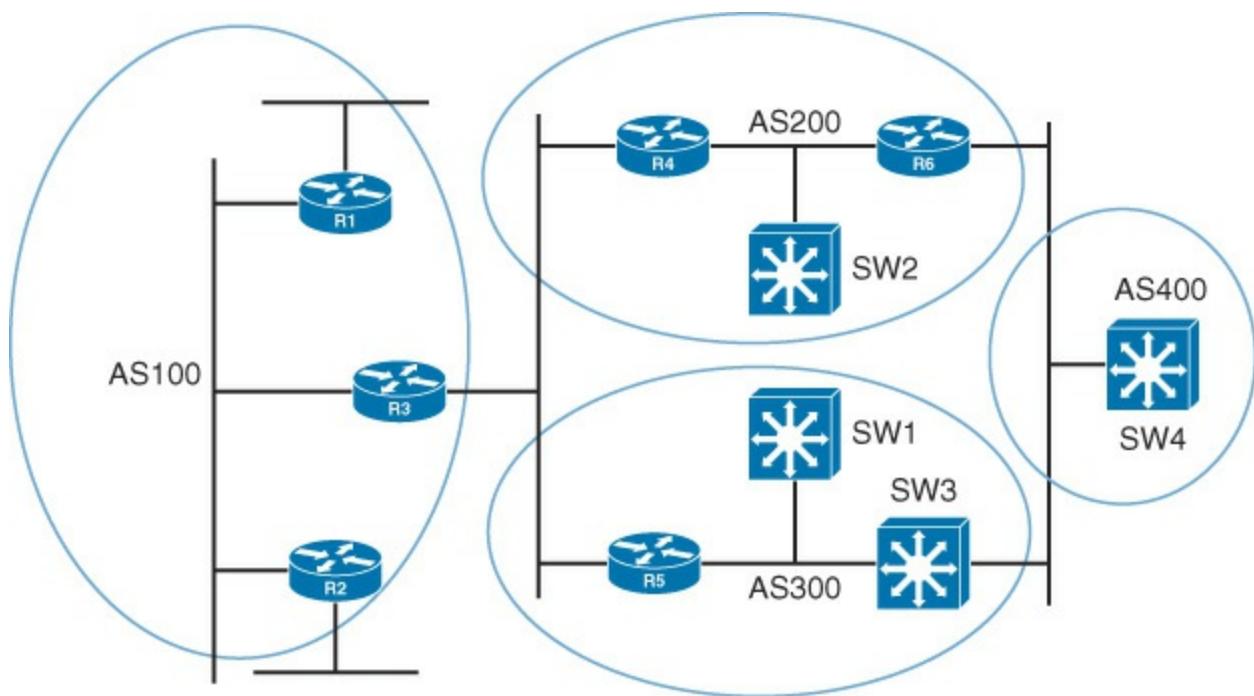


Figure 2-7 BGP Topology

- Routers R1 and R2 in AS100 should be made to only passively accept BGP sessions. R3 should be configured to only actively create BGP sessions to R1 and R2 within AS100. (3 points)
- Configure the following loopback interfaces on R3 and SW4; advertise these networks into BGP using the **network** command: (2 points)
 - R3 – Loopback interface 5 (152.100.100.1/24)
 - SW4 – Loopback interface 5 (152.200.32.1/24)
 - SW4 – Loopback interface 6 (152.200.33.1/24)
 - SW4 – Loopback interface 7 (152.200.34.1/24)
 - SW4 – Loopback interface 8 (152.200.35.1/24)
- Configure R3 to inform R4 that it does not want to receive routes advertised from SW4 for networks 152.200.33.0/24, 152.200.34.0/24, and 152.200.35.0/24. Achieve this in such a manner that R4 does not actually advertise these routes toward R3. You may also configure R4. (4 points)
- Configure a route map on R5 that prepends its local autonomous system an additional two times for network 152.200.32.0/24 when advertised to R3. The route map may contain multiple **permit** statements, but only one prepend is permitted per line. (3 points)

Section 4: IPv6 (12 Points)

- Configure IPv6 addresses on your network as follows:

```

2007:C15:C0:10::1/64 – R1 Gi0/1
2007:C15:C0:11::1/64 – R1 tunnel0
2007:C15:C0:11::3/64 – R3 tunnel0
2007:C15:C0:12::2/64 – R2 tunnel0
2007:C15:C0:12::3/64 – R3 tunnell1
  
```

2007:C15:C0:13::2/64 – R2 fe0/1
2007:C15:C0:14::3/64 – R3 Gi0/0
2007:C15:C0:14::4/64 – R4 Gi0/0
2007:C15:C0:14::5/64 – R5 Gi0/0
2007:C15:C0:15::4/64 – R4 Gi0/1
2007:C15:C0:15::6/64 – R6 Gi0/0

Section 4.1: EIGRPv6

- Configure EIGRPv6 with an autonomous system of 6 between R1, R2, and R3. EIGRPv6 should not be enabled directly under the interfaces of the routers. Build your tunnels from R1 to R3 and R2 to R3 with source interfaces from VLAN 132 to advertise IPv6 edge networks from each router using **ipv6ip** mode. (2 points)

Section 4.2: OSPFv3

- Configure OSPFv3 per [Figure 2-8](#); use an OSPFv3 process of 1 on each router. (2 points)

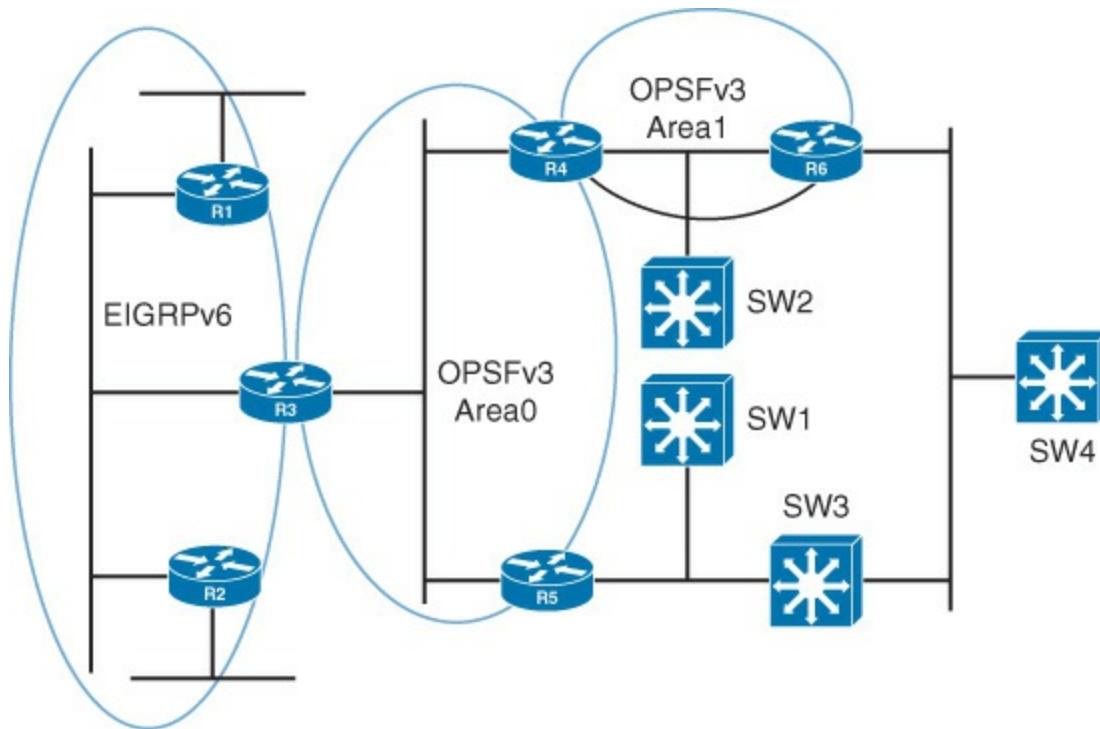


Figure 2-8 IPv6 Topology

- Configure Area 1 with IPsec authentication, use message digest 5, a security policy index of 500, and a key of DEC0DECC1E0DDBA11B0BB0BBEDB00B00. (2 points)
- Ensure the area router in Area 1 receives the following route. You may configure R4 to achieve this: (2 points)
I 2007::/16 [110/2]
via XXXX::XXXX:XXXX:XXXX:XXXX, GigabitEthernet0/0

Section 4.3: Redistribution

- Redistribute EIGRPv6 into OSPFv3 on R3. Redistributed EIGRPv6 routes should have a metric of 5000 associated with them, regardless of which area they are seen in within the OSPFv3 network. (2 points)
- Configure R3 so that both R1 and R2 have the following IPv6 EIGRPv6 route in place. Do not redistribute OSPF into EIGRPv6 to achieve this, and do ensure that all routers have full visibility: (2 points)

D 2007::/16 [90/XXXXXXXXXX]

via XXXX::XXXX:XXXX:XXXX:XXXX, Tunnel0

Section 5: QoS (6 Points)

- Two IP video conferencing units are to be installed onto Switch 2 ports Fast Ethernet 0/15 and 0/16 on VLAN 200. The devices use TCP ports 3230–3231 and UDP ports 3230–3235, and this traffic is unmarked from the devices as it enters the switch. Configure Switch 2 to assign a DSCP value of AF41 to video traffic from both of these devices. Ensure that the switch ports assigned to the devices do not participate in the usual spanning-tree checks, cannot form trunk links, and cannot be configured as EtherChannels. (3 points)
- Configure R2 to assign a strict-priority queue with a 40 percent reservation of the WAN bandwidth for the video conferencing traffic in the previous question. Maximize the available bandwidth by ensuring the RTP headers within the video stream are compressed. The remainder of the bandwidth should be guaranteed for a default queue with WRED enabled. (3 points)

Section 6: Multicast (9 Points)

- Configure routers R1, R2, R3, and R4 for IPv4 multicast. Each router should use PIM sparse dense mode. Both R1 and R2 should be configured to be candidate RPs specifically for the following multicast groups: 225.225.0.1, 225.225.0.2, 225.225.0.3, and 225.225.0.4 (by use of their Loopback 0 interfaces). You should limit the boundary of your multicast network so that it does propagate further into your network than R4. R3 should be configured as a mapping agent to announce the rendezvous points for the multicast network with the same boundary constraints. (3 points)
- Configure R3 to ensure R4 has a candidate RP as R1 for groups 225.225.0.1 and 225.225.0.2 and R2 for groups 225.225.0.3 and 225.225.0.3. (3 points)
- Configure R1 to monitor traffic forwarded through itself for traffic destined to the multicast group of 225.225.0.1. If no packet for this group is received within a single 10-second interval, ensure that an SNMP trap is sent to an SNMP management station on 120.100.100.100 using a community string of **public**. (3 points)

Section 7: Security (10 Points)

- Allow router R6 to passively watch the SYN connections that flow to only VLAN 63 for servers that might reside on this subnet. To prevent a potential denial-of-service (DoS) attack from a flood of SYN requests, the router should be configured to randomly drop SYN packets from any source to this VLAN that have not been correctly established within 20 seconds. (2 points)
- Configure an ACL on R1 to allow TCP sessions generated on this router and through its Ethernet interface and to block TCP sessions from entering on its VLAN 132 interface that were not initiated on it or through it originally. Do not use the established feature within standard ACLs to achieve this, and apply ACLs only on the VLAN 132 interface. The ACL should timeout after 100 seconds of locally initiated TCP inactivity; it should also enable ICMP traffic inbound for testing purposes. (3 points)
- Configure R1 so that it can perform SCP. The router should belong to a domain of toughest.co.uk. Use local authentication with a username and password of cisco, a key size of 768 bits, and an SSH timeout of 2 minutes and retry value of 2. (2 points)
- The network administrator has determined that IPv6 router advertisements are being sourced from routers on VLAN 34. Disable these advertisements from entering and propagating on VLAN 34. You may use an ACL applied in a single location in your solution. Do not use the RA guard solution with untrusted ports. (3 points)

“Ask the Proctor”

Note

This section should be used only if you require clues to complete the questions. In the actual CCIE lab, the proctor will not enter into any discussions about the questions or answers. He or she will be present to ensure that you do not have problems with the lab environment and to maintain the timing element of the exam.

Section 1: LAN Switching

- Q. Do you just want me to configure the root and secondary root bridges into 802.1w spanning tree?
- A. You should ensure that your network runs a consistent version of spanning tree.
- Q. Can I change the root bridge assignments of odd- and even-numbered VLANs to ensure that different interfaces are used on Switch 3 and Switch 4?
- A. No, the root bridge assignment should remain as per the first question.
- Q. If a copper Ethernet cable fails between Switch 1 and Switch 2, surely I wouldn't encounter spanning-tree issues, because there would not be any loops present. Am I correct in thinking this?
- A. Not entirely. Consider a partial failure rather than a complete breakage.
- Q. The switches are connected with Ethernet copper cables; wouldn't a feature like UDLD be

beneficial only if the connections are fiber?

A. UDLD can operate over copper Ethernet in the same manner as fiber.

Q. Would you like me to configure a native VLAN of 30 on trunks to the two new switches?

A. No, a native VLAN would not facilitate transportation of multiple VLANs over the single VLAN 30 between Switch 1 and Switch 2.

Q. Are you looking for a GRE type of tunnel between switches?

A. No, use a Layer 2 switch tunneling feature.

Q. I assume you require remote span configured for R2 traffic. Is it okay to send both TX and RX traffic to Switch 2?

A. Read the question carefully; this information has been provided.

Q. Would you like me to configure UDLD aggressive mode on Switch 2 to transition the required port to error-disable mode if a trunk failure occurs?

A. No, you must configure a feature that will place a nontrunk link into error-disable mode if all the trunks on Switch 2 fail.

Q. Can I just shut down ports 0/9 and 0/10 so that they can't communicate?

A. Nice try; look for a security feature to disable communication between these ports.

Section 2: IPv4 IGP Protocols

Section 2.1: EIGRP

Q. If I can't apply the summary statement directly under the interface can I apply it within the process instance?

A. Yes.

Q. If I configure a summary address on R1, this route overrides the VLAN 100 and Loopback 0 routes from R3 as received on R3. Is this correct?

A. Yes, this is the expected behavior of summarization; you need to enable a feature that enables the more specific routes to be received on R3.

Q. I think I can achieve this with multiple summary routes but the question restricts this. Can I use a new EIGRP process instead?

A. No, use a feature that enables your specific routes to leak from the summary route.

Q. Is it acceptable to adjust the hold time on the Ethernet interfaces to change the hello interval?

A. Yes.

Q. Can I manipulate the delay associated with network 150.101.1.0/24 because this advertisement leaves R2 rather than by changing an interface delay on R2?

A. Yes.

Section 2.2: OSPF

- Q.** Is it acceptable to provide tunnels between R4 and R5 to join Area 0?
- A.** No, this solution would involve a neighbor relationship being formed between the routers in Area 0.
- Q.** I'd normally use a virtual link to extend Area 0 into a transit area. Can I use this technique to stretch Area 0 between R4 and R5?
- A.** You can use virtual links in your solution; think about where the links need to be, though, to ensure that your topology operates correctly.

Section 2.3: Redistribution

- Q.** I've followed the redistribution instructions, but I don't receive the EIGRP AS2 route on R1 after redistribution.
- A.** You will have some underlying issues before receiving the route on R1. Use your troubleshooting skills to determine the problem.
- Q.** I've noticed that due to the preconfigured loopback interfaces on R1 and R3 both of these routers have the same EIGRP router ID. Can I manually change the router ID on one of the routers to see if this helps?
- A.** Yes.
- Q.** I've managed to get the EIGRP AS2 route redistributed from R3 into EIGRP on R1, but the next hop is showing as R3. Can I policy route on R1 so that the next hop for this route is directly via R2?
- A.** No, you must have the routing table reflect the next hop of this route via R2 and not R3.
- Q.** Can I use the EIGRP third-party next-hop feature to leave the next hop of the route unaltered from R2?
- A.** Yes.
- Q.** Can I modify the OSPF cost on the interface connecting R3 to the OSPF network to attempt to change the next hop for the subnet 120.100.63.0/24?
- A.** No, this would affect routes received on R3 from both R4 and R5 equally because R4 and R5 reside on the same subnet as R3.
- Q.** Can I use an offset list or similar feature on R4 to penalize the route 120.100.63.0/24 as it advertised to R3?
- A.** No, you are permitted to configure only R3.
- Q.** Is it acceptable to use a route map on R3 and match a route source to penalize the route to 120.100.63.0/24?
- A.** Yes.

Section 3: BGP

- Q.** If I can't use **ebgp-multipath** on my peering on R6, Switch 3, and Switch 4, will my peering fail because I am peering from my loopback interfaces?
- A.** Yes, it will. You must configure a feature that overrides this behavior.
- Q.** Can I try to use NAT to fix my peering?
- A.** No, use a specific BGP feature to disregard the TTL check.
- Q.** I'm experiencing peering issues between R1 and R3 and have BGP notifications displayed on the console. Is this expected behavior?
- A.** Yes, you had a similar issue within EIGRP; check your router ID.
- Q.** Do you want me to configure an ACL to limit BGP connections to purely inbound or outbound on TCP port 179?
- A.** No, an ACL would actually break the peering entirely. Use a BGP feature to force the peering to become directional.
- Q.** Can I just configure a filter on R4 to stop advertising specific routes to R3?
- A.** No, you must dynamically inform R4 to not advertise specific routes via R3.
- Q.** Can I use BGP ORF?
- A.** Yes.

Section 4: IPv6

- Q.** Do you want a tunnel between R1 and R2 also?
- A.** No, just from R1 to R3 and from R2 to R3. These tunnels will advertise the edge networks of each router within EIGRPv6.
- Q.** Would you like me to configure an additional IPv6 subnet on R4 to receive the 2007::/16 route?
- A.** No, investigate an alternative method to create this route from the preconfigured subnets you already have, ensuring that the route is received as illustrated in the question.
- Q.** Would you like me to redistribute routes into OSPFv3 as external type 1 or type 2?
- A.** The question provides you with sufficient information to determine the redistribution type to use.

Section 5: QoS

- Q.** Do the VC units use UDP Ports 3230 and 3235 or 3230 through 3235?
- A.** They use the range 3230 through 3235.
- Q.** Do you want me to trust the ports assigned to the VC units?
- A.** The VC devices are not marking the traffic, so there is a need to trust these ports.
- Q.** Would you like me to disable trunking, channeling, and spanning-tree checks on the ports

assigned to the VC units?

A. Yes, but remember there is a single command that will disable all these features.

Q. If I use the bandwidth percent command on R2 in my 40-percent guaranteed reservation, is this sufficient to answer the question?

A. No, the question dictates that a priority queue be used.

Section 6: Multicast

Q. If I configure R1 and R2 for the same multicast groups, won't R3 and R4 see both routers as RPs for the same groups?

A. Yes, you will address this behavior in the following question.

Q. To have R1 and R2 as candidate RPs for different groups, can I just configure group lists on R3?

A. If you were permitted to configure R1 and R2, group lists would achieve the desired results, but you are permitted to configure only R3. Group lists can assist in your solution on R3, but you need to find a method of assigning these specifically to R1 and R2.

Q. Do you want me to actually configure an IGMP join group on R1 for 225.225.0.1 for the SNMP question?

A. No, this isn't required; traffic destined to this group will be sent to R1 regardless because it is the candidate RP for this group.

Section 7: Security

Q. Do you want me to configure an ACL to block SYN packets coming into VLAN 63?

A. No, SYN packets should still enter into VLAN 63. You need to configure a feature that monitors the SYN packets and closes down any half-opened connections.

Q. Can I use a reflexive ACL to drop SYN packets that are not correctly established by the servers?

A. No, there is a specific TCP feature used to protect servers from a flood of SYN packets that could cause a DoS attack.

Q. Can I just use a standard ACL on R1 on the VLAN 132 interface to permit sessions outbound and deny everything else inbound?

A. No, this would block return path traffic initiated by R1.

Q. Can I use a reflexive ACL to dynamically permit the return traffic with a time limit of 100 seconds?

A. Yes.

Q. I have configured SCP with the required SSH parameters, but I am not confident of my configuration. Any suggestions?

A. If you have time, try to copy the IOS image from flash on R1 with RCP. If you are prompted for a password and gain access to the file, you have configured this feature correctly.

Q. To stop the RA, can I apply an ACL on each port that connects to each router?

A. No, the question stipulates the ACL can only be used in one location.

Q. So, can I just apply it to VLAN34?

A. Yes.

Q. I have applied the ACL blocking RA ICMPv6 from entering the switch, but I am still seeing the RAs when I debug IPv6 on the routers. Am I missing something?

A. You must consider that by default the switch would be completely transparent to IPv6 and you would need to make the switch understand what it has to filter.

Q. So, do I need to enable IPv6 on the switch?

A. Yes, but consider that it isn't just a case of enabling it; there is an additional step for VLAN 34.

Practice Lab Debrief

The section analyzes each question, showing you what was required and how to achieve the desired results. You should use this section to produce an overall score for the practice lab.

Section 1: LAN Switching (22 Points)

- Configure your switched network to use 802.1w spanning tree. Switch 1 should be the root bridge for VLANs 34, 46, 53, 63, 100, 132, and 200, with Switch 2 being the secondary root bridge for all listed VLANs. (3 points)

802.1w is a Rapid Spanning Tree; the switches will be in the default mode of standard Per-VLAN Spanning Tree (PVST) and require configuration to **rapid-pvst** mode. Switch 1 is required to be the root bridge and Switch 2 the secondary root bridge for VLANs 34, 46, 53, 63, 100, 132, and 200. If you have configured this correctly, as shown in [Example 2-1](#), you have earned 3 points. [Example 2-1](#) also shows confirmation of the root bridge and which interfaces are used to reach the root bridge from the neighboring switches. VLAN 34 is used as an example, but each VLAN would be identical in this configuration.

Example 2-1 SW1, SW2, SW3, and SW4 Configuration and Verification

[Click here to view code image](#)

```
SW1(config)# spanning-tree mode rapid-pvst
SW1(config)# spanning-tree vlan 34,46,53,63,100,132,200 root primary

SW2(config)# spanning-tree mode rapid-pvst
SW2(config)# spanning-tree vlan 34,46,53,63,100,132,200 root secondary

SW3(config)# spanning-tree mode rapid-pvst

SW4(config)# spanning-tree mode rapid-pvst

SW1# show spanning-tree vlan 34 | include root
      This bridge is the root
SW1# show spanning-tree vlan 46 | include root
```

```

        This bridge is the root
SW1# show spanning-tree vlan 53 | include root
        This bridge is the root
SW1# show spanning-tree vlan 63 | include root
        This bridge is the root
SW1# show spanning-tree vlan 100 | include root
        This bridge is the root
SW1# show spanning-tree vlan 132 | include root
        This bridge is the root

SW1# show spanning-tree vlan 200 | include root
        This bridge is the root

SW2# show spanning-tree vlan 34 | include Root FWD
Fa0/23          Root FWD 19          128.25    P2p

SW3# show spanning-tree vlan 34 | include Root FWD
Fa0/19          Root FWD 19          128.21    P2p

SW4# show spanning-tree vlan 34 | include Root FWD
Fa0/21          Root FWD 19          128.23    P2p

```

- Switch 3 should use its interface directly connecting to Switch 2 (Fast Ethernet 0/21) for traffic directed toward even-numbered VLANs (34, 46, 100, 132, 200) and the interface directly connecting to Switch 1 (Fast Ethernet 0/19) for odd-numbered VLANs (53, 63). (3 points)

This is a straightforward VLAN load-balancing question to ensure that trunk links are utilized efficiently and not logically disabled by spanning tree. Switch 3 uses the interface directly connecting to Switch 1 (Fast Ethernet 0/19) for all VLANs as the lowest root cost path by default. To adjust this behavior, this interface must effectively be penalized for the even-numbered VLANs to ensure a more attractive path is via Switch 2 (Fast Ethernet 0/21). If you have configured this correctly, as shown in [Example 2-2](#), you have scored 3 points.

Example 2-2 SW3 VLAN Load-Balancing Configuration and Verification

[Click here to view code image](#)

```

SW3(config)# interface fastethernet 0/19
SW3(config-if)# spanning-tree vlan 34,46,100,132,200 cost 100

SW3(config-if)# do show spanning-tree root

```

Vlan	Root ID	Root Cost	Hello Time	Max Age	Fwd Dly	Root Port
VLAN0001	32769 0013.806d.9400	19	2	20	15	Fa0/19
VLAN0034	24610 0013.806d.9400	38	2	20	15	Fa0/21
VLAN0046	24622 0013.806d.9400	38	2	20	15	Fa0/21
VLAN0053	24629 0013.806d.9400	19	2	20	15	Fa0/19
VLAN0063	24639 0013.806d.9400	19	2	20	15	Fa0/19
VLAN0100	24676 0013.806d.9400	38	2	20	15	Fa0/21
VLAN0132	24676 0013.806d.9400	38	2	20	15	Fa0/21
VLAN0200	24776 0013.806d.9400	38	2	20	15	Fa0/21

- Switch 4 should use its interface directly connecting to Switch 2 (Fast Ethernet0/19) for traffic

destined toward even-numbered VLANs (34, 46, 100, 132, 200) and the interface directly connected to Switch 1 (Fast Ethernet 0/21) for odd-numbered VLANs (53, 63). (3 points)

Following from the previous question, to ensure a balanced access topology for VLAN load balancing, Switch 4 uses the interface directly connecting to Switch 1 (Fast Ethernet 0/21) for all VLANs as the lowest root cost path by default, rendering the second trunk connecting to Switch 2 unused unless a failover condition occurs. As per the previous question, the directly connected interface to Switch 1 needs to be penalized for the even-numbered VLANs. If you have configured this correctly, as shown in [Example 2-3](#), you have scored 3 points.

Example 2-3 SW4 VLAN Load-Balancing Configuration and Verification

[Click here to view code image](#)

```
SW4(config)# interface fastethernet 0/21
SW4(config-if)# spanning-tree vlan 34,46,100,132,200 cost 100
SW4(config-if)# do show spanning-tree root
```

Vlan	Root ID	Root Cost	Hello Time	Max Age	Fwd Dly	Root Port
VLAN0001	32769 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0034	24610 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0046	24622 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0053	24629 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0063	24639 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0100	24676 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0132	24676 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0200	24776 0013.806d.9400	38	2	20	15	Fa0/19

- Ensure that a cable fault between Switches 1 and 2 could not result in one-way traffic between the two switches, resulting in spanning-tree issues. (2 points)

Unidirectional Link Detection (UDLD) detects unidirectional links on fiber-optic connections, in aggressive mode. UDLD also detects unidirectional links because of one-way traffic on twisted-pair links. If you configure the ports between Switch 1 and Switch 2 into aggressive mode, the switches become UDLD neighbors, can detect one-way links, and shut down the link if this condition arises to mitigate spanning-tree issues. If you have configured this correctly, as shown in [Example 2-4](#), you have scored 2 points.

Example 2-4 SW1 and SW2 UDLD Configuration and Verification

[Click here to view code image](#)

```
SW1(config)# interface fastethernet 0/23
SW1(config-if)# udld port aggressive
```

```
SW2(config)# interface fastethernet 0/23
SW2(config-if)# udld port aggressive
```

```
SW1# show udld fastethernet 0/23
```

Interface Fa0/23

```
---  
Port enable administrative configuration setting: Enabled / in aggressive mode  
Port enable operational state: Enabled / in aggressive mode  
Current bidirectional state: Bidirectional  
Current operational state: Advertisement - Single neighbor detected  
Message interval: 15  
Time out interval: 5
```

Entry 1

```
---  
Expiration time: 44  
Cache Device index: 1  
Current neighbor state: Bidirectional  
Device ID: CAT0935N2GQ  
Port ID: Fa0/23  
Neighbor echo 1 device: CAT0911X17K  
Neighbor echo 1 port: Fa0/23
```

```
Message interval: 15  
Time out interval: 5  
CDP Device name: SW2
```

-
- Configure Switch 1 and Switch 2 to allow connectivity of two further switches in the future to be connected to ports Fast Ethernet 0/18 on each switch. The new switches should be able to tunnel their own configured VLANs through a new VLAN (30) between Switch 1 and Switch 2. There is no requirement to configure a root bridge or VLAN load balancing for the new VLAN between Switch 1 and Switch 2. (4 points)

This is a service provider requirement whereby customers tunnel their own VLANs through the providers network; To mitigate any VLAN overlaps from other customers, a unique service provider VLAN is used to transport the customer VLANs. [Example 2-5](#) shows VLAN 30 being used to transport VLANs over a dot1q-tunnel. Use the **show dot1q-tunnel** command to verify your tunnel configuration on your switches. If your ports are shutdown by initial configuration, it would be worth enabling them to protect your points. If you have configured this correctly, as shown in [Example 2-5](#), you have scored 4 points.

Example 2-5 SW1 and SW2 Q in Q Configuration

[Click here to view code image](#)

```
SW1(config)# vlan 30  
SW1(config-vlan)# exit  
SW1(config)# interface fastethernet 0/18  
SW1(config-if)# switchport access vlan 30  
SW1(config-if)# switchport mode dot1q-tunnel
```

```
SW2(config)# vlan 30  
SW2(config-vlan)# exit  
SW2(config)# interface fastethernet 0/18  
SW2(config-if)# switchport access vlan 30  
SW2(config-if)# switchport mode dot1q-tunnel
```

-
- Configure your switched network to monitor the VLAN 200 interface associated with R2

(Switch 2 Fast Ethernet 0/1) and send only traffic destined to R2 on this switch port across your network to Switch 3 port Fast Ethernet 0/17; use a new VLAN (20) to assist in this configuration. There is no requirement to configure a root bridge or VLAN load balancing for the new VLAN. (3 points)

This is a remote span question. The only complexity is based around the question statement of where you actually need to monitor: “traffic destined to R2.” This means that you need to configure the span parameters to only send the traffic transmitted out of the switch port toward R2, which is configured by the **tx** parameter. If this optional parameter is not configured, both transmit and receive traffic is monitored. Remote span requires a VLAN to propagate the span traffic between switches, which is why you need to configure VLAN 20 on both Switches 1 and 2. If your ports are shut down by initial configuration, it would be worth enabling them to protect your points. If you have configured this correctly, as shown in [Example 2-6](#), you have scored 3 points.

Example 2-6 SW2 and SW3 Remote Span Configuration and Verification

[Click here to view code image](#)

```
SW2(config)# vlan 20
SW2(config-vlan)# remote-span
SW2(config-vlan)# exit
SW2(config)# monitor session 1 source interface fastethernet 0/1 tx
SW2(config)# monitor session 1 destination remote vlan 20
SW2(config)# do show monitor session 1
Session 1
-----
Type : Remote Source Session
Source Ports :
    TX Only : Fa0/1
Dest RSPAN VLAN : 20

SW3(config)# vlan 20
SW3(config-vlan)# exit
SW3(config)# monitor session 1 source remote vlan 20
SW3(config)# monitor session 1 destination interface fast 0/17
SW3(config)# do show monitor session 1
Session 1
-----
Type : Remote Destination Session
Source RSPAN VLAN : 20
Destination Ports : Fa0/17

Encapsulation : Native
Ingress : Disabled
```

- Configure the interface on Switch 2, which connects to R5 VLAN 53 (Fast Ethernet 0/5) in such a way that if all the trunks on Switch 2 connecting to Switch 1, Switch 3, and Switch 4 should fail, this Ethernet port transitions into error-disable state. (3 points)

The question requires link-state tracking to be configured. This feature provides redundancy in the network when used with server NIC adapter teaming. If a link is lost on the primary interface, connectivity is transparently switched to the secondary interface. Ports connected to servers are

configured as downstream ports, and ports connected to other switches are configured as upstream ports. If the upstream trunk ports on Switch 2 fail, link-state tracking automatically puts the downstream port connected to R5 into error-disable state. [Example 2-7](#) shows the associated configuration and testing by shutting down the trunk ports on Switch 2, which connects to Switch 1, Switch 3, and Switch 4, which forces Fast Ethernet downstream port into error-disable state. If you have configured this correctly, as shown in [Example 2-7](#), you have scored 3 points.

Example 2-7 SW2 Link-State Tracking Configuration and Verification

[Click here to view code image](#)

```
SW2(config)# link state track 1
SW2(config)# interface fast0/5
SW2(config-if)# link state group 1 downstream
SW2(config-if)# interface fastethernet 0/19
SW2(config-if)# link state group 1 upstream
SW2(config-if)# interface fastethernet 0/21
SW2(config-if)# link state group 1 upstream
SW2(config-if)# interface fastethernet 0/23
SW2(config-if)# link state group 1 upstream

SW2# show interface fastethernet 0/5 | include connected
fastethernet0/5 is up, line protocol is up (connected)
```

```
SW2(config-if)# int fast 0/19
SW2(config-if)# shut
SW2(config-if)# int fast 0/21
SW2(config-if)# shut
SW2(config-if)# int fast 0/23
SW2(config-if)# shut
```

```
SW2# show interface fastethernet 0/5 | include err-disabled
fastethernet0/5 is down, line protocol is down (err-disabled)
```

- Configure interfaces Fast Ethernet 0/9 and 0/10 on Switch 1 so that even if they are configured to belong to the same VLAN they cannot forward unicast, broadcast, or multicast traffic to one another. Do not use any form of ACL or configure the ports to belong to a PVLAN. (1 point)

You are required to configure the interfaces with the command **switchport protected** to ensure that no traffic is forwarded between these ports. Traffic is forwarded as normal between a protected and an unprotected port. If you have configured this correctly, you have scored 1 point.

Section 2: IPv4 IGP Protocols (26 Points)

Section 2.1: EIGRP

- Configure EIGRP per [Figure 2-5](#) using an instance name of CCIE and autonomous system of 1. Each EIGRP router should have its Loopback 0 interface configured and advertised within EIGRP. (2 points)

Use vanilla EIGRP with a virtual instance configuration in preparation for the following questions. If you have configured this correctly, as shown in [Example 2-8](#), you have scored 2 points.

Example 2-8 EIGRP Configuration and Verification

[Click here to view code image](#)

```
R1(config)# router eigrp CCIE
R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# net 120.100.1.0 0.0.0.255
R1(config-router-af)# net 120.100.123.0 0.0.0.255
R1(config-router-af)# net 120.100.100.0 0.0.0.255

R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router-af)# network 120.100.2.0 0.0.0.255
R2(config-router-af)# network 120.100.123.0 0.0.0.255
R2(config-router-af)# network 120.100.200.0 0.0.0.255

R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# network 120.100.3.0 0.0.0.255
R3(config-router-af)# network 120.100.123.0 0.0.0.255
R3(config-router-af)# network 120.100.34.0 0.0.0.255

R4(config-router)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# network 120.100.4.0 0.0.0.255
R4(config-router-af)# network 120.100.34.0 0.0.0.255

R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# network 120.100.5.0 0.0.0.255
R5(config-router-af)# network 120.100.34.0 0.0.0.255

R1# sh ip route eigrp
  120.0.0.0/24 is subnetted, 9 subnets
D    120.100.4.0 [90/158720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D    120.100.5.0 [90/158720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D    120.100.2.0 [90/156160] via 120.100.123.2, 00:23:32, GigabitEthernet0/0
D    120.100.3.0 [90/156160] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D    120.100.34.0 [90/30720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D    120.100.200.0 [90/30720] via 120.100.123.2, 00:23:32, GigabitEthernet0/0
```

- Configure R1 to advertise a summary route of 120.100.0.0/16 outbound on its VLAN132 interface. R3 should see the original VLAN 100 and Loopback 0 individual routes in addition to the summary route. You can only use one summary route in your configuration. Do not apply the summary configuration directly to the interface. (3 points)

Summarization will by default block all longer prefixes covered by the supernet configured on an interface; therefore, the VLAN 100 and Loopback 0 route from R1 would not be seen by R3. Allowing specific routes to be advertised with summary routes can be a valid requirement. One method used to achieve this is by configuring multiple summary routes, but the question does not permit this approach. To facilitate the specific routes with the summary, a leak map should be configured to match the VLAN 100 and Loopback 0 interfaces on R1. The leak map, which is configured per a normal route map, is then applied to the standard summary route statement on R1.

Because you cannot apply the summary configuration directly to the interface as per earlier EIGRP configuration, you must apply it to the address family af-interface within the Enhanced Interior Gateway Routing Protocol (EIGRP) instance. If you have configured this correctly, as shown in [Example 2-9](#), you have scored 3 points.

Example 2-9 R1 Leak Map Configuration and Verification

[Click here to view code image](#)

```
R1(config)# route-map LEAK-VLAN-100-LOOP0 permit 10
R1(config-route-map)# match ip address 1
R1(config-route-map)# exit
R1(config)# access-list 1 permit 120.100.100.0
R1(config)# access-list 1 permit 120.100.1.0
R1(config)# router eigrp CCIE
R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# af-interface Gigabit0/0
R1(config-router-af-interface)# summary-address 120.100.0.0 255.255.0.0 leak-map
LEAK-VLAN-100-LOOP0

R3# show ip route eigrp

R3# show ip route eigrp
  120.0.0.0/8 is variably subnetted, 10 subnets, 2 masks
D      120.100.4.0/24 [90/156160] via 120.100.34.4, 00:23:32, GigabitEthernet0/0
D      120.100.5.0/24 [90/156160] via 120.100.34.5, 00:23:32, GigabitEthernet0/0
D      120.100.0.0/16 [90/30720] via 120.100.123.1, 00:00:53, GigabitEthernet0/1
D      120.100.1.0/24 [90/156160] via 120.100.123.1, 00:23:32, GigabitEthernet0/1
D      120.100.2.0/24 [90/156160] via 120.100.123.2, 00:23:32, GigabitEthernet0/1
D      120.100.100.0/24 [90/30720] via 120.100.123.1, 00:23:32, GigabitEthernet0/1
D      120.100.200.0/24 [90/30720] via 120.100.123.2, 00:23:32, GigabitEthernet0/1
```

- Ensure that the length of time that EIGRP considers neighbors to be valid without receiving a hello packet on the VLAN 132 network between R1, R2, and R3 is 200 seconds. Do not change the hello-interval parameter. (2 points)

EIGRP considers neighbors to be valid up to three times the hello interval, the VLAN 132 network is a high-speed link, and hello packets will be sent every 5 seconds. You could usually tune the hold time by manipulating the hello intervals on an interface, but this question ensures that you can achieve the desired result only by manually changing the hold time to 200 under the VLAN 132 interfaces of routers R1, R2, and R3. [Example 2-10](#) shows the required configuration and verification of hold time by displaying the neighbors' statistics as seen by R3. If you have configured this correctly, as shown in [Example 2-10](#) (either directly under the interfaces or within the EIGRP address family af-interface), you have scored 2 points.

Example 2-10 EIGRP Hold-Time Configuration and Verification

[Click here to view code image](#)

```
R1(config)# interface GigabitEthernet0/0
R1(config-if)# ip hold-time eigrp 1 200
R1(config-if)
```

Enter configuration commands, one per line. End with CNTL/Z.

```
R2(config)# interface fastethernet0/0
R2(config-if)# ip hold-time eigrp 1 200
R2(config-if)
```

```
R3(config)# interface GigabitEthernet0/0
```

```
R3(config-if)# ip hold-time eigrp 1 200
```

```
R3(config-if)# do sh ip eigrp neighbors
```

IP-EIGRP neighbors for process 1

H	Address	Interface	Hold (sec)	Uptime (ms)	SRTT	RTO	Q	Seq Cnt	Num
3	120.100.123.1	Gi0/1	198	00:00:57	3	200	0	25	
2	120.100.123.2	Gi0/1	199	00:01:00	3	200	0	18	
1	120.100.34.5	Gi0/0	12	00:23:32	1	200	0	21	
0	120.100.34.4	Gi0/0	12	00:23:35	35	210	0	22	

- Configure new loopback interfaces on R1 and R2 using a Loopback 2 interface with an identical IP address of 150.101.1.1/24 on both routers; advertise this network into EIGRP on each router. Ensure that R3 prefers the route from R1 by manipulating the delay associated with this route. Do not manually adjust the delay associated with the interface by use of the **delay** command, and you are permitted to configure only R2 to influence the delay. (3 points)

R3 will receive identical routes from both R1 and R2 for network 150.101.1.0/24; therefore, both routes will be stored in the topology and routing table. R2 could influence the metric calculated by R3 by manipulating the delay of the new loopback interface or of the Ethernet interface connecting to R3, but this is not permitted. Because configuration is required solely on R2, the only method available is to create an offset list, which enables you to match specific routes and append further delay to them as they are advertised on R2 toward R3. If the offset list is not applied to the VLAN 132 interface, it would affect the whole process and not just advertisements toward R3. [Example 2-11](#) shows the configuration required to advertise the new routes and the routes as they are received on R3. Initial delay is shown to be 5100μS. Post configuration of the offset list on R2, the delay is seen to increase to 5103μS for the route received from R2; therefore, the route installed into the routing table of R3 is then the original advertised from R1 with the more appealing value of 5100μS. If you have configured this correctly, as shown in [Example 2-11](#), you have scored 3 points.

Example 2-11 EIGRP Configuration and Verification

[Click here to view code image](#)

```
R1(config)# interface Loopback2
R1(config-if)# ip address 150.101.1.1 255.255.255.0
R1(config-if)# router eigrp CCIE
R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# net 150.101.1.0 0.0.0.255
```

```
R2(config)# interface Loopback2
R2(config-if)# ip address 150.101.1.1 255.255.255.0
R2(config-if)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router-af)# net 150.101.1.0 0.0.0.255
```

```
R3# show ip route 150.101.1.0
Routing entry for 150.101.1.0/24
Known via "eigrp 1", distance 90, metric 156160, type internal
Redistributing via eigrp 1
Last update from 120.100.123.2 on fastethernet1/1, 00:00:23 ago
Routing Descriptor Blocks:
  120.100.123.2, from 120.100.123.2, 00:00:23 ago, via fastethernet1/1
    Route metric is 156160, traffic share count is 1
    Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
* 120.100.123.1, from 120.100.123.1, 00:00:23 ago, via fastethernet1/1
  Route metric is 156160, traffic share count is 1
  Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
  Reliability 255/255, minimum MTU 1500 bytes
  Loading 1/255, Hops 1
```

```
R3# show ip eigrp topology 150.101.1.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 150.101.1.0/24
State is Passive, Query origin flag is 1, 2 Successor(s), FD is 156160
Routing Descriptor Blocks:
  120.100.123.1 (GigabitEthernet0/1), from 120.100.123.1, Send flag is 0x0
    Composite metric is (156160/128256), Route is Internal
    Vector metric:
      Minimum bandwidth is 100000 Kbit
      Total delay is 5100 microseconds
      Reliability is 255/255
      Load is 1/255
      Minimum MTU is 1500
      Hop count is 1
  120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
    Composite metric is (156160/128256), Route is Internal
    Vector metric:
      Minimum bandwidth is 100000 Kbit
      Total delay is 5100 microseconds
      Reliability is 255/255
      Load is 1/255
      Minimum MTU is 1500
      Hop count is 1
```

```
R2# show interface Fast0/0 | include DLY
MTU 1500 bytes, BW 100000 Kbit/sec, DLY 100 usec,
R2(config)# access-list 1 permit 150.101.1.0
R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router-af)# topology base
R2(config-router-af)# offset-list 1 out 100 fastethernet0/0

R3# show ip route 150.101.1.0
Routing entry for 150.101.1.0/24
Known via "eigrp 1", distance 90, metric 156160, type internal
Redistributing via eigrp 1
Last update from 120.100.123.1 on GigabitEthernet0/1, 00:00:17 ago
Routing Descriptor Blocks:
* 120.100.123.1, from 120.100.123.1, 00:00:17 ago, via GigabitEthernet0/1
```

```
Route metric is 156160, traffic share count is 1
Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
Reliability 255/255, minimum MTU 1500 bytes
Loading 1/255, Hops 1
```

```
R3# show ip eigrp topology 150.101.1.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 150.101.1.0/24
  State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
  Routing Descriptor Blocks:
    120.100.123.1 (GigabitEthernet0/1), from 120.100.123.1, Send flag is 0x0
      Composite metric is (156160/128256), Route is Internal
      Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
    120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
      Composite metric is (156260/128356), Route is Internal
      Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5103 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
```

- Configure EIGRP with a new instance name of CCIE2 between R2 and R3 over VLAN 132 with an autonomous system of 2 and 256-bit encryption with a password of lake2aho3. Any additional connections to AS2 should be encrypted using the same password without further configuration on R2 and R3. Configure a new loopback interface on R2 (Loopback 3) with an IP address of 150.101.2.1/24, and advertise this and only this network to R3 from R2. (2 points)

This straightforward configuration within a new EIGRP instance facilitates subsequent redistribution between EIGRP AS1 to AS2. The only twist to the question is to perform authentication without the need for further configuration should there be additional peering to AS2. The simple fix to this is to apply authentication to all interfaces using the **af-interface default** command. [Example 2-16](#) shows the basic EIGRP configuration on R2 and R3 with HMAC authentication. If you have configured this correctly, as shown in [Example 2-12](#), you have scored 2 points.

Example 2-12 R2 and R3 EIGRP AS2 Configuration and Verification

[Click here to view code image](#)

```
R2(config)# interface Loopback3
R2(config-if)# ip add 150.101.2.1 255.255.255.0
R2(config-if)# router eigrp CCIE2
R2(config-router)# address-family ipv4 unicast autonomous-system 2
R2(config-router-af)# af-interface default
R2(config-router-af-interface)# authentication mode hmac-sha-256 0 lake2aho3
R2(config-router-af-interface)# exit
R2(config-router-af)# network 150.101.2.0 0.0.0.255
R2(config-router-af)# network 120.100.123.0 0.0.0.255
```

```
R3(config)# router eigrp CCIE2
R3(config-router)# address-family ipv4 unicast autonomous-system 2
R3(config-router-af)# af-interface default
R3(config-router-af-interface)# authentication mode hmac-sha-256 0 lake2aho3
R3(config-router-af-interface)# exit
R3(config-router-af)# network 120.100.123.0
R3(config-router-af)# sh ip route eigrp 2
    150.101.0.0/24 is subnetted, 2 subnets
D      150.101.2.0 [90/156160] via 120.100.123.2, 00:00:25, GigabitEthernet0/1
```

Section 2.2: OSPF

- Configure OSPF per [Figure 2-6](#) using a process ID of 1; all OSPF configuration where possible should not be configured under the process ID. Each OSPF router should also have its Loopback 0 interface configured and advertised within OSPF as follows: (2 points)

R4 Loopback 0 – Area 0
R5 Loopback 0 – Area 0
R6 Loopback 0 – Area 1
SW1 Loopback 0 – Area 2
SW2 Loopback 0 – Area 1
SW3 Loopback 0 – Area 2
SW4 Loopback 0 – Area 3

As per Lab 1, the question directs you to configure OSPF directly under the interfaces of the routers; the switches still require configuration under the OSPF process running this version of IOS. Did you notice that Area 0 is partitioned? If you have configured this correctly, as shown in [Example 2-13](#), you have scored 2 points. Consider using the **show ip ospf interface** command to verify your configuration.

Example 2-13 Initial OSPF Configuration

[Click here to view code image](#)

```
R4(config)# interface Loopback 0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# exit
R4(config)# interface GigabitEthernet 0/1
R4(config-if)# ip ospf 1 area 1

R5(config)# interface Loopback 0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# exit
R5(config)# interface GigabitEthernet 0/1
R5(config-if)# ip ospf 1 area 2

R6(config)# interface Loopback 0
R6(config-if)# ip ospf 1 area 1
R6(config-if)# interface GigabitEthernet 0/0
R6(config-if)# ip ospf 1 area 1
```

```

R6(config-if)# interface GigabitEthernet 0/1
R6(config-if)# ip ospf 1 area 3

SW1(config)# ip routing
SW1(config)# router ospf 1
SW1(config-router)# network 120.100.7.1 0.0.0.0 area 2
SW1(config-router)# network 120.100.53.1 0.0.0.0 area 2

SW2(config)# ip routing
SW2(config-if)# router ospf 1
SW2(config-router)# net 120.100.8.1 0.0.0.0 area 1
SW2(config-router)# net 120.100.46.2 0.0.0.0 area 1

SW3(config)# ip routing
SW3(config)# router ospf 1
SW3(config-router)# network 120.100.53.3 0.0.0.0 area 2
SW3(config-router)# network 120.100.63.3 0.0.0.0 area 3
SW3(config-router)# network 120.100.9.1 0.0.0.0 area 2

SW4(config)# ip routing
SW4(config)# router ospf 1
SW4(config-router)# network 120.100.10.1 0.0.0.0 area 3
SW4(config-router)# network 120.100.63.4 0.0.0.0 area 3

```

- Area 0 is partitioned between R4 and R5; ensure that your network can accommodate this issue.
You are not permitted to form any Area 0 neighbor relationship directly between R4 and R5 to join Area 0. (4 points)

A fundamental rule of the Open Shortest Path First (OSPF) Protocol is not to design your network with a partitioned backbone Area 0 or partition if of a failure condition occurs. A virtual link between R4 and R5 would not work here because you would need to transit multiple OSPF areas. A tunnel between the two routers is also not permitted because this would form a direct neighbor relationship. You are required to configure a virtual link between R5 and Switch 3 to propagate Area 3 routes and similarly between R4 and R6. By then creating an additional virtual link between R6 and Switch 3, the two effective halves of the network have been joined at an Area 0 level. Remember to configure all virtual links to the router ID of the remote router as opposed to the physical IP address on the corresponding interface. [Example 2-14](#) shows the required configuration to create virtual links between R5-SW3, R4-R6, and R6-SW3. The resulting routing table verification on Switch 4 shows all networks are being learned correctly post configuration. If you have configured this correctly, as shown in [Example 2-14](#), you have scored 4 points.

Example 2-14 OSPF Virtual-Link Configuration and Routing Table Verification

[Click here to view code image](#)

```

R5(config)# router ospf 1
R5(config-router)# area 2 virtual-link 120.100.9.1

SW3(config-router)# router ospf 1
SW3(config-router)# area 2 virtual-link 120.100.5.1

R4(config)# router ospf 1
R4(config-router)# area 1 virtual-link 120.100.6.1

```

```

R6(config-if)# router ospf 1
R6(config-router)# area 1 virtual-link 120.100.4.1
R6(config-router)# area 3 virtual-link 120.100.9.1

SW3(config-if)# router ospf 1
SW3(config-router)# area 3 virtual-link 120.100.6.1

SW4# sh ip route ospf
    120.0.0.0/8 is variably subnetted, 10 subnets, 2 masks
O IA    120.100.9.1/32 [110/2] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.8.1/32 [110/3] via 120.100.63.6, 00:00:54, Vlan63
O IA    120.100.5.1/32 [110/3] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.4.1/32 [110/3] via 120.100.63.6, 00:00:54, Vlan63
O IA    120.100.7.1/32 [110/3] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.6.1/32 [110/2] via 120.100.63.6, 00:00:54, Vlan63
O IA    120.100.53.0/24 [110/2] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.46.0/24 [110/2] via 120.100.63.6, 00:00:55, Vlan63

```

Section 2.3: Redistribution

- Perform a one-way redistribution of EIGRP AS2 into EIGRP AS1 on R3 using the following default metric: 1544 20000 255 1 1500. Ensure that R1 shows a next hop for the AS1 advertised route of 150.101.2.0/24 of R2. Perform configuration only on R3 for this task. (3 points)

This is a simple redistribution question. Upon inspection, you would believe the only complexity would be that of modifying the next-hop attribute for R1, which would by default show as R3 for the AS2 route advertised by R2. In fact, you would find that the AS2 route would not be seen on R1 post redistribution from R3. This is due to an inherent safety mechanism within EIGRP that will cause redistribution issues with routers that have duplicate EIGRP router IDs. Pre-lab configuration ensured that both R1 and R3 have the same Loopback 255 IP address, which will force the router ID to be identical. [Example 2-15](#) shows the redistribution configuration on R3. The AS2 route of 150.101.2.0/24 is received on R3 but is absent on R1. Inspection of the EIGRP topology table for the route on R3 shows that it is being advertised into EIGRP and that the router ID of R3 is 200.200.200.200; similarly, the router ID of R1 is also 200.200.200.200. If you change the router ID of R3 to that of its Loopback 0 interface (120.100.3.1), the route is then accepted by R1, but of course a next hop is shown as R3, even though R2 resides on the same IP subnet as R1 and R2 and is the originating router. The EIGRP third-party next-hop feature can be used to modify the next-hop attribute with a router redistributing another routing protocol into EIGRP in a similar manner to that of BGP. If you have configured this correctly, as shown in [Example 2-15](#), you have scored 3 points.

Example 2-15 R3 EIGRP Redistribution Configuration and Verification

[Click here to view code image](#)

```

R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# topology base
R3(config-router-af-topology)# redistribute eigrp 2

```

```
R1# show ip route 150.101.2.0
```

```
% Subnet not in table
```

```
R3# show ip eigrp topology 150.101.2.0/24
```

```
IP-EIGRP (AS 1): Topology entry for 150.101.2.0/24
```

```
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
```

```
Routing Descriptor Blocks:
```

```
120.100.123.2, from Redistributed, Send flag is 0x0
```

```
Composite metric is (156160/0), Route is External
```

```
Vector metric:
```

```
Minimum bandwidth is 100000 Kbit
```

```
Total delay is 5100 microseconds
```

```
Reliability is 255/255
```

```
Load is 1/255
```

```
Minimum MTU is 1500
```

```
Hop count is 1
```

```
External data:
```

```
Originating router is 200.200.200.200 (this system)
```

```
AS number of route is 2
```

```
External protocol is EIGRP, external metric is 156160
```

```
Administrator tag is 0 (0x00000000)
```

```
IP-EIGRP (AS 2): Topology entry for 150.101.2.0/24
```

```
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
```

```
Routing Descriptor Blocks:
```

```
120.100.123.2 (GigaEthernet0/1), from 120.100.123.2, Send flag is 0x0
```

```
Composite metric is (156160/128256), Route is Internal
```

```
Vector metric:
```

```
Minimum bandwidth is 100000 Kbit
```

```
Total delay is 5100 microseconds
```

```
Reliability is 255/255
```

```
Load is 1/255
```

```
Minimum MTU is 1500
```

```
Hop count is 1
```

```
R3# show ip eigrp topology | include ID
```

```
IP-EIGRP Topology Table for AS(1)/ID(200.200.200.200)
```

```
IP-EIGRP Topology Table for AS(2)/ID(200.200.200.200)
```

```
R1# show ip eigrp topology | include ID
```

```
IP-EIGRP Topology Table for AS(1)/ID(200.200.200.200)
```

```
R1#
```

```
R3(config)# router eigrp CCIE
```

```
R3(config-router)# address-family ipv4 unicast autonomous-system 1
```

```
R3(config-router-af)# eigrp router-id 120.100.3.1
```

```
R3# show ip eigrp topology | include ID
```

```
IP-EIGRP Topology Table for AS(1)/ID(120.100.3.1)
```

```
R3# show ip eigrp topology 150.101.2.0/24
```

```
IP-EIGRP (AS 1): Topology entry for 150.101.2.0/24
```

```
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
```

```
Routing Descriptor Blocks:
```

```
120.100.123.2, from Redistributed, Send flag is 0x0
```

```
Composite metric is (156160/0), Route is External
```

```
Vector metric:
```

```
Minimum bandwidth is 100000 Kbit
```

```
Total delay is 5100 microseconds
```

```

Reliability is 255/255
Load is 1/255
Minimum MTU is 1500
Hop count is 1
External data:
    Originating router is 120.100.3.1 (this system)
    AS number of route is 2
    External protocol is EIGRP, external metric is 156160
    Administrator tag is 0 (0x00000000)
IP-EIGRP (AS 2): Topology entry for 150.101.2.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
Routing Descriptor Blocks:
120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
    Composite metric is (156160/128256), Route is Internal
Vector metric:
    Minimum bandwidth is 100000 Kbit
    Total delay is 5100 microseconds
    Reliability is 255/255
    Load is 1/255
    Minimum MTU is 1500
    Hop count is 1

```

```

R1# show ip route 150.101.2.0
Routing entry for 150.101.2.0/24
Known via "eigrp 1", distance 170, metric 158720, type external
Redistributing via eigrp 1
Last update from 120.100.123.3 on GigabitEthernet0/0, 00:03:06 ago
Routing Descriptor Blocks:
* 120.100.123.3, from 120.100.123.3, 00:03:06 ago, via GigabitEthernet0/0
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1

```

```

R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# af-interface GigabitEthernet0/1
R3(config-router-af-interface)# no next-hop-self

```

```

R1# show ip route 150.101.2.0
Routing entry for 150.101.2.0/24
Known via "eigrp 1", distance 170, metric 158720, type external
Redistributing via eigrp 1
Last update from 120.100.123.2 on Gigabit0/0, 00:00:24 ago
Routing Descriptor Blocks:
* 120.100.123.2, from 120.100.123.3, 00:00:24 ago, via GigabitEthernet0/0
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
Reliability 255/255, minimum MTU 1500 bytes
Loading 1/255, Hops 1

```

-
- Perform mutual redistribution of EIGRP AS1 and OSPF on R4 and R5. Use a metric of 5000 for redistributed routes into OSPF, which should appear as external type 2 routes and the following K values for OSPF routes redistributed into EIGRP: 1544 20000 255 1 1500. (2 points)

This is an unambiguous redistribution question that sets the scene for the question that follows.

Example 2-16 shows the required configuration on R4 and R5 with verification of external EIGRP

received routes on R3. Because the metrics are identical on R4 and R5, there are multiple routes with load-sharing potential. If you have configured this correctly, you have scored 2 points.

Example 2-16 R4 and R5 Redistribution Configuration and Verification on R3

[Click here to view code image](#)

```
R4(config-router)# router ospf 1
R4(config-router)# redistribute eigrp 1 subnets
R4(config-router)# default-metric 5000
R4(config-router)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topology)# redistribute ospf 1
R4(config-router-af-topology)# default-metric 1544 20000 255 1 1500

R5(config-router)# router ospf 1
R5(config-router)# redistribute eigrp 1 subnets
R5(config-router)# default-metric 5000
R5(config-router)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# topology base
R5(config-router-af-topology)# redistribute ospf 1
R5(config-router-af-topology)# default-metric 1544 20000 255 1 1500

R3# show ip route eigrp
      150.101.0.0/24 is subnetted, 2 subnets
D        150.101.1.0 [90/2297856] via 120.100.123.1, 00:05:05, Gigabit0/1
      120.0.0.0/8 is variably subnetted, 20 subnets, 3 masks
D EX    120.100.9.1/32
          [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
          [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.8.1/32
          [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
          [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.10.1/32
          [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
          [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.5.1/32
          [170/6780416] via 120.100.34.4, 00:01:51, GigabitEthernet0/0
D      120.100.4.0/24
          [90/156160] via 120.100.34.4, 00:07:17, GigabitEthernet0/0
D      120.100.5.0/24
          [90/156160] via 120.100.34.5, 00:07:17, GigabitEthernet0/0
D EX    120.100.4.1/32
          [170/6780416] via 120.100.34.5, 00:00:23, GigabitEthernet0/0
D EX    120.100.7.1/32
          [170/6780416] via 120.100.34.5, 00:00:23, GigabitEthernet0/0
          [170/6780416] via 120.100.34.4, 00:00:23, GigabitEthernet0/0
D EX    120.100.6.1/32
          [170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
          [170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0
D      120.100.0.0/16 [90/2172416] via 120.100.123.1, 00:05:07, GigabitEthernet0/1
D      120.100.1.0/24 [90/2297856] via 120.100.123.1, 00:05:07, GigabitEthernet0/1
D      120.100.2.0/24 [90/2297856] via 120.100.123.2, 00:05:07, GigabitEthernet0/1
D EX    120.100.63.0/24
          [170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
```

```

[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0
D EX 120.100.53.0/24
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0
D EX 120.100.46.0/24
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0
D 120.100.100.0/24 [90/2172416] via 120.100.123.1, 00:05:07,
GigabitEthernet0/1
D 120.100.200.0/24 [90/2172416] via 120.100.123.2, 00:05:08,
GigabitEthernet0/1

```

- R3 will have equal-cost external EIGRP routes to the redistributed OSPF subnet 120.100.63.0/24 (VLAN 63). Configure only R3 to ensure that R3 routes via a next hop of R5 (120.100.34.5) for this destination subnet. If this route fails, the route advertised from R4 (120.100.34.4) should be used dynamically. (3 points)

[Example 2-20](#) shows both routes for 120.100.63.0/24 received on R3 from R4 and R5. Because all routers share a common media, the interface connecting to R4 or R5 cannot be modified on R3 because this would affect both routes. Similarly, an offset list to manipulate delay would be of no use because you are permitted to configure only R3. You are therefore required to penalize the route received from R4 only to ensure that the R5-generated route is preferred on R3. By configuring a route map on R3 to match only the route source of R4, you can increase the metric for the required route (120.100.63.0/24). This simply enables the original route received from R5 to take precedence. [Example 2-17](#) shows the required configuration and verification that the route is preferred via the R5. The topology table shows that the R4 route is also present and that R4 is effectively the feasible successor for this network on this router. If the route from R5 is withdrawn, the route from R5 would enter the routing table automatically. You will need a second **permit** statement on the route map (**permit 20**) to enable all other routes inbound to R3 to enter unaltered. [Example 2-17](#) also details the routing tables of each device to confirm redistribution from EIGRP into OSPF or vice versa. If you have configured this correctly, as shown in [Example 2-17](#), you have scored 3 points.

Example 2-17 R3 OSPF Redistribution Configuration and Verification

[Click here to view code image](#)

```

R3# show ip route 120.100.63.0
Routing entry for 120.100.63.0/24
Known via "eigrp 1", distance 170, metric 6780416, type external
Redistributing via eigrp 1
Last update from 120.100.34.5 on GigabitEthernet0/0, 00:01:59 ago
Routing Descriptor Blocks:
  120.100.34.5, from 120.100.34.5, 00:01:59 ago, via GigabitEthernet0/0
    Route metric is 6780416, traffic share count is 1
    Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
* 120.100.34.4, from 120.100.34.4, 00:01:59 ago, via GigabitEthernet0/0
  Route metric is 6780416, traffic share count is 1
  Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
  Reliability 255/255, minimum MTU 1500 bytes
  Loading 1/255, Hops 1

```

```
R3(config)# access-list 1 permit 120.100.34.4
R3(config)# access-list 2 permit 120.100.63.0
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# topology base
R3(config-router-af-topology)# distribute-list route-map PENALISE-VLAN63 in
    GigabitEthernet0/0
R3(config-router-af-topology)# exit
R3(config-router-af)# exit
R3(config-router)# exit
R3(config)# route-map PENALISE-VLAN63 permit 10
R3(config-route-map)# match ip address 2
R3(config-route-map)# match ip route-source 1
R3(config-route-map)# set metric +500000
R3(config-route-map)# route-map PENALISE-VLAN63 permit 20
```

```
R3# show ip route 120.100.63.0
Routing entry for 120.100.63.0/24
Known via "eigrp 1", distance 170, metric 6780416, type external
Redistributing via eigrp 1
Last update from 120.100.34.5 on GigabitEthernet0/0, 00:00:21 ago
Routing Descriptor Blocks:
* 120.100.34.5, from 120.100.34.5, 00:00:21 ago, via GigabitEthernet0/0
    Route metric is 6780416, traffic share count is 1
    Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R3# show ip eigrp topology 120.100.63.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 120.100.63.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 6780416
Routing Descriptor Blocks:
120.100.34.5 (GigabitEthernet0/0), from 120.100.34.5, Send flag is 0x0
    Composite metric is (6780416/6777856), Route is External
    Vector metric:
        Minimum bandwidth is 1544 Kbit
Total delay is 200100 microseconds
    Reliability is 255/255
    Load is 1/255
    Minimum MTU is 1500
    Hop count is 1
    External data:
        Originating router is 120.100.5.1
        AS number of route is 1
        External protocol is OSPF, external metric is 2
        Administrator tag is 0 (0x00000000)
120.100.34.4 (GigabitEthernet0/0), from 120.100.34.4, Send flag is 0x0
    Composite metric is (128000000/6777856), Route is External
    Vector metric:
        Minimum bandwidth is 20 Kbit
        Total delay is 0 microseconds
        Reliability is 0/255
        Load is 0/255
        Minimum MTU is 0
        Hop count is 1
    External data:
        Originating router is 120.100.4.1
        AS number of route is 1
        External protocol is OSPF, external metric is 2
        Administrator tag is 0 (0x00000000)
```

Note

The full IP routing tables of each device are provided within the accompanying configurations to verify your redistributed routes.

Section 3: BGP (15 Points)

- Configure BGP peering per [Figure 2-7](#) as follows: iBGP R1-R3, R2-R3, R4-R6, R4-SW2, R5-SW1, R5-SW3, eBGP R3-R4, R3-R5, SW4-SW3, R6-SW4. Use loopback interfaces to peer on all routers with the exception of peering between R3-R4 and R3-R5. Do not use the command **ebgp-multihop** within your configurations. (3 points)

The restrictions within the internal Border Gateway Protocol (iBGP) peering require you to configure R3, R4, and R5 as route reflectors within their own autonomous system. Autosummarization is disabled to ensure BGP does not summarize routes, and synchronization is disabled because the internal gateway protocol (IGP) will not be synchronized to BGP within this lab. The question does not dictate that you must configure peer groups, but it is considered good practice when you have more than one peer with a similar peering configuration. The question does, however, dictate that you must not use **ebgp- multihop**. This feature would, of course, be required for the peering from AS400 to AS300 and AS400 to AS200 because loopback interfaces are used for the external peering here, unlike AS100 to AS200 and AS300, which peer from connected interfaces. Without **ebgp-multihop**, the peering fails inbound and outbound from AS400. The only way to fix this is to use a feature that disables connection verification to establish an external BGP (eBGP) peering session with a single-hop peer that uses a loopback interface. Use of the command **neighbor disable-connected-check** on R6, SW3, and SW4 for the required peering allows the peering to be formed successfully. [Example 2-18](#) shows the basic peering configuration for BGP, the eBGP failure condition observed on peering to and from AS400, and the required configuration to rectify the condition. If you have configured this correctly, you have scored 3 points.

Example 2-18 BGP Peering Configuration and Verification

[Click here to view code image](#)

```
R1(config)# router bgp 100
R1(config-router)# no auto-summary
R1(config-router)# no synchronization
R1(config-router)# neighbor 120.100.3.1 remote-as 100
R1(config-router)# neighbor 120.100.3.1 update-source Loopback0

R2(config)# router bgp 100
R2(config-router)# no auto-summary
R2(config-router)# no synchronization
R2(config-router)# neighbor 120.100.3.1 remote-as 100
R2(config-router)# neighbor 120.100.3.1 update-source Loopback0

R3(config)# router bgp 100
R3(config-router)# no auto-summary
R3(config-router)# no synchronization
R3(config-router)# neighbor AS100 peer-group
```

```
R3(config-router)# neighbor AS100 remote-as 100
R3(config-router)# neighbor AS100 update-source Loopback0
R3(config-router)# neighbor 120.100.1.1 peer-group AS100
R3(config-router)# neighbor 120.100.2.1 peer-group AS100
R3(config-router)# neighbor AS100 route-reflector-client
R3(config-router)# neighbor 120.100.34.4 remote-as 200
R3(config-router)# neighbor 120.100.34.5 remote-as 300
```

```
R4(config)# router bgp 200
R4(config-router)# router bgp 200
R4(config-router)# no auto-summary
R4(config-router)# no synchronization
R4(config-router)# neighbor AS200 peer-group
R4(config-router)# neighbor AS200 remote-as 200
R4(config-router)# neighbor AS200 update-source Loopback0
R4(config-router)# neighbor AS200 route-reflector-client
R4(config-router)# neighbor 120.100.6.1 peer-group AS200
R4(config-router)# neighbor 120.100.8.1 peer-group AS200
R4(config-router)# neighbor 120.100.34.3 remote-as 100
```

```
R5(config)# router bgp 300
R5(config-router)# no auto-summary
R5(config-router)# no synchronization
R5(config-router)# neighbor AS300 peer-group
R5(config-router)# neighbor AS300 remote-as 300
R5(config-router)# neighbor AS300 update-source Loopback0
R5(config-router)# neighbor AS300 route-reflector-client
R5(config-router)# neighbor 120.100.7.1 peer-group AS300
R5(config-router)# neighbor 120.100.9.1 peer-group AS300
R5(config-router)# neighbor 120.100.34.3 remote-as 100
```

```
R6(config)# router bgp 200
R6(config-router)# no auto-summary
R6(config-router)# no synchronization
R6(config-router)# neighbor 120.100.4.1 remote-as 200
R6(config-router)# neighbor 120.100.4.1 update-source Loopback0
R6(config-router)# neighbor 120.100.10.1 remote-as 400
R6(config-router)# neighbor 120.100.10.1 update-source Loopback0
```

```
SW1(config)# router bgp 300
SW1(config-router)# no auto-summary
SW1(config-router)# no synchronization
SW1(config-router)# neighbor 120.100.5.1 remote-as 300
SW1(config-router)# neighbor 120.100.5.1 update-source Loopback0
```

```
SW2(config)# router bgp 200
SW2(config-router)# no auto-summary
SW2(config-router)# no synchronization
SW2(config-router)# neighbor 120.100.4.1 remote-as 200
SW2(config-router)# neighbor 120.100.4.1 update-source Loopback0
```

```
SW3(config)# router bgp 300
SW3(config-router)# no auto-summary
SW3(config-router)# no synchronization
SW3(config-router)# neighbor 120.100.5.1 remote-as 300
SW3(config-router)# neighbor 120.100.5.1 update-source Loopback0
SW3(config-router)# neighbor 120.100.10.1 remote-as 400
SW3(config-router)# neighbor 120.100.10.1 update-source Loopback0
```

```

SW4(config)# router bgp 400
SW4(config-router)# no auto-summary
SW4(config-router)# no synchronization
SW4(config-router)# neighbor 120.100.6.1 remote-as 200
SW4(config-router)# neighbor 120.100.6.1 update-source Loopback0
SW4(config-router)# neighbor 120.100.9.1 remote-as 300
SW4(config-router)# neighbor 120.100.9.1 update-source Loopback0

SW4# sh ip bgp neigh 120.100.6.1 | include External
External BGP neighbor not directly connected.
SW4# show ip bgp neighbors 120.100.9.1 | include External
External BGP neighbor not directly connected.
SW4#

SW4# sh ip bgp neighbors 120.100.6.1 | include active
No active TCP connection
SW4# sh ip bgp neighbors 120.100.9.1 | include active
No active TCP connection

SW4(config-router)# neighbor 120.100.6.1 disable-connected-check
SW4(config-router)# neighbor 120.100.9.1 disable-connected-check

R6(config-router)# neighbor 120.100.10.1 disable-connected-check

SW3(config-router)# neighbor 120.100.10.1 disable-connected-check

SW4# show ip bgp neighbors 120.100.6.1 | include Established
BGP state = Established, up for 00:02:01
SW4# show ip bgp neighbors 120.100.9.1 | include Established
BGP state = Established, up for 00:02:05

```

You will also find peering issues between R1 and R3. [Example 2-19](#) shows the routers are informing each other they have an incorrect BGP identifier. This is simply because both routers have identical loopback interface address of 200.200.200.200, which is used as the BGP identifier. By changing the ID of one router, the peering is established. It does not matter what you change the ID to, but it needs to be unique; the Loopback 0 interface would be a good choice. No extra points for this task because this is part of the original peering.

Example 2-19 R1 and R3 Peering Issue Configuration and Verification

[Click here to view code image](#)

```

R1# * 19:30:13.287: %BGP-3-NOTIFICATION: sent to neighbor 120.100.3.1 2/3 (BGP
identifier wrong) 4 bytes C8C8C8C8

R3# * 19:25:30.043: %BGP-3-NOTIFICATION: received from neighbor 120.100.1.1 2/
3 (BGP identifier wrong) 4 bytes C8C8C8C8

R1# show ip bgp summary | include identifier
BGP router identifier 200.200.200.200, local AS number 100

R3# show ip bgp summary | include identifier
BGP router identifier 200.200.200.200, local AS number 100

```

```
R1(config-router)# bgp router-id 120.100.1.1
*19:34:45.467: %BGP-5-ADJCHANGE: neighbor 120.100.3.1 Up
```

- Routers R1 and R2 in AS100 should be made to passively accept only BGP sessions. R3 should be configured to actively create only BGP sessions to R1 and R2 within AS100. (3 points)

A BGP speaker by default will attempt to open a session on TCP port 179 with a configured peer, because such a normal peering arrangement will see two sessions being established to build a successful neighbor relationship. This behavior can be modified to effectively allow sessions to be established only either inbound or outbound. The solution to the question is achieved by configuring the **neighbor transport connection-mode** to passive (only inbound connections will be established) on R1 and R2 and active (only outbound sessions will be established) on R3. You must manually activate each neighbor on each router for the solution to work effectively. If you have configured this correctly, as shown in [Example 2-20](#), you have scored 3 points. Consider using the **show ip bgp summary** command to verify your configuration.

Example 2-20 R1, R2, and R3 Connection Mode Configuration

[Click here to view code image](#)

```
R1(config)# router bgp 100
R1(config-router)# neighbor 120.100.3.1 transport connection-mode passive
R1(config-router)# neighbor 120.100.3.1 activate

R2(config)# router bgp 100
R2(config-router)# neighbor 120.100.3.1 transport connection-mode passive
R2(config-router)# neighbor 120.100.3.1 activate

R3(config)# router bgp 100
R3(config-router)# neighbor AS100 transport connection-mode active
R3(config-router)# neighbor 120.100.1.1 activate
R3(config-router)# neighbor 120.100.2.1 activate
```

- Configure the following loopback interfaces on R3 and SW4; advertise these networks into BGP using the **network** command: (2 points)

R3 – Loopback interface 5 (152.100.100.1/24)
SW4 – Loopback interface 5 (152.200.32.1/24)
SW4 – Loopback interface 6 (152.200.33.1/24)
SW4 – Loopback interface 7 (152.200.34.1/24)
SW4 – Loopback interface 8 (152.200.35.1/24)

This simple question creates BGP routes for the following task. If you have configured this correctly, as shown in [Example 2-21](#), you have scored 2 points.

Example 2-21 R3 and SW4 Network Advertisement Configuration and Verification

[Click here to view code image](#)

```

R3(config)# interface Loopback5
R3(config-if)# ip address 152.100.100.1 255.255.255.0
R3(config-if)# router bgp 100
R3(config-router)# network 152.100.100.0 mask 255.255.255.0

SW4(config)# interface Loopback5
SW4(config-if)# ip address 152.200.32.1 255.255.255.0
SW4(config-if)# interface Loopback6
SW4(config-if)# ip address 152.200.33.1 255.255.255.0
SW4(config-if)# interface Loopback7
SW4(config-if)# ip address 152.200.34.1 255.255.255.0
SW4(config-if)# interface Loopback8
SW4(config-if)# ip address 152.200.35.1 255.255.255.0
SW4(config-if)# router bgp 400
SW4(config-router)# network 152.200.32.0 mask 255.255.255.0
SW4(config-router)# network 152.200.33.0 mask 255.255.255.0
SW4(config-router)# network 152.200.34.0 mask 255.255.255.0
SW4(config-router)# network 152.200.35.0 mask 255.255.255.0

```

```

R3# show ip bgp
BGP table version is 10, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 152.100.100.0/24	0.0.0.0	0		32768	i
* 152.200.32.0/24	120.100.34.4			0	200 400 i
*>	120.100.34.5			0	300 400 i
* 152.200.33.0/24	120.100.34.4			0	200 400 i
*>	120.100.34.5			0	300 400 i
* 152.200.34.0/24	120.100.34.4			0	200 400 i
*>	120.100.34.5			0	300 400 i
* 152.200.35.0/24	120.100.34.4			0	200 400 i
*>	120.100.34.5			0	300 400 i

- Configure R3 to inform R4 that it does not want to receive routes advertised from SW4 for networks 152.200.33.0/24, 152.200.34.0/24, and 152.200.35.0/24. Achieve this in such a manner that R4 does not actually advertise these routes toward R3. You may also configure R4. (4 points)

BGP has a prefix-based outbound route filtering (ORF) mechanism that can send and receive capabilities to minimize BGP updates sent between BGP peers. Advertisement of ORF capability indicates that a peer will accept a prefix list from a neighbor and apply the prefix list received from a neighbor locally to avoid the unnecessary sending of routes that would be blocked by the receiver anyway. R3 is therefore configured with a prefix list that blocks the required routes generated from SW4, which is sent via ORF to R4. R4 is configured to receive this prefix list via ORF, and the routes are blocked outbound at R4. [Example 2-2](#) shows the required ORF and prefix-list filtering with the resulting outbound advertisement on R4. The BGP table on R3 is also displayed showing the routes are no longer being received from R4 and solely from R5. If you have configured this correctly, as shown in [Example 2-22](#), you have scored 4 points.

Example 2-22 BGP ORF Configuration and Verification

```
R3(config)# router bgp 100
R3(config-router)# neighbor 120.100.34.4 capability orf prefix-list send
R3(config-router)# neighbor 120.100.34.4 prefix-list FILTER in
R3(config)# ip prefix-list FILTER seq 5 deny 152.200.33.0/24
R3(config)# ip prefix-list FILTER seq 10 deny 152.200.34.0/24
R3(config)# ip prefix-list FILTER seq 15 deny 152.200.35.0/24
R3(config)# ip prefix-list FILTER seq 20 permit 0.0.0.0/0 le 32

R4(config)# router bgp 200
R4(config-router)# neighbor 120.100.34.3 capability orf prefix-list receive
R4(config-router)# exit
R4(config)# exit
R4# show ip bgp neighbors 120.100.34.3 advertised-routes
BGP table version is 17, local router ID is 120.100.4.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*->i152.200.32.0/24  120.100.10.1        0       100      0 400 i

Total number of prefixes 1

R3# clear ip bgp *
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*-> 152.100.100.0/24 0.0.0.0             0       32768   i
*-> 152.200.32.0/24  120.100.34.4        0       200     400 i
*          120.100.34.5        0       300     400 i
*> 152.200.33.0/24  120.100.34.5        0       300     400 i
*> 152.200.34.0/24  120.100.34.5        0       300     400 i
*> 152.200.35.0/24  120.100.34.5        0       300     400 i
```

- Configure a route map on R5 that prepends its local autonomous system an additional two times for network 152.200.32.0/24 when advertised to R3. The route map may contain multiple **permit** statements, but only one prepend is permitted per line. (3 points)

This is a simple autonomous system path prepend question, or so it seems. Normally you would prepend the same autonomous system number multiple times within the same **permit** statement, but the question restricts this, so you are forced to use multiple **permit** statements with the same autonomous system **prepend** statement. [Example 2-22](#) shows the route 152.200.32.0/24 as received initially on R3 from R5 with an autonomous system path of 300-400. After configuration of the route map to prepend the route on R5 twice, the network is received on R3 with an autonomous system path of 300-300-400. This might look like the route has indeed been prepended twice, but the question requests an “additional” two times; in fact, the route has been prepended only once. The problem is that the route map **permit 10** statement on R3 has been executed, and the route map will then not evaluate any additional route map entries and simply drops out, so the **permit 20** statement is never

actually executed. By configuring a **continue 20** statement within the **permit 10** line, the router is forced to evaluate the permit 20 line. Rather than dropping out of the route map after successful execution of the **permit 10** statement, the final verification within [Example 2-23](#) shows the route received on R3 with successful prepend applied by R5. If you have configured this correctly, as shown in [Example 2-23](#), you have scored 3 points.

Example 2-23 R5 Prepend Configuration and Verification

[Click here to view code image](#)

```
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*> 152.100.100.0/24 0.0.0.0                  0        32768  i
*> 152.200.32.0/24 120.100.34.4             0        200    400  i
*          120.100.34.5                         0        300    400  i
*> 152.200.33.0/24 120.100.34.5             0        300    400  i
*> 152.200.34.0/24 120.100.34.5             0        300    400  i
*> 152.200.35.0/24 120.100.34.5             0        300    400  i

R5(config)# router bgp 300
R5(config-router)# neighbor 120.100.34.3 route-map PREPEND out
R5(config-router)# exit
R5(config)# access-list 1 permit 152.200.32.0
R5(config)# route-map PREPEND permit 10
R5(config-route-map)# match ip address 1
R5(config-route-map)# set as-path prepend 300
R5(config-route-map)# route-map PREPEND permit 20
R5(config-route-map)# match ip address 1
R5(config-route-map)# set as-path prepend 300
R5(config-route-map)# route-map PREPEND permit 30

R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*> 152.100.100.0/24 0.0.0.0                  0        32768  i
*> 152.200.32.0/24 120.100.34.4             0        200    400  i
*          120.100.34.5                         0        300    300    400  i
*> 152.200.33.0/24 120.100.34.5             0        300    400  i
*> 152.200.34.0/24 120.100.34.5             0        300    400  i
*> 152.200.35.0/24 120.100.34.5             0        300    400  i

R5(config)# route-map PREPEND permit 10
R5(config-route-map)# continue 20

R3# clear ip bgp *
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
```

r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 152.100.100.0/24	0.0.0.0	0		32768	i
*> 152.200.32.0/24	120.100.34.4		0	200	400 i
*	120.100.34.5		0	300	300 300 400 i
*> 152.200.33.0/24	120.100.34.5		0	300	400 i
*> 152.200.34.0/24	120.100.34.5		0	300	400 i
*> 152.200.35.0/24	120.100.34.5		0	300	400 i

Section 4: IPv6 (12 Points)

- Configure IPv6 addresses on your network as follows:

2007:C15:C0:10::1/64 – R1 Gi0/1
2007:C15:C0:11::1/64 – R1 tunnel0
2007:C15:C0:11::3/64 – R3 tunnel0
2007:C15:C0:12::2/64 – R2 tunnel0
2007:C15:C0:12::3/64 – R3 tunnell1
2007:C15:C0:13::2/64 – R2 fe0/1
2007:C15:C0:14::3/64 – R3 Gi0/0
2007:C15:C0:14::4/64 – R4 Gi0/0
2007:C15:C0:14::5/64 – R5 Gi0/0
2007:C15:C0:15::4/64 – R4 Gi0/1
2007:C15:C0:15::6/64 – R6 Gi0/0

The prerequisite to the following questions is configuration of the IPv6 addresses and tunnel interfaces. [Example 2-24](#) shows the initial IPv6 configuration; tunnel specifics are provided in later questions, so just creating the tunnel interfaces and configuring an IPv6 address is required at this point. No points are on offer here for this task, unfortunately. Consider using the **show ipv6 interfaces brief** command for a quick check of your interface configuration.

Example 2-24 IPv6 Initial Configuration

[Click here to view code image](#)

```
R1(config)# ipv6 unicast-routing
R1(config)# interface GigabitEthernet0/1
R1(config-if)# ipv6 address 2007:C15:C0:10::1/64
R1(config-if)# interface tunnel0
R1(config-if)# ipv6 address 2007:C15:C0:11::1/64

R2(config)# ipv6 unicast-routing
R2(config)# interface fastethernet 0/1
R2(config-if)# ipv6 address 2007:C15:C0:13::2/64
R2(config-if)# interface tunnel0
R2(config-if)# ipv6 address 2007:C15:C0:12::2/64

R3(config)# ipv6 unicast-routing
```

```

R3(config)# int GigabitEthernet0/0
R3(config-if)# ipv6 address 2007:C15:C0:14::3/64
R3(config-if)# interface tunnel0
R3(config-if)# ipv6 address 2007:C15:C0:11::3/64
R3(config-if)# interface tunnel1
R3(config-if)# ipv6 address 2007:C15:C0:12::3/64

R4(config)# ipv6 unicast-routing
R4(config)# interface GigabitEthernet0/0
R4(config-if)# ipv6 address 2007:C15:C0:14::4/64
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ipv6 address 2007:C15:C0:15::4/64

R5(config)# ipv6 unicast-routing
R5(config)# interface GigabitEthernet0/0
R5(config-if)# ipv6 address 2007:C15:C0:14::5/64

R6(config)# ipv6 unicast-routing
R6(config)# interface GigabitEthernet0/0
R6(config-if)# ipv6 address 2007:C15:C0:15::6/64

```

Section 4.1: EIGRPv6

- Configure EIGRPv6 with an autonomous system of 6 between R1, R2, and R3. EIGRPv6 should not be enabled directly under the interfaces of the routers. Build your tunnels from R1 to R3 and R2 to R3 with source interfaces from VLAN 132 to advertise IPv6 edge networks from each router using **ipv6ip** mode. (2 points)

This is a straightforward EIGRPv6 configuration that requires the autonomous system number of 6 enabled by the **address-family ipv6** command under the existing EIGRP process as opposed to enabling EIGRPv6 under each interface.

The tunnel mode of **ipv6ip** is supplied within the question for the manually configured IPv6 tunnel, which provides connectivity from R3 to R2 and R1. The source interfaces of each tunnel are the VLAN 132 Ethernet interfaces. You should ensure that you make the IPv6-enabled interface on R3, which will actually belong to the OSPFv3 domain passive within EIGRPv6 as a matter of good practice. If you have configured this correctly, as shown in [Example 2-25](#), you have scored 2 points.

Example 2-25 EIGRPv6 Configuration and Verification

[Click here to view code image](#)

```

R1(config-if)# interface Tunnel0
R1(config-if)# tunnel source Gigabit0/0
R1(config-if)# tunnel destination 120.100.123.3
R1(config-if)# tunnel mode ipv6ip
R1(config-if)# router eigrp CCIE
R1(config-router)# address-family ipv6 unicast autonomous-system 6
R1(config-router-af)# af-interface Tunnel0
R1(config-router-af-interface)# no shutdown
R1(config-router-af-interface)# af-interface Gigabit0/1
R1(config-router-af-interface)# no shutdown

R2(config-if)# interface Tunnel0

```

```
R2(config-if)# tunnel source fastethernet0/0
R2(config-if)# tunnel destination 120.100.123.3
R2(config-if)# tunnel mode ipv6ip
R2(config-if)# router eigrp CCIE
R2(config-router)# address-family ipv6 unicast autonomous-system 6
R2(config-router-af)# af-interface Tunnel0
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface fastethernet0/1
R2(config-router-af-interface)# no shutdown
```

```
R3(config-if)# tunnel source Gigabit0/1
R3(config-if)# tunnel destination 120.100.123.1
R3(config-if)# tunnel mode ipv6ip
R3(config-if)# interface Tunnel1

R3(config-if)# tunnel source Gigabit0/1
R3(config-if)# tunnel destination 120.100.123.2
R3(config-if)# tunnel mode ipv6ip
R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 6
R3(config-router-af)# af-interface GigabitEthernet0/0
R3(config-router-af-interface)# passive-interface
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface Tunnel0
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface Tunnel1
R3(config-router-af-interface)# no shutdown
```

```
R1# show ipv6 route eigrp
IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007:C15:C0:12::/64 [90/310044416]
    via FE80::7864:7B03, Tunnel0
D  2007:C15:C0:13::/64 [90/310070016]
    via FE80::7864:7B03, Tunnel0
```

```
R2# show ipv6 route eigrp
IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007:C15:C0:10::/64 [90/310070016]
    via FE80::7864:7C03, Tunnel0
D  2007:C15:C0:11::/64 [90/310044416]
    via FE80::7864:7C03, Tunnel0
```

```
R3# show ipv6 route eigrp
IPv6 Routing Table - 9 entries
```

Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
U - Per-user Static route, M - MIPv6
I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
D - EIGRP, EX - EIGRP external

```
D 2007:C15:C0:10::/64 [90/297270016]
  via FE80::7864:7B01, Tunnel0
D 2007:C15:C0:13::/64 [90/297270016]
  via FE80::7864:7C02, Tunnel1
```

Section 4.2: OSPFv3

- Configure OSPFv3 per [Figure 2-8](#); use an OSPFv3 process of 1 on each router. (2 points)

Use vanilla OSPFv3 configuration between R3, R4, R5, and R6. If you have configured this correctly, as shown in [Example 2-26](#), you have scored 2 points.

Example 2-26 OSPFv3 Configuration and Verification

[Click here to view code image](#)

```
R3(config)# interface GigabitEthernet 0/0
R3(config-if)# ipv6 ospf 1 area 0

R4(config)# interface GigabitEthernet0/0
R4(config-if)# ipv6 ospf 1 area 0
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ipv6 ospf 1 area 1

R5(config)# interface GigabitEthernet0/0
R5(config-if)# ipv6 ospf 1 area 0

R6(config)# interface GigabitEthernet0/0
R6(config-if)# ipv6 ospf 1 area 1

R3# show ipv6 route ospf
IPv6 Routing Table - 11 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI 2007:C15:C0:15::/64 [110/2]
  via FE80::213:C3FF:FE7B:E4A0, GigabitEthernet0/0

R5# show ipv6 route ospf
IPv6 Routing Table - 5 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI 2007:C15:C0:15::/64 [110/2]
  via FE80::213:C3FF:FE7B:E4A0, GigabitEthernet0/0
```

```
R6# show ipv6 route ospf
IPv6 Routing Table - 5 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI 2007:C15:C0:14::/64 [110/2]
  via FE80::213:C3FF:FE7B:E4A1, GigabitEthernet0/0
```

- Configure Area 1 with IPsec authentication, use message digest 5, a security policy index of 500, and a key of DEC0DECC1E0DDBA11B0BB0BBEDB00B00. (2 points)

Authentication is required on R4 and R6 because they both belong to Area 1. The question explicitly states the specific parameters required, and you should not encounter any issues unless you incorrectly enter one of the keys. At 32 hex characters long, this could easily be done while under a time constraint. If you have configured this correctly, as shown in [Example 2-27](#), you have scored 2 points.

Example 2-27 Area 1 Authentication Configuration

[Click here to view code image](#)

```
R4(config)# ipv6 router ospf 1
R4(config-router)# area 1 authentication ipsec spi 500 md5
DEC0DECC1E0DDBA11B0BB0BBEDB00B00

R6(config)# ipv6 router ospf 1
R6(config-router)# area 1 authentication ipsec spi 500 md5
DEC0DECC1E0DDBA11B0BB0BBEDB00B00
```

- Ensure that the area router in Area 1 receives the following route; you may configure R4 to achieve this: (2 points)

OI 2007::/16 [110/2]
via XXXX::XXXX:XXXX:XXXX:XXXX, GigabitEthernet0/0

The only area router within Area 1 is R6. R4 is the area border router within this area. OI within the routing table is an OSPF interarea route, so this route must be generated from another area. Because Area 0 is the only other area within the OSPFv3 network, the route must be generated from this area as opposed to a redistributed route, which would show as an external route. A summary route generated on the area border router R4 of 2007::/16 within Area 0 will provide the required route to be received on R6. If you have configured this correctly, as shown in [Example 2-28](#), you have scored 2 points.

Example 2-28 OSPFv3 Configuration and Verification

[Click here to view code image](#)

```
R4(config)# ipv6 router ospf 1
R4(config-rtr)# area 0 range 2007::/16
```

```
R6# show ipv6 route ospf | include OI
OI 2007::/16 [110/2]
via FE80::213:C3FF:FE7B:E4A1, GigabitEthernet0/0
```

Section 4.3: Redistribution

- Redistribute EIGRPv6 into OSPFv3 on R3. Redistributed EIGRPv6 routes should have a metric of 5000 associated with them, regardless of which area they are seen in within the OSPFv3 network. (2 points)

A one-way redistribution of EIGRPv6 to OSPFv3 is required on R3. The default redistribution behavior ensures that external routes are advertised as external type 2, which have a fixed cost associated with them regardless of which area or location of the OSPFv3 network they are seen in. You simply require the metric set to 5000 on the OSPFv3 process. You must remember to advertise connected routes also; otherwise, the OSPFv3 network will not see the directly connected tunnel interfaces on R3. If you have configured this correctly, as shown in [Example 2-29](#), you have scored 2 points.

Example 2-29 R3 Ipv6 Redistribution Configuration and Verification

[Click here to view code image](#)

```
R3(config)# ipv6 router ospf 1
R3(config-rtr)# redistribute eigrp 6 include-connected metric 5000

R4# show ipv6 route ospf
IPv6 Routing Table - 11 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
O 2007::/16 [110/0]
  via ::, Null0
OE2 2007:C15:C0:10::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:11::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:12::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:13::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
```

- Configure R3 so that both R1 and R2 have the following IPv6 EIGRPv6 route in place. Do not redistribute OSPF into EIGRPv6 to achieve this, and do ensure that all routers have full visibility: (2 points)

D 2007::/16 [90/XXXXXXXXXX]
via XXXX::XXXX:XXXX:XXXX:XXXX, Tunnel0

You should have noticed in the previous question that mutual redistribution was not required; therefore, the EIGRPv6 network would not have reachability of the OSPFv3 network. This question

ensures the EIGRPv6 network sends traffic to R3 for the summarized network of 2007::/16. Because you are not permitted to redistribute OSPFv3 with a summary address, you need to configure EIGRPv6 summarization on the tunnel interfaces on R3 toward R1 and R2; this will provide the correct route and hop count as per the question. [Example 2-33](#) shows the required configuration and verification of the route, in addition to ICMP reachability to the remote OSPFv3 Area 1 network on R6. This test clearly demonstrates full end-to-end reachability from EIGRPv6 to OSPFv3. If you have configured this correctly, as shown in [Example 2-30](#), you have scored 2 points.

Example 2-30 R3 Ipv6 Summarization Configuration and Verification

[Click here to view code image](#)

```
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 6
R3(config-router-af)# af-interface Tunnel0
R3(config-router-af-interface)# summary-address 2007::/16
R3(config-router-af-interface)# af-interface Tunnel1
R3(config-router-af-interface)# summary-address 2007::/16

R1# show ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007::/16 [90/310044416]
    via FE80::7864:7B03, Tunnel0
R1# ping ipv6 2007:C15:C0:15::6

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:15::6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/7/8 ms

R2# show ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007::/16 [90/310044416]
    via FE80::7864:7C03, Tunnel0

R2# ping ipv6 2007:C15:C0:15::6

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:15::6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/7/8 ms
```

Section 5: QoS (6 Points)

- Two IP video conferencing units are to be installed onto Switch 2 ports Fast Ethernet 0/15 and 0/16 on VLAN 200. The devices use TCP ports 3230–3231 and UDP ports 3230–3235, and this traffic is unmarked from the devices as it enters the switch. Configure Switch 2 to assign a DSCP value of AF41 to video traffic from both of these devices. Ensure that the switch ports assigned to the devices do not participate in the usual spanning-tree checks, cannot form trunk links, and cannot be configured as EtherChannels. (3 points)

This is a differentiated services code point (DSCP) coloring of application traffic question. The TCP and UDP port information is provided so that access lists matching these ports within a class map are required for identification of the video traffic, and a policy map colors the traffic to a DSCP value of 41. The overall quality of service (QoS) service policy is applied to the video conferencing ports of Fast Ethernet 0/15 and 0/16 on Switch 2. The ports are required to be set to VLAN 200 with spanning-tree checks disabled, and trunking and channeling disabled using the command **switchport host**. The ports can also be explicitly configured to disable each feature individually, but the **switchport host** command does all this for you. If you have configured this correctly, as shown in [Example 2-31](#), you have scored 3 points. Use the **show policy-map** command to verify your configuration.

Example 2-31 QoS Configuration

[Click here to view code image](#)

```
SW2(config)# interface range fastethernet 0/15-16
SW2(config-if-range)# switchport access vlan 200
SW2(config-if-range)# switchport host
SW2(config-if-range)# exit
SW2(config)# mls qos
SW2(config)# class-map VIDEO
SW2(config-cmap)# match access-group 100
SW2(config-cmap)# exit
SW2(config)# access-list 100 permit tcp any any range 3230 3231
SW2(config)# access-list 100 permit udp any any range 3230 3235
SW2(config)# policy-map VIDEO-MARK
SW2(config-pmap)# class VIDEO
SW2(config-pmap-c)# set dscp AF41
SW2(config-pmap-c)# exit
SW2(config)# interface range fastethernet 0/15-16
SW2(config-if-range)# service-policy input VIDEO-MARK
```

- Configure R2 to assign a strict-priority queue with a 40-percent reservation of the WAN bandwidth for the video conferencing traffic in the previous question. Maximize the available bandwidth by ensuring the RTP headers within the video stream are compressed. The remainder of the bandwidth should be guaranteed for a default queue with WRED enabled. (3 points)

Following from the previous question, R2 is required to provide QoS on the Ethernet link toward the rest of the network. A class map matches the precolored video traffic of DSCP 41; a policy map is then required to call the class map and assign a strict 40-percent priority queue with the command **priority percent 40**. RTP compression is configured within the policy map for the video traffic. The

default queue has a guaranteed bandwidth reservation with the command **bandwidth percent 60**, and weighted random early detection (WRED) is enabled within this queue. If you have configured this correctly, as shown in [Example 2-32](#), you have scored 3 points.

Example 2-32 R2 QoS Configuration and Verification

[Click here to view code image](#)

```
R2(config)# class-map match-all VIDEO
R2(config-cmap)# match dscp af41
R2(config-cmap)# policy-map VIDEO-QOS
R2(config-pmap)# class VIDEO
R2(config-pmap-c)# priority percent 40
R2(config-pmap-c)# compression header ip rtp

R2(config-pmap-c)# class class-default
R2(config-pmap-c)# bandwidth percent 60
R2(config-pmap-c)# random-detect
R2(config-pmap-c)# exit
R2(config)# interface fastethernet0/0
R2(config-if)# service-policy output VIDEO-QOS
```

Section 6: Multicast (9 Points)

- Configure routers R1, R2, R3, and R4 for IPv4 multicast. Each router should use PIM sparse dense mode. Both R1 and R2 should be configured to be candidate RPs specifically for the following multicast groups: 225.225.0.1, 225.225.0.2, 225.225.0.3, and 225.225.0.4 (by use of their Loopback 0 interfaces). You should limit the boundary of your multicast network so that it does propagate further into your network than R4. R3 should be configured as a mapping agent to announce the rendezvous points for the multicast network with the same boundary constraints. (3 points)

The question dictates that R1 and R2 be rendezvous points (RPs) and advertise the same groups to the multicast network. R3 is required to announce the rendezvous points, and R4 will by default elect R2 as the RP for each group because it has the higher loopback address compared to R1 for the same groups. TTL scoping is used within the configuration to limit the boundary of advertisements on both the candidate RPs and the discovery agent up to R4. [Example 2-33](#) shows the required configuration and RP mappings as received on R4. If you have configured this correctly, as shown in [Example 2-33](#), you have scored 3 points.

Example 2-33 R1, R2, R3, and R4 Multicast Configuration and Verification

[Click here to view code image](#)

```
R1(config)# ip multicast-routing
R1(config)# interface Loopback0
R1(config-if)# ip pim sparse-dense-mode
R1(config-if)# interface GigabitEthernet0/0
R1(config-if)# ip pim sparse-dense-mode
R1(config-if)# ip pim send-rp-announce Loopback0 scope 3 group-list GROUPS
R1(config)# ip access-list standard GROUPS
```

```

R1(config-std-nacl)# permit 225.225.0.1
R1(config-std-nacl)# permit 225.225.0.2
R1(config-std-nacl)# permit 225.225.0.3
R1(config-std-nacl)# permit 225.225.0.4

R2(config)# ip multicast-routing
R2(config)# interface Loopback0
R2(config-if)# ip pim sparse-dense-mode
R2(config-if)# interface fastethernet0/0
R2(config-if)# ip pim sparse-dense-mode
R2(config-if)# ip pim send-rp-announce Loopback0 scope 3 group-list GROUPS
R2(config)# ip access-list standard GROUPS
R2(config-std-nacl)# permit 225.225.0.1
R2(config-std-nacl)# permit 225.225.0.2
R2(config-std-nacl)# permit 225.225.0.3
R2(config-std-nacl)# permit 225.225.0.4

R3(config)# ip multicast-routing
R3(config)# interface Loopback0
R3(config-if)# ip pim sparse-dense-mode
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-dense-mode
R3(config-if)# interface GigabitEthernet0/1
R3(config-if)# ip pim sparse-dense-mode
R3(config-if)# exit
R3(config)# ip pim send-rp-discovery lo0 scope 2
R4(config-if)# ip multicast-routing
R4(config-if)# interface GigabitEthernet0/0
R4(config-if)# ip pim sparse-dense-mode

R4# show ip pim rp mapping
PIM Group-to-RP Mappings

Group(s) 225.225.0.1/32
  RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:52
Group(s) 225.225.0.2/32
  RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:56
Group(s) 225.225.0.3/32
  RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:55
Group(s) 225.225.0.4/32
  RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:55

```

-
- Configure R3 to ensure that R4 has a candidate RP as R1 for groups 225.225.0.1 and 225.225.0.2 and R2 for groups 225.225.0.3 and 225.225.0.4. (3 points)

As detailed in the previous example, R2 will by default become the candidate RP as selected by the discovery agent (R3) because of having a higher loopback IP address as used in the PIM announcements compared to R1. By configuring a group list on the discovery agent, RP announcements can be filtered. Configuring two filter lists with each candidate RP associated with

them allows the discovery agent to announce two different RPs. [Example 2-34](#) shows the required configuration, a debug of the auto-RP announcements on R3 to detail the filtering and the resulting RP mappings on R4. If you have configured this correctly, as shown in [Example 2-34](#), you have scored 3 points.

Example 2-34 R3 RP Multicast Configuration and Verification

[Click here to view code image](#)

```
R3(config)# ip pim rp-announce-filter rp-list R1 group-list R1-GROUPS
R3(config)# ip pim rp-announce-filter rp-list R2 group-list R2-GROUPS
R3(config)# ip access-list standard R1
R3(config-std-nacl)# permit 120.100.1.1
R3(config-std-nacl)# exit
R3(config)# ip access-list standard R2
R3(config-std-nacl)# permit 120.100.2.1
R3(config-std-nacl)# exit
R3(config# ip access-list standard R1-GROUPS
R3(config-std-nacl)# permit 225.225.0.1
R3(config-std-nacl)# permit 225.225.0.2
R3(config-std-nacl)# exit
R3(config)# ip access-list standard R2-GROUPS
R3(config-std-nacl)# permit 225.225.0.3
R3(config-std-nacl)# permit 225.225.0.4

R3# debug ip pim auto-rp
PIM Auto-RP debugging is on
Auto-RP(0): Received RP-announce, from 120.100.1.1, RP_cnt 1, ht 181
Auto-RP(0): Update (225.225.0.1/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Update (225.225.0.2/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Filtered 225.225.0.3/32 for RP 120.100.1.1
Auto-RP(0): Filtered 225.225.0.4/32 for RP 120.100.1.1
Auto-RP(0): Received RP-announce, from 120.100.1.1, RP_cnt 1, ht 181
Auto-RP(0): Update (225.225.0.1/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Update (225.225.0.2/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Filtered 225.225.0.3/32 for RP 120.100.1.1
Auto-RP(0): Filtered 225.225.0.4/32 for RP 120.100.1.1

R4# show ip pim rp mapping
PIM Group-to-RP Mappings

Group(s) 225.225.0.1/32
RP 120.100.1.1 (?), v2v1
Info source: 120.100.34.3 (?), elected via Auto-RP
Uptime: 00:00:08, expires: 00:02:52
Group(s) 225.225.0.2/32
RP 120.100.1.1 (?), v2v1
Info source: 120.100.34.3 (?), elected via Auto-RP
Uptime: 00:00:08, expires: 00:02:51
Group(s) 225.225.0.3/32
RP 120.100.2.1 (?), v2v1
Info source: 120.100.34.3 (?), elected via Auto-RP
Uptime: 00:00:47, expires: 00:02:12
Group(s) 225.225.0.4/32
RP 120.100.2.1 (?), v2v1
Info source: 120.100.34.3 (?), elected via Auto-RP
Uptime: 00:00:47, expires: 00:02:09
```

- Configure R1 to monitor traffic forwarded through itself for traffic destined to the multicast group of 225.225.0.1. If no packet for this group is received within a single 10-second interval, ensure that an SNMP trap is sent to an SNMP management station on 120.100.100.100 using a community string of **public**. (3 points)

The IP multicast heartbeat feature facilitates the monitoring of the delivery of IP multicast packets and failure notification based on configurable parameters. By configuring R1 to enable the heartbeat monitoring for the group 225.255.0.1 with the subparameters of 1 and 10, the router monitors a packet lost within 1 interval of 10 seconds and will send an Simple Network Management Protocol (SNMP) trap to the SNMP host 120.100.100.100, which is required to be configured within the basic SNMP trap configuration. [Example 2-35](#) details the required multicast heartbeat configuration and verification of the SNMP trap by issue of a ping to 225.225.0.1 from R3. Even though R1 does not have a valid IGMP join group for this group, traffic is still directed to it, and the heartbeat process is activated. If you have configured this correctly, as shown in [Example 2-35](#), you have scored 3 points.

Example 2-35 R1 Multicast Heartbeat Configuration

[Click here to view code image](#)

```
R1(config)# snmp-server host 120.100.100.100 traps public
R1(config)# snmp-server enable traps ipmulticast
R1(config)# ip multicast heartbeat 225.225.0.1 1 1 10

R1# debug snmp packets

R3# ping 225.225.0.1

R1# SNMP: Queuing packet to 120.100.100.100
SNMP: V1 Trap, ent ciscoExperiment.2.3.1, addr 120.100.100.1, gentrap 6, spectrap 1
  ciscoIpMRouteHeartBeatEntry.2.225.225.0.1 = 120.100.123.3
  ciscoIpMRouteHeartBeatEntry.3.225.225.0.1 = 10
  ciscoIpMRouteHeartBeatEntry.4.225.225.0.1 = 1
  ciscoIpMRouteHeartBeatEntry.5.225.225.0.1 = 0
```

Section 7: Security (10 Points)

- Allow router R6 to passively watch the SYN connections that flow to only VLAN 63 for servers that might reside on this subnet. To prevent a potential DoS attack from a flood of SYN requests, the router should be configured to randomly drop SYN packets from any source to this VLAN that have not been correctly established within 20 seconds. (2 points)

The question requires that the TCP intercept feature be configured on R6. This protects TCP servers from TCP SYN-flooding attacks with a wave of half-opened connections overwhelming the server's CPU, the result of which can effectively cause a DoS attack. The default behavior of the feature is to intercept the SYN connections to a server and effectively proxy the connection until it has been correctly established. Because you are requested to passively monitor the connection, you must configure the feature into watch mode by use of the global **ip tcp intercept mode watch** command. You are also requested to ensure that the feature is enabled only on VLAN 63 from any source, so an access list is required to which the intercept features restricts its monitoring. The default behavior of

the feature is to drop SYN connections based on the oldest first, but the question dictated that random connections must be dropped. This is achieved with the global command **ip tcp intercept drop-mode random**. To ensure that the 20-second limit is met as opposed to the default 30 second, adjustment of the timers is required with the global command **ip tcp intercept watch-timeout 20**. If you have configured this correctly, as shown in [Example 2-36](#), you have scored 2 points. Use the **show tcp intercept connections** command to verify your configuration.

Example 2-36 R6 TCP Intercept Configuration

[Click here to view code image](#)

```
R6(config)# ip tcp intercept list 100
R6(config)# access-list 100 permit tcp any 120.100.63.0 0.0.0.255
R6(config)# ip tcp intercept mode watch
R6(config)# ip tcp intercept drop-mode random
R6(config)# ip tcp intercept watch-timeout 20
```

- Configure an ACL on R1 to allow TCP sessions generated on this router and through its Ethernet interface and to block TCP sessions from entering on its VLAN 132 interface that were not initiated on it or through it originally. Do not use the established feature within standard ACLs to achieve this, and only apply ACLs on the VLAN 132 interface. The ACL should timeout after 100 seconds of locally initiated TCP inactivity; it should also enable ICMP traffic inbound for testing purposes. (3 points)

The question requires that a reflexive access control list (ACL) be configured on R1. This enables TCP traffic for sessions originating from within the network but denies TCP traffic for sessions originating from outside the network. The reflexive ACL contains only temporary entries, which are automatically created when a new TCP session is initiated. The entries are simply removed, by default, 300 seconds after the session ends. However, the question requires this to be modified to 100 seconds. To facilitate the reflexive ACL, you must configure a standard ACL inbound on the VLAN 132 Ethernet interface, which permits the required traffic inbound to R1 and only returns traffic matching the reflexive ACL. Required traffic is, of course, EIGRP, PIM, IPv6 tunneling, and as directed, ICMP for testing. It's a cruel question because if you forget to permit any of the required traffic inbound, you will lose points from a previous section in which you might have otherwise scored the total possible points. If you did not know what protocol IPv6 uses, you can simply use the **log** option on your inbound ACL on a final **deny** statement. This would show you that the tunneling from R3 inbound to R1 uses IP protocol 41, which must be included in your inbound ACL.

[Example 2-37](#) shows the required configuration and verification of the reflexive ACL. Because traffic is evaluated only by the ACL as it passes through the router, Switch 1 has been configured to belong to VLAN 100 to telnet through R1 to R3 in the example. When initiated by Switch 1, the Telnet session passes through the ACL **FILTER-OUT** on R1 and creates an entry in the reflexive ACL **DYNAMIC-TCP**. Real-time details can be seen by issuing the **show access-lists** command on R1. The reflexive ACL permits return traffic to the Telnet session inbound for the configured inactivity interval of 100 seconds. If you have configured this correctly, as shown in [Example 2-37](#), you have scored 3 points.

Note

The reflexive ACL is valid only for traffic flowing through the router. Therefore, you might experience connectivity issues if you initiate a Telnet session from R1 without manipulating the Telnet source option. This behavior has no bearing on points scored and should be considered a by-product of the solution. If you face a similar question in the actual exam and Telnet connectivity was required from the router you are configuring, you would specifically be instructed to ensure the correct operation of Telnet on that router.

Example 2-37 R1 Reflexive ACL Configuration and Verification

[Click here to view code image](#)

```
R1(config-if)# ip access-list extended FILTER-IN
R1(config-ext-nacl)# permit icmp any any
R1(config-ext-nacl)# permit eigrp any any
R1(config-ext-nacl)# permit pim any any
R1(config-ext-nacl)# permit tcp host 120.100.3.1 host 120.100.1.1 eq bgp
R1(config-ext-nacl)# permit 41 host 120.100.123.3 host 120.100.123.1
R1(config-ext-nacl)# evaluate DYNAMIC-TCP
R1(config-ext-nacl)# ip access-list extended FILTER-OUT
R1(config-ext-nacl)# permit tcp any any reflect DYNAMIC-TCP
R1(config-ext-nacl)# exit
R1(config)# ip reflexive-list timeout 100
R1(config)# interface GigabitEthernet0/0
R1(config-if)# ip access-group FILTER-IN in
R1(config-if)# ip access-group FILTER-OUT out
```

```
SW1(config)# interface vlan 100
SW1(config-if)# ip add 120.100.100.100 255.255.255.0
SW1(config-if)# exit
SW1(config)# ip route 120.100.3.1 255.255.255.255 120.100.100.1
SW1(config)# exit
SW1# trace 120.100.3.1
```

Type escape sequence to abort.
Tracing the route to 120.100.3.1

```
1 120.100.100.1 0 msec 4 msec 0 msec
2 120.100.100.1 !A * !A
```

```
SW1# telnet 120.100.3.1
```

```
Trying 120.100.3.1 ... Open
```

User Access Verification

Password:

R3>enable

Password:

R3#

```
R1# show access-lists
```

```
Standard IP access list 1
```

```

10 permit 120.100.1.0 (3 matches)
20 permit 120.100.100.0 (3 matches)
Standard IP access list GROUPS
10 permit 225.225.0.1
20 permit 225.225.0.2
30 permit 225.225.0.3
40 permit 225.225.0.4
Reflexive IP access list DYNAMIC-TCP
    permit tcp host 120.100.3.1 eq telnet host 120.100.100.100 eq 11034 (34
        matches) (time left 90)
Extended IP access list FILTER-IN
5 permit icmp any any (150 matches)
10 permit eigrp any any (1710 matches)
20 permit pim any any (92 matches)
25 permit tcp host 120.100.3.1 host 120.100.1.1 eq bgp (126 matches)
30 evaluate DYNAMIC-TCP
Extended IP access list FILTER-OUT
10 permit tcp any any reflect DYNAMIC-TCP (18 matches)

```

- Configure R1 so that it is capable of performing SCP. The router should belong to a domain of toughtest.co.uk. Use local authentication with a username and password of cisco, a key size of 768 bits, and an SSH timeout of 2 minutes and retry value of 2. (2 points)

SCP is Secure Copy Protocol. It is similar to Remote Copy but requires Secure Shell (SSH) to be running on the router for security purposes. It is a tough question because this is the kind of feature for which you will need to check the documentation. You will need to realize aspects of SSH are considered prerequisites to enable SCP. Even if you hadn't configured SSH or SCP previously, you should realize that you would need to configure a domain ID, local authentication with a username and password, a key of some form, and some SSH timeout and retry values based on the directions. Be careful on the values and remember to enter the timeout in seconds, not minutes. Your username and password combination requires a privilege level of 15 set for SCP. If you have configured this correctly, as shown in [Example 2-38](#), you have scored 2 points.

Example 2-38 R1 RCP Configuration

[Click here to view code image](#)

```

R1(config)# ip domain-name toughtest.co.uk
R1(config)# crypto key generate rsa modulus 768
The name for the keys will be: R1.toughtest.co.uk

% The key modulus size is 768 bits
% Generating 768 bit RSA keys, keys will be non-exportable... [OK]

R1(config)# aaa new-model
R1(config)# aaa authentication login default local
R1(config)# aaa authorization exec default local
R1(config)# username cisco privilege 15 password 0 cisco
R1(config)# ip ssh time-out 120
R1(config)# ip ssh authentication-retries 2
R1(config)# ip scp server enable
R1(config)#
00:57:29.343: %SSH-5-ENABLED: SSH 1.99 has been enabled

```

- The network administrator has determined that IPv6 router advertisements are being sourced from routers on VLAN 34. Disable these advertisements from entering and propagating on VLAN 34. You may use an ACL applied in a single location in your solution. Do not use the RA guard solution with untrusted ports. (3 points)

Routers R3, R4, and R5 will begin to send RAs as soon as they are configured with an IPv6 address. A simple solution is to enable RA guard on the switch, whereby you could set the switch ports connecting to the routers as untrusted, but this is not permitted. Because you are permitted to use an ACL in only a single location, this needs to be applied to the VLAN 34 interface. The ACL needs to deny router advertisements, which, of course, use ICMPv6. The ACL then needs to permit everything else; otherwise, you have just broken your IPv6 network. You need to remember that for the switch to process IPv6 packets, it needs to be running IPv6 and have a valid IPv6 address assigned to VLAN 34, something you might have overlooked under the time constraints and pressure of the practice exam. If your switch was not previously enabled for IPv6, remember to use the command **sdm prefer dual-ipv4-and-ipv6 routing** (and reboot the device for this to take effect). If you have configured this correctly, as shown in [Example 2-39](#), you have scored 3 points.

Example 2-39 SW1 RA ACL Configuration

[Click here to view code image](#)

```
SW1(config)# ipv6 unicast-routing
SW1(config)# ipv6 access-list RA
SW1(config-ipv6-acl)# deny icmp any any router-advertisement
SW1(config-ipv6-acl)# permit ipv6 any any
SW1(config-ipv6-acl)# exit
SW1(config)# int vlan 34
SW1(config-if)# ipv6 traffic-filter RA in
SW1(config-if)# ipv6 address 2007:C15:C0:15::10/64

SW1# show log
*Oct 4 17:58:23: %IPV6-6-ACCESSLOGDP: list RA/10 denied icmpv6
  FE80::219:AAFF:FEBA:BE40 -> FF02::1 (134/0), 1 packet
*Oct 4 17:58:23: %IPV6-6-ACCESSLOGDP: list RA/10 denied icmpv6
  FE80::218:18FF:FEA2:3250 -> FF02::1 (134/0), 1 packet
```

Lab Wrap-Up

So, how did it go? Did you run out of time? Did you manage to finish but miss what was actually required? If you scored over 80, well done. If you accomplished this within 8 hours or less, you will be prepared for any scenario that you are likely to face during the 5.5 hours of the Configuration section of the actual exam. Remember that the Troubleshooting section on the v5.0 exam is a separate section from the Configuration section and has a different scenario; you will have 2 hours to complete the Troubleshooting section. This lab was designed to ensure that you troubleshoot your own work as you progress through the questions.

Did you manage to configure items such as EIGRP third-party next hop and the **continue** statement within your BGP prepending? Items such as these might seem inconsequential, but they can make or break your lab.

Practice Lab 3

Equipment List

Practice Lab 3 follows an identical format to Lab 1 and 2 with timings and also consists of 100 points. You need the following hardware and software components to begin this practice lab:

- Six routers loaded with Cisco IOS Software Release 15.3T Advanced Enterprise image and the minimum interface configuration, as documented in [Table 3-1](#)

Router	Model	Ethernet I/F
R1	3925	2
R2	3925	2
R3	3925	2
R4	3925	2
R5	3925	2
R6	3925	2

Table 3-1 Hardware Required per Router

- Four 3560X switches with IOS 15.0S IP Services

Setting Up the Lab

You can use any combination of routers as long as you fulfill the requirements within the topology diagram, as shown in [Figure 3-1](#). However, you should use the same model of routers because this makes life easier if you load configurations directly from those supplied into your own devices.

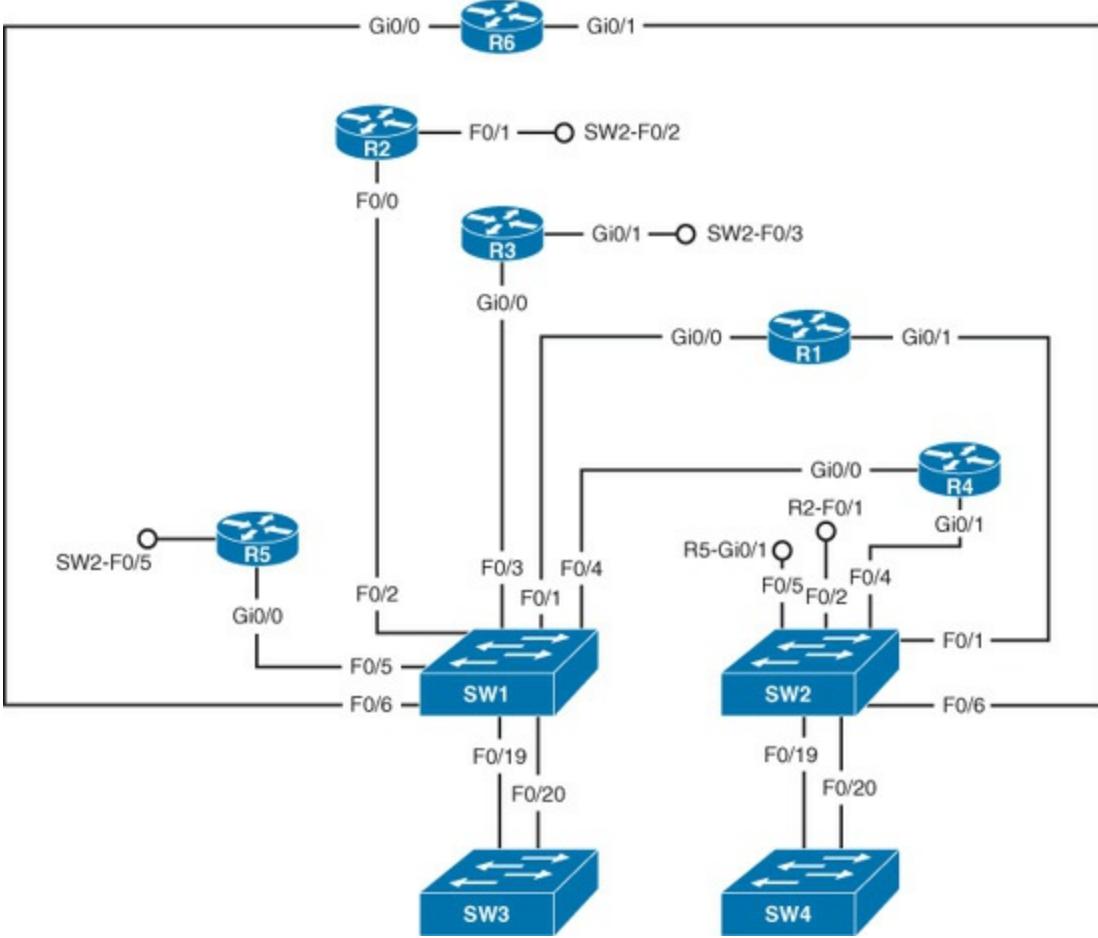


Figure 3-1 Practice Lab 3 Network Topology

Note

Notice in the initial configurations supplied that some interfaces do not have IP addresses preconfigured. This is because either you do not use that interface or you need to configure this interface from default within the exercise. The initial configurations supplied should be used to preconfigure your routers and switches before the lab starts.

If your routers have different interface speeds than those used within this book, adjust the **bandwidth** statements on the relevant interfaces to keep all interface speeds in line. This ensures that you do not get unwanted behavior because of differing IGP metrics.

Lab Topology

This practice lab uses the topology as outlined in [Figure 3-1](#), which you must re-create with your own equipment.

Switch Instructions

Configure VLAN assignments from the configurations supplied on the CD-ROM or from [Table 3-2](#).

VLAN	Switch 1	Switch 2	Switch 3	Switch 4
25		Fa0/2, Fa0/5		
32	Fa0/1, Fa0/2, Fa0/3			
34		Fa0/3, Fa0/4		
45	Fa0/4, Fa0/5, Fa0/6	—	—	—
200	Fa0/19	—	—	—
400	—	Fa0/19	—	—
Trunk	Fa0/1	Fa0/6	—	—
Trunk	Fa0/20	Fa0/20	Fa0/20	Fa0/20

Table 3-2 VLAN Assignment

Connect your switches with RJ-45 Ethernet cross-over cables, as shown in [Figure 3-2](#).

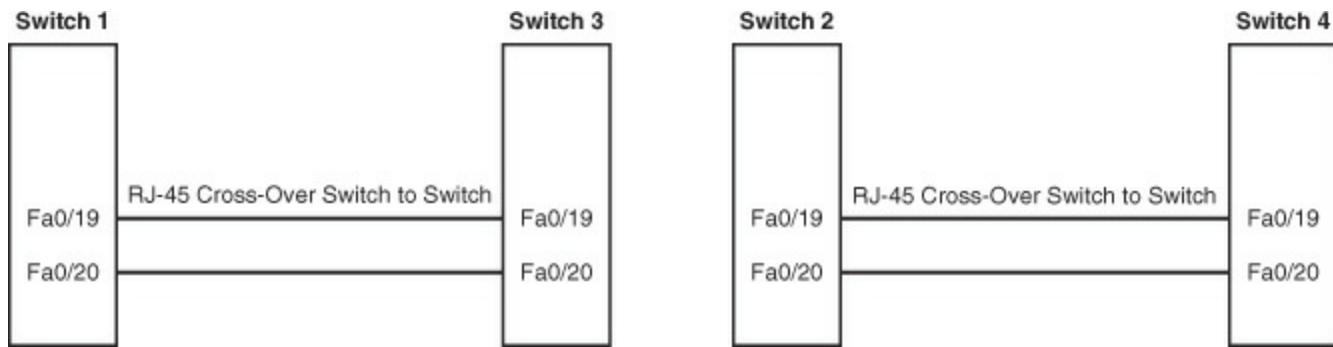


Figure 3-2 Switch-to-Switch Connectivity

IP Address Instructions

In the actual CCIE lab, you find that the majority of your IP addresses are preconfigured. For this exercise, you are required to configure your IP addresses as shown in [Figure 3-3](#) or to load the initial router configurations supplied. If you are manually configuring your equipment, be sure you include the following loopback addresses:

R1 Lo0 120.100.1.1/32

R2 Lo0 120.100.2.1/32

R3 Lo0 120.100.3.1/32

R4 Lo0 120.100.4.1/32

R5 Lo0 120.100.5.1/32

R6 Lo0 120.100.6.1/32 SW1 Lo0 10.1.1.1/24

Lo1 10.1.2.1/24

Lo2 10.1.3.1/24

SW2 Lo0 10.2.2.1/24

Lo1 10.2.3.1/24

Lo2 10.2.4.1/24

SW3 Lo0 10.33.33.1/24

Lo1 10.33.34.1/24

Lo2 10.33.35.1/24

SW4 Lo0 10.44.44.1/24

Lo1 10.44.45.1/24

Lo2 10.44.46.1/24

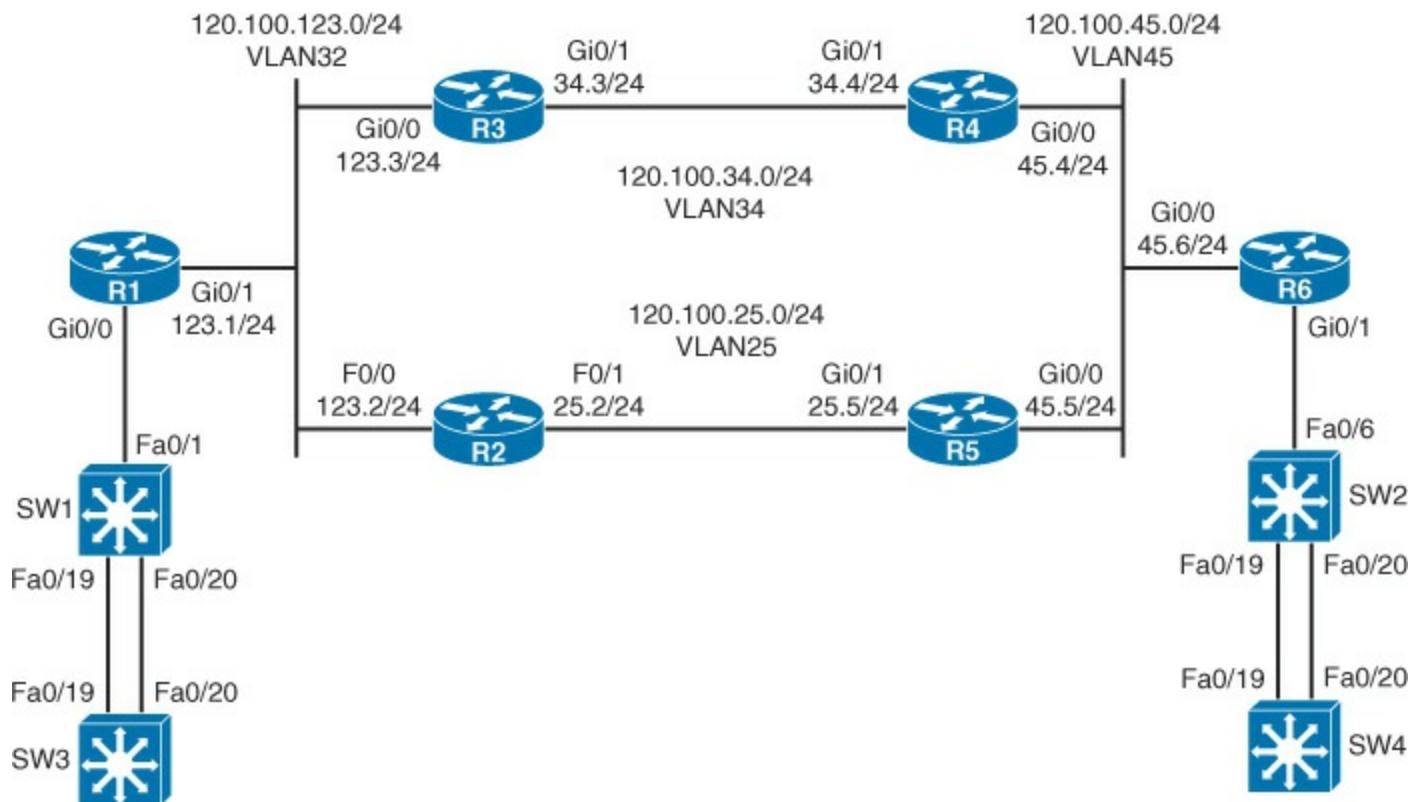


Figure 3-3 IP Addressing Diagram

Pre-Lab Tasks

- Build the lab topology per [Figure 3-1](#) and [Figure 3-2](#).
- Configure the IP addresses on each router as shown in [Figure 3-3](#) and add the loopback addresses. Alternatively, you can load the initial configuration files supplied if your router is compatible with those used to create this exercise.

General Guidelines

- Read the whole lab before you start.
- Do not configure any static/default routes unless otherwise specified.
- Ensure full IP visibility between routers for ping testing/Telnet access to your devices.
- If you are running out of time, choose questions that you are confident you can answer. Failing this, choose questions with a higher point rating to maximize your potential score.
- Get into a comfortable and quiet environment where you can focus for the next 8 hours.
- Take a 30-minute break midway through the exercise.
- Have available a Cisco documentation CD-ROM, or access online the latest documentation from the following URL:

Note that access to this URL is likely to be restricted within the real exam.

Note

Access only these URLs, not the whole [Cisco.com](#) website (because if you are permitted to use documentation during your CCIE lab exam, it will be restricted). To save time during your lab, consider opening several windows with the pages you are likely to look at.

Practice Lab Three

You will now be answering questions in relation to the network topology, as shown in [Figure 3-4](#).

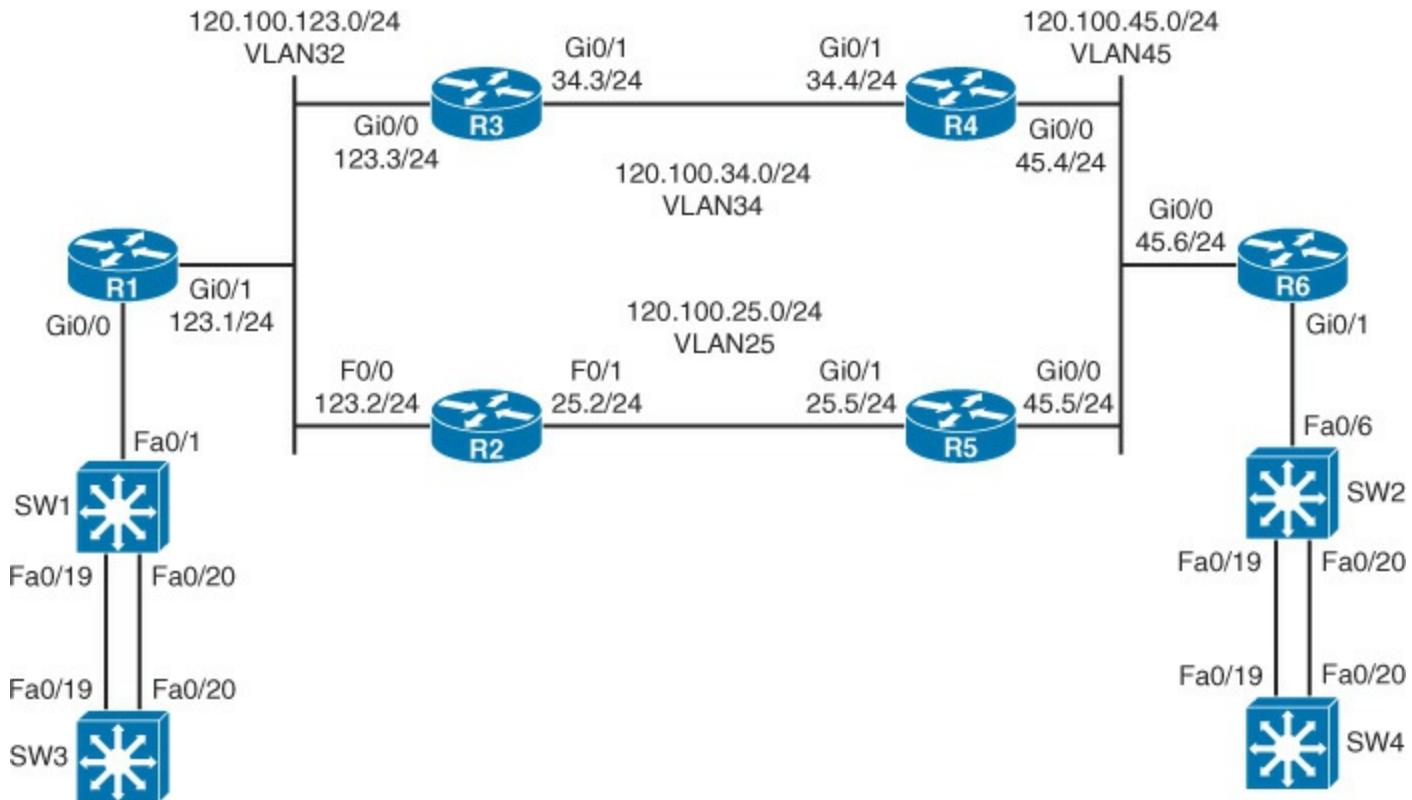


Figure 3-4 Lab Topology Diagram

Section 1: LAN Switching (4 Points)

- Configure your switched network per [Figure 3-5](#). Your switched network is physically nonlooped and therefore does not require any STP root bridge configuration. Configure SW1 Fa0/19 to belong to VLAN 200 and SW2 Fa0/19 to belong to VLAN 400. Configure interface Fa0/1 on SW1 to become a trunk port toward R1 and Fa0/6 on SW2 to become a trunk port toward R6. Ports should use 802.1Q encapsulation. Restrict the VLANs permissible to use the trunk on Switch 1 Fa0/1 to VLAN 10, 50, and 200 and VLAN 20, 100, and 400 on Switch 2 Fa0/6. Interface Fa0/20 of each switch has been preconfigured to be a trunk port. You should also configure R1 and R6 to terminate the VLANs on each router. Connectivity between switches will be provided via R1 and R6 later in the lab. (3 points)

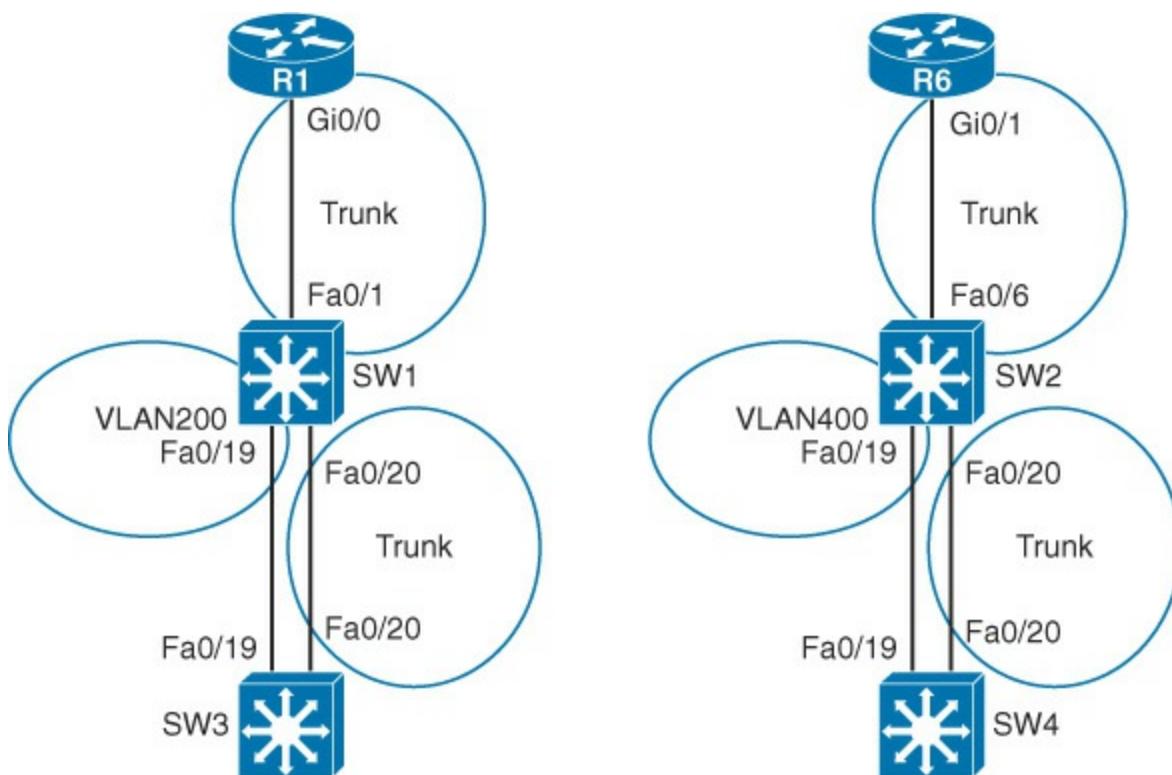


Figure 3-5 Switched Network Topology

- SW3 interface Fa0/19 and SW4 interface Fa0/19 are required to communicate with each other on the same IP subnet of 1.1.1.0/24; configure these interfaces with IP addresses 1.1.1.1/24 and 1.1.1.2/24, respectively. The interfaces should be configured to communicate as if connected directly as a point-to-point link. (Actual IP end-to-end connectivity will be achieved in a later section.) (1 point)

Section 2: MPLS and OSPF (27 Points)

- Configure OSPF on your routers per [Figure 3-6](#) to enable your network to transport MPLS and MP-BGP. All required interfaces (including Loopback 0) should be configured to belong to Area 0. Ensure that all OSPF configuration is entered under the interfaces. (3 points)

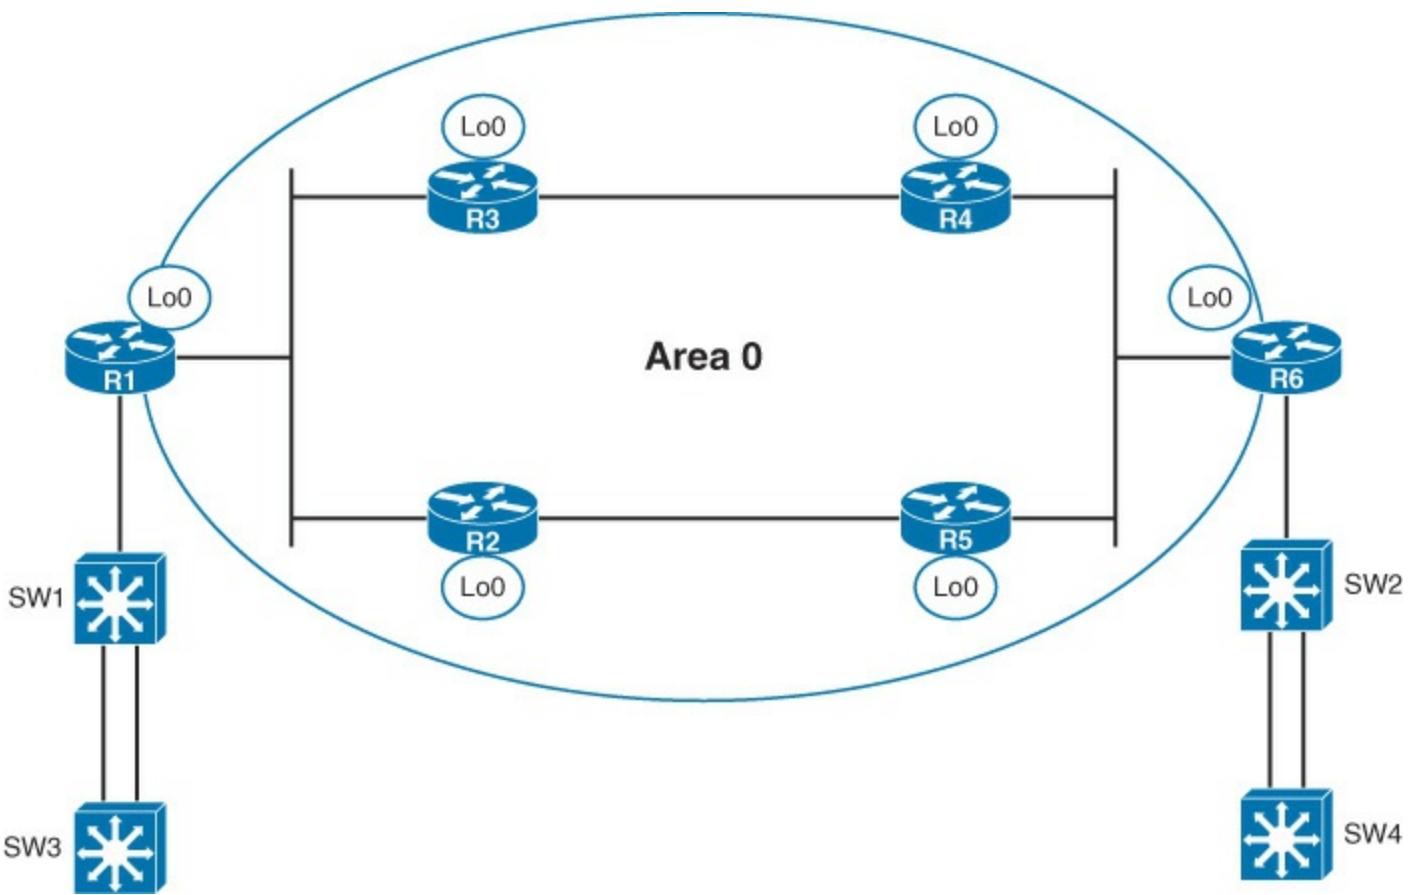


Figure 3-6 MPLS/OSPF Topology

- Configure MPLS on all routers within the OSPF domain; use LDP, ensuring that TDP can be used on unused interfaces without specifically configuring these interfaces for TDP. Routers R1 and R6 will become your PE routers, whereas R2, R3, R4, and R5 will become P routers. (4 points)
- You will be configuring two VPNs over your MPLS networks per [Figure 3-7](#) between PE routers of BLUE and RED. At this point, assign the following interfaces on each PE router into separate routing instances within the routers:
 - PE R1 interface Gi0/0 VLAN10 connection into VPN BLUE
 - PE R1 interface Gi0/0 VLAN 50 connection into VPN RED
 - PE R6 interface Gi0/1 VLAN 20 connection into VPN BLUE
 - PE R6 interface Gi0/1 VLAN 100 connection into VPN RED
 Configure VPN BLUE to use an RD of 100 and VPN RED to use an RD of 200 for both importing and exporting routes into your BGP network, which will be configured later with an autonomous system of AS65001. (4 points)

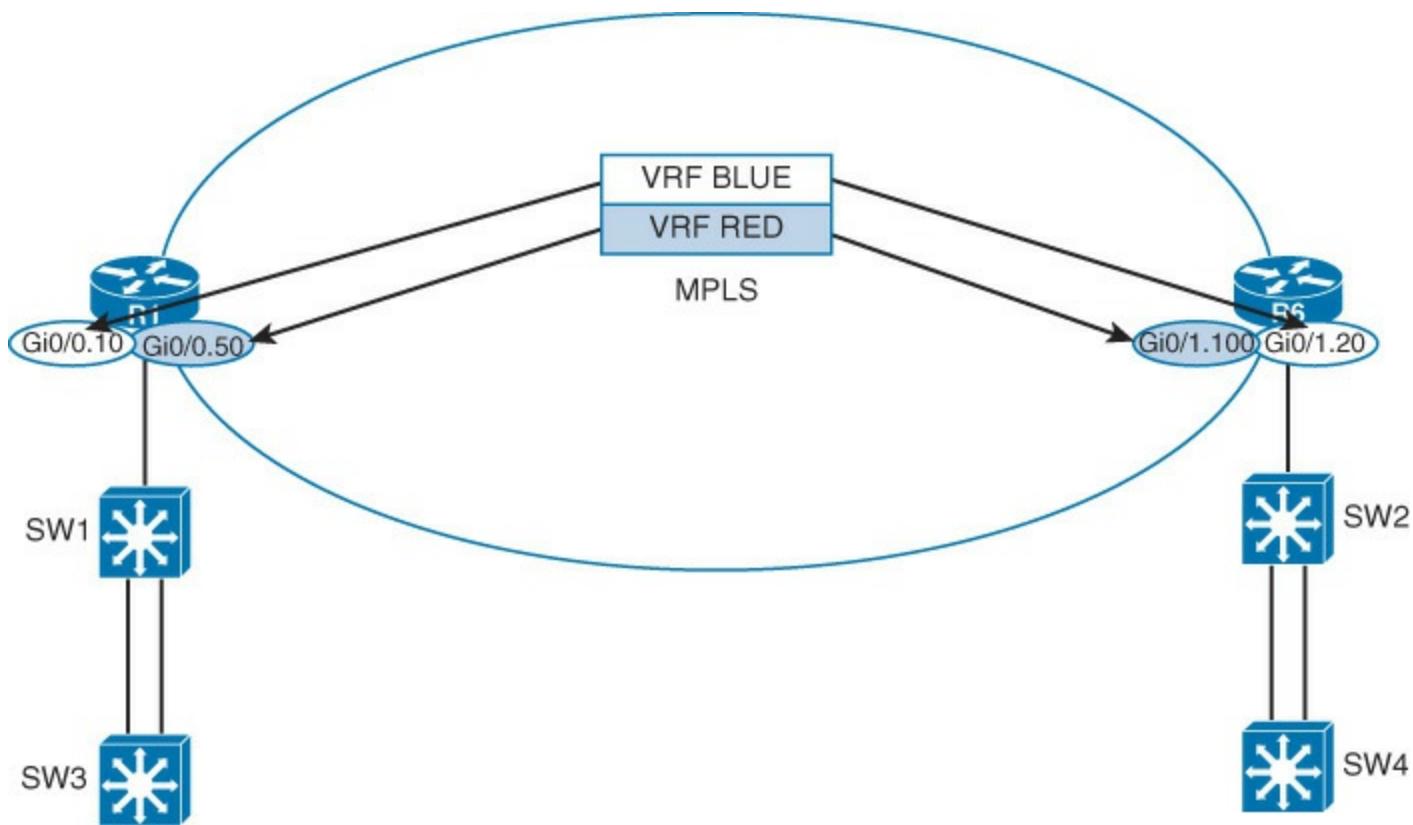


Figure 3-7 MPLS VPN Topology

- Create a network between PE router R1 and CE device SW1 using a VLAN 10 interface on SW1 that can be trunked toward R1; this network will reside in the BLUE VPN. Use a subnet of 10.10.10.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)
- Create a network between PE router R6 and CE device SW2 using a VLAN 20 interface on SW2 that can be trunked toward R6; this network will reside in the BLUE VPN. Use a subnet of 10.10.20.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)
- Create a network between PE router R1 and CE device SW3 using a VLAN 50 interface on SW3 that can be trunked toward R1; this network will reside in the RED VPN. Use a subnet of 130.50.50.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)
- Create a network between PE router R6 and CE device SW4 using a VLAN 100 interface on SW4 that can be trunked toward R6; this network will reside in the RED VPN. Use a subnet of 130.100.100.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)

Section 3: BGP (5 Points)

- Configure MP-BGP between your PE routers, per [Figure 3-8](#), to enable your network to transport the VPNV4 addresses of your configured VPNs (BLUE and RED). Use loopback interfaces for peering between your PE routers. You will configure the actual VPN routing in later questions. (4 points)

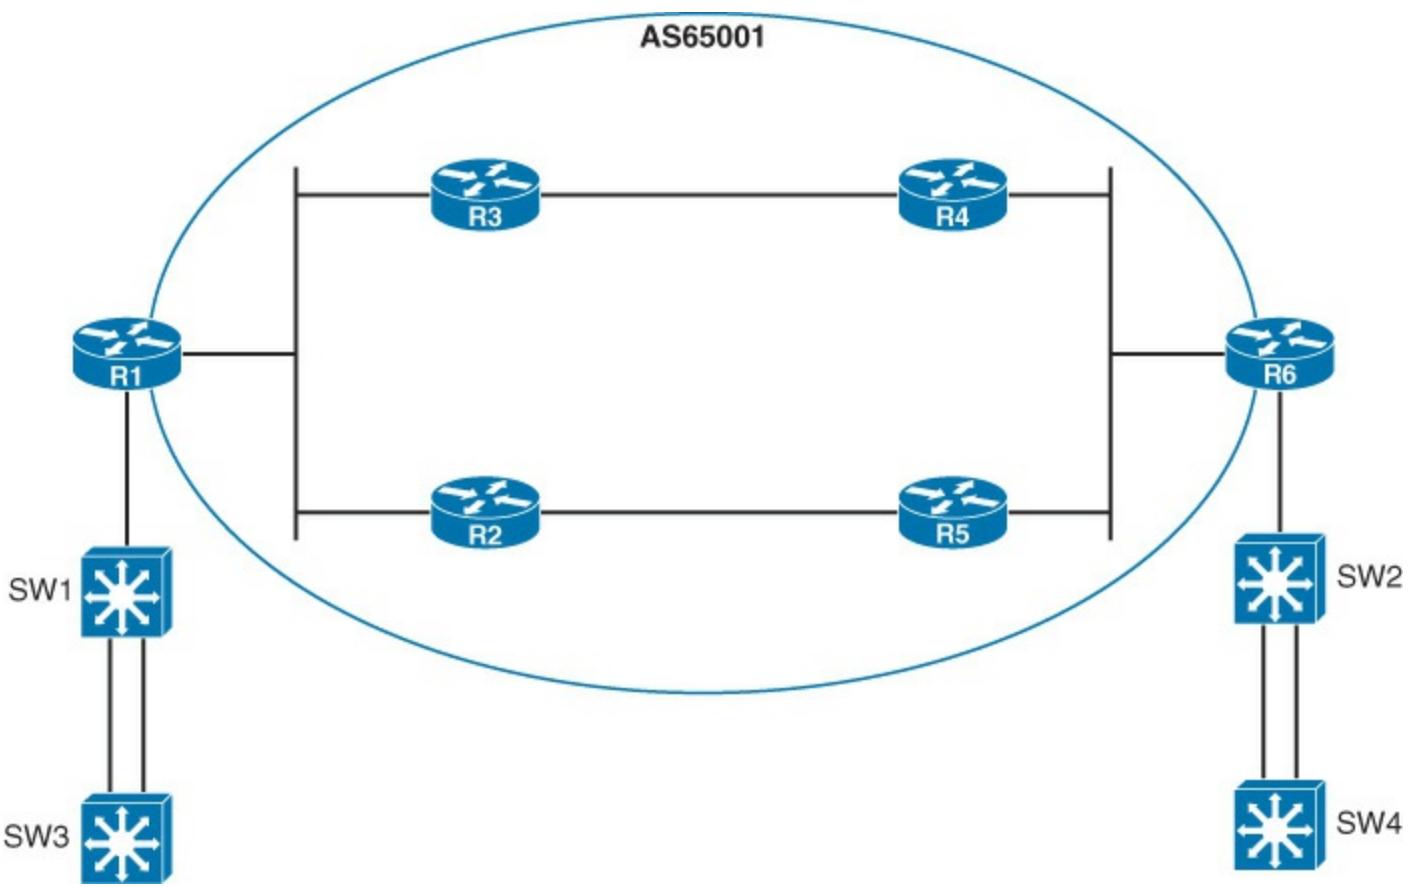


Figure 3-8 BGP Topology

Section 4: EIGRP and MP-BGP (3 Points)

- Configure EIGRP per [Figure 3-9](#) between your PE router R6 and CE Switch SW2. Use an EIGRP virtual instance name of VPN on R6 and a process number of 10 on SW2. Use VLAN 20 for EIGRP connectivity between R6 and SW2. Advertise all preconfigured loopback networks on SW2 to R6 for the BLUE VPN. (1 point)

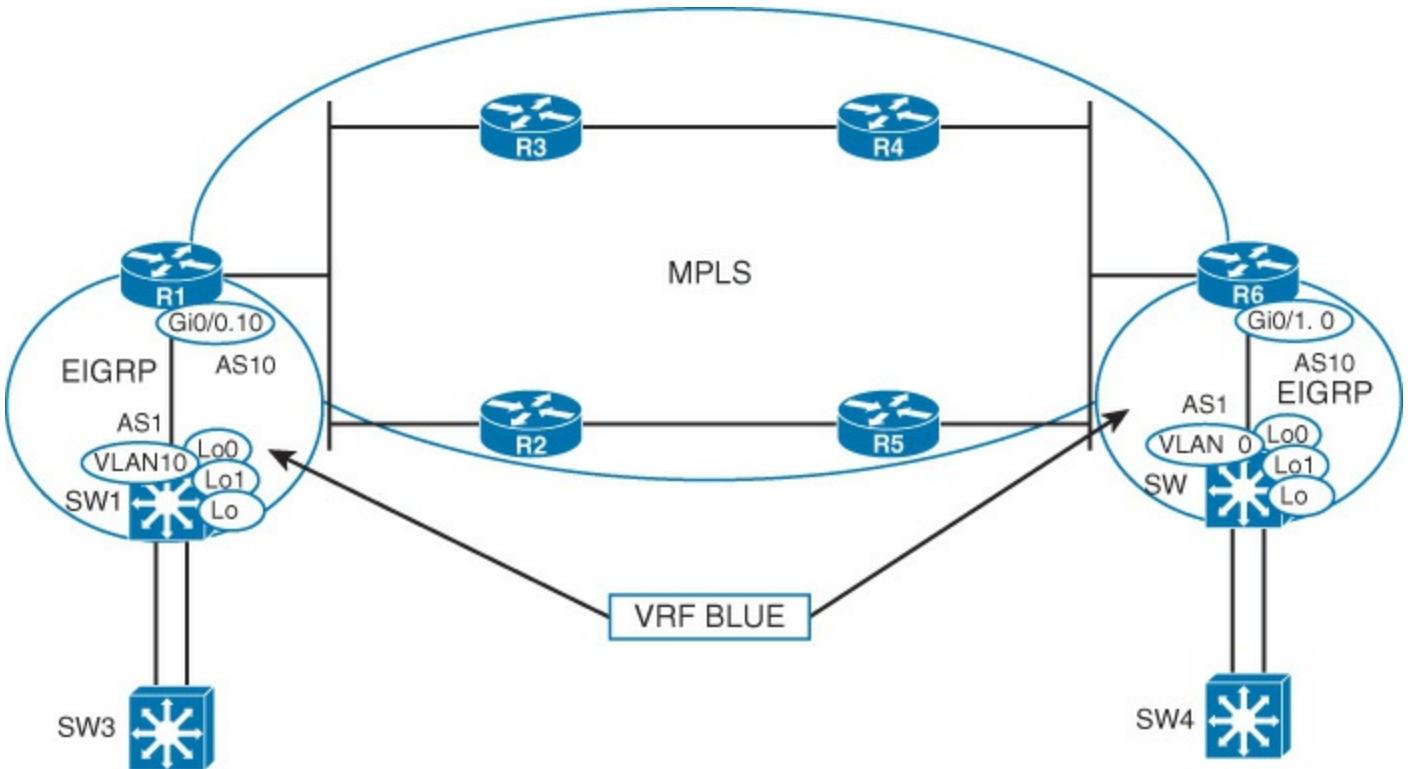


Figure 3-9 EIGRP/MP-BGP Topology

- Configure EIGRP per [Figure 3-9](#) between your PE router R1 and CE switch SW1. Use an EIGRP virtual instance name of VPN on R1 and a process number of 10 on SW1. Use VLAN 10 for EIGRP connectivity between R1 and SW1. Advertise all preconfigured loopback networks on SW1 to R1 for the BLUE VPN. (1 point)
- Configure your PE routers R1 and R6 to transport EIGRP routes from your CE devices between the BLUE VPN using MP-BGP. EIGRP networks residing on SW1 should be seen as internal EIGRP routes on SW2 and vice versa. Ensure that all EIGRP routes have a MED of 50 assigned to them within MP-BGP. Use a default metric of 10000 100 255 1 1500 for BGP routes when redistributed into EIGRP. (1 point)

Section 5: OSPF and MP-BGP (6 Points)

- Configure OSPF per [Figure 3-10](#) for your VRF RED with a process number of 3 on PE router R1 and SW3 using VLAN 50 for connectivity. Use a process ID of 2 on PE router R6 and CE device SW4 using VLAN 100 for connectivity. You should permit only internal OSPF routes to be advertised across your VPN and ensure that the redistribution of BGP routes into OSPF are assigned as type 1 external routes with no manually adjusted cost associated with them. It is acceptable for these routes to come through as /32 routes because of default OSPF behavior of loopback interfaces. (2 points)

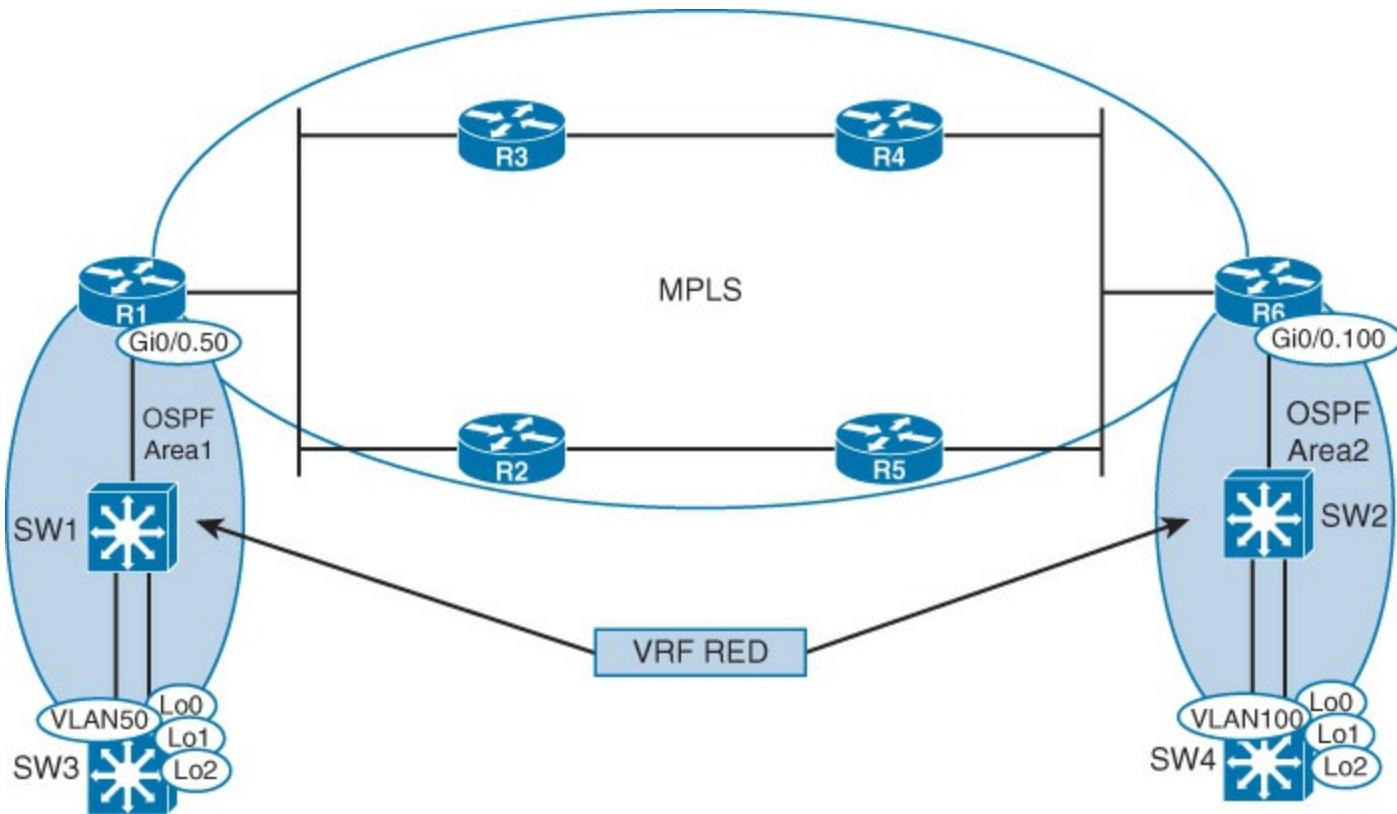


Figure 3-10 OSPF Topology

- You will notice that your OSPF IA (intra-area) routes between CE devices SW3 and SW4 appear as type 1 external routes. Configure your OSPF network appropriately to ensure that the routes are displayed correctly as IA routes. You are not permitted to adjust the OSPF redistribution into BGP as directed in the previous question. Maintain the OSPF process IDs are previously directed. You are permitted to configure only router R1. (4 points)

Section 6: MPLS (7 Points)

- Leak network 10.1.1.0/24 from SW1 VRF BLUE on PE R1 into the VRF RED on PE1; similarly, leak 10.44.44.0/24 from VRF RED into VRF BLUE on R6. Both Switch 1 and Switch 4 should receive the following routes:

SW1# **show ip route | include 10.44.44.0**

D EX 10.44.44.0/24 [170/XXXXXX] via 10.10.10.1, 00:00:27, Vlan10

SW1#

SW4# **show ip route | include 10.1.1.0**

O E1 10.1.1.0/24 [110/XX] via 130.100.100.1, 00:03:04, Vlan100

SW4#

Verify your configuration by pinging from VRF RED SW4 10.44.44.1 to VRF BLUE SW1 10.1.1.1 SW1. (5 points)

- Configure your PE routers R1 and R6 to ensure that the MPLS P routers are not listed as intermediate hops when a trace route is performed on your CE devices. (2 points)

Section 7: VPLS Simulation (10 Points)

- Switches 3 and 4 will have been configured to belong to the subnet of 1.1.1.0/24 within a previous question. Create an L2TPv3 Xconnect attachment circuit on your PE routers R1 and R6 for your CE devices (SW3 Fast Ethernet 0/19 1.1.1.1/24 and SW4 Fast Ethernet 0/19 1.1.1.2/24) to communicate using a Layer 2 tunneling solution (use Version 3) across your Layer 3 network. You should use existing loopback interfaces on your PE routers for peering over your MPLS network. Be aware that the SW3 resides in VLAN 200 and that SW4 resides in VLAN 400 in respective PE router subinterfaces. (10 points)

Section 8: Multicast (10 Points)

- Configure your MPLS network for multicast support of the RED VRF using PIM sparse mode. PE routers R1 and R6 should be configured to tunnel multicast traffic using an MDT address of 232.0.0.11 from CE device Switch 3 VLAN 50 to CE device SW4 VLAN 100 over the RED VRF. Switch 4 should be configured to reply to an ICMP ping on its VLAN 100 interface directed to 226.2.2.2 from Switch 3 VLAN 50. It can be assumed that the mVRF bandwidth requirement is low; configure MDT appropriately. Ensure that PE router R6's associated VLAN 100 IP address is used as the rendezvous point for the RED VRF multicast traffic. (10 points)

Section 9: IPv6 (6 Points)

- Configure the following IPv6 address on the PE routers R1 and R6, and implement IPv6 over MPLS between the 6PE routers to advertise the prefixes between 6PEs. Make sure that your loopback IPv6 addresses are used to source any locally generated IPv6 traffic. (6 points)

R1 Lo0 2010:C15:C0:1::1/64

R1 Gi0/0.10 2010:C15:C0:11::1/64

R6 Lo0 2010:C15:C0:6::1/64

R6 Gi0/1.20 2010:C15:C0:62::1/64

Section 10: QoS (7 Points)

- Create the following QoS profile on your PE router R1 for traffic egressing to your CE device connected to the BLUE VRF. The total bandwidth between the PE to CE should be shaped to 1 Mbps. Ensure that voice traffic is assigned to an LLQ. Use an appropriate method of prioritizing DSCP traffic so that AF31 packets are statistically dropped more frequently than AF32 during congestion, and reduce the effects of TCP global synchronization within your Mission-Critical class, and solely reduce the effect of TCP global synchronization within the Default class. (4 points)

Class	DSCP Value	% of Bandwidth Assigned
Voice	EF, CS5	35
Mission-Critical	CS6, AF31, AF32, CS3	40
Default	Any	25

- Create the following QoS profile on your PE router R1 for traffic ingressing from your CE device connected to the BLUE VRF into the MPLS network. The total aggregate speed from the CE to PE should be restricted to 1 Mbps.

Class	CIR (bps)
Voice	350,000
Mission-Critical	400,000
Default	250,000

- Traffic in the Voice class within the detailed CIR should have the MPLS EXP set to 5 and above discarded. Traffic in the Mission-Critical class within the detailed CIR should have the MPLS EXP set to 3 and above set to 7. Traffic in the Default class within the detailed CIR should have the MPLS EXP set to 0 and above set to 4. (3 points)

Section 11: Security (15 Points)

- Create three new loopback IP addresses of loopback1 on R4, R5, and R6; use IP addresses of 4.4.4.4/24, 5.5.5.5/24, and 6.6.6.6/24, respectively. Use EIGRP with a named virtual instance of VPN and autonomous system of 1 to advertise the loopback networks between routers over a common GRE tunnel network of 100.100.100.X/24 (X = router number) sourced from each router's common Ethernet interface, using IPsec to encrypt all traffic between the loopback networks using a preshared ISAKMP key of CCIE. Use an IPsec transform set of esp-des esp-md5-hmac on each router. R6 is to be a hub router, with R4 and R5 being effectively spoke routers in your solution. You are not permitted to enable EIGRP on your Ethernet interfaces between routers. Spoke routers must communicate with each other directly using dynamic IPsec connections with the aid of NHRP at the hub, whereas hub-to-spoke IPsec connections should be permanent. The hub router should provide all necessary direct next-hop information to the spoke routers when they are required to communicate between themselves. NHRP should be authenticated with a password of SECRET. Use an MTU of 1416 for your secure traffic, an NHRP timeout of 100 seconds for spoke replies, and a delay of 2 microseconds on the tunnel network. Test your solution by extended pings sourced from the configured loopback interfaces. (10 points)
- Following on from the previous question, add R2 into the common GRE tunnel network as a spoke router using identical security parameters as used on R4 and R5, ensuring that it receives routes from R4, R5, and R6 using the same common EIGRP parameters. The source interface for the tunnel configuration on R2 should be Fast Ethernet 1/1, and the destination should be the Gigabit Ethernet 0/0 interface of R6. Add new Loopback 2 identical IP addresses of 45.45.45.45/24 on both R4 and R5, and advertise this identical network from R4 and R5 to the hub router R6 on the common GRE tunnel interface. Configure R6 to advertise both destinations (R4 and R5) to spoke router R2 for network 45.45.45.0/24 in EIGRP over the common GRE tunnel network. (3 points)
- The network manager of your network cannot justify a full security implementation, but wants to implement a solution that provides a password prompt from R1 only when the keyboard entry 1 is entered on the console port (as opposed to the normal CR/Enter key). Configure R1 appropriately. (2 points)

Practice Lab 3: “Ask the Proctor”

Note

This section should be used only if you require clues to complete the questions. In the actual CCIE lab, the proctor will not enter into any discussions about the questions or answers. He or she will be present to ensure that you do not have problems with the lab environment and to maintain the timing element of the exam.

Section 1: LAN Switching and Frame Relay

- Q. Do you want me to configure Layer 2 between Switch 3 and Switch 4 so that they can communicate on the subnet 1.1.1.0/24?

A. No, simply configure the switches as directed in the question and Layer 2 connectivity will be provisioned later within the lab when your core network is configured.

Q. With my Frame Relay, I can only reach my spoke routers from the hub. Is this acceptable?

A. No, the question states that each device must be reachable over the Frame Relay network; this includes spoke-to-spoke communication.

Section 2: MPLS and OSPF

Q. Do you require OSPF for any interfaces on R1 and R6 that connect to the switches?

A. No, just configure OSPF per the figure; this is required to advertise your loopback addresses for MPLS.

Q. Does it matter what OSPF process ID I use on my routers?

A. No, the question doesn't direct you to use a specific process ID, so you can use an ID of your choice.

Q. To protect 66.66.66.66/32, is this related to MPLS TE and is a tunnel required between R1 and R2?

A. No, this is a pure IP solution designed to speed convergence in the event of a failure without the need to tune convergence timers.

Q. Can I configure OSPFv2 Fast Reroute for the 6.6.6.6/32 prefix?

A. Yes.

Q. Do you want the OSPF from the core routers extended into the RED VRF I created so that I run end-to-end OSPF between CE Switch 1 and CE Switch 2?

A. No, you will ultimately achieve this connectivity through an MPLS VPN and not by simply extending OSPF through your core devices.

Q. Do you want me to configure my RED VRF with a route descriptor of 100 and 200 for the BLUE VRF?

A. You have been provided with additional information in the question that enables you to facilitate use of MP-BGP extended communities.

Q. So, just add in the MP-BGP autonomous system number to the RD?

A. A combination of the two will achieve the desired results.

Q. I can't ping to my VLAN 10 interface on Switch 1 from R1. Do I need to perform any further configuration to make this work?

A. No, just remember that R1 is now a PE router with multiple VRF routing tables. You need to ensure that you source your ping correctly; otherwise, R1 would use its default routing table (which is used for the MPLS connectivity).

Section 3: BGP

Q. Do you want me to configure a full mesh of BGP between all routers?

A. No, MP-BGP is simply required between the PE routers.

Q. Do you need me to configure the PEs to send community values to each other?

A. You must remember how MPLS works and ensure that the route targets are propagated to successfully configure your VPNs.

Q. I usually configure next-hop self on my BGP configurations. Is this acceptable here?

A. You haven't been instructed not to use this command at this point even though this is an iBGP configuration.

Section 4: EIGRP and MP-BGP

Q. EIGRP requires the same autonomous system number on neighbor routers to peer successfully. If I use a different number on R6 and Switch 2, they cannot peer correctly.

A. Correct. Look for a method of making the autonomous system number the same within your VRF specific configuration on R6.

Section 5: OSPF and MP-BGP

Q. Do you want me to configure OSPF, MPLS, and BGP initially within the OSPF section?

A. No, just initially as directed OSPF; this will enable your network to transport MPLS and BGP within later questions.

Q. Changing the process ID on OSPF peers wouldn't affect any adjacency. Why would I need to do this?

A. You are correct, but you have been directed to do so in the question. It will become evident why you have been asked to do this in a later question.

Q. Why would I want to advertise the OSPF routes as external type 1 routes within BGP; surely the routes should appear as standard interarea routes through the VPN.

A. Correct, this question is a little misleading. The routes will come out as type 1 external routes on your CE devices, and it would appear that you have modified this behavior with your redistribution configuration. This behavior should become apparent why in the following question.

Q. I think if I change the redistribution of OSPF into BGP, I can make the OSPF routes appear as intra-area routes. Do I score any points if I change the redistribution?

A. No, by all means try to change the redistribution, though; it might help you understand the issue.

Q. I changed the redistribution, but the routes remain identical. This must have something to do with the different OSPF process ID I had to configure. I can't adjust this, so I am stuck.

A. You had a similar issue with EIGRP autonomous system numbers; just investigate what is possible within your VRF configuration.

Q. If I change the domain ID on R1, is that acceptable?

A. Find an appropriate value and try it out.

Section 6: MPLS

Q. I can manage to leak routes between VRFs but my route comes out as a host route. Can I modify my loopback interface with the OSPF **network** command on Switch 4 so that it is advertised with the correct mask?

A. Yes.

Section 7: VPLS Simulation

Q. Is this MPLS specific, or could I do this over a standard Layer 3 network?

A. You could achieve the same result over a standard Layer 3 network. Just exercise caution where you configure your parameters to achieve the correct results in the appropriate VRF.

Q. I have my L2TPv3 tunnel up end to end, yet I cannot ping between switches. I suspect a spanning-tree type of issue if the question states VLAN differences when I need to provide Layer 2 adjacency. Am I at liberty to manipulate spanning tree?

A. Yes.

Section 8: Multicast

Q. Do you want me to enable PIM over my P routers or just PE routers?

A. The question states “MPLS network.” To provide end-to-end multicast support, you might find that configuring PIM end to end is required.

Q. Do you want PIM on my MPLS router loopback interfaces?

A. You might find it is required at certain points within your MPLS network.

Q. I have a Multicast Distribution Tree tunnel between PE routers, but I don’t understand what the low-bandwidth requirement is.

A. MDT has differing requirements for high- and low-bandwidth sources; you might or might not require a Data MDT.

Q. To get Switch 4 to reply to a ping to 226.2.2.2, can I just configure an IGMP join group appropriately on its VLAN 100 interface?

A. You can.

Section 9: IPv6

Q. Do you want me to run IPv6 down to my CE switches and redistribute anything over MPLS?

A. Your switches are currently not capable of running IPv6.

Q. Should I just advertise my IPv6 prefixes with the BGP **network** command?

A. Yes, because there is no redistribution to be configured.

Section 10: QoS

Q. Do you want the first QoS policy outbound on the BLUE VRF interface on PE router R1?

A. Yes.

Q. To prioritize DSCP traffic, do you want me to configure some priority queuing within a class for AF32 flows?

A. No, use a common technique whereby traffic is dropped randomly as queues fill. AF31 packets should be dropped more frequently than AF32, though.

Q. Are you looking for random early detect?

A. You're almost there; this wouldn't offer the inherent drop preference, though.

Q. The second QoS policy limits traffic to 1 Mbps, yet the first will be line rate at 1 Gbps. Is this correct?

A. Yes, I appreciate that this isn't the real world; it just provides you with two different configuration exercises.

Q. Do I use the same packet-marking classes in each question?

A. Yes.

Q. Is this DiffServ, whereby you want me to modify the topmost bits in the EXP field?

A. Yes.

Q. Do you want the policy applied to the CE-facing VRF BLUE interface as an input service policy?

A. Yes, this would then modify the traffic as it flows into the MPLS network.

Section 11: Security

Q. Don't I need an ACL to mark all traffic that should be encrypted?

A. No, your solution will not require an ACL, and all traffic flowing from the new subnets you created should automatically be encrypted.

Q. The clues in the question suggest this is a DMVPN question. I have configured my solution correctly, yet I don't get spoke routes on the spoke routers. Is this acceptable?

A. No, you need full network visibility from all devices and not just the hub.

Q. This sounds like a split-horizon issue; can I disable this behavior?

A. Yes.

Q. I still show a next hop of the hub between spoke networks. Is this okay?

A. No, the question specifically states that spoke routers must be able to communicate with each other directly.

Q. Can I modify the next hop from the hub?

A. Yes.

Q. I have added R2 as a spoke to the DMVPN network, and it receives a single route to network 45.45.45.0/24 via the hub router. Is this acceptable?

A. No, you must configure R6 to advertise both spokes (R4 and R5) as valid next hops for this

destination.

Q. Can I configure **max-paths** on R2?

A. No, use a similar feature on R6 hub to actually advertise both spokes rather than just one as a valid next hop.

Q. Do you want me to get R1 to somehow translate a CR into a 1 to then provide a password prompt?

A. No, just make the router provide a prompt when it receives an ASCII 1, rather than a CR on the line con 0 port.

Practice Lab 3 Debrief

This section now analyzes each question showing you what was required and how to achieve the desired results. You should use this section to produce an overall score for Practice Lab 3.

Section 1: LAN Switching (4 Points)

■ Configure your switched network per [Figure 3-6](#). Your switched network is physically nonlooped and therefore does not require any STP root bridge configuration. Configure SW1 Fa0/19 to belong to VLAN 200 and SW2 Fa0/19 to belong to VLAN 400. Configure Interface Fa0/1 on SW1 to become a trunk port toward R1 and Fa0/6 on SW2 to become a trunk port toward R6; ports should use 802.1Q encapsulation. Restrict the VLANs permissible to use the trunk on Switch 1 Fa0/1 to VLAN10, 50, and 200 and VLAN 20, 100, and 400 on Switch 2 Fa0/6. Interface Fa0/20 of each switch has been preconfigured to be a trunk port. You should also configure R1 and R6 to terminate the VLANs on each router. Connectivity between switches will be provided via R1 and R6 later in the lab. (3 points)

This is a simple question, but you are required to complete multiple configuration items to gain your points. The configuration enables connectivity between switches when the MPLS section has been completed later in the lab. To begin, Ports Fa0/19 of Switch 1 and Switch 2 should be assigned the correct VLAN. (The actual VLANs would have been created previously in the initial configuration.) Next, the trunking is configured as directed with allowed VLANs of 10, 50, and 200 for Switch 1 and 20, 100, and 400 for Switch 2. R1 and R6 are configured with the corresponding VLAN numbers as sub interfaces to terminate the trunk connections from Switch 1 and Switch 2 using an identical reference for the dot1q encapsulation. If you have configured this correctly, as shown in [Example 3-1](#), you have scored 3 points.

Note

R1 and R6 use the VLAN number for the encapsulation and the subinterface number. Your subinterface number does not need to match the VLAN number, but it is considered good practice to do so.

Example 3-1 SW1, SW2, R1, and R6 Configuration

[Click here to view code image](#)

```
Switch1# show run interface fastethernet 0/19
!
interface fastethernet0/19
  switchport access vlan 200
  switchport mode access
```

```
Switch1# show run interface fastethernet 0/1
!
interface fastethernet0/1
  switchport trunk encapsulation dot1q
  switchport trunk allowed vlan 10,50,200
  switchport mode trunk
```

```
Switch2# show run interface fastethernet 0/19
!
interface fastethernet0/19
  switchport access vlan 400
  switchport mode access
```

```
Switch2# show run interface fastethernet 0/6
!
interface fastethernet0/6
  switchport trunk encapsulation dot1q
  switchport trunk allowed vlan 20,100,400
  switchport mode trunk
```

```
R1# show run | begin interface GigabitEthernet0/0
!
interface GigabitEthernet0/0
  no ip address
!
interface GigabitEthernet0/0.10
  encapsulation dot1Q 10
!
interface GigabitEthernet0/0.50
  encapsulation dot1Q 50
!
interface GigabitEthernet0/0.200
  encapsulation dot1Q 200
```

```
R6# show run | begin interface GigabitEthernet0/1
!
interface GigabitEthernet0/1
no ip address
!
interface GigabitEthernet0/1.20
  encapsulation dot1Q 20
!
interface GigabitEthernet0/0.100
  encapsulation dot1Q 100
!
interface GigabitEthernet0/1.400
```

-
- SW3 interface Fa0/19 and SW4 interface Fa0/19 are required to communicate with each other on the same IP subnet of 1.1.1.0/24. Configure these interfaces with IP addresses 1.1.1.1/24 and 1.1.1.2/24, respectively. The interfaces should be configured to communicate as if connected directly as a point-to-point link. (Actual IP end-to-end connectivity will be achieved in a later

section.) (1 point)

This is a straightforward configuration task to change the operation of the ports to non-switch-port Layer 3 mode where an IP address can be configured. End-to-end connectivity is achieved through the IP network at a later stage. If you have configured this correctly, as shown in [Example 3-2](#), you have scored 1 point.

Example 3-2 SW3 and SW4 Configuration

[Click here to view code image](#)

```
Switch3# show run interface fastethernet 0/19
!
interface fastethernet0/19
  no switchport
  ip address 1.1.1.1 255.255.255.0

Switch4# show run interface fastethernet 0/19
!
interface fastethernet0/19
  no switchport
  ip address 1.1.1.2 255.255.255.0
```

Section 2: MPLS and OSPF (27 Points)

- Configure OSPF on your routers, per [Figure 3-6](#), to enable your network to transport MPLS and MP-BGP. All required interfaces (including Loopback 0) should be configured to belong to Area 0. Ensure that all OSPF configuration is entered under the interfaces. (3 points)

OSPF is used as the IGP in which to advertise the router loopback addresses, which will, of course, be used for the MPLS connectivity. The question directs you to configure OSPF directly under the interfaces of the routers. [Example 3-3](#) shows the loopback interfaces of each router from R1's perspective advertised as host routes as required for MPLS. If you have configured this correctly, as shown in [Example 3-3](#), you have scored 3 points. Consider using the **show ip ospf interface** command to verify your configuration.

Example 3-3 OSPF Configuration and Verification

[Click here to view code image](#)

```
R1(config-if)# int lo0
R1(config-if)# ip ospf 1 area 0
R1(config-if)# int Gi0/1
R1(config-if)# ip ospf 1 area 0

R2(config-if)# int lo0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# int Fa0/0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# int Fa0/1
R2(config-if)# ip ospf 1 area 0

R3(config-if)# int lo0
```

```

R3(config-if)# ip ospf 1 area 0
R3(config-if)# int Gi0/0
R3(config-if)# ip ospf 1 area 0
R3(config-if)# int Gi0/1
R3(config-if)# ip ospf 1 area 0

R4(config-if)# int lo0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# int gi0/0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# int Gi0/1
R4(config-if)# ip ospf 1 area 0

R5(config-if)# int lo0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# int gi0/0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# int Gi0/1
R5(config-if)# ip ospf 1 area 0

R6(config-if)# int lo0
R6(config-if)# ip ospf 1 area 0
R6(config-if)# int gi0/0
R6(config-if)# ip ospf 1 area 0

```

```

R1# show ip route ospf
    120.0.0.0/8 is variably subnetted, 11 subnets, 2 masks
O      120.100.2.1/32 [110/2] via 120.100.132.2, 00:05:00, GigabitEthernet0/1
O      120.100.3.1/32 [110/2] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.4.1/32 [110/12] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.5.1/32 [110/22] via 120.100.132.3, 00:02:36, GigabitEthernet0/1
O      120.100.6.1/32 [110/22] via 120.100.132.3, 00:01:19, GigabitEthernet0/1
O      120.100.25.0/24 [110/31] via 120.100.132.3, 00:02:26, GigabitEthernet0/1
O      120.100.34.0/24 [110/11] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.45.0/24 [110/21] via 120.100.132.3, 00:06:16, GigabitEthernet0/1

```

- Configure MPLS on all routers within the OSPF domain; use LDP, ensuring that TDP can be used on unused interfaces without specifically configuring these interfaces for TDP. Routers R1 and R6 will become your PE routers, whereas R2, R3, R4, and R5 will become P routers. (4 points)

Configuration is required on each router for them to become LSRs (label switch routers). The LSRs must have loopback interfaces with an address mask of 32 bits, and these interfaces must be reachable within the global IP routing table (which the previous question achieved). R1 and R6 are the PE (provider edge) routers, which will be used to connect to switches in later questions simulating CE (customer edge) devices. R2, R3, R4, and R5 become the P (provider) routers, which will be used to switch labeled packets between the PE routers. The question tells you to use LDP (Label Distribution Protocol) but facilitate the future use of TDP (Tag Distribution Protocol) without further configuration on unused interfaces. This is achieved by configuring TDP globally and LDP under each interface used for MPLS within this lab. (The default global and interface configuration is LDP.) The PE routers require only MPLS configured on their serial interfaces toward the P routers. If you have configured this correctly, as shown in [Example 3-4](#), you have scored 4 points.

Example 3-4 MPLS Configuration

[Click here to view code image](#)

```
R1(config)# mpls label protocol tdp
R1(config)# interface Gi0/1
R1(config-if)# mpls label protocol ldp
R1(config-if)# mpls ip

R2(config)# mpls label protocol tdp
R2(config)# interface Fa0/0
R2(config-if)# mpls label protocol ldp
R2(config-if)# mpls ip
R2(config)# interface Fa0/1
R2(config-if)# mpls label protocol ldp
R2(config-if)# mpls ip

R3(config)# mpls label protocol tdp
R3(config)# interface Gi0/0
R3(config-if)# mpls label protocol ldp
R3(config-if)# mpls ip
R3(config-if)# interface Gi0/1
R3(config-if)# mpls label protocol ldp
R3(config-if)# mpls ip

R4(config)# mpls label protocol tdp
R4(config)# interface GigabitEthernet0/0
R4(config-if)# mpls label protocol ldp
R4(config-if)# mpls ip
R4(config-if)# interface Gi0/1
R4(config-if)# mpls label protocol ldp
R4(config-if)# mpls ip

R5(config)# mpls label protocol tdp
R5(config)# interface Gi0/0
R5(config-if)# mpls label protocol ldp
R5(config-if)# mpls ip
R5(config-if)# interface Gi0/1
R5(config-if)# mpls label protocol ldp
R5(config-if)# mpls ip

R6(config)# mpls label protocol tdp
R6(config)# interface Gi0/0
R6(config-if)# mpls label protocol ldp
R6(config-if)# mpls ip
```

[Example 3-5](#) shows verification of the configuration with the LDP peering between each router. Notice that the loopback addresses are used for LDP peer identification.

Example 3-5 MPLS Configuration Verification

[Click here to view code image](#)

```
R1# show mpls ldp neighbor
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.1.1:0
```

```
TCP connection: 120.100.2.1.40418 - 120.100.1.1.646
State: Oper; Msgs sent/rcvd: 69/71; Downstream
Up time: 00:47:20
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.123.2
Addresses bound to peer LDP Ident:
    120.100.123.2    120.100.25.2    120.100.2.1
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.1.1:0
TCP connection: 120.100.3.1.51369 - 120.100.1.1.646
State: Oper; Msgs sent/rcvd: 68/68; Downstream
Up time: 00:47:18
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.123.3
Addresses bound to peer LDP Ident:
    120.100.123.3    120.100.3.1    120.100.34.3
```

```
R2# show mpls ldp neighbor
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.3.1.16991 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 71/68; Downstream
Up time: 00:46:33
LDP discovery sources:
    fastethernet0/0, Src IP addr: 120.100.123.3
    fastethernet0/1, Src IP addr: 120.100.34.3
Addresses bound to peer LDP Ident:
    120.100.123.3    120.100.3.1    120.100.34.3
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.5.1.13826 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 73/76; Downstream
Up time: 00:46:24
LDP discovery sources:
    fastethernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
    120.100.25.5    120.100.5.1    5.5.5.5    120.100.45.5
    100.100.100.5
Peer LDP Ident: 120.100.1.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.1.1.646 - 120.100.2.1.40418
State: Oper; Msgs sent/rcvd: 69/68; Downstream
Up time: 00:46:07
LDP discovery sources:
    fastethernet0/0, Src IP addr: 120.100.123.1
Addresses bound to peer LDP Ident:
    120.100.123.1    120.100.1.1
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.4.1.47401 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 54/57; Downstream
Up time: 00:32:28
LDP discovery sources:
    fastethernet0/1, Src IP addr: 120.100.34.4
Addresses bound to peer LDP Ident:
    120.100.4.1    4.4.4.4    120.100.45.4    100.100.100.4
    120.100.34.4
```

```
R3# show mpls ldp neighbor
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.2.1.646 - 120.100.3.1.16991
State: Oper; Msgs sent/rcvd: 69/72; Downstream
Up time: 00:47:11
LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.123.2
    GigabitEthernet0/1, Src IP addr: 120.100.25.2
```

```

Addresses bound to peer LDP Ident:
  120.100.123.2    120.100.25.2    120.100.2.1
Peer LDP Ident: 120.100.1.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.1.1.646 - 120.100.3.1.51369
State: Oper; Msgs sent/rcvd: 67/67; Downstream
Up time: 00:46:43
LDP discovery sources:
  GigabitEthernet0/0, Src IP addr: 120.100.123.1
Addresses bound to peer LDP Ident:
  120.100.123.1    120.100.1.1
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.5.1.53107 - 120.100.3.1.646
State: Oper; Msgs sent/rcvd: 67/74; Downstream
Up time: 00:45:22
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
  120.100.25.5    120.100.5.1    5.5.5.5          120.100.45.5
  100.100.100.5
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.4.1.15940 - 120.100.3.1.646
State: Oper; Msgs sent/rcvd: 52/56; Downstream
Up time: 00:33:06
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.34.4
Addresses bound to peer LDP Ident:
  120.100.4.1      4.4.4.4        120.100.45.4    100.100.100.4
  120.100.34.4

```

```

R4# show mpls ldp neighbor
Peer LDP Ident: 120.100.6.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.6.1.55234 - 120.100.4.1.646
State: Oper; Msgs sent/rcvd: 74/76; Downstream
Up time: 00:43:52
LDP discovery sources:
  GigabitEthernet0/0, Src IP addr: 120.100.45.6
Addresses bound to peer LDP Ident:
  120.100.6.1      6.6.6.6        100.100.100.6   120.100.45.6
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.5.1.57689 - 120.100.4.1.646
State: Oper; Msgs sent/rcvd: 72/74; Downstream
Up time: 00:43:48
LDP discovery sources:
  GigabitEthernet0/0, Src IP addr: 120.100.45.5
  GigabitEthernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
  120.100.25.5    120.100.5.1    5.5.5.5          120.100.45.5
  100.100.100.5
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.2.1.646 - 120.100.4.1.47401
State: Oper; Msgs sent/rcvd: 55/52; Downstream
Up time: 00:30:52
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.25.2
Addresses bound to peer LDP Ident:
  120.100.123.2    120.100.25.2    120.100.2.1
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.3.1.646 - 120.100.4.1.15940
State: Oper; Msgs sent/rcvd: 54/50; Downstream
Up time: 00:30:52
LDP discovery sources:

```

```
GigabitEthernet0/1, Src IP addr: 120.100.34.3
Addresses bound to peer LDP Ident:
  120.100.123.3    120.100.3.1      120.100.34.3

R5# show mpls ldp neighbor
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.5.1:0
  TCP connection: 120.100.2.1.646 - 120.100.5.1.13826
  State: Oper; Msgs sent/rcvd: 80/77; Downstream
  Up time: 00:49:55
  LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.25.2
  Addresses bound to peer LDP Ident:
    120.100.123.2    120.100.25.2      120.100.2.1
Peer LDP Ident: 120.100.6.1:0; Local LDP Ident 120.100.5.1:0
  TCP connection: 120.100.6.1.18472 - 120.100.5.1.646
  State: Oper; Msgs sent/rcvd: 81/81; Downstream
  Up time: 00:48:58
  LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.6
  Addresses bound to peer LDP Ident:
    120.100.6.1      6.6.6.6          100.100.100.6    120.100.45.6
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.5.1:0
  TCP connection: 120.100.4.1.646 - 120.100.5.1.57689
  State: Oper; Msgs sent/rcvd: 80/78; Downstream
  Up time: 00:48:54
  LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.4
    GigabitEthernet0/1, Src IP addr: 120.100.34.4
  Addresses bound to peer LDP Ident:
    120.100.4.1      4.4.4.4          120.100.45.4      100.100.100.4
    120.100.34.4
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.5.1:0
  TCP connection: 120.100.3.1.646 - 120.100.5.1.53107
  State: Oper; Msgs sent/rcvd: 77/70; Downstream
  Up time: 00:48:17
  LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.34.3
  Addresses bound to peer LDP Ident:
    120.100.123.3    120.100.3.1      120.100.34.3

R6# show mpls ldp neighbor
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.6.1:0
  TCP connection: 120.100.5.1.646 - 120.100.6.1.18472
  State: Oper; Msgs sent/rcvd: 82/82; Downstream
  Up time: 00:49:31
  LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.5
  Addresses bound to peer LDP Ident:
    120.100.25.5    120.100.5.1      5.5.5.5          120.100.45.5
    100.100.100.5
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.6.1:0
  TCP connection: 120.100.4.1.646 - 120.100.6.1.55234
  State: Oper; Msgs sent/rcvd: 82/80; Downstream
  Up time: 00:49:31
  LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.4
  Addresses bound to peer LDP Ident:
    120.100.4.1      4.4.4.4          120.100.45.4      100.100.100.4
    120.100.34.4
```

- You will be configuring two VPNs over your MPLS networks per [Figure 3-8](#) between PE routers of BLUE and RED. At this point, assign the following interfaces on each PE router into separate routing instances within the routers:

PE R1 interface Gi0/0 VLAN10 connection into VPN BLUE

PE R1 interface Gi0/0 VLAN 50 connection into VPN RED

PE R6 interface Gi0/1 VLAN 20 connection into VPN BLUE

PE R6 interface Gi0/1 VLAN 100 connection into VPN RED

Configure VPN BLUE to use an RD of 100 and VPN RED to use an RD of 200 for both importing and exporting routes into your BGP network, which will be configured later with an autonomous system of AS65001. (4 points)

You are required to create virtual routing forwarding (VRF) instances on the PE routers and assign the subinterfaces on each PE router into these. This will ultimately provide end-to-end virtual private networking (VPN) connectivity over the MPLS network for your CE devices to communicate. You are directed to use a route descriptor (RD) of 100 for the BLUE VRF and 200 for the RED VRF and must combine this with the BGP autonomous system number of 65001 to import and export route target extended communities for the specified VRFs. The actual BGP configuration will be configured later in the lab. If you have configured this correctly, as shown in [Example 3-6](#), you have scored 4 points.

Example 3-6 VRF Configuration

[Click here to view code image](#)

```
R1(config)# ip vrf BLUE
R1(config-vrf)# rd 65001:100
R1(config-vrf)# route-target export 65001:100
R1(config-vrf)# route-target import 65001:100
R1(config-vrf)# !
R1(config-vrf)# ip vrf RED
R1(config-vrf)# rd 65001:200
R1(config-vrf)# route-target export 65001:200
R1(config-vrf)# route-target import 65001:200
R1(config-vrf)# exit
R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# ip vrf forwarding BLUE
R1(config-subif)# interface GigabitEthernet0/0.50
R1(config-subif)# ip vrf forwarding RED
```

```
R6(config)# ip vrf BLUE
R6(config-vrf)# rd 65001:100
R6(config-vrf)# route-target export 65001:100
R6(config-vrf)# route-target import 65001:100
R6(config-vrf)# !
R6(config-vrf)# ip vrf RED
R6(config-vrf)# rd 65001:200
R6(config-vrf)# route-target export 65001:200
R6(config-vrf)# route-target import 65001:200
R6(config-vrf)# exit
R6(config)# interface GigabitEthernet0/1.20
R6(config-subif)# ip vrf forwarding BLUE
```

```
R6(config)# interface GigabitEthernet0/1.100
R6(config-subif)# ip vrf forwarding RED
```

- Create a network between PE router R1 and CE device SW1 using a VLAN10 interface on SW1 that can be trunked toward R1. This network will reside in the BLUE VPN. Use a subnet of 10.10.10.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)

This is a simple configuration task to assign IP connectivity between the PE and CE devices for future routing between the devices and remote VPN connectivity via R6. The new VLAN10 must be created on SW1, and this VLAN should have already been permitted to flow through to R1 as an allowed VLAN. The subinterface of Gigabit 0/0.10 on R1 has been assigned to the BLUE VRF during the previous question, so connectivity between SW1 and R1 should now be possible (when IP addresses are assigned). When testing, remember that R1 must use the appropriate VRF to confirm connectivity, because a normal ping would be sourced from the global routing table and will fail. If you have configured this correctly, as shown in [Example 3-7](#), you have scored 2 points.

Example 3-7 BLUE VRF IP Addressing and Local Connectivity Testing

[Click here to view code image](#)

```
R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# ip add 10.10.10.1 255.255.255.252

Switch1(config)# vlan 10
Switch1(config-vlan)# exit
Switch1(config)# interface vlan 10
Switch1(config-if)# no shutdown
Switch1(config-if)# ip add 10.10.10.2 255.255.255.252

R1# ping vrf BLUE 10.10.10.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.10.10.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

- Create a network between PE router R6 and CE device SW2 using a VLAN 20 interface on SW2 that can be trunked toward R6. This network will reside in the BLUE VPN. Use a subnet of 10.10.20.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)

This is a simple configuration task as per the previous question to assign connectivity between the PE and CE devices for future routing between the devices and remote VPN connectivity via R1. The new VLAN 20 must be created on SW2, and this VLAN already should have been permitted to flow through to R6 as an allowed VLAN. The subinterface of Gigabit 0/1.20 on R6 has been assigned to the BLUE VRF during a previous question, so connectivity between SW2 and R6 should now be possible. When testing, remember that R6 must use the appropriate VRF to confirm connectivity. If you have configured this correctly, as shown in [Example 3-8](#), you have scored 2 points.

Example 3-8 BLUE VRF IP Addressing and Local Connectivity Testing

[Click here to view code image](#)

```
R6(config)# interface GigabitEthernet0/1.20
R6(config-subif)# ip add 10.10.20.1 255.255.255.252

Switch2(config)# vlan 20
Switch2(config-vlan)# exit
Switch2(config)# interface vlan 20
Switch2(config-if)# no shutdown
Switch2(config-if)# ip add 10.10.20.2 255.255.255.252

R6# ping vrf BLUE 10.10.20.2
```

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.10.20.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms

-
- Create a network between PE router R1 and CE device SW3 using a VLAN 50 interface on SW3 that can be trunked toward R1; this network will reside in the RED VPN. Use a subnet of 130.50.50.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)

Here's another simple configuration to assign connectivity between the PE and CE devices for future routing between the devices and remote VPN connectivity via R6. VLAN 50 has been previously created on SW3 and SW1 within the initial configuration. This VLAN should have already been permitted to flow through SW1 to R1 as an allowed VLAN. The subinterface of Gigabit 0/0.50 on R1 has been assigned to the RED VRF during a previous question, so connectivity between SW3 and R1 should now be possible. When testing, remember that R1 must use the appropriate VRF to confirm connectivity. If you have configured this correctly, as shown in [Example 3-9](#), you have scored 2 points.

Example 3-9 RED VRF IP Addressing and Local Connectivity Testing

[Click here to view code image](#)

```
R1(config)# interface GigabitEthernet0/0.50
R1(config-subif)# ip add 130.50.50.1 255.255.255.252

Switch3(config)# interface vlan 50
Switch3(config-if)# no shutdown
Switch3(config-if)# ip add 130.50.50.2 255.255.255.252

R1# ping vrf RED 130.50.50.2
```

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 130.50.50.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms

-
- Create a network between PE router R6 and CE device SW4 using a VLAN 100 interface on SW4 that can be trunked toward R6; this network will reside in the RED VPN. Use a subnet of 130.100.100.0/30 with .1/30 assigned to the PE and .2/30 assigned to the CE. (2 points)

This is the final configuration task to assign connectivity between the PE and CE devices for future routing between the devices and remote VPN connectivity via R1. VLAN 100 has been previously created on SW4 and SW2, within the initial configuration; this VLAN should have already been permitted to flow through SW2 to R6 as an allowed VLAN. The subinterface of Gigabit 0/1.100 on R6 has been assigned to the RED VRF during a previous question, so connectivity between SW4 and R6 should now be possible. When testing, remember that R6 must use the appropriate VRF to confirm connectivity. If you have configured this correctly, as shown in [Example 3-10](#), you have scored 2 points.

Example 3-10 RED VRF IP Addressing and Local Connectivity Testing

[Click here to view code image](#)

```
R6(config)# interface GigabitEthernet0/1.100
R6(config-subif)# ip add 130.100.100.1 255.255.255.252

Switch4(config)# interface vlan 100
Switch4(config-if)# no shutdown
Switch4(config-if)# ip add 130.100.100.2 255.255.255.252

R6# ping vrf RED 130.100.100.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 130.100.100.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

Section 3: BGP (5 Points)

- Configure MP-BGP between your PE routers, per [Figure 3-9](#), to enable your network to transport the VPNv4 addresses of your configured VPNs (BLUE and RED). Use loopback interfaces for peering between your PE routers. You will configure the actual VPN routing in later questions. (4 points)

MPLS requires the use of Multiprotocol BGP (MP-BGP) between the PE routers to exchange VPNv4 addresses in addition to IPv4 addresses. The VPNs will be mapped into the configuration later, so this question is a straightforward peering and VPNv4 setup task. The configuration requires you to peer from your loopback interfaces, which are advertised via your P routers within OSPF and that extended communities are used between PE routers to advertise your VPNv4 addresses successfully. You should be aware that route targets (RTs) are implemented by the use of the BGP extended community (64 bits) and therefore the **send-community both** value must be configured within MP-BGP. The **next-hop-self** command is optional and strictly required only when you have an eBGP configuration to preserve the next-hop information to peers; you will not lose any points if you added this or left it out. The actual VPN portion of MP-BGP will be configured later within the IPv4 address family for VRF-specific advertisements. This is a simple MP-BGP network with only two PE routers; additional PE routers would require a full mesh of iBGP peering or configuration of route-reflection to aid scalability.

If you have configured this correctly, as shown in [Example 3-11](#), you have scored 4 points.

Example 3-11 MP-BGP Configuration

[Click here to view code image](#)

```
R1(config)# router bgp 65001
R1(config-router)# no synchronization
R1(config-router)# no auto-summary
R1(config-router)# neighbor 120.100.6.1 remote-as 65001
R1(config-router)# neighbor 120.100.6.1 update-source Loopback0
R1(config-router)# address-family vpng4
R1(config-router-af)# neighbor 120.100.6.1 activate
R1(config-router-af)# neighbor 120.100.6.1 next-hop-self
R1(config-router-af)# neighbor 120.100.6.1 send-community both

R6(config)# router bgp 65001
R6(config-router)# no sync
R6(config-router)# no auto-summary
R6(config-router)# neighbor 120.100.1.1 remote-as 65001
R6(config-router)# neighbor 120.100.1.1 update-source Loopback0

R6(config-router)# address-family vpng4
R6(config-router-af)# neighbor 120.100.1.1 activate
R6(config-router-af)# neighbor 120.100.1.1 next-hop-self
R6(config-router-af)# neighbor 120.100.1.1 send-community both
```

Section 4: EIGRP and MP-BGP (3 Points)

- Configure EIGRP per [Figure 3-9](#) between your PE router R6 and CE switch SW2. Use an EIGRP virtual instance name of VPN on R6 and a process number of 10 on SW2. Use VLAN 20 for EIGRP connectivity between R6 and SW2. Advertise all preconfigured loopback networks on SW2 to R6 for the BLUE VPN. (1 point)

Until now, the questions have merely dealt with setting up the infrastructure for MPLS connectivity. Now you are requested to advertise routes from your CE switch SW2 to PE router R6, which will ultimately be advertised throughout the BLUE VPN to the remote PE router R1 and CE switch SW1. You'll realize that to peer successfully with EIGRP you would need to be operating within the same autonomous system number, yet the question enforces you to run differing autonomous system numbers. PE routers would normally connect to multiple customers, so it is unreasonable to expect that each EIGRP domain should run the same autonomous system number. Therefore, the autonomous system is assigned with the **address-family vrf**-specific command. [Example 3-12](#) details the EIGRP configuration and resulting neighbor relationship and route propagation between R6 and SW2. If you have configured this correctly, as shown in [Example 3-12](#), you have scored 1 point.

Note

The IP addressing for VLAN 20 on SW2 and associated subinterfaces on R6 has previously been configured. The BLUE VRF has also been associated to the R6 subinterface previously.

Example 3-12 R6 and Switch 2 EIGRP Configuration and Verification

[Click here to view code image](#)

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R6(config-router-af)# network 10.10.20.0 0.0.0.3
```

```
Switch2(config)# ip routing
Switch2(config)# router eigrp 10
Switch2(config-router)# no auto-summary
Switch2(config-router)# network 10.10.20.0 0.0.0.3
Switch2(config-router)# network 10.2.2.0 0.0.0.255
Switch2(config-router)# network 10.2.3.0 0.0.0.255
Switch2(config-router)# network 10.2.4.0 0.0.0.255
```

```
R6# show ip eigrp vrf BLUE neighbors
IP-EIGRP neighbors for process 10
      H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
      0   10.10.20.2       Gi0/1.20        (sec)   (ms)          Cnt Num
      0   10.10.20.2       Gi0/1.20        11 00:04:18   1    200  0   1
R6#
R6# show ip route vrf BLUE eigrp
  10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks
D     10.2.2.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
D     10.2.3.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
D     10.2.4.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
```

Note

The IP addressing for VLAN 10 on SW1 and associated subinterfaces on R1 has previously been configured. The BLUE VRF has also been associated to the R1 subinterface previously.

- Configure EIGRP per [Figure 3-9](#) between your PE router R1 and CE switch SW1. Use an EIGRP virtual instance name of VPN on R1 and a process number of 10 on SW1. Use VLAN10 for EIGRP connectivity between R1 and SW1. Advertise all preconfigured loopback networks on SW1 to R1 for the BLUE VPN. (1 point)

As per the previous question, you are requested to advertise routes from your CE switch SW1 to PE router R1, which will ultimately be advertised throughout the BLUE VPN to the remote PE router R6 and CE switch SW2. [Example 3-13](#) details the EIGRP configuration and resulting neighbor relationship and route propagation between R1 and SW1. If you have configured this correctly, as shown in [Example 3-13](#), you have scored 1 point.

Example 3-13 R1 and Switch 1 EIGRP Configuration and Verification

[Click here to view code image](#)

```
R1(config)# router eigrp VPN
R1(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
```

```
R1(config-router-af) # network 10.10.10.0 0.0.0.3

Switch1(config) # ip routing
Switch1(config) # router eigrp 10
Switch1(config-router) # no auto-summary
Switch1(config-router) # network 10.10.10.0 0.0.0.3
Switch1(config-router) # network 10.1.1.0 0.0.0.255
Switch1(config-router) # network 10.1.2.0 0.0.0.255
Switch1(config-router) # network 10.1.3.0 0.0.0.255
```

R1# show ip eigrp vrf BLUE neighbors

IP-EIGRP neighbors for process 10

H	Address	Interface	Hold (sec)	Uptime	SRTT (ms)	RTO	Q Cnt	Seq Num
0	10.10.10.2	Gi0/0.10	13	00:00:24	1	200	0	1

R1# show ip eigrp vrf BLUE neighbors

IP-EIGRP neighbors for process 10

H	Address	Interface	Hold (sec)	Uptime	SRTT (ms)	RTO	Q Cnt	Seq Num
0	10.10.10.2	Gi0/0.10	13	00:00:24	1	200	0	1

R1# show ip route vrf BLUE eigrp

10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks

D	10.1.3.0/24	[90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10
D	10.1.2.0/24	[90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10
D	10.1.1.0/24	[90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10

- Configure your PE routers R1 and R6 to transport EIGRP routes from your CE devices between the BLUE VPN using MP-BGP. EIGRP networks residing on SW1 should be seen as internal EIGRP routes on SW2 and vice versa. Ensure that all EIGRP routes have a MED of 50 assigned to them within MP-BGP. Use a default metric of 10000 100 255 1 1500 for BGP routes when redistributed into EIGRP. (1 point)

The full end-to-end VPN routing is achieved at this point by redistributing EIGRP into the appropriate address family for the VRF. The question dictates the metrics you should use. In reality, the metrics are not required because the extended community values of MP-BGP previously configured will effectively transport the internal metrics of EIGRP and ensure that the routes are shown as internal EIGRP routes at the remote location, even though they have been redistributed via another routing protocol. The question is just looking for accuracy and giving you the opportunity to view routes with the metrics and later without if you choose to. [Example 3-14](#) details the configuration required on the PE routers and resulting routes on the CE devices SW1 and SW2. If you have configured this correctly, as shown in [Example 3-14](#), you have scored 1 point.

Example 3-14 PE and CE MP-BGP Redistribution Configuration and Verification

[Click here to view code image](#)

```
R1(config) # router eigrp VPN
R1(config-router) # address-family ipv4 vrf BLUE autonomous-system 10
R1(config-router) # topology base
```

```
R1(config-router-af-topology)# redistribute bgp 65001 metric 10000 100 255 1 1500
R1(config-router-af-topology)# router bgp 65001
R1(config-router)# address-family ipv4 vrf BLUE
R1(config-router-af)# redistribute eigrp 10 metric 50
```

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R6(config-router)# topology base
R6(config-router-af-topology)# redistribute bgp 65001 metric 10000 100 255 1 1500
R6(config-router-af-topology)# router bgp 65001
R6(config-router)# address-family ipv4 vrf BLUE
R6(config-router-af)# redistribute eigrp 10 metric 50
```

```
SW1# show ip route eigrp
D      10.2.2.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.2.3.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.2.4.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.10.20.0/30 [90/28416] via 10.10.10.1, 00:32:05, Vlan10
```

```
SW2# show ip route eigrp
D      10.1.3.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.1.2.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.1.1.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.10.10.0/30 [90/26112] via 10.10.20.1, 00:33:07, Vlan20
```

[Example 3-15](#) details the BGP routes received on the PE routers with the assigned MED value of 50; it also details the MPLS forwarding table for the BLUE VRF. Notice the iBGP routes on the PE routers from the remote PE router with the MED of 50; these are the routes that are propagated to EIGRP CE devices. If you have configured this correctly, as shown in [Example 3-15](#), you have scored 3 points.

Example 3-15 PE MP-BGP and MPLS Verification

[Click here to view code image](#)

```
R6# show ip bgp vpnv4 vrf BLUE
BGP table version is 17, local router ID is 120.100.6.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*>i10.1.1.0/24	120.100.1.1	50	100	0	?
*>i10.1.2.0/24	120.100.1.1	50	100	0	?
*>i10.1.3.0/24	120.100.1.1	50	100	0	?
*> 10.2.2.0/24	10.10.20.2	50		32768	?
*> 10.2.3.0/24	10.10.20.2	50		32768	?
*> 10.2.4.0/24	10.10.20.2	50		32768	?
*>i10.10.10.0/30	120.100.1.1	0	100	0	?
*> 10.10.20.0/30	0.0.0.0	0		32768	?

```
R1# show ip bgp vpnv4 vrf BLUE
BGP table version is 17, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
```

Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*> 10.1.1.0/24	10.10.10.2	50		32768	?
*> 10.1.2.0/24	10.10.10.2	50		32768	?
*> 10.1.3.0/24	10.10.10.2	50		32768	?
*>i10.2.2.0/24	120.100.6.1	50	100	0	?
*>i10.2.3.0/24	120.100.6.1	50	100	0	?
*>i10.2.4.0/24	120.100.6.1	50	100	0	?
*> 10.10.10.0/30	0.0.0.0		0	32768	?
*>i10.10.20.0/30	120.100.6.1	0	100	0	?

R1# show mpls forwarding-table vrf BLUE

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes switched	Outgoing interface	Next Hop
26	Untagged	10.1.3.0/24[V]	0	Gi0/0.10	10.10.10.2
27	Untagged	10.1.2.0/24[V]	0	Gi0/0.10	10.10.10.2
28	Aggregate	10.10.10.0/30[V]	0		
29	Untagged	10.1.1.0/24[V]	0	Gi0/0.10	10.10.10.2

R6# show mpls forwarding-table vrf BLUE

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes switched	Outgoing interface	Next Hop
26	Untagged	10.2.2.0/24[V]	0	Gi0/1.20	10.10.20.2
27	Untagged	10.2.3.0/24[V]	0	Gi0/1.20	10.10.20.2
28	Untagged	10.2.4.0/24[V]	0	Gi0/1.20	10.10.20.2
29	Aggregate	10.10.20.0/30[V]	0		

Section 5: OSPF and MP-BGP (6 Points)

- Configure OSPF per [Figure 3-10](#) for your VRF RED with a process number of 3 on PE router R1 and SW3 using VLAN 50 for connectivity. Use a process ID of 2 on PE router R6 and CE device SW4 using VLAN 100 for connectivity. You should permit only internal OSPF routes to be advertised across your VPN and ensure that the redistribution of BGP routes into OSPF are assigned as type 1 external routes with no manually adjusted cost associated with them. It is acceptable for these routes to come through as /32 routes because of default OSPF behavior of loopback interfaces. (2 points)

You are requested to configure OSPF over your MPLS network between CE devices SW3 and SW4 via your PE routers R1 and R6. [Figure 3-10](#) indicates that all loopback interfaces are to be included in OSPF on both CE devices. You should be aware that OSPF will advertise these as host routes, but the question states that this is acceptable behavior. As with the EIGRP question, you are requested to manipulate the redistribution of the IGP into BGP, but in reality the routes would appear to have not been redistributed through another routing protocol by default. This direction is actually a red herring for the next question when the routes at the CE devices appear as external routes when they should in fact be internal routes. You are requested to permit only internal OSPF routes to be redistributed into BGP, which is a simple **match internal** parameter on the redistribution configuration. [Example 3-16](#) details the required configuration and verification. If you have configured this correctly, as shown in [Example 3-16](#), you have scored 2 points.

Example 3-16 VRF RED OSPF Configuration and Verification

```
SW3(config)# ip routing
SW3(config)# router ospf 3
SW3(config-router)# network 130.50.50.0 0.0.0.3 area 0
SW3(config-router)# network 10.33.33.0 0.0.0.255 area 1
SW3(config-router)# network 10.33.34.0 0.0.0.255 area 1
SW3(config-router)# network 10.33.35.0 0.0.0.255 area 1

SW4(config)# ip routing
SW4(config)# router ospf 2
SW4(config-router)# network 130.100.100.0 0.0.0.3 area 0
SW4(config-router)# network 10.44.44.0 0.0.0.255 area 2
SW4(config-router)# network 10.44.45.0 0.0.0.255 area 2
SW4(config-router)# network 10.44.46.0 0.0.0.255 area 2

R1(config)# router ospf 3 vrf RED
R1(config-router)# network 130.50.50.0 0.0.0.3 area 0
R1(config-router)# redistribute bgp 65001 subnets metric-type 1
R1(config-router)# router bgp 65001
R1(config-router)# address-family ipv4 vrf RED
R1(config-router-af)# redistribute ospf 3 match internal

R6(config)# router ospf 2 vrf RED
R6(config-router)# net 130.100.100.0 0.0.0.3 area 0
R6(config-router)# redistribute bgp 65001 subnets metric-type 1
R6(config-router)# router bgp 65001
R6(config-router)# address-family ipv4 vrf RED
R6(config-router-af)# redistribute ospf 2 match internal

R1# show ip route vrf RED ospf

Routing Table: RED

    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.33.34.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.35.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.33.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50

R6# show ip route vrf RED ospf

Routing Table: RED

    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.44.46.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.45.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.44.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100

SW3# show ip route ospf
    130.100.0.0/30 is subnetted, 1 subnets
O E1    130.100.100.0 [110/2] via 130.50.50.1, 00:06:08, Vlan50
        10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.44.46.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.45.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.44.1/32 [110/3] via 130.50.50.1, 00:02:55, Vlan50

SW4# show ip route ospf
    130.50.0.0/30 is subnetted, 1 subnets
O E1    130.50.50.0 [110/2] via 130.100.100.1, 00:03:37, Vlan100
```

```
10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.33.34.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.35.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.33.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
```

Note

The IP addressing for VLAN50 on SW3 and associated subinterface on R1 and VLAN 100 on SW4 and associated subinterface on R6 has previously been configured. The RED VRF has also been associated to the R1 and R6 subinterfaces previously.

- You will notice that your OSPF IA (intra-area) routes between CE devices SW3 and SW4 appear as type 1 external routes. Configure your OSPF network appropriately to ensure that the routes are displayed correctly as IA routes. You are not permitted to adjust the OSPF redistribution into BGP as directed in the previous question. Maintain the OSPF process IDs are previously directed. You are permitted to configure only router R1. (4 points)

This is a tricky question and one that will really eat into your time (the kind of question that if the answer doesn't jump out at you and the points don't look appealing enough, it's one to park and come back to). Because you have your routes in place and following questions do not build from this one, you can confidently leave questions like this for later. As stated previously, the redistribution into type 1 is actually somewhat misleading. When you look at the routes in [Example 3-17](#) for the PE routers, you will see that they are actually IA routes at this point. So, it is only when these routes are advertised to the CE devices that the type 1 external route change occurs.

Example 3-17 VRF RED OSPF Routes

[Click here to view code image](#)

```
R1# show ip route vrf RED ospf
```

Routing Table: RED

```
10.0.0.0/32 is subnetted, 6 subnets
O IA    10.33.34.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.35.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.33.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
```

```
R6# show ip route vrf RED ospf
```

Routing Table: RED

```
10.0.0.0/32 is subnetted, 6 subnets
O IA    10.44.46.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.45.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.44.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
```

```
SW3# show ip route ospf
```

```
130.100.0.0/30 is subnetted, 1 subnets
O E1    130.100.100.0 [110/2] via 130.50.50.1, 00:06:08, Vlan50
10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.44.46.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.45.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
```

```
O E1      10.44.44.1/32 [110/3] via 130.50.50.1, 00:02:55, Vlan50
SW4# show ip route ospf
  130.50.0.0/30 is subnetted, 1 subnets
O E1      130.50.50.0 [110/2] via 130.100.100.1, 00:03:37, Vlan100
  10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1      10.33.34.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1      10.33.35.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1      10.33.33.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
```

The clue is actually in the question “Maintain the OSPF process IDs as previously directed.” Statements such as this should make you think, “Okay, so if I did change the process ID, it would most likely work. Why would that do it, and how else can I achieve that?” OSPF has a domain ID by default. This is the same as the process ID. If the process IDs differ on PE routers that form the VPN, the LSA is changed to a type 5 and the routes become external. You might not have known that, but it is the kind of thing that you gain through research and rack time. Because you are not permitted to change the process ID, you are left with only the option of changing the domain ID. [Example 3-18](#) details the domain ID information on your PE routers, the configuration required to change the domain ID on one of your PE’s Router R1, and the resulting IA routes received on your CE devices. If you have configured this correctly, as shown in [Example 3-18](#), you have scored 4 points.

Example 3-18 Domain ID Configuration and OSPF Route Verification

[Click here to view code image](#)

```
R1# show ip ospf 3 | include Domain
  Domain ID type 0x0005, value 0.0.0.3

R6# show ip ospf 2 | include Domain
  Domain ID type 0x0005, value 0.0.0.2

R1(config)# router ospf 3 vrf RED
R1(config-router)# domain-id 0.0.0.2

SW3# show ip route ospf
  130.100.0.0/30 is subnetted, 1 subnets
O IA      130.100.100.0 [110/2] via 130.50.50.1, 00:00:09, Vlan50
  10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O IA      10.44.46.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50
O IA      10.44.45.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50
O IA      10.44.44.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50
SW3# 

SW4# show ip route ospf
  130.50.0.0/30 is subnetted, 1 subnets
O IA      130.50.50.0 [110/2] via 130.100.100.1, 00:00:07, Vlan100
  10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O IA      10.33.34.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100
O IA      10.33.35.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100
O IA      10.33.33.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100
```

Section 6: MPLS (7 Points)

- Leak network 10.1.1.0/24 from SW1 VRF BLUE on PE R1 into the VRF RED on PE1. Similarly, leak 10.44.44.0/24 from VRF RED into VRF BLUE on R6. Both Switch 1 and Switch 4 should receive the following routes:

SW1# **show ip route | include 10.44.44.0**

D EX 10.44.44.0/24 [170/XXXXXX] via 10.10.10.1, 00:00:27, Vlan10

SW1#

SW4# **show ip route | include 10.1.1.0**

O E1 10.1.1.0/24 [110/XX] via 130.100.100.1, 00:03:04, Vlan100

SW4#

Verify your configuration by pinging from VRF RED SW4 10.44.44.1 to VRF BLUE SW1 10.1.1.1 SW1. (5 points)

This is a straightforward VRF export question with a slight twist for the attentive in that the OSPF route 10.44.44.0/24 originates from a loopback interface on Switch 4, so OSPF must be manipulated to treat this interface as a point-to-point network to advertise the /24 mask. The route leaking is achieved by creation of export maps on the PE routers R1 and R6, permitting the required routes from each VRF to the existing BLUE and RED VRF advertisements by adding them to the appropriate route target (RT) within MP-BGP by use of the **set extcommunity rt XXXXX:XXX additive** command.

[Example 3-19](#) details the required configuration on PE routers R1, R6, and the CE device SW4; the resulting verification of the route advertisements and testing are also shown. If you have configured this correctly, as shown in [Example 3-19](#), you have scored 5 points.

Example 3-19 Selective VRF Export Configuration and Verification

[Click here to view code image](#)

```
Sw4(config)# interface Loopback0
Sw4(config-if)# ip ospf network point-to-point

R1(config)# ip vrf BLUE
R1(config-vrf)# export map SW1
R1(config-vrf)# access-list 10 permit 10.1.1.0 0.0.0.255
R1(config-vrf)# exit
R1(config)# route-map SW1 permit 10
R1(config-route-map)# match ip address 10
R1(config-route-map)# set extcommunity rt 65001:200 additive

R6(config)# ip vrf RED
R6(config-vrf)# export map SW4
R6(config-vrf)# access-list 10 permit 10.44.44.0 0.0.0.255
R6(config-vrf)# exit
R6(config)# route-map SW4 permit 10
R6(config-route-map)# match ip address 10
R6(config-route-map)# set extcommunity rt 65001:100 additive

! R1 is now sending 10.1.1.0 into VRF RED and R6 10.44.44.0 into VRF BLUE

R1# show ip bgp vpnv4 vrf RED
```

```
BGP table version is 33, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:200 (default for vrf RED)					
*> 10.33.33.1/32	130.50.50.2	2		32768	?
*> 10.33.34.1/32	130.50.50.2	2		32768	?
*> 10.33.35.1/32	130.50.50.2	2		32768	?
*>i10.44.44.1/32	120.100.6.1	2	100	0	?
*>i10.44.45.1/32	120.100.6.1	2	100	0	?
*>i10.44.46.1/32	120.100.6.1	2	100	0	?
*> 130.50.50.0/30	0.0.0.0		0		32768 ?
*>i130.100.100.0/30	120.100.6.1	0	100	0	?

! No sign of the 10.1.1.0 route, clear the BGP session to kick start the export map

```
R1# clear ip bgp *
R1# show ip bgp vpng4 vrf RED
BGP table version is 34, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:200 (default for vrf RED)					
*> 10.1.1.0/24	10.10.10.2	50		32768	?
*> 10.33.33.1/32	130.50.50.2	2		32768	?
*> 10.33.34.1/32	130.50.50.2	2		32768	?
*> 10.33.35.1/32	130.50.50.2	2		32768	?
*>i10.44.44.1/32	120.100.6.1	2	100	0	?
*>i10.44.45.1/32	120.100.6.1	2	100	0	?
*>i10.44.46.1/32	120.100.6.1	2	100	0	?
*> 130.50.50.0/30	0.0.0.0		0		32768 ?
*>i130.100.100.0/30	120.100.6.1	0	100	0	?

```
R6# show ip bgp vpng4 vrf BLUE
BGP table version is 35, local router ID is 120.100.6.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*>i10.1.1.0/24	120.100.1.1	50	100	0	?
*>i10.1.2.0/24	120.100.1.1	50	100	0	?
*>i10.1.3.0/24	120.100.1.1	50	100	0	?
*> 10.2.2.0/24	10.10.20.2	50		32768	?
*> 10.2.3.0/24	10.10.20.2	50		32768	?
*> 10.2.4.0/24	10.10.20.2	50		32768	?
*>i10.10.10.0/30	120.100.1.1	0	100	0	?
*> 10.10.20.0/30	0.0.0.0		0		32768 ?
*> 10.44.44.1/32	130.100.100.2	2		32768	?

! Notice the 10.44.44.0 route is actually listed as a host route, change the loopback interface on Sw4 to a point-to-point for OSPF to advertise it correctly

```
SW4(config)# interface lo0
SW4(config-if)# ip ospf network point-to-point
```

```

R6# show ip bgp vpng4 vrf BLUE | include 10.44.44.0
*> 10.44.44.0/24      130.100.100.2          2      32768 ?
Switch1# show ip route | include 10.44.44.0
D EX 10.44.44.0/24 [170/281856] via 10.10.10.1, 00:00:51, Vlan10
Switch1#
SW4# show ip route | include 10.1.1.0
O E1 10.1.1.0/24 [110/51] via 130.100.100.1, 00:02:45, Vlan100

! Now test with an extended ping to ensure that the loopback interface is used as the
source

SW1# ping
Protocol [ip]:
Target IP address: 10.44.44.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 10.1.1.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.44.44.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/9/12 ms

R1# show mpls forwarding-table vrf BLUE
Local  Outgoing      Prefix           Bytes tag  Outgoing      Next Hop
tag    tag or VC    or Tunnel Id    switched    interface
34     Untagged     10.1.3.0/24[V]   0          Gi0/0.10    10.10.10.2
35     Untagged     10.1.2.0/24[V]   0          Gi0/0.10    10.10.10.2
36     Aggregate    10.10.10.0/30[V] 0          Gi0/0.10    10.10.10.2
37     Untagged     10.1.1.0/24[V]   590        Gi0/0.10    10.10.10.2

R1# show mpls forwarding-table vrf RED
Local  Outgoing      Prefix           Bytes tag  Outgoing      Next Hop
tag    tag or VC    or Tunnel Id    switched    interface
38     Aggregate    130.50.50.0/30[V] 0          Gi0/0.50    130.50.50.2
39     Untagged     10.33.34.1/32[V]  0          Gi0/0.50    130.50.50.2
40     Untagged     10.33.35.1/32[V]  0          Gi0/0.50    130.50.50.2
41     Untagged     10.33.33.1/32[V]  0          Gi0/0.50    130.50.50.2

! Note the Routes are not leaked within the MPLS forwarding-table

R6# show mpls forwarding-table vrf BLUE
Local  Outgoing      Prefix           Bytes tag  Outgoing      Next Hop
tag    tag or VC    or Tunnel Id    switched    interface
34     Untagged     10.2.2.0/24[V]   0          Gi0/1.20    10.10.20.2
35     Untagged     10.2.3.0/24[V]   0          Gi0/1.20    10.10.20.2
36     Untagged     10.2.4.0/24[V]   0          Gi0/1.20    10.10.20.2
37     Aggregate    10.10.20.0/30[V] 0

R6# show mpls forwarding-table vrf RED

```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes switched	Outgoing interface	Next Hop
38	Aggregate	130.100.100.0/30 [V]	0		
39	Untagged	10.44.46.1/32 [V]	0	Gi0/1.100	130.100.100.2
40	Untagged	10.44.45.1/32 [V]	0	Gi0/1.100	130.100.100.2
42	Untagged	10.44.44.0/24 [V]	1534	Gi0/1.100	130.100.100.2

! Note the Routes are not leaked within the MPLS forwarding-table

- Configure your PE routers R1 and R6 to ensure that the MPLS P routers are not listed as intermediate hops when a trace route is performed on your CE devices. (2 points)

By default, the MPLS network will be shown when a traceroute is performed. This can be changed, with the **no mpls ip propagate-ttl** global command within your PE routers, so that only PE routers are shown as next hops. [Example 3-20](#) shows the default behavior and modified behavior after configuration from a **traceroute** command issued on CE device SW1. If you have configured this correctly, as shown in [Example 3-20](#), you have scored 2 points.

Example 3-20 MPLS Traceroute Configuration and Testing

[Click here to view code image](#)

```
SW1# traceroute 10.2.2.1

Type escape sequence to abort.
Tracing the route to 10.2.2.1

1 10.10.10.1 0 msec 0 msec 0 msec
2 120.100.123.2 12 msec 12 msec 16 msec
3 120.100.25.5 8 msec 12 msec 8 msec
4 10.10.20.1 8 msec 8 msec 8 msec
5 10.10.20.2 8 msec * 4 msec
```

```
R1(config)# no mpls ip propagate-ttl
```

```
R6(config)# no mpls ip propagate-ttl
```

```
SW1# traceroute 10.2.2.1
```

```
Type escape sequence to abort.
Tracing the route to 10.2.2.1
```

```
1 10.10.10.1 4 msec 0 msec 0 msec
2 10.10.20.1 12 msec 8 msec 12 msec
3 10.10.20.2 4 msec * 4 msec
```

Section 7: VPLS Simulation (10 Points)

- Switches 3 and 4 will have been configured to belong to the subnet of 1.1.1.0/24 in a previous question. Create an L2TPv3 Xconnect attachment circuit on your PE routers R1 and R6 for your CE devices (SW3 Fast Ethernet 0/19 1.1.1.1/24 and SW4 Fast Ethernet 0/19 1.1.1.2/24) to communicate using a Layer 2 tunneling solution (use Version 3) across your Layer 3 network. You should use existing loopback interfaces on your PE routers for peering over your MPLS network. Be aware that the SW3 resides in VLAN 200 and that SW4 resides in VLAN 400 in respective PE router subinterfaces. (10 points)

This question simulates VPLS and requires that L2TPv3 (Layer 2 Tunneling Protocol Version 3) is configured between your PE routers connecting the two subinterfaces that connect to SW3 and SW4 interfaces via SW1 and SW4 (VLAN 200 and VLAN 400, respectively). Strictly speaking, L2TPv3 is not covered in the current blueprint, but the simple solution is included here to create a switching issue that will enable you to hone your troubleshooting skills in this area and apply a relevant solution based on your findings. SW3 and SW4 will use a pseudowire to communicate over the IP network and logically will connect in the same Layer 2 domain. The pseudowire class PW-CLASS configures the encapsulation to L2TPv3 and sets the loopback interfaces of the PE routers to be used for peering. The **xconnect** subinterface command binds the local PE interface to the remote PE loopback with a VC ID (virtual channel ID), which in the example matches the subinterface number of the specific PE router. (You could have used any ID here.) Note that Cisco Express Forwarding (CEF) must be enabled for the L2TPv3 feature to function correctly. [Example 3-21](#) details the required PE configuration on routers R1 and R2.

Example 3-21 PE L2TPv3 Configuration

[Click here to view code image](#)

```
R1(config)# pseudowire-class PW-CLASS
R1(config-pw-class)# encapsulation l2tpv3
R1(config-pw-class)# protocol l2tpv3
R1(config-pw-class)# ip local interface Loopback0
R1(config-pw-class)# interface GigabitEthernet0/0.200
R1(config-subif)# xconnect 120.100.6.1 200 pw-class PW-CLASS

R6(config)# pseudowire-class PW-CLASS
R6(config-pw-class)# encapsulation l2tpv3
R6(config-pw-class)# protocol l2tpv3
R6(config-pw-class)# ip local interface Loopback0
R6(config-pw-class)# interface GigabitEthernet0/1.400
R6(config-subif)# xconnect 120.100.1.1 200 pw-class PW-CLASS
```

[Example 3-22](#) shows the successful L2TPv3 session established between PE R1 to PE R6, yet the ping test from SW3 to 1.1.1.2 fails. As the session is up, you can safely assume that there is a connectivity type issue between either SW3 and PE R1 or SW4 and PE R6, or possibly between both connections. The question does bring your attention to the fact that both CE devices reside in different VLANs, so this should give you a starting point in your investigation. When logging is enabled on SW1 and SW2 (these CE devices bring SW3 and SW4 Fast Ethernet 0/19 interfaces into VLAN 200 and VLAN 400, respectively), you can see spanning-tree inconsistencies exist between VLAN 200

being “bridged” to VLAN 400 via your L2TPv3 solution. Closer inspection reveals that spanning tree has actually blocked ports on SW1 and SW2 from PE routers R1 and R6, respectively, even though you have previously allowed the local VLAN 200 and 400 through the trunk on PE routers R1 and R6, respectively. The problem is actually resolved by enabling BPDU filtering on SW1 with the **spanning-tree bpdufilter enable** command on the trunk interface toward the PE router R1. Enabling BPDU filtering on an interface is equivalent to disabling the spanning tree on an interface; it is possible to create bridging loops if this command is not correctly used. If you have configured this correctly, per [Examples 3-22](#) and [3-23](#), you have scored 10 points.

Example 3-22 PE and CE L2TPv3 Verification Testing and Configuration

[Click here to view code image](#)

```
R1# show l2tp session

L2TP Session Information Total tunnels 1 sessions 1

LocID RemID Remote Name      State   Remote Address  Port   Sessions L2TP Class/
                                         VPDN Group
51446  36190    R6           est     120.100.6.1      0       1

LocID      RemID      TunID      Username, Intf/
                                         Vcid, Circuit
51003      9619       51446     200, Gi0/0.200:200  est     00:24:40  1

R6# show l2t session

L2TP Tunnel and Session Information Total tunnels 1 sessions 1

LocID RemID Remote Name      State   Remote Address  Port   Sessions L2TP Class/
                                         VPDN Group
36190  51446    R1           est     120.100.1.1      0       1

LocID      RemID      TunID      Username, Intf/
                                         Vcid, Circuit
9619       51003     36190     200, Gi0/1.400:400  est     00:25:26  1

SW3# ping 1.1.1.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:
.....
!Make sure you are logging on your CE devices

SW1(config)# logging console

SW1#
03:22:19: %SPANTREE-2-RECV_PVID_ERR: Received BPDU with inconsistent peer vlan id
400 on fastethernet0/1 VLAN200.
03:22:19: %SPANTREE-2-BLOCK_PVID_LOCAL: Blocking fastethernet0/1 on VLAN0200.
Inconsistent local vlan.

SW1# show spanning-tree blockedports

Name          Blocked Interfaces List
```

VLAN0200

Fa0/1

Number of blocked ports (segments) in the system : 1

SW2#03:22:21: %SPANTREE-2-RECV_PVID_ERR: Received BPDU with inconsistent peer vlan id 200 on fastethernet0/6 VLAN400.
03:22:21: %SPANTREE-2-BLOCK_PVID_PEER: Blocking fastethernet0/6 on VLAN0200.
Inconsistent peer vlan.

SW2# **show spanning-tree blockedports**

Name	Blocked Interfaces List
VLAN0200	Fa0/6
VLAN0400	Fa0/6

Number of blocked ports (segments) in the system : 2

SW3# **ping 1.1.1.2**

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:
.....
Success rate is 0 percent (0/5)

SW1# **show spanning-tree blockedports**

Name	Blocked Interfaces List
VLAN0200	Fa0/1

Number of blocked ports (segments) in the system : 1

SW1(config)# **int fast 0/1**
SW1(config-if)# **spanning-tree bpdufilter enable**
SW1(config-if)#03:33:57: %SPANTREE-2-UNBLOCK_CONSIST_PORT: Unblocking fastethernet0/1 on VLAN0200. Port consistency restored.

SW3# **ping 1.1.1.2**

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:
.!!!!
Success rate is 80 percent (4/5), round-trip min/avg/max = 8/12/17 ms

Section 8: Multicast (10 Points)

- Configure your MPLS network for multicast support of the RED VRF using PIM sparse mode. PE routers R1 and R6 should be configured to tunnel multicast traffic using an MDT address of 232.0.0.11 from CE device Switch 3 VLAN 50 to CE device SW4 VLAN 100 over the RED VRF. Switch 4 should be configured to reply to an ICMP ping on its VLAN 100 interface directed to 226.2.2.2 from Switch 3 VLAN 50. It can be assumed that the mVRF bandwidth requirement is low; configure MDT appropriately. Ensure that PE router R6's associated VLAN 100 IP address is used as the rendezvous point (RP) for the RED VRF multicast traffic. (10 points)

Multicast support for MPLS VPNs is provided by configuring multicast routing within the core network. As directed, PIM sparse mode is required in your solution and should be enabled on all P router MPLS interfaces and P-facing PE router MPLS interfaces. PIM sparse mode is also configured on the CE interfaces on VLAN 50 and VLAN 100 on Switches 3 and 4, respectively, and corresponding PE terminating interfaces on the PE routers R1 and R6. PIM sparse mode is finally configured on the loopback interfaces of the PE routers R1 and R6 because Multicast Distribution Tree (MDT) will tunnel between these interfaces.

Don't forget that multicast routing is enabled on the CE switches with the command **ip multicast-routing distributed** and on the routers with **ip multicast-routing**. The **mdt default group-address** is configured to 232.0.0.11 on PE routers R1 and R6 within the RED VRF. Source Specific Multicast (SSM) is enabled on all MPLS routers with the command **ip pim ssm default** to allow transport of multicast information between all P and PE routers.

The question states that the mVRF (multicast VRF) bandwidth requirement is low, which simply means that a Data MDT is not required in this solution. (These are used for high-bandwidth sources and limit the traffic received to the routers' part of the multicast tree.) You should also realize that a Data MDT is not required because there was no mention of threshold values or access-lists within the question, which are required for Data MDT configurations.

The address of 130.100.100.1 (R6 VRF RED) is used as the RP for the mVRF, and this is configured on both CE (Switch 3 and Switch 4) devices and both PE routers (R1 and R6) within the RED VRF.

CE device Switch 4 is finally configured with **ip igmp join-group 226.2.2.2** under its VLAN 100 interface for it to reply to a multicast ping from CE device Switch 3 over the MPLS VPN.

The question is comprehensive as to the number of items that require configuration, and it would be an easy mistake to miss tasks such as enabling PIM on the PE loopback interfaces, where you might not immediately assume that it is required. As with all questions, testing is key. [Example 3-23](#) details the required configuration for the solution.

Example 3-23 Multicast Configuration

[Click here to view code image](#)

```
! Initial Multicast Setup for the MPLS Core Routers

R1(config)# ip multicast-routing
R1(config-vrf)# interface Loopback0
R1(config-if)# ip pim sparse-mode
R1(config-if)# interface GigabitEthernet0/1
R1(config-if)# ip pim sparse-mode

R2(config)# ip multicast-routing
R2(config)# interface fastethernet0/0
R2(config-if)# ip pim sparse-mode
R2(config-if)# interface fastethernet0/1
R2(config-if)# ip pim sparse-mode

R3(config)# ip multicast-routing
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-mode
R3(config-if)# interface GigabitEthernet0/1
```

```

R3(config-if)# ip pim sparse-mode

R4(config)# ip multicast-routing
R4(config)# interface GigabitEthernet0/0
R4(config-if)# ip pim sparse-mode
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ip pim sparse-mode

R5(config)# ip multicast-routing
R5(config)# interface GigabitEthernet0/0
R5(config-if)# ip pim sparse-mode
R5(config-if)# interface GigabitEthernet0/1
R5(config-if)# ip pim sparse-mode

R6(config)# ip multicast-routing
R6(config)# interface Loopback0
R6(config-if)# ip pim sparse-mode
R6(config)# interface GigabitEthernet0/0
R6(config-if)# ip pim sparse-mode

! PE Specific mVRF and MDT Configuration

R1(config)# ip multicast-routing vrf RED
R1(config)# ip vrf RED
R1(config-vrf)# mdt default 232.0.0.11
R1(config-vrf)# interface GigabitEthernet0/0.50
R1(config-subif)# ip pim sparse-mode
R1(config-subif)# exit
R1(config)# ip pim vrf RED rp-address 130.100.100.1
R1(config)# ip pim ssm default

R6(config)# ip vrf RED
R6(config-vrf)# mdt default 232.0.0.11
R6(config-vrf)# interface GigabitEthernet0/1.100
R6(config-subif)# ip pim sparse-mode
R6(config-subif)# exit
R6(config)# ip pim vrf RED rp-address 130.100.100.1
R6(config)# ip pim ssm default

! CE Specific Configuration

SW3(config)# ip multicast-routing distributed
SW3(config)# int vlan 50
SW3(config-if)# ip pim sparse-mode
SW3(config-if)# exit
SW3(config)# ip pim rp-address 130.100.100.1

SW4(config)# ip multicast-routing distributed
SW4(config)# interface vlan 100
SW4(config-if)# ip pim sparse-mode
SW4(config-if)# ip igmp join-group 226.2.2.2
SW4(config-if)# exit
SW4(config)# ip pim rp-address 130.100.100.1

```

[Example 3-24](#) details the testing for the solution; the MDT tunnel is detailed and shown as an interface used for PIM adjacency between the PE routers. If you have configured your solution per [Example 3-24](#) and can successfully ping between Switch 3 and Switch 4, you have scored 10 points.

Example 3-24 Multicast Testing

[Click here to view code image](#)

```
R6# show ip pim vrf RED neigh
PIM Neighbor Table
Mode: B - Bidir Capable, DR - Designated Router, N - Default DR Priority,
      S - State Refresh Capable
Neighbor           Interface          Uptime/Expires   Ver   DR
Address
Prio/Mode
130.100.100.2    GigabitEthernet0/1.100 00:02:08/00:01:34 v2    1 / DR S
120.100.1.1       Tunnel1                00:00:05/00:01:39 v2    1 / S

R1# ping vrf RED 226.2.2.2
Type escape sequence to abort.
Sending 1, 100-byte ICMP Echos to 226.2.2.2, timeout is 2 seconds:
Reply to request 0 from 130.100.100.2, 12 ms

SW3# ping 226.2.2.2
Type escape sequence to abort.
Sending 1, 100-byte ICMP Echos to 226.2.2.2, timeout is 2 seconds:
Reply to request 0 from 130.100.100.2, 9 ms
SW3# show ip pim rp
Group: 226.2.2.2, RP: 130.100.100.1, v2, uptime 00:00:37, expires never
Group: 224.0.1.40, RP: 130.100.100.1, v2, uptime 01:01:24, expires never

R1# show ip pim mdt bgp
Peer (Route Distinguisher + IPv4)          Next Hop
MDT group 232.0.0.11
2:65001:200:120.100.6.1                  120.100.6.1

R6# show ip pim mdt bgp
Peer (Route Distinguisher + IPv4)          Next Hop
MDT group 232.0.0.11
2:65001:200:120.100.1.1                  120.100.1.1
```

Section 9: IPv6 (6 Points)

- Configure the following IPv6 address on the PE routers R1 and R6, and implement IPv6 over MPLS between the six PE routers to advertise the prefixes between six PEs. Make sure that your loopback IPv6 addresses are used to source any locally generated IPv6 traffic. (6 points)

R1 Lo0 2010:C15:C0:1::1/64

R1 Gi0/0.10 2010:C15:C0:11::1/64

R6 Lo0 2010:C15:C0:6::1/64

R6 Gi0/1.20 2010:C15:C0:62::1/64

In this relatively straightforward IPv6 question, you must deal with no IPv6 redistribution or complex issues. The question directs you to configure IPv6 onto your VRF BLUE interfaces of the PE routers. You would usually extend this IPv6 domain into your CE devices. IPv6 over MPLS backbones

enables isolated IPv6 domains to communicate with each other over an MPLS IPv4 core network. To ensure that the loopback IPv6 addresses of the PE routers are used to source locally generated IPv6 traffic, the PE routers are configured with **mpls ipv6 source-interface Loopback0**. MP-BGP is used to advertise the IPv6 prefixes between PE routers, and the configuration is nearly identical to that of IPv4. Aggregate label binding and advertisement is enabled for IPv6 prefixes using the **neighbor send-label** command. Connected IPV6 routes are redistributed using BGP with the **network** command under the IPv6 address family, and IPv6 routing and IPv6 CEF must be enabled on your PE routers. If you have configured your routers correctly, per [Example 3-25](#), you have scored 6 points.

Example 3-25 PE IPv6 Configuration and Verification

[Click here to view code image](#)

```

R1(config)# ipv6 unicast-routing
R1(config)# ipv6 cef
R1(config)# mpls ipv6 source-interface Loopback0
R1(config)# interface loopback0
R1(config-if)# ipv6 add 2010:C15:C0:1::1/64
R1(config-if)# interface GigabitEthernet0/0.10
R1(config-subif)# ipv6 address 2010:C15:C0:11::1/64
R1(config-subif)# router bgp 65001
R1(config-router)# no bgp default ipv4-unicast
R1(config-router)# address-family ipv6
R1(config-router-af)# neighbor 120.100.6.1 activate
R1(config-router-af)# neighbor 120.100.6.1 send-label
R1(config-router-af)# network 2010:C15:C0:11::0/64
R1(config-router-af)# network 2010:C15:C0:1::/64
R1(config-router-af)# exit-address-family

R6(config)# ipv6 unicast-routing
R6(config)# ipv6 cef
R6(config)# mpls ipv6 source-interface Loopback0
R6(config)# interface loopback0
R6(config-if)# ipv6 add 2010:C15:C0:6::1/64
R6(config-if)# interface GigabitEthernet0/1.20
R6(config-subif)# ipv6 address 2010:C15:C0:62::1/64
R6(config-subif)# router bgp 65001
R6(config-router)# no bgp default ipv4-unicast
R6(config-router)# address-family ipv6
R6(config-router-af)# neighbor 120.100.1.1 activate
R6(config-router-af)# neighbor 120.100.1.1 send-label
R6(config-router-af)# network 2010:C15:C0:62::/64
R6(config-router-af)# network 2010:C15:C0:6::/64
R6(config-router-af)# exit-address-family

R1# show ip bgp ipv6 unicast
BGP table version is 5, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*> 2010:C15:C0:1::/64
                           ::                               0        32768  i
*>i2010:C15:C0:6::/64

```

	0	100	0	i
*> 2010:C15:C0:11::/64	0	32768	i	
*>i2010:C15:C0:62::/64	0	100	0	i

R6# **show ip bgp ipv6 unicast**

BGP table version is 5, local router ID is 120.100.6.1
 Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
 r RIB-failure, S Stale
 Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
*>i2010:C15:C0:1::/64	::FFFF:120.100.1.1	0	100	0	i
*> 2010:C15:C0:6::/64	::	0	32768	i	
*>i2010:C15:C0:11::/64	::FFFF:120.100.1.1	0	100	0	i
*> 2010:C15:C0:62::/64	::	0	32768	i	

R1# **ping ipv6 2010:C15:C0:62::1**

Type escape sequence to abort.
 Sending 5, 100-byte ICMP Echos to 2010:C15:C0:62::1, timeout is 2 seconds:
 !!!!!
 Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms
 R1# **ping ipv6 2010:C15:C0:6::1**

Type escape sequence to abort.
 Sending 5, 100-byte ICMP Echos to 2010:C15:C0:6::1, timeout is 2 seconds:
 !!!!!
 Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms

R6# **ping ipv6 2010:C15:C0:11::1**

Type escape sequence to abort.
 Sending 5, 100-byte ICMP Echos to 2010:C15:C0:11::1, timeout is 2 seconds:
 !!!!!
 Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms
 R6# **ping ipv6 2010:C15:C0:1::1**

Type escape sequence to abort.
 Sending 5, 100-byte ICMP Echos to 2010:C15:C0:1::1, timeout is 2 seconds:
 !!!!!
 Success rate is 100 percent (5/5), round-trip min/avg/max = 8/9/12 ms

R1# **show ipv6 route**

IPv6 Routing Table - 8 entries
 Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
 U - Per-user Static route
 I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
 O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
 ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
 D - EIGRP, EX - EIGRP external
 C 2010:C15:C0:1::/64 [0/0]

```

    via ::, Loopback0
L 2010:C15:C0:1::1/128 [0/0]
    via ::, Loopback0
B 2010:C15:C0:6::/64 [200/0]
    via ::FFFF:120.100.6.1, IPv6-mpls
C 2010:C15:C0:11::/64 [0/0]
    via ::, GigabitEthernet0/0.10
L 2010:C15:C0:11::1/128 [0/0]
    via ::, GigabitEthernet0/0.10
B 2010:C15:C0:62::/64 [200/0]
    via ::FFFF:120.100.6.1, IPv6-mpls
L FE80::/10 [0/0]
    via ::, Null0
L FF00::/8 [0/0]
    via ::, Null0

```

R6# show ipv6 route

```

IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
        U - Per-user Static route
        I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
        O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
        ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
        D - EIGRP, EX - EIGRP external
B 2010:C15:C0:1::/64 [200/0]
    via ::FFFF:120.100.1.1, IPv6-mpls
C 2010:C15:C0:6::/64 [0/0]
    via ::, Loopback0
L 2010:C15:C0:6::1/128 [0/0]
    via ::, Loopback0
B 2010:C15:C0:11::/64 [200/0]
    via ::FFFF:120.100.1.1, IPv6-mpls
C 2010:C15:C0:62::/64 [0/0]
    via ::, GigabitEthernet0/1.10
L 2010:C15:C0:62::1/128 [0/0]
    via ::, GigabitEthernet0/1.20
L FE80::/10 [0/0]
    via ::, Null0
L FF00::/8 [0/0]
    via ::, Null0

```

Section 10: QoS (7 Points)

- Create the following QoS profile on your PE router R1 for traffic egressing to your CE device connected to the BLUE VRF. The total bandwidth between the PE to CE should be shaped to 1Mbps. Ensure that voice traffic is assigned to an LLQ. Use an appropriate method of prioritizing DSCP traffic so that AF31 packets are statistically dropped more frequently than AF32 during congestion and reduce the effects of TCP global synchronization within your Mission-Critical class and solely reduce the effect of TCP global synchronization within the Default class. (4 points)

Class	DSCP Value	% of Bandwidth Assigned
Voice	EF, CS5	35
Mission-Critical	CS6, AF31, AF32, CS3	40
Default	Any	25

This is a three-class PE-to-CE QoS question that requires assigning traffic to queues based on DSCP values into the listed classes and assignment of bandwidth on a per-class basis. Voice traffic is assigned into the LLQ by configuration of a priority queue with the command **priority percent 35**. DSCP prioritization is achieved in the Mission-Critical class by enabling WRED with the **random-detect dscp-based** command, whereby lower-priority DSCP traffic will be dropped more aggressively than higher priority under congestion, thus reducing the effect of global synchronization. A similar non-DSCP-based effect is achieved within the Default class by use of the **random-detect** command. HQF Multiple Policy Support is required for the question with a parent policy shaping the output of the PE to the CE at 1Mbps. The child policy map is called from within the parent policy to provide the QoS for Voice, Mission-Critical, and Default traffic. The parent policy map is applied outbound on the PE interface connecting to the BLUE VRF CE device. [Example 3-26](#) details the required configuration on PE router R1. If you have configured this correctly, you have scored 4 points.

Example 3-26 PE to CE QoS Configuration

[Click here to view code image](#)

```
R1(config)# class-map match-any VOICE
R1(config-cmap)# match ip dscp ef
R1(config-cmap)# match ip dscp cs5
R1(config-cmap)# class-map match-any MISSION-CRITICAL
R1(config-cmap)# match ip dscp cs6
R1(config-cmap)# match ip dscp af31
R1(config-cmap)# match ip dscp af32
R1(config-cmap)# match ip dscp cs3
R1(config-cmap)# policy-map PE-CE-CHILD
R1(config-pmap)# class VOICE
R1(config-pmap-c)# priority percent 35
R1(config-pmap-c)# class MISSION-CRITICAL
R1(config-pmap-c)# bandwidth percent 40
R1(config-pmap-c)# random-detect dscp-based
R1(config-pmap-c)# class class-default
R1(config-pmap-c)# bandwidth percent 25
R1(config-pmap-c)# random-detect
R1(config-pmap-c)# exit
R1(config-cmap)# policy-map PE-CE-PARENT
R1(config-pmap-c)# class class-default
R1(config-pmap-c)# shape average 1000000
R1(config-pmap-c)# service-policy PE-CE-CHILD
R1(config-pmap-c)# exit
R1(config-pmap)# exit

R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# service-policy output PE-CE-PARENT
```

- Create the following QoS profile on your PE router R1 for traffic ingressing from your CE device connected to the BLUE VRF into the MPLS network. The total aggregate speed from the CE to PE should be restricted to 1 Mbps.

Class	CIR (bps)
Voice	350,000
Mission-Critical	400,000
Default	250,000

Traffic in the Voice class within the detailed CIR should have the MPLS EXP set to 5 and above discarded. Traffic in the Mission-Critical class within the detailed CIR should have the MPLS EXP set to 3 and above set to 7. Traffic in the Default class within the detailed CIR should have the MPLS EXP set to 0 and above set to 4. (3 points)

This DiffServ tunneling question requires that the classes you have configured in the previous question be policed to an aggregate of 1 Mbps and have their MPLS EXP values adjusted. The policy map is applied to the input interface of the PE router, which connects to the BLUE VRF CE device and affects the traffic as it flows through the MPLS network. [Example 3-27](#) details the required configuration on PE router R1. If you have configured this correctly, you have scored 3 points.

Example 3-27 CE to PE QoS Configuration

[Click here to view code image](#)

```
R1(config)# policy-map CE-PE-SHAPE
R1(config-pmap)# class VOICE
R1(config-pmap-c)# police cir 350000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 5
R1(config-pmap-c-police)# exceed-action drop
R1(config-pmap-c-police)# class MISSION-CRITICAL
R1(config-pmap-c)# police cir 400000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 3
R1(config-pmap-c-police)# exceed-action set-mpls-exp-topmost-transmit 7
R1(config-pmap-c-police)# class class-default
R1(config-pmap-c)# police cir 250000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 0
R1(config-pmap-c-police)# exceed-action set-mpls-exp-topmost-transmit 4
R1(config-pmap-c-police)# interface GigabitEthernet0/0.10
R1(config-subif)# service-policy input CE-PE-SHAPE
```

Section 11: Security (15 Points)

- Create three new loopback IP addresses of loopback1 on R4, R5, and R6; use IP addresses of 4.4.4.4/24, 5.5.5.5/24, and 6.6.6.6/24, respectively. Use EIGRP with a named virtual instance of VPN and autonomous system of 1 to advertise the loopback networks between routers over a common GRE tunnel network of 100.100.100.X/24 (X = router number) sourced from each router's common Ethernet interface, using IPsec to encrypt all traffic between the loopback networks using a preshared ISAKMP key of CCIE. Use an IPsec transform set of esp-des esp-md5-hmac on each router. R6 is to be a hub router, with R4 and R5 being effectively spoke routers in your solution. You are not permitted to enable EIGRP on your Ethernet interfaces between routers. Spoke routers must communicate with each other directly using dynamic IPsec connections with the aid of NHRP at the hub, whereas hub-to-spoke IPsec connections should be permanent. The hub router should provide all necessary direct next-hop information to the spoke routers when they are required to communicate between themselves. NHRP should be authenticated with a password of SECRET. Use an MTU of 1416 for your secure traffic, an NHRP timeout of 100 seconds for spoke replies, and a delay of 2 milliseconds on the tunnel network. Test your solution by extended pings sourced from the configured loopback interfaces. (10 points)

This is a classic Dynamic Multipoint VPN (DMVPN) question in which a hub-and-spoke design is used with Next Hop Resolution Protocol (NHRP) for the spoke routers to communicate with each other. You have numerous tasks to perform, so this could be the kind of question that is best saved until later and tackled if you have time. The question dictates that you configure a tunnel network 100.100.100.0/24 in which to advertise each router's new loopback network over GRE and EIGRP sourced from the common Ethernet interfaces, which is uncomplicated; the complexity begins when you enable IPsec and NHRP. The **crypto isakmp policy** command configures the preshared key to **CCIE** and sets the transform set with the required parameters of **esp-des esp-md5-hmac**, which are applied to the tunnel interface by the use of the **tunnel protection ipsec profile IPSEC** command. The MTU is fixed at 1416 as directed within the question on the tunnel interfaces to allow for overhead of the VPN connection.

A delay of 2000 is configured on each tunnel interface as directed in the question, which is 2 milliseconds, so be aware of the unit values, which are microseconds. The tunnel source of each router is the common Ethernet network 120.100.45. Because the spoke routers will terminate their connection to the hub on the same interface, the tunnel mode must be set to **tunnel mode gre multipoint**. NHRP is enabled on the tunnel interface of each router with an identical network ID to match the broadcast domain for all three routers, and the authentication password is set to SECRET as directed within the question. The command **ip nhrp map multicast dynamic** permits the registration of the multicast address for EIGRP during boot or during initiation of spoke-to-hub sessions. The **ip nhrp holdtime 100** command sets the NHRP time for a spoke to keep the NHRP reply to 100 seconds and is configured on the hub-and-spoke routers.

The required configuration for the loopback and tunnel interfaces and the DMVPN is detailed in [Example 3-28](#).

Example 3-28 DMVPN Configuration

[Click here to view code image](#)

```
R4(config)# interface loopback1
R4(config-if)# ip add 4.4.4.4 255.255.255.0
R4(config-if)# router eigrp VPN
R4(config-router)# address-family ipv4 autonomous-system 1
R4(config-router-af)# network 100.100.100.0 0.0.0.255
R4(config-router-af)# network 4.4.4.0 0.0.0.255

R5(config)# interface loopback1
R5(config-if)# ip address 5.5.5.5 255.255.255.0
R5(config-if)# router eigrp VPN
R5(config-router)# address-family ipv4 autonomous-system 1
R5(config-router-af)# network 100.100.100.0 0.0.0.255
R5(config-router-af)# network 5.5.5.0 0.0.0.255

R6(config)# interface loopback1
R6(config-if)# ip address 6.6.6.6 255.255.255.0
R6(config-if)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
R6(config-router-af)# network 100.100.100.0 0.0.0.255
R6(config-router-af)# network 6.6.6.0 0.0.0.255

R6(config)# crypto isakmp policy 1
R6(config-isakmp)# authentication pre-share
R6(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R6(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R6(crypto-trans)# crypto ipsec profile IPSEC
R6(ipsec-profile)# set transform-set DMVPN
R6(ipsec-profile)# interface Tunnel1
R6(config-if)# ip address 100.100.100.6 255.255.255.0
R6(config-if)# ip mtu 1416
R6(config-if)# ip nhrp authentication SECRET
R6(config-if)# ip nhrp map multicast dynamic
R6(config-if)# ip nhrp network-id 10
R6(config-if)# ip nhrp holdtime 100
R6(config-if)# delay 2000
R6(config-if)# tunnel source gig 0/0
R6(config-if)# tunnel mode gre multipoint
R6(config-if)# tunnel key 1
R6(config-if)# tunnel protection ipsec profile IPSEC

R4(config)# crypto isakmp policy 1
R4(config-isakmp)# authentication pre-share
R4(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R4(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R4(crypto-trans)# crypto ipsec profile IPSEC
R4(ipsec-profile)# set transform-set DMVPN
R4(ipsec-profile)# interface Tunnel0
R4(config-if)# ip address 100.100.100.4 255.255.255.0
R4(config-if)# ip mtu 1416
R4(config-if)# ip nhrp authentication SECRET
R4(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R4(config-if)# ip nhrp map multicast 120.100.45.6
R4(config-if)# ip nhrp network-id 10
R4(config-if)# ip nhrp holdtime 100
R4(config-if)# ip nhrp nhs 100.100.100.6
R4(config-if)# delay 2000
R4(config-if)# tunnel source gig 0/0
```

```

R4(config-if)# tunnel mode gre multipoint
R4(config-if)# tunnel key 1
R4(config-if)# tunnel protection ipsec profile IPSEC

R5(config)# crypto isakmp policy 1
R5(config-isakmp)# authentication pre-share
R5(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R5(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R5(cfg-crypto-trans)# crypto ipsec profile IPSEC
R5(ipsec-profile)# set transform-set DMVPN
R5(ipsec-profile)# interface Tunnel0
R5(config-if)# ip address 100.100.100.5 255.255.255.0
R5(config-if)# ip mtu 1416
R5(config-if)# ip nhrp authentication SECRET
R5(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R5(config-if)# ip nhrp map multicast 120.100.45.6
R5(config-if)# ip nhrp network-id 10
R5(config-if)# ip nhrp holdtime 100
R5(config-if)# ip nhrp nhs 100.100.100.6
R5(config-if)# delay 2000
R5(config-if)# tunnel source gig 0/0
R5(config-if)# tunnel mode gre multipoint
R5(config-if)# tunnel key 1
R5(config-if)# tunnel protection ipsec profile IPSEC

```

[Example 3-29](#) details the EIGRP routes received on all routers. As you can see, the hub router shows both spoke networks, yet each spoke router discovers only the hub network; this is a classic split-horizon issue. The hub router R6 must be configured to disable the split-horizon behavior to ensure that the spoke routers receive each other's routes. However, the question dictates that spoke routers should be able to communicate "directly." As shown in [Example 3-29](#), the next hop for spoke networks show as the hub router 100.100.100.6 for each spoke network. The command **no next-hop-self** on the hub router R6 ensures that the spoke routers are used as next hops when spoke-to-spoke communication is required, and this will enable the dynamic IPsec peering between spokes as directed in the question.

Example 3-29 DMVPN Spoke-to-Spoke Routing

[Click here to view code image](#)

```

R4# show ip route eigrp
    6.0.0.0/24 is subnetted, 1 subnets
D        6.6.6.0 [90/285084416] via 100.100.100.6, 00:02:42, Tunnel0

R5# show ip route eigrp
    6.0.0.0/24 is subnetted, 1 subnets
D        6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:50, Tunnel0

R6# show ip route eigrp
    4.0.0.0/24 is subnetted, 1 subnets
D        4.4.4.0 [90/285084416] via 100.100.100.4, 00:03:06, Tunnel0
    5.0.0.0/24 is subnetted, 1 subnets
D        5.5.5.0 [90/285084416] via 100.100.100.5, 00:01:02, Tunnel0

```

!R6 has both spoke routes yet each spoke (R4 and R5) only have the hub network

```
route, !a classic split horizon issue.
```

```
R6(config)# router eigrp VPN  
R6(config-router)# address-family ipv4 autonomous-system 1
```

```
R6(config-router-af)# af-interface Tunnel0  
R6(config-router-af-interface)# no split-horizon
```

```
R4# show ip route eigrp  
 5.0.0.0/24 is subnetted, 1 subnets  
D      5.5.5.0 [90/285596416] via 100.100.100.6, 00:00:22, Tunnel0  
 6.0.0.0/24 is subnetted, 1 subnets  
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:04:14, Tunnel0
```

```
R5# show ip route eigrp  
 4.0.0.0/24 is subnetted, 1 subnets  
D      4.4.4.0 [90/285596416] via 100.100.100.6, 00:00:33, Tunnel0  
 6.0.0.0/24 is subnetted, 1 subnets  
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:02:20, Tunnel0  
R5#
```

```
! The next-hop for spoke to spoke routes shows as the hub router (100.100.100.6) yet  
! the question states traffic must flow directly between spokes so the next-hop must  
! be modified
```

```
R6(config)# router eigrp VPN  
R6(config-router)# address-family ipv4 autonomous-system 1  
R6(config-router-af)# af-interface Tunnel1  
R6(config-router-af-interface)# no next-hop-self
```

```
R4# show ip route eigrp  
 5.0.0.0/24 is subnetted, 1 subnets  
D      5.5.5.0 [90/285596416] via 100.100.100.5, 00:00:28, Tunnel0  
 6.0.0.0/24 is subnetted, 1 subnets  
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:29, Tunnel0
```

```
R5# show ip route eigrp  
 4.0.0.0/24 is subnetted, 1 subnets  
D      4.4.4.0 [90/285596416] via 100.100.100.4, 00:00:39, Tunnel0  
 6.0.0.0/24 is subnetted, 1 subnets  
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:39, Tunnel0
```

[Example 3-30](#) shows the ISAKMP IPsec connection on spoke Router R5 to the hub. To bring up a dynamic ISAKMP IPsec connection to the other spoke router R4, an extended ping is required from loopback interface to loopback interface.

This question was extremely complex and is the reason why it was weighted so heavily. You had multiple items to configure within the standard DMVPN solution, such as split horizon. It should make you realize the importance of reading the question a number of times and taking the time to test your configurations to ensure that you have successfully answered the question. If you have configured your routers correctly, as detailed in [Examples 3-29](#) and [3-30](#), congratulations, and you have earned a hefty 10 points.

Example 3-30 DMVPN Spoke-to-Spoke Testing

[Click here to view code image](#)

```
R5# show crypto map
Crypto Map "Tunnel0-head-0" 65536 ipsec-isakmp
    Profile name: IPSEC
    Security association lifetime: 4608000 kilobytes/3600 seconds
    PFS (Y/N): N
    Transform sets={
        DMVPN,
    }

Crypto Map "Tunnel0-head-0" 65537 ipsec-isakmp
    Map is a PROFILE INSTANCE.
    Peer = 120.100.45.6
    Extended IP access list
        access-list permit gre host 120.100.45.5 host 120.100.45.6
    Current peer: 120.100.45.6
    Security association lifetime: 4608000 kilobytes/3600 seconds
    PFS (Y/N): N
    Transform sets={
        DMVPN,
    }
    Interfaces using crypto map Tunnel0-head-0:
        Tunnel0
```

```
R5# show crypto isakmp sa
```

dst	src	state	conn-id	slot	status
120.100.45.6	120.100.45.5	QM_IDLE	4001	0	ACTIVE

```
IPv6 Crypto ISAKMP SA
```

R5 spoke router only has a connection to the hub router. An extended ping sourced from the loopback interface of one spoke to another is required to bring up the dynamic spoke to spoke connection.

```
R5# ping
Protocol [ip]:
Target IP address: 4.4.4.4
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 5.5.5.5
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 4.4.4.4, timeout is 2 seconds:
Packet sent with a source address of 5.5.5.5
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms
```

```
R5# show crypto isakmp sa
```

dst	src	state	conn-id	slot	status
-----	-----	-------	---------	------	--------

120.100.45.5	120.100.45.4	QM_IDLE	4002	0	ACTIVE
120.100.45.6	120.100.45.5	QM_IDLE	4001	0	ACTIVE

IPv6 Crypto ISAKMP SA

R5# **show crypto isakmp sa**

IPv4 Crypto ISAKMP SA

dst	src	state	conn-id	slot	status
120.100.45.5	120.100.45.4	QM_IDLE	4002	0	ACTIVE
120.100.45.6	120.100.45.5	QM_IDLE	4001	0	ACTIVE

IPv6 Crypto ISAKMP SA

- Following on from the previous question, add R2 into the common GRE tunnel network as a spoke router using identical security parameters as used on R4 and R5, ensuring it receives routes from R4, R5, and R6 using the same common EIGRP parameters. The source interface for the tunnel configuration on R2 should be Fast Ethernet 1/1, and the destination should be the Gigabit Ethernet 0/0 interface of R6. Add new Loopback 2 identical IP addresses of 45.45.45.45/24 on both R4 and R5 and advertise this identical network from R4 and R5 to the hub router R6 on the common GRE tunnel interface. Configure R6 to advertise both destinations (R4 and R5) to spoke router R2 for network 45.45.45.0/24 in EIGRP over the common GRE tunnel network. (3 points)

Adding R2 as an additional spoke router into the DMVPN network is a relatively simple task if you were successful with the previous question; it is simply a spoke repetition task. R4 and R5 are configured with a new Loopback 2 interface with an identical IP address of 45.45.45.45/24. This network is then advertised within EIGRP over the DMVPN toward the preconfigured hub router R6. [Example 3-31](#) shows the required configuration on R2, R4, and R5 and the resulting route advertisements for the new network on R4 and R5 successfully received on R6 and R2.

Example 3-31 DMVPN R2, R4, and R5 Configuration and Verification

[Click here to view code image](#)

```
R2(config-if)# router eigrp VPN
R2(config-router)# address-family ipv4 autonomous-system 1
R2(config-router-af)# network 100.100.100.0 0.0.0.255
R2(config-router-af)# exit-address-family
R2(config-router)# crypto isakmp policy 1
R2(config-isakmp)# authentication pre-share
R2(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R2(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R2(crypto-trans)# crypto ipsec profile IPSEC
R2(ipsec-profile)# set transform-set DMVPN
R2(ipsec-profile)# interface Tunnel0
R2(config-if)# ip address 100.100.100.2 255.255.255.0
R2(config-if)# ip mtu 1416
R2(config-if)# ip nhrp authentication SECRET
R2(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R2(config-if)# ip nhrp map multicast 120.100.45.6
R2(config-if)# ip nhrp network-id 10
R2(config-if)# ip nhrp holdtime 100
R2(config-if)# ip nhrp nhs 100.100.100.6
```

```

R2(config-if)# delay 2000
R2(config-if)# tunnel source fastethernet0/1
R2(config-if)# tunnel mode gre multipoint
R2(config-if)# tunnel key 1
R2(config-if)# tunnel protection ipsec profile IPSEC
R4(config)# interface loopback2
R4(config-if)# ip add 45.45.45.45 255.255.255.0
R4(config-if)# router eigrp VPN
R4(config-router)# address-family ipv4 autonomous-system 1
R4(config-router-af)# network 45.45.45.0 0.0.0.255

R5(config)# interface loopback2
R5(config-if)# ip add 45.45.45.45 255.255.255.0
R5(config-if)# router eigrp VPN
R5(config-router)# address-family ipv4 autonomous-system 1
R5(config-router-af)# network 45.45.45.0 0.0.0.255

R6# show ip route eigrp
 4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/61440640] via 100.100.100.4, 00:00:16, Tunnel0
 5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      5.5.5.0/24 [90/61440640] via 100.100.100.5, 00:00:16, Tunnel0
 45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      45.45.45.0/24 [90/61440640] via 100.100.100.5, 00:01:10, Tunnel0
                                         [90/61440640] via 100.100.100.4, 00:01:10, Tunnel0

R2# show ip route eigrp
 4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/71680640] via 100.100.100.4, 00:01:40, Tunnel0
 5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      5.5.5.0/24 [90/71680640] via 100.100.100.5, 00:01:40, Tunnel0
 6.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      6.6.6.0/24 [90/61440640] via 100.100.100.6, 00:07:05, Tunnel0
 45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      45.45.45.0/24 [90/71680640] via 100.100.100.4, 00:03:39, Tunnel0

```

[Example 3-32](#) shows that only a single route for network 45.45.45.0/24 is shown in the routing table of R2. This is the default behavior of the hub router R6 when a hub has more than one path (with the same metric but through different spokes) to reach the same network; EIGRP advertises only one path as the best path to connected spokes. With Add Path Support in EIGRP, hubs can advertise up to four additional best paths to connected spokes, thereby allowing load balancing and path redundancy. In this instance, the command **add-paths 2** under the Tunnel 0 interface of the EIGRP af-interface section ensures that the spoke router R2 receives both paths to network 45.45.45.0/24 through R4 and R5. If you have configured this correctly, as shown in [Example 3-32](#), you have scored 3 points.

Example 3-32 DMVPN R2, R4, and R5 Configuration and Verification

[Click here to view code image](#)

```

R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
R6(config-router-af)# af-interface Tunnel0
R6(config-router-af-interface)# add-paths 2

```

```
R2# show ip route eigrp
 4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/71680640] via 100.100.100.4, 00:01:16, Tunnel0
 5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      5.5.5.0/24 [90/71680640] via 100.100.100.5, 00:01:14, Tunnel0
 6.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      6.6.6.0/24 [90/61440640] via 100.100.100.6, 00:01:22, Tunnel0
 45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D      45.45.45.0/24 [90/71680640] via 100.100.100.4, 00:01:16, Tunnel0
                                         [90/61440640] via 100.100.100.5, 00:01:16, Tunnel0
```

- The network manager of your network cannot justify a full security implementation but wants to implement a solution that provides only a password prompt from R1 when the keyboard entry 1 is entered on the console port (as opposed to the normal CR/Enter key). Configure R1 appropriately. (2 points)

This question makes use of the **activation-character** command on the console port. This is a nasty question because the CLI entry requires an ASCII entry. You would need to search to discover that ASCII numeric figures (0 to 9) are prefixed by the binary value of 0011, so a value of 1 (0001) would be 00110001. Therefore, the decimal conversion is $32 + 16 + 1 = 49$. This is good question on which to use the (?) on the CLI for clues and your documentation CD or search facility in the lab if you were not aware of this feature. If you have configured this correctly per [Example 3-33](#), you have scored 2 points.

Example 3-33 R1 Console Activation-Character Configuration

[Click here to view code image](#)

```
R1(config)# line con 0
R1(config-line)# activation-character ?
CHAR or <0-127> Activation character or its decimal equivalent
R1(config-line)# activation-character 49
```

Lab 3 Wrap-Up

So, how did it go? Did you run out of time? Did you manage to finish but miss what was actually required? If you scored over 80, well done. If you accomplished this within 8 hours or less, you will be prepared for any scenario that you are likely to face during the 5.5 hours of the Configuration section of the actual exam. Remember that the Troubleshooting section on the v5.0 exam is a separate section from the Configuration section and has a different scenario; you will have 2 hours to complete the Troubleshooting section. This lab was designed to ensure that you troubleshoot your own work as you progress through the questions.

Did you manage to configure items such as disabling split horizon for DMVPN and the area ID for OSPF? This attention to detail and complete understanding of the protocols will ultimately earn you your number.

Summary

Are You Ready?

This question became a well-known [Cisco.com](#) slogan that identified the Internet revolution. By the end of these practice exams, you should have a good idea of whether you are ready. Did you feel confident working through the questions, or was it a complete shock to the system? Are you more used to being spoon-fed solitary scenarios than actually having to analyze questions and piece together parts of a complex network jigsaw?

Life is full of challenges. During your education and career, the CCIE certification is as tough as it gets. The exam is designed to test your technical skills, your understanding and analysis of complex topologies, and your capacity to build and troubleshoot a network with IP routing protocols and features. You need to achieve a minimum score of 80 percent to pass.

Further Reading

The following Cisco Press titles are on topics appearing on the CCIE exam blueprint. These books are not required study resources, but you can use them to build knowledge in certain areas:

CCIE Routing and Switching Exam Certification Guide, Fourth Edition

CCIE Routing and Switching Exam Quick Reference, Second Edition

CCIE Routing and Switching Troubleshooting Practice Labs

Routing TCP/IP, Volume I, 2/e

Routing TCP/IP, Volume II

Troubleshooting IP Routing Protocols

Inside Cisco IOS Software Architecture

Cisco LAN Switching

Cisco OSPF Command and Configuration Handbook

Cisco BGP-4 Command and Configuration Handbook

Cisco Router Configuration Handbook, Second Edition

Cisco LAN Switching Configuration Handbook, Second Edition

Developing IP Multicast Networks, Volume I

Internet Routing Architectures, Second Edition

MPLS and VPN Architectures

MPLS and VPN Architectures, Volume II

Cisco Catalyst QoS

End-to-End QoS Network Design

Deploying IPv6 Networks

Network Security Technologies and Solutions

Help and Advice

- Look at http://www.cisco.com/web/learning/le3/ccie/rs/lab_exam.html for the latest information about the CCIE certification, which includes suggested training and reading.
- Keep your schedule flexible during your rack time. Include time for breaks and relaxation; you will often find that 5 minutes away from the keyboard can help you consider possible solutions. Most important, do not forget the people you care for and make time for them, too.
- Build your study plan based on a balance between theory and practice. You need to understand the concepts through the theory; then consolidate this during your rack time.
- Begin with simple topics in isolation; then work up to complex lab scenarios. Spend as much time repeating your configurations as possible to improve your speed and ability to perform basic configurations with your eyes shut. This will save you time for where you need it during the exam.
- Explore the Cisco CD documentation or the URL <http://www.cisco.com/univercd/home/home.htm>. This will be your research lifeline during the exam, where you can find information, concepts, and samples related to all technologies involved in the exam.
- Start to plan for your exam at least 6 months before the lab date.
- If you find these practice labs have highlighted weak areas, do not be afraid to postpone your lab date.

How Can I Schedule My CCIE Lab Exam?

Go to http://www.cisco.com/web/learning/le3/ccie/rs/lab_exam.html to find all the information on how to schedule your exam, including locations, start times, and more. You must have a CCO user ID, your CCIE written exam date, and score to be able to view your profile and schedule your exam.

The Day Before

If you are traveling to take your exam, try to arrive the day before to familiarize yourself with the area. Take a tour to the lab location so that you won't be late on the day; the last thing you need is to arrive flustered. The day before is a day to be relaxed, not a time to attempt any last-minute studying. Have a light dinner and try to get a good night's sleep. Most important, save the beer until after the exam; pass or fail you will feel like one or two for sure. The CCIE exam might be the reason why Stella Artois is so popular in Brussels!

The Day of the Exam

On the day of the exam, plan to arrive for registration at least 15 minutes before the exam begins. The proctor will walk you to the lab and give you a briefing before the exam starts, telling you about the lab environment, on which rack or station you will be working, and the general guidelines for the day. The proctor will not discuss solutions or possible solutions for a given question with you. The proctor will be available to help you understand the wording or meaning of the questions, to make sure that the backbone routers are working properly, and to make sure that the hardware and software on your rack are working perfectly so that your exam runs smoothly. Ask the proctor for any assistance or verification; the worst he or she can say is, "Sorry, everything looks okay from my side;

please check your configuration.” Read the entire exam before you start to get the bigger picture, ensuring that you fully understand each question and its requirements. Begin by performing easier tasks, leaving the most difficult for later. Take some small breaks during the morning and the afternoon to refresh yourself and relieve the stress.

Pass or Fail, What Next?

If you pass, you certainly have something to celebrate; you have just joined an elite club that will in no doubt enhance your career. You have achieved the highest level of certification in the networking world and should aim to continue your thirst for knowledge that sets you apart from your peers. However, take a break before starting your next CCIE track!

If you failed, don’t worry, and don’t take it personally; most people fail on their first attempt. You will have to put it down to experience and get back on the keyboard as soon as you can to work out what went wrong. You will more than likely be successful the next time and will ultimately become a better engineer for your extra rack time.

I hope these practice exams and tips prove helpful and guide you to take your exam with success.

Where are the Companion Content Files?

Register this digital version of Cisco CCIE Routing and Switching v5.0 Configuration Practice Labs, Third Edition to access important downloads.

Register this eBook to unlock the companion configuration files for each practice lab. Follow these steps:

1. Go to www.ciscopress.com/register and log in or create a new account.
2. Enter the ISBN: **9780133786316**
3. Answer the challenge question as proof of purchase.
4. Click on the “Access Bonus Content” link in the Registered Products section of your account page, to be taken to the page where your downloadable content is available.

The Professional and Personal Technology Brands of Pearson



Cisco Press

IBM
Press.

informIT

PEARSON IT Certification



QUE

SAMS

VMWARE PRESS

```
SW3 (config)# interface range fastethernet 0/23-24  
SW3 (config-if-range)# shut
```

```
SW4 (config)# interface range fastethernet 0/23-24  
SW4 (config-if-range)# shut
```

```
SW1(config)# spanning-tree mode rapid-pvst
```

```
SW2(config)# spanning-tree mode rapid-pvst
```

```
SW1(config)# spanning-tree vlan 1 root primary
SW1(config)# spanning-tree vlan 300 root primary
SW1(config-if)# interface fastethernet 0/19
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/20
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/21
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface fastethernet 0/22
SW1(config-if)# spanning-tree guard root
```

```
SW2(config)# spanning-tree vlan 1 root secondary
SW2(config)# spanning-tree vlan 300 root secondary
SW2(config-if)# interface fastethernet 0/19
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/20
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/21
SW2(config-if)# spanning-tree guard root
SW2(config-if)# interface fastethernet 0/22
SW2(config-if)# spanning-tree guard root
```

```
SW1(config-if)# interface range fastethernet0/19-20
SW1(config-if)# channel-group 1 mode on
SW1(config-if)# interface range fastethernet0/21-22
SW1(config-if)# channel-group 2 mode on
SW1(config-if)# interface range fastethernet0/23-24
SW1(config-if)# channel-group 3 mode on
SW1(config-if)# interface Port-channel1
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface Port-channel2
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
SW1(config-if)# spanning-tree guard root
SW1(config-if)# interface Port-channel3
SW1(config-if)# switchport trunk encapsulation dot1q
SW1(config-if)# switchport mode trunk
```

```
SW2(config-if)# interface range fastethernet0/19-20
SW2(config-if)# channel-group 1 mode on
SW2(config-if)# interface range fastethernet0/21-22
SW2(config-if)# channel-group 2 mode on
SW2(config-if)# interface range fastethernet0/23-24
SW2(config-if)# channel-group 3 mode on
SW2(config-if)# interface Port-channel1
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk
SW2(config-if)# interface Port-channel2
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk
SW2(config-if)# interface Port-channel3
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport mode trunk
```

```
SW3(config-if)# interface range fastethernet0/19-20
SW3(config-if)# channel-group 1 mode on
SW3(config-if)# interface range fastethernet0/21-22
SW3(config-if)# channel-group 2 mode on
SW3(config-if)# interface Port-channel1
SW3(config-if)# switchport trunk encapsulation dot1q
```

```
SW3(config-if)# switchport mode trunk
SW3(config-if)# interface Port-channel2
SW3(config-if)# switchport trunk encapsulation dot1q
SW3(config-if)# switchport mode trunk
```

```
SW4(config-if)# interface range fastethernet0/19-20
SW4(config-if)# channel-group 1 mode on
SW4(config-if)# interface range fastethernet0/21-22
SW4(config-if)# channel-group 2 mode on
SW4(config-if)# interface Port-channel1
SW4(config-if)# switchport trunk encapsulation dot1q
SW4(config-if)# switchport mode trunk
SW4(config-if)# interface Port-channel2
SW4(config-if)# switchport trunk encapsulation dot1q
SW4(config-if)# switchport mode trunk
```

```
SW1# show interfaces port-channel 1 status
```

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

```
SW1# show interfaces port-channel 2 status
```

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

```
SW1# show interfaces port-channel 3 status
```

Port	Name	Status	Vlan	Duplex	Speed	Type
Po3		connected	trunk	a-full	a-100	

```
SW1# show etherchannel summary
```

Number of channel-groups in use: 3

Number of aggregators: 3

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)
3	Po3 (SU)	-	Fa0/23 (P) Fa0/24 (P)

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)
3	Po3 (SU)	-	Fa0/23 (P) Fa0/24 (P)

```
SW2# show interfaces port-channel 1 status
```

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

```
SW2# show interfaces port-channel 2 status
```

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW2# show interfaces port-channel 3 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po3		connected	trunk	a-full	a-100	

SW2# show etherchannel summary

Number of channel-groups in use: 3

Number of aggregators: 3

Group Port-channel Protocol Ports

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)
3	Po3 (SU)	-	Fa0/23 (P) Fa0/24 (P)

SW3# show interface port-channel 1 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW3# show interface port-channel 2 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

SW3# show etherchannel summary

Number of channel-groups in use: 2

Number of aggregators: 2

Group Port-channel Protocol Ports

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)

SW4# show interface port-channel 1 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po1		connected	trunk	a-full	a-100	

SW4# show interface port-channel 2 status

Port	Name	Status	Vlan	Duplex	Speed	Type
Po2		connected	trunk	a-full	a-100	

```
SW4# show etherchannel summary
```

```
Number of channel-groups in use: 2
```

```
Number of aggregators: 2
```

Group	Port-channel	Protocol	Ports
-------	--------------	----------	-------

Group	Port-channel	Protocol	Ports
1	Po1 (SU)	-	Fa0/19 (P) Fa0/20 (P)
2	Po2 (SU)	-	Fa0/21 (P) Fa0/22 (P)

```
SW1(config)# port-channel load-balance dst-mac
```

```
SW2(config)# port-channel load-balance dst-mac
```

```
SW3(config)# port-channel load-balance dst-mac
```

```
SW4(config)# port-channel load-balance dst-mac
```

```
SW1# show etherchannel load-balance
```

```
EtherChannel Load-Balancing Operational State (dst-mac) :
```

```
Non-IP: Destination MAC address
```

```
IPv4: Destination MAC address
```

```
IPv6: Destination IP address
```

```
SW1(config)# errdisable recovery cause link-flap
SW1(config)# errdisable recovery interval 35
SW1(config)# interface fastethernet 0/10
SW1(config-if)# storm-control multicast level 0
SW1(config-if)# storm-control action shutdown

SW2(config)# errdisable recovery cause link-flap
SW2(config)# errdisable recovery interval 35
SW2(config)# interface fastethernet 0/10
SW2(config-if)# storm-control multicast level 0
SW2(config-if)# storm-control action shutdown

SW3(config)# errdisable recovery cause link-flap
SW3(config)# errdisable recovery interval 35
SW3(config)# interface fastethernet 0/10
SW3(config-if)# storm-control multicast level 0
SW3(config-if)# storm-control action shutdown

SW4(config)# errdisable recovery cause link-flap
SW4(config)# errdisable recovery interval 35
SW3(config)# interface fastethernet 0/10
SW3(config-if)# storm-control multicast level 0
SW3(config-if)# storm-control action shutdown
```

```
SW1(config)# ip dhcp snooping
SW1(config)# ip dhcp snooping vlan 300
SW1(config)# ip dhcp snooping information option
SW1(config)# int fastethernet 0/18
SW1(config-if)# ip dhcp snooping trust
SW1(config)# interface range fastethernet 0/11-17
SW1(config-if-range)# ip dhcp snooping limit rate 10
SW1(config)# interface range fastethernet 0/11-18
SW1(config-if-range)# switchport host
SW1(config-if-range)# switchport access vlan 300

SW2(config)# ip dhcp snooping
SW2(config)# ip dhcp snooping vlan 300
SW2(config)# ip dhcp snooping information option
SW2(config)# interface range fastethernet 0/11-17
SW2(config-if-range)# ip dhcp snooping limit rate 10
SW2(config-if-range)# switchport host
SW2(config-if-range)# switchport access vlan 300

SW3(config)# ip dhcp snooping
SW3(config)# ip dhcp snooping vlan 300
SW3(config)# ip dhcp snooping information option
SW3(config)# interface range fastethernet 0/11-17
SW3(config-if-range)# ip dhcp snooping limit rate 10
SW3(config-if-range)# switchport host
SW3(config-if-range)# switchport access vlan 300
```

```
SW4(config)# ip dhcp snooping
SW4(config)# ip dhcp snooping vlan 300
SW4(config)# ip dhcp snooping information option
SW4(config)# interface range fastethernet 0/11-17
SW4(config-if-range)# ip dhcp snooping limit rate 10
SW4(config-if-range)# switchport host
SW4(config-if-range)# switchport access vlan 300
```

```
SW1# sh ip dhcp snooping
Switch DHCP snooping is enabled
DHCP snooping is configured on following VLANs:
300
```

```
Insertion of option 82 is enabled
  circuit-id format: vlan-mod-port
  remote-id format: MAC
```

```
Option 82 on untrusted port is not allowed
```

```
Verification of hwaddr field is enabled
```

Interface	Trusted	Rate limit (pps)
fastethernet0/11	no	10
fastethernet0/12	no	10
fastethernet0/13	no	10
fastethernet0/14	no	10
fastethernet0/15	no	10
fastethernet0/16	no	10
fastethernet0/17	no	10
fastethernet0/18	yes	unlimited

```
SW1(config)# interface range fast 0/11-17
SW1(config-if-range)# ip verify source

SW2(config)# interface range fast 0/11-17
SW2(config-if-range)# ip verify source

SW3(config)# interface range fast 0/11-17
SW3(config-if-range)# ip verify source

SW4(config)# interface range fast 0/11-17
SW4(config-if-range)# ip verify source
```

```
R4(config)# interface GigabitEthernet0/1.45
R4(config-if)# encapsulation dot1Q 45
R4(config-if)# ip address 120.100.45.4 255.255.255.0
R4(config-if)# interface GigabitEthernet0/1.46
R4(config-if)# encapsulation dot1Q 46
R4(config-if)# ip address 120.100.46.4 255.255.255.0

SW2(config)# vlan 45-46
SW2(config)# interface fastethernet0/4
SW2(config-if)# switchport trunk encapsulation dot1q
SW2(config-if)# switchport trunk allowed vlan 45,46
SW2(config-if)# switchport mode trunk
```

```
R1(config)# interface GigabitEthernet 0/0
R1(config-if)# ip ospf 1 area 100
R1(config)# interface GigabitEthernet 0/1
R1(config-if)# ip ospf 1 area 0
R1(config-if)# interface Loopback 0
R1(config-if)# ip ospf 1 area 0

R2(config)# interface Loopback 0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# interface fastethernet 0/0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# interface Serial 0/1
R2(config-if)# ip ospf 1 area 5
R2(config-if)# interface fastethernet 0/1
R2(config-if)# ip ospf 1 area 200

R3(config)# interface loopback 0
R3(config-if)# ip ospf 1 area 0
R3(config-if)# interface GigabitEthernet 0/1
R3(config-if)# ip ospf 1 area 0
R3(config-if)# interface GigabitEthernet 0/0
R3(config-if)# ip ospf 1 area 34

R4(config)# interface Loopback 0
R4(config-if)# ip ospf 1 area 34
R4(config-if)# interface GigabitEthernet 0/0
R4(config-if)# ip ospf 1 area 34
R4(config-if)# interface GigabitEthernet 0/1.45
R4(config-if)# ip ospf 1 area 5

R5(config)# interface Loopback 0
R5(config-if)# ip ospf 1 area 5
R5(config-if)# interface GigabitEthernet 0/0
R5(config-if)# ip ospf 1 area 5
R5(config-if)# interface Serial 0/0/1
R5(config-if)# ip ospf 1 area 5
```

```
R2# sh ip route | include /32
O      120.100.5.1/32 [110/65] via 120.100.25.5, 00:04:34, Serial0/1
O IA    120.100.4.1/32 [110/66] via 120.100.123.3, 00:00:42, Serial0/0
O      120.100.1.1/32 [110/129] via 120.100.123.3, 00:01:00, Serial0/0
O      120.100.3.1/32 [110/65] via 120.100.123.3, 00:01:00, Serial0/0

R2# sh ip route | include /32
C      120.100.25.5/32 is directly connected, Serial0/1
O      120.100.5.1/32 [110/65] via 120.100.25.5, 00:39:59, Serial0/1
O IA    120.100.4.1/32 [110/3] via 120.100.123.3, 00:47:32, GigabitEthernet0/1
O      120.100.1.1/32 [110/2] via 120.100.123.1, 00:50:56, GigabitEthernet0/1
O      120.100.3.1/32 [110/2] via 120.100.123.3, 00:49:20, GigabitEthernet0/1

R1# conf t
R1(config)# int Loopback 0
R1(config-if)# ip ospf network point-to-point

R2# conf t
R2(config)# interface Loopback 0
R2(config-if)# ip ospf network point-to-point

R3# conf t
R3(config)# int Loopback 0
R3(config-if)# ip ospf network point-to-point

R4# conf t
R4(config)# int Loopback 0
R4(config-if)# ip ospf network point-to-point

R5# conf t
R4(config)# int Loopback 0
R4(config-if)# ip ospf network point-to-point

R2# sh ip route ospf 1 | include /24
O IA    120.100.4.0/24 [110/3] via 120.100.123.3, 01:42:09, fastethernet0/0
O      120.100.5.0/24 [110/65] via 120.100.25.5, 00:17:09, Serial0/1
O      120.100.1.0/24 [110/2] via 120.100.123.1, 01:43:00, fastethernet0/0
O      120.100.3.0/24 [110/2] via 120.100.123.3, 01:42:26, fastethernet0/0
O      120.100.45.0/24 [110/65] via 120.100.25.5, 02:52:46, Serial0/1
O IA    120.100.34.0/24 [110/2] via 120.100.123.3, 01:43:00, fastethernet0/0
O IA    120.100.100.0/24 [110/2] via 120.100.123.1, 00:00:04, fastethernet0/0
```

```
R1# show ip ospf int GigabitEthernet 0/1
GigabitEthernet0/1 is up, line protocol is up
  Internet Address 150.100.1.1/24, Area 100
  Process ID 1, Router ID 120.100.1.1, Network Type BROADCAST, Cost: 1
  Enabled by interface config, including secondary ip addresses
  Transmit Delay is 1 sec, State DR, Priority 1
  Designated Router (ID) 120.100.1.1, Interface address 150.100.1.1
  No backup designated router on this network
  Timer intervals configured, Hello 10, Dead 40, Wait 40, Retransmit 5
    oob-resync timeout 40
    Hello due in 00:00:00
  Supports Link-local Signaling (LLS)
  Cisco NSF helper support enabled
  IETF NSF helper support enabled
  Index 1/1, flood queue length 0
  Next 0x0(0)/0x0(0)
  Last flood scan length is 0, maximum is 0
  Last flood scan time is 0 msec, maximum is 0 msec
  Neighbor Count is 0, Adjacent neighbor count is 0
  Suppress hello for 0 neighbor(s)
```

```
R1(config)# interface GigabitEthernet 0/1
R1(config-if)# ip ospf 1 area 100 secondaries none
```

```
R2# sh ip route 120.100.100.0
% Subnet not in table
```

```
R5(config)# ip sla 1
R5(config-ip-sla)# icmp-echo 120.100.25.5
R5(config-ip-sla-echo)# exit
R5(config)# ip sla schedule 1 life forever start-time now
R5(config)# track 1 rtr 1 reachability
```

```
R5# show ip sla statistics
```

```
Round Trip Time (RTT) for Index 1
    Latest RTT: 4 milliseconds
Latest operation start time: *21:17:10.683 UTC Mon Aug 05 2013Latest operation
return code: OK
Number of successes: 2
Number of failures: 0
Operation time to live: Forever
```

```
R4(config)# interface GigabitEthernet0/1.45
R4(config-if)# ip ospf network non-broadcast
R4(config-if)# router ospf 1
R4(config-router)# neighbor 120.100.45.5
```

```
R5(config)# interface GigabitEthernet0/0
R5(config-if)# ip ospf network non-broadcast
R5(config-if)# router ospf 1
R5(config-router)# neighbor 120.100.45.4
R5(config-router)# exit
R5(config)# access-list 100 permit ospf host 120.100.45.5 host 120.100.45.4
R5(config)# route-map TEST permit 10
R5(config-route-map)# match ip address 100
R5(config-route-map)# set ip next-hop verify-availability 120.100.25.2 10 track 1
R5(config-route-map)# interface GigabitEthernet0/0
R5(config-if)# ip policy route-map TEST
R5(config-if)# exit
R5(config)# ip local policy route-map TEST
```

```
R2# debug ip icmp
ICMP packet debugging is on
R2#
*Feb 26 22:17:12.847: ICMP: time exceeded (time to live) sent to 120.100.45.5 (d
est was 120.100.45.4)
R2#
```

```
R5# show ip ospf neigh
```

Neighbor ID	Pri	State	Dead Time	Address	Interface
120.100.2.1	0	FULL/-	00:00:37	120.100.25.2	Serial0/0/1
120.100.4.1	1	INIT/DROTHER	00:01:45	120.100.45.4	GigabitEthernet0/0

```

R3(config)# router ospf 1
R3(config-router)# area 34 virtual-link 120.100.4.1

R4(config)# router ospf 1
R4(config-router)# area 34 virtual-link 120.100.3.1

R5(config)# interface s0/0/1
R5(config-if)# shut
R5(config-if)#
*Jan  2 21:58:16.811: %OSPF-5-ADJCHG: Process 1, Nbr 120.100.2.1 on Serial0/0/1 from
  FULL to DOWN, Neighbor Down: Interface down or detached
*Jan  2 21:58:18.807: %LINK-5-CHANGED: Interface Serial0/0/1, changed state to
  administratively down
*Jan  2 21:58:19.807: %LINEPROTO-5-UPDOWN: Line protocol on Interface Serial0/0/1,
  changed state to down
R5(config-if)# do show ip ospf neigh

Neighbor ID      Pri   State            Dead Time       Address          Interface
N/A              0     ATTEMPT/DROTHER 00:00:33    120.100.45.4 GigabitEthernet0/0
R5(config-if)#
*Jan  2 21:59:43.547: %OSPF-5-ADJCHG: Process 1, Nbr 0.0.0.0 on GigabitEthernet0/0
  from ATTEMPT to DOWN, Neighbor Down: Dead timer expired
R5(config-if)#
*Jan  2 22:00:08.135: %OSPF-5-ADJCHG: Process 1, Nbr 120.100.4.1 on
  GigabitEthernet0/0 from LOADING to FULL, Loading Done
R5(config-if)#

R5# sh ip route ospf
  150.100.0.0/24 is subnetted, 2 subnets
O IA    150.100.2.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
        120.0.0.0/24 is subnetted, 9 subnets
O IA    120.100.4.0 [110/2] via 120.100.45.4, 00:04:49, GigabitEthernet0/0
O IA    120.100.1.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.2.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.3.0 [110/3] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.34.0 [110/2] via 120.100.45.4, 00:04:49, GigabitEthernet0/0
O IA    120.100.123.0 [110/3] via 120.100.45.4, 00:00:12, GigabitEthernet0/0
O IA    120.100.100.0 [110/4] via 120.100.45.4, 00:00:12, GigabitEthernet0/0

```

```
R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# network 120.100.4.0 0.0.0.255
R4(config-router-af)# network 120.100.45.0 0.0.0.255
R4(config-router-af)# network 120.100.46.0 0.0.0.255
```

```
R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# network 120.100.5.0 0.0.0.255
R5(config-router-af)# network 120.100.45.0 0.0.0.255
R5(config-router-af)# network 120.100.3.0 0.0.0.255
```

```
R6(config)# router eigrp CCIE
R6(config-router)# address-family ipv4 unicast autonomous-system 1
R6(config-router-af)# network 120.100.6.0 0.0.0.255
R6(config-router-af)# network 120.100.46.0 0.0.0.255
R6(config-router-af)# network 120.100.3.0 0.0.0.255
```

```
SW1(config)# ip routing
SW1(config)# exit
SW1# sh run | beg eigrp
router eigrp 1
network 120.100.7.1 0.0.0.0
network 150.100.3.7 0.0.0.0
no auto-summary
```

```
SW2(config)# ip routing
SW2(config)# exit
SW2# sh run | beg eigrp
router eigrp 1
network 120.100.8.1 0.0.0.0
network 150.100.3.8 0.0.0.0
no auto-summary
```

```
SW3(config)# ip routing
SW3(config)# exit
SW3# sh run | beg eigrp
router eigrp 1
network 120.100.9.1 0.0.0.0
```

```
network 150.100.3.9 0.0.0.0
no auto-summary

SW4(config)# ip routing
SW4(config)# exit
SW4# sh run | beg eigrp
router eigrp 1
network 120.100.10.1 0.0.0.0
network 150.100.3.10 0.0.0.0
no auto-summary
```

```
R4# show ip route eigrp
  150.100.0.0/24 is subnetted, 3 subnets
D      150.100.3.0
        [90/30720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/30720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
  120.0.0.0/8 is variably subnetted, 16 subnets, 2 masks
D      120.100.8.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.9.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.10.0/24
        [90/158720] via 120.100.46.6, 00:01:07, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:01:07, GigabitEthernet0/1.45
D      120.100.5.0/24
        [90/156160] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
D      120.100.6.0/24
        [90/156160] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
D      120.100.7.0/24
        [90/158720] via 120.100.46.6, 00:00:10, GigabitEthernet0/1.46
        [90/158720] via 120.100.45.5, 00:00:10, GigabitEthernet0/1.45
```

```
R4# show ip route 120.100.8.0
Routing entry for 120.100.8.0/24
Known via "eigrp 1", distance 90, metric 158720, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:15 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:15 ago, via GigabitEthernet0/1.46
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 2
```

```
R4# show ip route 120.100.9.0
Routing entry for 120.100.9.0/24
Known via "eigrp 1", distance 90, metric 158720, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:25 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:25 ago, via GigabitEthernet0/1.46
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 2
```

```
R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topolgy)# metric maximum-hops 1
R4(config-router-af-topology)# do show ip route eigrp
    150.100.0.0/24 is subnetted, 3 subnets
D        150.100.3.0
            [90/30720] via 120.100.46.6, 00:00:04, GigabitEthernet0/1.46
            [90/30720] via 120.100.45.5, 00:00:04, GigabitEthernet0/1.45
    120.0.0.0/8 is variably subnetted, 13 subnets, 2 masks
D        120.100.5.0/24
            [90/156160] via 120.100.45.5, 00:00:04, GigabitEthernet0/1.45
D        120.100.6.0/24
            [90/156160] via 120.100.46.6, 00:00:04, GigabitEthernet0/1.46
```

```
R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 30720, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.45.5 on GigabitEthernet0/1.45, 00:25:40 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:25:40 ago, via GigabitEthernet0/1.46
    Route metric is 30720, traffic share count is 1
    Total delay is 200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 254/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
120.100.45.5, from 120.100.45.5, 00:25:40 ago, via GigabitEthernet0/1.45
    Route metric is 30720, traffic share count is 1
    Total delay is 200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 252/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R4(config)# route-map CHANGEMETRIC permit 10
R4(config-route-map)# match interface gigabitEthernet 0/1.45
R4(config-route-map)# set metric 2000 10 255 1 1500
R4(config-route-map)# route-map CHANGEMETRIC permit 20
R4(config-route-map)# set metric 1000 10 255 1 1500
R4(config-route-map)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topolgy)# distribute-list route-map CHANGEMETRIC in
R4(config-router-af-topolgy)# ^z
R4# clear ip route *
R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 1282560, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.45.5 on GigabitEthernet0/1.45, 00:03:10 ago
Routing Descriptor Blocks:
* 120.100.45.5, from 120.100.45.5, 00:03:10 ago, via GigabitEthernet0/1.45
    Route metric is 1282560, traffic share count is 1
    Total delay is 100 microseconds, minimum bandwidth is 2000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R5(config)# int gig0/0
R5(config-if)# shutdown

R4# sh ip route 150.100.3.0
Routing entry for 150.100.3.0/24
Known via "eigrp 1", distance 90, metric 2562560, type internal
Redistributing via ospf 1, eigrp 1
Advertised by ospf 1 metric 5000 subnets
Last update from 120.100.46.6 on GigabitEthernet0/1.46, 00:00:10 ago
Routing Descriptor Blocks:
* 120.100.46.6, from 120.100.46.6, 00:00:10 ago, via GigabitEthernet0/1.46
    Route metric is 2562560, traffic share count is 1
    Total delay is 100 microseconds, minimum bandwidth is 1000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R4(config-route-map)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topology)# redistribute ospf 1
R4(config-router-af-topology)# default-metric 10000 100 255 1 1500
R4(config-router-af-topology)# router ospf 1
R4(config-router)# redistribute eigrp 1 subnets
R4(config-router)# default-metric 5000

R1# show ip route ospf | include E2
O E2      150.100.3.0 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet0/0
O E2      120.100.6.0/24 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet 0/0
O E2      120.100.46.0/24 [110/5000] via 120.100.123.3, 00:00:46, GigabitEthernet 0/0

SW1# show ip route eigrp | include EX
D EX      150.100.2.0 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      150.100.1.0 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.25.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.1.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.2.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.3.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.34.0/24 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.123.3/32 [170/284416] via 150.100.3.6, 00:01:43, Vlan300
D EX      120.100.123.0/24 [170/284416] via 150.100.3.6, 00:01:44, Vlan300
```

```
R4(config)# router ospf 1
R4(config-router)# redistribute maximum-prefix 5 80
```

```
R1# sh run | begin bgp
router bgp 10
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.3.1 update-source Loopback0
no auto-summary
```

```
R2# sh run | begin bgp
router bgp 10
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.5.1 remote-as 300
neighbor 120.100.5.1 ebgp-multihop 2
neighbor 120.100.5.1 update-source Loopback0
no auto-summary
```

```
R3# sh run | begin bgp
router bgp 10
no synchronization
neighbor IBGP peer-group
neighbor IBGP remote-as 10
neighbor IBGP update-source Loopback0
neighbor IBGP route-reflector-client
neighbor 120.100.1.1 peer-group IBGP
neighbor 120.100.2.1 peer-group IBGP
neighbor 120.100.4.1 remote-as 200
neighbor 120.100.4.1 ttl-security hops 2
neighbor 120.100.4.1 update-source Loopback0
no auto-summary
```

```
R4# sh run | begin bgp
router bgp 200
no synchronization
neighbor 120.100.3.1 remote-as 10
neighbor 120.100.3.1 ebgp-multihop 2
neighbor 120.100.3.1 update-source Loopback0
neighbor 120.100.6.1 remote-as 300
neighbor 120.100.6.1 ebgp-multihop 2
neighbor 120.100.6.1 update-source Loopback0
neighbor 120.100.45.5 remote-as 300
no auto-summary
```

```
R3(config)# access-list 100 permit ip host 120.100.4.1 host 120.100.3.1
R3(config)# exit
R3# debug ip packet 100 detail dump
IP packet debugging is on (detailed) (dump) for access list 100
```

```
R3# TCP src=42692, dst=179, seq=2600279946, ack=0, win=163
84 SYN
0F400C00: C204 07400000 B..@..
0F400C10: C20211E0 00100800 45C0002C 6A870000 B..`....E@.,j...
0F400C20: 0106467E 01010101 03030303 A6C400B3 ..F~.....&D.3
0F400C30: 9AFD1F8A 00000000 60024000 F1BB0000 .}.....`.@.q;..
0F400C40: 02040218 .....
```

```
! The TTL from R4 is decremented to 01 Hex = 01 decimal as R4 has ebgp-multihop 2
! configured and the BGP session will not be established as R3 has the TTL security
! check enabled, from R3's perspective R4 could be 254 hops away!
! Configure R4 so the TTL value will read 253 decimal (FD hex) by configuring an
! ebgp multihop value of 255 (this value will decrement down to 253 when it is
! processed by R3).
```

```
R4(config)# router bgp 200
R4(config)# neighbor 120.100.3.1 ebgp-multihop 255
```

```
R3# TCP src=44109, dst=179, seq=3925370469, ack=3209854606
, win=16263 ACK
0F7CBB60: C204 07400000 B..@..
0F7CBB70: C20211E0 00100800 45C00028 8C9A0000 B..`....E@.(....
0F7CBB80: FD06286E 01010101 03030303 AC4D00B3 }.(n.....,M.3
0F7CBB90: E4028565 BF527E8E 50103F87 13FC0000 d..e?R~.P.?...|..
0F7CBBA0:
```

```
! Now a hex value of FD (253 Decimal) can be seen at R3 from R4, this shows that R4
! can not be further than 2 hops away from R3 and the security check passes and BGP
! is established.
```

```
R3# sh ip bgp neighbor | include hops | TTL
External BGP neighbor may be up to 2 hops away.
Connection is ECN Disabled, Minimum incoming TTL 253, Outgoing TTL 255
```

```
R5# sh run | begin bgp
router bgp 300
```

```
no synchronization
neighbor 120.100.2.1 remote-as 10
neighbor 120.100.2.1 ebgp-multipath 2
neighbor 120.100.2.1 update-source Loopback0
neighbor 120.100.6.1 remote-as 300
neighbor 120.100.6.1 update-source Loopback0
neighbor 120.100.45.6 remote-as 200
neighbor 120.100.7.1 remote-as 300
neighbor 120.100.7.1 update-source Loopback0
no auto-summary
```

```
R6# sh run | beg bgp
router bgp 300
no synchronization
neighbor IBGP peer-group
neighbor IBGP remote-as 300
neighbor IBGP update-source Loopback0
neighbor 120.100.4.1 remote-as 200
neighbor 120.100.4.1 ebgp-multipath 2
neighbor 120.100.4.1 update-source Loopback0
neighbor 120.100.5.1 peer-group IBGP
neighbor 120.100.7.1 peer-group IBGP
no auto-summary
```

```
SW1# sh run | begin bgp
router bgp 300
no synchronization
neighbor IBGP peer-group
neighbor IBGP remote-as 300
neighbor IBGP update-source Loopback0
neighbor 120.100.5.1 peer-group IBGP
neighbor 120.100.6.1 peer-group IBGP
no auto-summary
```

```
R5(config)# int s0/0/1
R5(config-if)# shut

R5# trace 120.100.2.1

Type escape sequence to abort.
Tracing the route to 120.100.2.1

1 120.100.45.4 0 msec 0 msec 0 msec
2 120.100.34.3 0 msec 4 msec 0 msec
3 120.100.123.2 4 msec * 4 msec

R5# debug ip icmp
ICMP packet debugging is on
R5#
*Jan 17 21:32:32.455: ICMP: time exceeded rcvd from 120.100.34.3
R5#
*Jan 17 21:32:34.179: ICMP: time exceeded rcvd from 120.100.34.3
R5#

R2# debug ip icmp
ICMP packet debugging is on
R2#
Jan 17 21:26:11.310: ICMP: time exceeded rcvd from 120.100.34.4
R2#
Jan 17 21:26:13.306: ICMP: time exceeded rcvd from 120.100.34.4
```

```

R5# sh ip bgp
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*->i130.100.200.0/24 120.100.4.1                  0       100      0 200 10 i

R2(config)# interface Loopback2
R2(config-if)# ip address 130.100.200.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 130.100.200.0 mask 255.255.255.0
R2(config-router)# neighbor 120.100.3.1 route-map NO-EXPORT out
R2(config-router)# neighbor 120.100.3.1 send-community
R2(config-router)# exit
R2(config)# access-list 5 permit 130.100.200.0
R2(config)# route-map NO-EXPORT permit 10
R2(config-route-map)# match ip address 5
R2(config-route-map)# set community no-export
R2(config-route-map)# route-map NO-EXPORT permit 20

R3# sh ip bgp 130.100.200.1
BGP routing table entry for 130.100.200.0/24, version 4
Paths: (1 available, best #1, table Default-IP-Routing-Table, not advertised to
EBGP peer)
    Advertised to update-groups:
        2
    Local, (Received from a RR-client)
        120.100.2.1 (metric 65) from 120.100.2.1 (130.100.200.1)
            Origin IGP, metric 0, localpref 100, valid, internal, best
            Community: no-export

R5# conf t
Enter configuration commands, one per line. End with CNTL/Z.
R5(config)# int s0/0/1
R5(config-if)# shut
R5(config-if)# ^Z
R5# show ip bgp

R5#

```

```
R5(config)# track 2 ip route 130.100.200.0 255.255.255.0 reachability
R5(config-track)# interface GigabitEthernet0/1
R5(config-if)# standby 1 ip 150.100.3.1
R5(config-if)# standby 1 preempt
R5(config-if)# standby 1 track 2 decrement 20
```

```
R6(config)# interface GigabitEthernet0/1
R6(config-if)# standby 1 ip 150.100.3.1
R6(config-if)# standby 1 priority 90
R6(config-if)# standby 1 preempt
```

```
R5# sh standby gigabitEthernet 0/1
GigabitEthernet0/1 - Group 1
  State is Active
    23 state changes, last state change 00:20:11
  Virtual IP address is 150.100.3.1
  Active virtual MAC address is 0000.0c07.ac01
    Local virtual MAC address is 0000.0c07.ac01 (v1 default)
  Hello time 3 sec, hold time 10 sec
    Next hello sent in 0.460 secs
  Preemption enabled
  Active router is local
  Standby router is 150.100.3.6, priority 90 (expires in 8.472 sec)
  Priority 100 (default 100)
    Track object 2 state Up decrement 20
  IP redundancy name is "hsrp-Gi0/1-1" (default)
```

```
R5#
```

```
R5# conf t
R5(config)# int s0/0/1
R5(config-if)# shut
R5(config-if)#

```

```
R5#%BGP-3-NOTIFICATION: sent to neighbor 120.100.2.1 4/0 (hold time expired) 0 bytes
```

```
R5#%HSRP-6-STATECHANGE: GigabitEthernet0/1 Grp 1 state Active -> Speak
```

```
R5#%HSRP-6-STATECHANGE: GigabitEthernet0/1 Grp 1 state Speak -> Standby
```

```
R5# sh standby gigabitEthernet 0/1
```

```
GigabitEthernet0/1 - Group 1
```

```
  State is Standby
```

```
25 state changes, last state change 00:00:10
Virtual IP address is 150.100.3.1
Active virtual MAC address is 0000.0c07.ac01
  Local virtual MAC address is 0000.0c07.ac01 (v1 default)
Hello time 3 sec, hold time 10 sec
  Next hello sent in 1.880 secs
Preemption enabled
Active router is 150.100.3.6, priority 90 (expires in 8.980 sec)
Standby router is local
Priority 80 (default 100)
  Track object 2 state Down decrement 20
IP redundancy name is "hsrp-Gi0/1-1" (default)
```

```

R1(config)# interface Loopback1
R1(config-if)# ip address 126.1.1.1 255.255.255.0
R1(config-if)# router bgp 10
R1(config-router)# network 126.1.1.0 mask 255.255.255.0

R2(config)# interface Loopback1
R2(config-if)# ip address 130.1.1.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 130.1.1.0 mask 255.255.255.0

R3(config)# access-list 1 permit 0.0.0.0 127.255.255.255
R3(config)# route-map UPTO128 permit 10
R3(config-route-map)# match ip add 1

R3(config)# route-map ABOVE128 deny 10
R3(config-route-map)# match ip add 1
R3(config-route-map)# route-map ABOVE128 permit 20

R3(config)# router bgp 10
R3(config-router)# neighbor 120.100.1.1 route-map UPTO128 in
R3(config-router)# neighbor 120.100.2.1 route-map ABOVE128 in

R3# sh ip bgp
BGP table version is 8, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop            Metric LocPrf Weight Path
*->i126.1.1.0/24    120.100.1.1        0      100      0 i
*->i130.1.1.0/24    120.100.2.1        0      100      0 i
*->i130.100.200.0/24 120.100.2.1        0      100      0 i
R3#

```

```

R1(config)# interface Loopback3
R1(config-if)# ip address 132.1.1.1 255.255.255.0
R1(config-if)# router bgp 10
R1(config-router)# network 132.1.1.0 mask 255.255.255.0
R1(config-router)# ^Z
R1# sh ip bgp neighbors 120.100.3.1 advertised
BGP table version is 7, local router ID is 126.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 126.1.1.0/24	0.0.0.0	0		32768	i
*> 132.1.1.0/24	0.0.0.0	0		32768	i

Total number of prefixes 2

```

R3# sh ip bgp
BGP table version is 4, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

```

Network	Next Hop	Metric	LocPrf	Weight	Path
*>i126.1.1.0/24	120.100.1.1	0	100	0	i
*>i130.1.1.0/24	120.100.2.1	0	100	0	i
*>i130.100.200.0/24	120.100.2.1	0	100	0	i

```

R2# conf t
R2(config)# int Loopback3
R2(config-if)# ip add 100.1.1.1 255.255.255.0
R2(config-if)# router bgp 10
R2(config-router)# network 100.1.1.0 mask 255.255.255.0
R2(config-router)# ^Z
R2# sh ip bgp neighbor 120.100.3.1 advertised
BGP table version is 5, local router ID is 130.100.200.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 100.1.1.0/24	0.0.0.0	0	32768	i	
*> 130.1.1.0/24	0.0.0.0	0	32768	i	
*> 130.100.200.0/24	0.0.0.0	0	32768	i	

Total number of prefixes 3

R3# sh ip bgp

BGP table version is 4, local router ID is 120.100.3.1
 Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
 r RIB-failure, S Stale
 Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
*>i126.1.1.0/24	120.100.1.1	0	100	0	i
*>i130.1.1.0/24	120.100.2.1	0	100	0	i
*>i130.100.200.0/24	120.100.2.1	0	100	0	i

```
R1(config)# ipv6 unicast-routing
R1(config)# interface gigabitEthernet 0/1
R1(config-if)# ipv6 address 2007:C15:C0:10::1/64
R1(config-if)# gigabitEthernet 0/0
R1(config-if)# ipv6 address 2007:C15:C0:11::1/64
```

```
R2(config)# ipv6 unicast-routing
R2(config)# interface fastethernet 0/1
R2(config-if)# ipv6 address 2007:C15:C0:12::2/64
R2(config-if)# interface fastethernet 0/0
R2(config-if)# ipv6 address 2007:C15:C0:11::2/64
R2(config-if)# interface serial 0/1
R2(config-if)# ipv6 address 2007:C15:C0:14::2/64
```

```
R3(config)# ipv6 unicast-routing
R3(config)# interface gigabitEthernet 0/0
R3(config-if)# ipv6 address 2007:C15:C0:15::3/64
R3(config-if)# gigabitEthernet 0/1
R3(config-if)# ipv6 address 2007:C15:C0:11::3/64
```

```
R4(config)# ipv6 unicast-routing
R4(config)# interface gigabitEthernet 0/0
R4(config-if)# ipv6 address 2007:C15:C0:15::4/64
```

```
R5(config)# ipv6 unicast-routing
R5(config)# interface gigabitEthernet 0/1
R5(config)# ipv6 address 2007:C15:C0:16::5/64
R5(config-if)# interface Serial0/0/1
R5(config-if)# ipv6 address 2007:C15:C0:14::5/64
```

```
R6(config)# ipv6 unicast-routing
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 address 2007:C15:C0:16::6/64
```

```
R1(config)# router eigrp CCIE
R1(config-router)# address-family ipv6 unicast autonomous-system 2
R1(config-router-af)# af-interface GigabitEthernet0/0
R1(config-router-af-interface)# no shutdown
R1(config-router-af-interface)# af-interface GigabitEthernet0/1
R1(config-router-af-interface)# no shutdown

R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv6 unicast autonomous-system 1
R2(config-router-af)# af-interface fastethernet0/1
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface fastethernet0/0
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface Serial0/1
R2(config-router-af-interface)# no shutdown
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 1
R3(config-router-af)# af-interface GigabitEthernet0/0
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface GigabitEthernet0/1
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# exit
R3(config-router-af)# exit
R3(config-router)# exit
R3(config)# interface GigabitEthernet0/1
R3(config-if)# ipv6 eigrp 2
```

```
R3(config-if)# ipv6 summary-address eigrp 2 ::/0
R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 1
R3(config-router)# topology base
R3(config-router-topology)# redistribute eigrp 2 route-map EIGRPv6-2-1
R3(config-router-topology)# exit
R3(config-router-af)# exit
R3(config-router)# ipv6 router eigrp 2
R3(config-rtr)# no shut
R3(config-rtr)# exit
R3(config)# route-map EIGRPv6-2-1 permit 10
R3(config-route-map)# match ipv6 address EIGRPv6-2
R3(config-route-map)# route-map EIGRPv6-2-1 deny 20
R3(config-route-map)# exit
R3(config)# ipv6 access-list EIGRPv6-2
R3(config-ipv6-acl)# permit ipv6 2007:C15:C0:10::/64 any
```

```
R4(config)# router eigrp CCIE
R4(config-router)# address-family ipv6 unicast autonomous-system 1
R4(config-router-af)# af-interface GigabitEthernet0/0
R4(config-router-af-interface)# no shutdown
```

```
R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
R5(config-router-af)# af-interface Serial0/0/1
R5(config-router-af-interface)# no shutdown

R1# sh ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D    ::/0 [90/30720]
      via FE80::216:47FF:FEBB:1E12, GigabitEthernet0/0

R2# sh ipv6 route
EX  2007:C15:C0:10::/64 [170/33280]
      via FE80::216:47FF:FEBB:1E12, fastethernet0/0
D   2007:C15:C0:15::/64 [90/30720]
      via FE80::216:47FF:FEBB:1E12, fastethernet0/0

R3# sh ipv6 route eigrp
```

```
D  ::/0 [5/28160]
    via ::, Null0
D  2007:C15:C0:10::/64 [90/30720]
    via FE80::214:69FF:FE61:5EF0, GigabitEthernet0/1
D  2007:C15:C0:12::/64 [90/30720]
    via FE80::215:C6FF:FEF2:ABF1, GigabitEthernet0/1
D  2007:C15:C0:14::/64 [90/2172416]
    via FE80::215:C6FF:FEF2:ABF1, GigabitEthernet0/1
```

```
R4# sh ipv6 route eigrp
EX  2007:C15:C0:10::/64 [170/33280]
    via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D   2007:C15:C0:11::/64 [90/30720]
    via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D   2007:C15:C0:12::/64 [90/33280]
    via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
D   2007:C15:C0:14::/64 [90/2174976]
    via FE80::216:47FF:FEBB:1E11, GigabitEthernet0/0
```

```
R5# sh ipv6 route eigrp
EX  2007:C15:C0:10::/64 [170/2177536]
    via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D   2007:C15:C0:11::/64 [90/2172416]
    via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D   2007:C15:C0:12::/64 [90/2172416]
    via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
D   2007:C15:C0:15::/64 [90/2174976]
    via FE80::215:C6FF:FEF2:ABE0, Serial0/0/1
```

```
R5(config)# interface gigabitEthernet 0/1
R5(config-if)# ipv6 ospf 1 area 0
```

```
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 ospf 1 area 0
```

```
R5# show ipv6 ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
120.100.6.1	1	FULL/DR	00:00:30	3	GigabitEthernet0/1

```
R6# show ipv6 ospf neighbor
```

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
120.100.5.1	1	FULL/BDR	00:00:39	3	GigabitEthernet0/1

```
R5(config)# interface gigabitEthernet 0/1
R5(config-if)# ipv6 ospf flood-reduction
```

```
R6(config)# interface gigabitEthernet 0/1
R6(config-if)# ipv6 ospf flood-reduction
```

```
R5(config)# ipv6 router ospf 1
R5(config-router)# redistribute eigrp 1 metric 5000

R6# sh ipv6 route ospf
IPv6 Routing Table - 10 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OE2  2007:C15:C0:10::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2  2007:C15:C0:11::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2  2007:C15:C0:12::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2  2007:C15:C0:13::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
OE2  2007:C15:C0:15::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1

R5(config)# ipv6 router ospf 1
R5(config-rtr)# redistribute eigrp 1 metric 5000 include-connected
R6# show ipv6 route 2007:C15:C0:14::
IPv6 Routing Table - 10 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OE2  2007:C15:C0:14::/64 [110/5000]
    via FE80::214:6AFF:FEFC:F131, GigabitEthernet0/1
```

```
R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
R5(config-router-af)# af-interface Serial0/0/1
R5(config-router-af-interface)# summary-address 2007::/16
```

```
R3# sh ipv6 route | include /16
D 2007::/16 [90/2684416]
```

```
R3# ping ipv6 2007:C15:C0:16::5
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 2007:C15:C0:16::5, timeout is 2 seconds:

!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/12/16 ms

```
R3# ping ipv6 2007:C15:C0:16::6
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 2007:C15:C0:16::6, timeout is 2 seconds:

!!!!!

Success rate is 100 percent (5/5), round-trip min/avg/max = 12/15/16 ms

```
R4(config)# interface Tunnel0
R4(config-if)# ipv6 address 2007:C15:C0:17::4/64
R4(config-if)# tunnel source GigabitEthernet0/1.45
R4(config-if)# tunnel destination 120.100.45.5
R4(config-if)# tunnel mode ipv6ip
R4(config-if)# router eigrp CCIE
R4(config-router)# address-family ipv6 unicast autonomous-system 1
R4(config-router-af)# af-interface Tunnel0
R4(config-router-af-interface)# no shutdown

R5(config)# interface Tunnel0
R5(config-if)# ipv6 address 2007:C15:C0:17::5/64
R5(config-if)# ipv6 eigrp 1
R5(config-if)# tunnel source GigabitEthernet0/0
R5(config-if)# tunnel destination 120.100.45.4
R5(config-if)# tunnel mode ipv6ip
R5(config-if)# router eigrp CCIE
R5(config-router)# address-family ipv6 unicast autonomous-system 1
R5(config-router-af)# af-interface Tunnel0
R5(config-router-af-interface)# no shutdown
R5(config-router-af-interface)# summary-address 2007::/16

R5# sh ipv6 route eigrp
```

```
D 2007::/16 [5/2169856]
  via ::, Null0
EX 2007:C15:C0:10::/64 [170/2177536]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:11::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:12::/64 [90/2172416]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
D 2007:C15:C0:15::/64 [90/2174976]
  via FE80::215:C6FF:FEF2:ABE0, Serial0/1/0
```

```
R5(config)# int s0/1/0
R5(config-if)# shut
R5(config-if)# do sh ipv6 route eigrp
D 2007::/16 [5/297244416]
  via ::, Null0
EX 2007:C15:C0:10::/64 [170/297252096]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:11::/64 [90/297249536]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:12::/64 [90/297252096]
  via FE80::7864:2D04, Tunnel0
D 2007:C15:C0:15::/64 [90/297246976]
  via FE80::7864:2D04, Tunnel0
```

```
SW1(config)# mls qos
SW1(config)# mls qos map policed-dscp 48 46 34 32 24 28 16 10 to 8
SW1(config)# access-list 1 permit any
SW1(config)# class-map POLICE
SW1(config-cmap)# match access-group 1
SW1(config-cmap)# exit
SW1(config)# policy-map RE-MARK
SW1(config-pmap)# class POLICE
SW1(config-pmap-c)# trust dscp
SW1(config-pmap-c)# police 5000000 8000 exceed-action policed-dscp-transmit
SW1(config-pmap-c)# exit
SW1(config-pmap)# exit
SW1(config)# interface range fastethernet 0/1-24
SW1(config-if-range)# service-policy input RE-MARK
```

```
SW1# show policy-map RE-MARK
Policy Map RE-MARK
  Class POLICE
    police 5000000 8000 exceed-action policed-dscp-transmit
    trust dscp
```

```
SW1(config)# mls qos map dscp-mutation AF43-TO-AF42 38 to 36
SW1(config)# interface Gig0/1
SW1(config-if)# mls qos trust dscp
SW1(config-if)# mls qos dscp-mutation AF43-TO-AF42
```

```
R2# sh run class-map
!
class-map match-all VOIP
  match ip dscp ef
class-map match-all BULK-DATA
  match ip dscp af11
class-map match-all NET-MAN
  match ip dscp cs2
class-map match-all VIDEO
  match ip dscp af41
class-map match-all ROUTING
  match ip dscp cs6
class-map match-all SCAVENGER
  match ip dscp cs1
class-map match-all TRANS-DATA
  match ip dscp af21
class-map match-all MISSION-CRIT
  match ip dscp af31
class-map match-all CALL-SIG
  match ip dscp cs3
!
end
```

```
R2# sh run policy-map
!
```

```
policy-map QOS
  class VOIP
    bandwidth percent 16
  class VIDEO
    bandwidth percent 16
  class BULK-DATA
    bandwidth percent 3
    random-detect
  class TRANS-DATA
    bandwidth percent 14
  class NET-MAN
    bandwidth percent 3
  class ROUTING
    bandwidth percent 3
  class SCAVENGER
    bandwidth percent 1
  class MISSION-CRIT
    bandwidth percent 16

  class CALL-SIG
    bandwidth percent 3
  class class-default
    bandwidth percent 25
!
end
```

```
R2# sh run int s0/1 | begin max-reserved-bandwidth 100
max-reserved-bandwidth 100
service-policy output QOS
end

R2# show policy-map QOS
Policy Map QOS
  Class VOIP
    Bandwidth 16 (%) Max Threshold 64 (packets)
  Class VIDEO
    Bandwidth 16 (%) Max Threshold 64 (packets)
  Class BULK-DATA
    Bandwidth 3 (%)
      exponential weight 9
      class      min-threshold      max-threshold      mark-probability
      -----
      0          -                  -                  1/10
      1          -                  -                  1/10
      2          -                  -                  1/10
      3          -                  -                  1/10
      4          -                  -                  1/10
      5          -                  -                  1/10
      6          -                  -                  1/10
      7          -                  -                  1/10
      rsvp       -                  -                  1/10
```

```

Class TRANS-DATA
  Bandwidth 14 (%) Max Threshold 64 (packets)
Class NET-MAN
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class ROUTING
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class SCAVENGER
  Bandwidth 1 (%) Max Threshold 64 (packets)
Class MISSION-CRIT
  Bandwidth 16 (%)

    exponential weight 9

    class      min-threshold      max-threshold      mark-probability
    -----
    0          -                  -                  1/10
    1          -                  -                  1/10
    2          -                  -                  1/10
    3          -                  -                  1/10
    4          -                  -                  1/10
    5          -                  -                  1/10
    6          -                  -                  1/10
    7          -                  -                  1/10
    rsvp       -                  -                  1/10

Class CALL-SIG
  Bandwidth 3 (%) Max Threshold 64 (packets)
Class class-default
  Bandwidth 25 (%)

    exponential weight 9

    class      min-threshold      max-threshold      mark-probability
    -----
    0          -                  -                  1/10
    1          -                  -                  1/10
    2          -                  -                  1/10
    3          -                  -                  1/10
    4          -                  -                  1/10
    5          -                  -                  1/10
    6          -                  -                  1/10
    7          -                  -                  1/10
    rsvp       -                  -                  1/10

```

```
R3(config)# ip nbar custom Hastings_Beer 2 ascii Hastings_Beer udp range 11664 11666
R3(config)# class-map match-all VIRUS
R3(config-cmap)# match protocol Hastings_Beer
R3(config-cmap)# policy-map BLOCK-VIRUS
R3(config-pmap)# class VIRUS
R3(config-pmap-c)# drop
R3(config-pmap-c)# interface gigabit0/0
R3(config-if)# Service-policy input BLOCK-VIRUS
```

```
R2(config)# ip route 192.0.2.1 255.255.255.255 null0
R2(config)# ip route 150.100.2.100 255.255.255.255 Null0 Tag 10
R2(config)# router bgp 10
R2(config-router)# redistribute static route-map BLACKHOLE
R2(config-router)# route-map BLACKHOLE permit 10
R2(config-route-map)# match tag 10
R2(config-route-map)# set ip next-hop 192.0.2.1
R2(config-route-map)# set community no-export
R2(config-route-map)# exit

R2(config)# do show ip bgp neigh 120.100.3.1 advertised
BGP table version is 6, local router ID is 130.100.200.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
Network          Next Hop            Metric LocPrf Weight Path
*> 130.1.1.0/24    0.0.0.0                  0        32768  i
*> 130.100.200.0/24 0.0.0.0                0        32768  i
*> 150.100.2.100/32 192.0.2.1                0        32768  i
Total number of prefixes 3

R2# show ip route 150.100.2.100
Routing entry for 150.100.2.100/32
Known via "static", distance 1, metric 0 (connected)
```

```
Tag 10
Redistributing via bgp 10
Advertised by bgp 10 route-map BLACKHOLE
Routing Descriptor Blocks:
* directly connected, via Null0
Route metric is 0, traffic share count is 1
Route tag 10

R3(config)# ip route 192.0.2.1 255.255.255.255 null0
R3(config)# do show ip bgp
BGP table version is 14, local router ID is 120.100.3.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
Network          Next Hop          Metric LocPrf Weight Path
*>i126.1.1.0/24    120.100.1.1          0     100      0 i
*>i130.1.1.0/24    120.100.2.1          0     100      0 i
*>i130.100.200.0/24 120.100.2.1          0     100      0 i
* i150.100.2.100/32 192.0.2.1          0     100      0 i
```

```
R1(config)# ip route 192.0.2.1 255.255.255.255 null0
R1(config)# interface Gigabit0/1
R1(config-if)# no icmp unreachable
R1(config-if)# do show ip bgp
BGP table version is 8, local router ID is 126.1.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
r RIB-failure, S Stale
```

Origin codes: i - IGP, e - EGP, ? - incomplete

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 126.1.1.0/24	0.0.0.0	0		32768	i
*>i130.1.1.0/24	120.100.2.1	0	100	0	i
*>i130.100.200.0/24	120.100.2.1	0	100	0	i
* i150.100.2.100/32	192.0.2.1	0	100	0	i

R1# show ip route 150.100.2.100

Routing entry for 150.100.2.100/32

Known via "bgp 10", distance 200, metric 0, type internal

Last update from 192.0.2.1 00:00:02 ago

Routing Descriptor Blocks:

* 192.0.2.1, from 120.100.3.1, 00:00:02 ago

Route metric is 0, traffic share count is 1

AS Hops 0

R1# show ip route 192.0.2.1

Routing entry for 192.0.2.1/32

Known via "static", distance 1, metric 0 (connected)

Routing Descriptor Blocks:

* directly connected, via Null0

Route metric is 0, traffic share count is 1

```
R6(config)# ip access-list extended TTL
R6(config-ext-nacl)# deny eigrp any any
R6(config-ext-nacl)# deny tcp any any eq bgp
R6(config-ext-nacl)# deny tcp any eq bgp any
R6(config-ext-nacl)# permit ip any any ttl eq 0 1
R6(config-ext-nacl)# class-map DROP-TTL-0/1
R6(config-cmap)# match access-group name TTL
R6(config-cmap)# policy-map CoPP-TTL
R6(config-pmap)# class DROP-TTL-0/1
R6(config-pmap-c)# drop
R6(config-pmap-c)# control-plane
R6(config-cp)# service-policy input CoPP-TTL
```

```
R3(config)# ip multicast-routing
R3(config)# ntp master
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-mode
R3(config-if)# ntp multicast ttl 2
R3(config-if)# GigabitEthernet0/1

R3(config-if)# ip pim sparse-mode
R3(config-if)# ip pim send-rp-announce GigabitEthernet0/0 scope 2 group-list 4
R3(config)# ip pim send-rp-discovery GigabitEthernet0/0 scope 2
R3(config)# access-list 4 permit 224.0.1.1
```

```
R3# show ntp status
Clock is synchronized, stratum 8, reference is 127.127.7.1
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1E61.2AE19310 (21:17:21.167 UTC Tue Feb 27 2007)
clock offset is 0.0000 msec, root delay is 0.00 msec
root dispersion is 0.02 msec, peer dispersion is 0.02 msec
```

```
R1(config)# ip multicast-routing
R1(config-if)# interface
GigabitEthernet0/0

R1(config-if)# ip pim sparse-mode
R1(config-if)# ntp multicast client
```

```
R1# show ntp status
Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1E79.9FB2321D (21:17:45.623 UTC Tue Feb 27 2007)
clock offset is 0.0157 msec, root delay is 3.88 msec
root dispersion is 0.06 msec, peer dispersion is 0.02 msec
R1(config-if)#
R1# show ip igmp group
IGMP Connected Group Membership
Group Address      Interface          Uptime    Expires   Last Reporter
224.0.1.1          Serial0/0/0       00:40:12  00:02:50  120.100.123.1
224.0.1.39         Serial0/0/0       00:07:21  00:02:51  120.100.123.3
224.0.1.40         Serial0/0/0       00:40:13  00:02:52  120.100.123.1
CHANGE interface
R2(config)# ip multicast-routing
R2(config-if)# interface fastethernet 0/0
R2(config-if)# ip pim sparse-mode
R2(config-if)# ntp multicast client

R2# show ntp status
Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1E73.83B73E68 (21:17:39.514 UTC Tue Feb 27 2007)
clock offset is 0.0182 msec, root delay is 4.14 msec
root dispersion is 15875.06 msec, peer dispersion is 15875.02 msec
R2# show ip igmp group
```

IGMP Connected Group Membership

Group Address	Interface	Uptime	Expires	Last Reporter
224.0.1.1	Serial0/0	00:41:08	00:02:59	120.100.123.2
224.0.1.39	Serial0/0	00:08:12	00:02:57	120.100.123.3
224.0.1.40	Serial0/0	00:41:09	00:01:59	120.100.123.2

Change IF

```
R4(config)# ip multicast-routing
R4(config-if)# interface GigabitEthernet0/0
R4(config-if)# ip pim sparse-mode
R4(config-if)# ntp multicast client
```

```
R4# show ntp status
```

```
Clock is synchronized, stratum 9, reference is 120.100.34.3
nominal freq is 250.0000 Hz, actual freq is 250.0000 Hz, precision is 2**18
reference time is C98F1EF1.2B7DB1F2 (21:19:45.169 UTC Tue Feb 27 2007)
clock offset is -0.6937 msec, root delay is 1.37 msec
root dispersion is 7877.08 msec, peer dispersion is 7876.34 msec
```

```
R4# show ip igmp group
```

IGMP Connected Group Membership

Group Address	Interface	Uptime	Expires	Last Reporter
224.0.1.1	GigabitEthernet0/0	00:41:29	00:02:42	120.100.34.4
224.0.1.39	GigabitEthernet0/0	00:08:35	00:02:42	120.100.34.3
224.0.1.40	GigabitEthernet0/0	00:41:07	00:02:42	120.100.34.4

```
R1(config)# aaa new-model
R1(config)# logging buffered
R1(config)# logging 120.100.99.1
R1(config)#
R1(config)# event manager applet CCIE-QUESTION
R1(config-applet)# event cli pattern "^no (aaa|logging).*" sync no skip yes
R1(config-applet)# action 1.0 syslog msg "UNAUTHORIZED-COMMAND-ENTERED"
R1(config-applet)# action 2.0 cli command "show user"
R1(config-applet)# action 3.0 mail server "120.100.99.2" to "security@lab-exam.net"
   from "eem@lab-exam.net" subject "User-Issue" body "$_cli_result"

R1(config-applet)# no aaa new-model
%HA_EM-6-LOG: CCIE-QUESTION: UNAUTHORISED-COMMAND-ENTERED
%HA_EM-3-FMPD_SMTP_CONNECT: Unable to connect to SMTP server: 120.100.99.2
%HA_EM-3-FMPD_ERROR: Error executing applet CCIE-QUESTION statement 3.0
R1(config)# no logging buffered
%HA_EM-6-LOG: CCIE-QUESTION: UNAUTHORISED-COMMAND-ENTERED
%HA_EM-3-FMPD_SMTP_CONNECT: Unable to connect to SMTP server: 120.100.99.2
%HA_EM-3-FMPD_ERROR: Error executing applet CCIE-QUESTION statement 3.0
R1(config)# do show run | include aaa new-model
aaa new-model
R1(config)# do show run | include logging buffered
logging buffered 4096 debugging
```

```
SW1(config)# spanning-tree mode rapid-pvst
SW1(config)# spanning-tree vlan 34,46,53,63,100,132,200 root primary

SW2(config)# spanning-tree mode rapid-pvst
SW2(config)# spanning-tree vlan 34,46,53,63,100,132,200 root secondary

SW3(config)# spanning-tree mode rapid-pvst

SW4(config)# spanning-tree mode rapid-pvst

SW1# show spanning-tree vlan 34 | include root
      This bridge is the root
SW1# show spanning-tree vlan 46 | include root
      This bridge is the root
SW1# show spanning-tree vlan 53 | include root
      This bridge is the root
SW1# show spanning-tree vlan 63 | include root
      This bridge is the root
SW1# show spanning-tree vlan 100 | include root
      This bridge is the root
SW1# show spanning-tree vlan 132 | include root
      This bridge is the root

SW1# show spanning-tree vlan 200 | include root
      This bridge is the root

SW2# show spanning-tree vlan 34 | include Root FWD
Fa0/23          Root FWD 19          128.25    P2p

SW3# show spanning-tree vlan 34 | include Root FWD
Fa0/19          Root FWD 19          128.21    P2p

SW4# show spanning-tree vlan 34 | include Root FWD
Fa0/21          Root FWD 19          128.23    P2p
```

```
SW3(config)# interface fastethernet 0/19
SW3(config-if)# spanning-tree vlan 34,46,100,132,200 cost 100

SW3(config-if)# do show spanning-tree root
```

Vlan	Root ID	Cost	Root		Hello	Max	Fwd	Root Port
			Time	Age	Dly			
VLAN0001	32769 0013.806d.9400	19	2	20	15			Fa0/19
VLAN0034	24610 0013.806d.9400	38	2	20	15			Fa0/21
VLAN0046	24622 0013.806d.9400	38	2	20	15			Fa0/21
VLAN0053	24629 0013.806d.9400	19	2	20	15			Fa0/19
VLAN0063	24639 0013.806d.9400	19	2	20	15			Fa0/19
VLAN0100	24676 0013.806d.9400	38	2	20	15			Fa0/21
VLAN0132	24676 0013.806d.9400	38	2	20	15			Fa0/21
VLAN0200	24776 0013.806d.9400	38	2	20	15			Fa0/21

```
SW4(config)# interface fastethernet 0/21
SW4(config-if)# spanning-tree vlan 34,46,100,132,200 cost 100
SW4(config-if)# do show spanning-tree root
```

Vlan	Root ID	Root	Hello	Max	Fwd	Root Port
		Cost	Time	Age	Dly	
VLAN0001	32769 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0034	24610 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0046	24622 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0053	24629 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0063	24639 0013.806d.9400	19	2	20	15	Fa0/21
VLAN0100	24676 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0132	24676 0013.806d.9400	38	2	20	15	Fa0/19
VLAN0200	24776 0013.806d.9400	38	2	20	15	Fa0/19

```
SW1(config)# interface fastethernet 0/23
SW1(config-if)# udld port aggressive

SW2(config)# interface fastethernet 0/23
SW2(config-if)# udld port aggressive

SW1# show udld fastethernet 0/23

Interface Fa0/23
---
Port enable administrative configuration setting: Enabled / in aggressive mode
Port enable operational state: Enabled / in aggressive mode
Current bidirectional state: Bidirectional
Current operational state: Advertisement - Single neighbor detected
Message interval: 15
Time out interval: 5

Entry 1
---
Expiration time: 44
Cache Device index: 1
Current neighbor state: Bidirectional
Device ID: CAT0935N2GQ
Port ID: Fa0/23
Neighbor echo 1 device: CAT0911X17K
Neighbor echo 1 port: Fa0/23

Message interval: 15
Time out interval: 5
CDP Device name: SW2
```

```
SW1(config)# vlan 30
SW1(config-vlan)# exit
SW1(config)# interface fastethernet 0/18
SW1(config-if)# switchport access vlan 30
SW1(config-if)# switchport mode dot1q-tunnel
```

```
SW2(config)# vlan 30
SW2(config-vlan)# exit
SW2(config)# interface fastethernet 0/18
SW2(config-if)# switchport access vlan 30
SW2(config-if)# switchport mode dot1q-tunnel
```

```
SW2(config)# vlan 20
SW2(config-vlan)# remote-span
SW2(config-vlan)# exit
SW2(config)# monitor session 1 source interface fastethernet 0/1 tx
SW2(config)# monitor session 1 destination remote vlan 20
SW2(config)# do show monitor session 1
Session 1
-----
Type : Remote Source Session
Source Ports :
    TX Only : Fa0/1
Dest RSPAN VLAN : 20

SW3(config)# vlan 20
SW3(config-vlan)# exit
SW3(config)# monitor session 1 source remote vlan 20
SW3(config)# monitor session 1 destination interface fast 0/17
SW3(config)# do show monitor session 1
Session 1
-----
Type : Remote Destination Session
Source RSPAN VLAN : 20
Destination Ports : Fa0/17

Encapsulation : Native
Ingress : Disabled
```

```
SW2(config)# link state track 1
SW2(config)# interface fast0/5
SW2(config-if)# link state group 1 downstream
SW2(config-if)# interface fastethernet 0/19
SW2(config-if)# link state group 1 upstream
SW2(config-if)# interface fastethernet 0/21
SW2(config-if)# link state group 1 upstream
SW2(config-if)# interface fastethernet 0/23
SW2(config-if)# link state group 1 upstream
```

```
SW2# show interface fastethernet 0/5 | include connected
fastethernet0/5 is up, line protocol is up (connected)
```

```
SW2(config-if)# int fast 0/19
SW2(config-if)# shut
SW2(config-if)# int fast 0/21
SW2(config-if)# shut
SW2(config-if)# int fast 0/23
SW2(config-if)# shut
```

```
SW2# show interface fastethernet 0/5 | include err-disabled
fastethernet0/5 is down, line protocol is down (err-disabled)
```

```
R1(config)# router eigrp CCIE

R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# net 120.100.1.0 0.0.0.255
R1(config-router-af)# net 120.100.123.0 0.0.0.255
R1(config-router-af)# net 120.100.100.0 0.0.0.255

R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router-af)# network 120.100.2.0 0.0.0.255
R2(config-router-af)# network 120.100.123.0 0.0.0.255
R2(config-router-af)# network 120.100.200.0 0.0.0.255

R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# network 120.100.3.0 0.0.0.255
R3(config-router-af)# network 120.100.123.0 0.0.0.255
R3(config-router-af)# network 120.100.34.0 0.0.0.255

R4(config-router)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# network 120.100.4.0 0.0.0.255
R4(config-router-af)# network 120.100.34.0 0.0.0.255

R5(config)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# network 120.100.5.0 0.0.0.255
R5(config-router-af)# network 120.100.34.0 0.0.0.255

R1# sh ip route eigrp
    120.0.0.0/24 is subnetted, 9 subnets
D      120.100.4.0 [90/158720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D      120.100.5.0 [90/158720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D      120.100.2.0 [90/156160] via 120.100.123.2, 00:23:32, GigabitEthernet0/0
D      120.100.3.0 [90/156160] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D      120.100.34.0 [90/30720] via 120.100.123.3, 00:23:32, GigabitEthernet0/0
D      120.100.200.0 [90/30720] via 120.100.123.2, 00:23:32, GigabitEthernet0/0
```

```
R1(config)# route-map LEAK-VLAN-100-LOOP0 permit 10
R1(config-route-map)# match ip address 1
R1(config-route-map)# exit
R1(config)# access-list 1 permit 120.100.100.0
R1(config)# access-list 1 permit 120.100.1.0
R1(config)# router eigrp CCIE
R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# af-interface Gigabit0/0
R1(config-router-af-interface)# summary-address 120.100.0.0 255.255.0.0 leak-map
LEAK-VLAN-100-LOOP0
```

```
R3# show ip route eigrp
```

```
R3# show ip route eigrp
```

```
    120.0.0.0/8 is variably subnetted, 10 subnets, 2 masks
D        120.100.4.0/24 [90/156160] via 120.100.34.4, 00:23:32, GigabitEthernet0/0
D        120.100.5.0/24 [90/156160] via 120.100.34.5, 00:23:32, GigabitEthernet0/0
D        120.100.0.0/16 [90/30720] via 120.100.123.1, 00:00:53, GigabitEthernet0/1
D        120.100.1.0/24 [90/156160] via 120.100.123.1, 00:23:32, GigabitEthernet0/1
D        120.100.2.0/24 [90/156160] via 120.100.123.2, 00:23:32, GigabitEthernet0/1
D        120.100.100.0/24 [90/30720] via 120.100.123.1, 00:23:32, GigabitEthernet0/1
D        120.100.200.0/24 [90/30720] via 120.100.123.2, 00:23:32, GigabitEthernet0/1
```

```
R1(config)# interface GigabitEthernet0/0
R1(config-if)# ip hold-time eigrp 1 200
R1(config-if)
```

Enter configuration commands, one per line. End with CNTL/Z.

```
R2(config)# interface fastethernet0/0
R2(config-if)# ip hold-time eigrp 1 200
R2(config-if)
```

```
R3(config)# interface GigabitEthernet0/0
```

```
R3(config-if)# ip hold-time eigrp 1 200
```

```
R3(config-if)# do sh ip eigrp neighbors
```

IP-EIGRP neighbors for process 1

H	Address	Interface	Hold	Uptime	SRTT	RTO	Q	Seq
			(sec)	(ms)		Cnt	Num	
3	120.100.123.1	Gi0/1	198	00:00:57	3	200	0	25
2	120.100.123.2	Gi0/1	199	00:01:00	3	200	0	18
1	120.100.34.5	Gi0/0	12	00:23:32	1	200	0	21
0	120.100.34.4	Gi0/0	12	00:23:35	35	210	0	22

```
R1(config)# interface Loopback2
R1(config-if)# ip address 150.101.1.1 255.255.255.0
R1(config-if)# router eigrp CCIE
R1(config-router)# address-family ipv4 unicast autonomous-system 1
R1(config-router-af)# net 150.101.1.0 0.0.0.255
```

```
R2(config)# interface Loopback2
R2(config-if)# ip address 150.101.1.1 255.255.255.0
R2(config-if)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router)# net 150.101.1.0 0.0.0.255
```

```
R3# show ip route 150.101.1.0
Routing entry for 150.101.1.0/24
```

Known via "eigrp 1", distance 90, metric 156160, type internal
Redistributing via eigrp 1

Last update from 120.100.123.2 on fastethernet1/1, 00:00:23 ago

Routing Descriptor Blocks:

```
120.100.123.2, from 120.100.123.2, 00:00:23 ago, via fastethernet1/1
  Route metric is 156160, traffic share count is 1
  Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
  Reliability 255/255, minimum MTU 1500 bytes
```

```
    Loading 1/255, Hops 1
* 120.100.123.1, from 120.100.123.1, 00:00:23 ago, via fastethernet1/1
  Route metric is 156160, traffic share count is 1
  Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
  Reliability 255/255, minimum MTU 1500 bytes
  Loading 1/255, Hops 1
```

```
R3# show ip eigrp topology 150.101.1.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 150.101.1.0/24
  State is Passive, Query origin flag is 1, 2 Successor(s), FD is 156160
  Routing Descriptor Blocks:
    120.100.123.1 (GigabitEthernet0/1), from 120.100.123.1, Send flag is 0x0
      Composite metric is (156160/128256), Route is Internal
      Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
    120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
      Composite metric is (156160/128256), Route is Internal
      Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
```

```
Reliability is 255/255
Load is 1/255
Minimum MTU is 1500
Hop count is 1
```

```
R2# show interface Fast0/0 | include DLY
MTU 1500 bytes, BW 100000 Kbit/sec, DLY 100 usec,
R2(config)# access-list 1 permit 150.101.1.0
R2(config)# router eigrp CCIE
R2(config-router)# address-family ipv4 unicast autonomous-system 1
R2(config-router-af)# topology base
R2(config-router-af)# offset-list 1 out 100 fastethernet0/0

R3# show ip route 150.101.1.0
Routing entry for 150.101.1.0/24
Known via "eigrp 1", distance 90, metric 156160, type internal
Redistributing via eigrp 1
Last update from 120.100.123.1 on GigabitEthernet0/1, 00:00:17 ago
Routing Descriptor Blocks:
* 120.100.123.1, from 120.100.123.1, 00:00:17 ago, via GigabitEthernet0/1
    Route metric is 156160, traffic share count is 1
    Total delay is 5100 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R3# show ip eigrp topology 150.101.1.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 150.101.1.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
Routing Descriptor Blocks:
120.100.123.1 (GigabitEthernet0/1), from 120.100.123.1, Send flag is 0x0
    Composite metric is (156160/128256), Route is Internal
    Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
    Composite metric is (156260/128356), Route is Internal
    Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5103 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
```

```
R2(config)# interface Loopback3
R2(config-if)# ip add 150.101.2.1 255.255.255.0
R2(config-if)# router eigrp CCIE2
R2(config-router)# address-family ipv4 unicast autonomous-system 2
R2(config-router-af)# af-interface default
R2(config-router-af-interface)# authentication mode hmac-sha-256 0 lake2aho3
R2(config-router-af-interface)# exit
R2(config-router-af)# network 150.101.2.0 0.0.0.255
R2(config-router-af)# network 120.100.123.0 0.0.0.255

R3(config)# router eigrp CCIE2
R3(config-router)# address-family ipv4 unicast autonomous-system 2
R3(config-router-af)# af-interface default
R3(config-router-af-interface)# authentication mode hmac-sha-256 0 lake2aho3
R3(config-router-af-interface)# exit
R3(config-router-af)# network 120.100.123.0
R3(config-router-af)# sh ip route eigrp 2
    150.101.0.0/24 is subnetted, 2 subnets
D        150.101.2.0 [90/156160] via 120.100.123.2, 00:00:25, GigabitEthernet0/1
```

```
R4(config)# interface Loopback 0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# exit
R4(config)# interface GigabitEthernet 0/1
R4(config-if)# ip ospf 1 area 1

R5(config)# interface Loopback 0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# exit
R5(config)# interface GigabitEthernet 0/1
R5(config-if)# ip ospf 1 area 2

R6(config)# interface Loopback 0
R6(config-if)# ip ospf 1 area 1
R6(config-if)# interface GigabitEthernet 0/0
R6(config-if)# ip ospf 1 area 1
R6(config-if)# interface GigabitEthernet 0/1
R6(config-if)# ip ospf 1 area 3

SW1(config)# ip routing
SW1(config)# router ospf 1
SW1(config-router)# network 120.100.7.1 0.0.0.0 area 2
SW1(config-router)# network 120.100.53.1 0.0.0.0 area 2

SW2(config)# ip routing
SW2(config-if)# router ospf 1
SW2(config-router)# net 120.100.8.1 0.0.0.0 area 1
SW2(config-router)# net 120.100.46.2 0.0.0.0 area 1

SW3(config)# ip routing
SW3(config)# router ospf 1
SW3(config-router)# network 120.100.53.3 0.0.0.0 area 2
SW3(config-router)# network 120.100.63.3 0.0.0.0 area 3
SW3(config-router)# network 120.100.9.1 0.0.0.0 area 2

SW4(config)# ip routing
SW4(config)# router ospf 1
SW4(config-router)# network 120.100.10.1 0.0.0.0 area 3
SW4(config-router)# network 120.100.63.4 0.0.0.0 area 3
```

```
R5(config)# router ospf 1
R5(config-router)# area 2 virtual-link 120.100.9.1

SW3(config-router)# router ospf 1
SW3(config-router)# area 2 virtual-link 120.100.5.1

R4(config)# router ospf 1
R4(config-router)# area 1 virtual-link 120.100.6.1

R6(config-if)# router ospf 1
R6(config-router)# area 1 virtual-link 120.100.4.1
R6(config-router)# area 3 virtual-link 120.100.9.1

SW3(config-if)# router ospf 1
SW3(config-router)# area 3 virtual-link 120.100.6.1

SW4# sh ip route ospf
    120.0.0.0/8 is variably subnetted, 10 subnets, 2 masks
O IA    120.100.9.1/32 [110/2] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.8.1/32 [110/3] via 120.100.63.6, 00:00:54, Vlan63
O IA    120.100.5.1/32 [110/3] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.4.1/32 [110/3] via 120.100.63.6, 00:00:54, Vlan63
O IA    120.100.7.1/32 [110/3] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.6.1/32 [110/2] via 120.100.63.6, 00:00:54, Vlan63

O IA    120.100.53.0/24 [110/2] via 120.100.63.3, 00:00:54, Vlan63
O IA    120.100.46.0/24 [110/2] via 120.100.63.6, 00:00:55, Vlan63
```

```
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# topology base
R3(config-router-af-topology)# redistribute eigrp 2
```

```
R1# show ip route 150.101.2.0
% Subnet not in table
```

```
R3# show ip eigrp topology 150.101.2.0/24
IP-EIGRP (AS 1): Topology entry for 150.101.2.0/24
  State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
  Routing Descriptor Blocks:
    120.100.123.2, from Redistributed, Send flag is 0x0
      Composite metric is (156160/0), Route is External
      Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
```

```
External data:  
    Originating router is 200.200.200.200 (this system)  
    AS number of route is 2  
    External protocol is EIGRP, external metric is 156160  
    Administrator tag is 0 (0x00000000)  
IP-EIGRP (AS 2): Topology entry for 150.101.2.0/24  
    State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160  
    Routing Descriptor Blocks:  
        120.100.123.2 (GigaEthernet0/1), from 120.100.123.2, Send flag is 0x0  
            Composite metric is (156160/128256), Route is Internal  
            Vector metric:  
                Minimum bandwidth is 100000 Kbit  
                Total delay is 5100 microseconds  
                Reliability is 255/255  
                Load is 1/255  
                Minimum MTU is 1500  
                Hop count is 1  
  
R3# show ip eigrp topology | include ID  
IP-EIGRP Topology Table for AS(1)/ID(200.200.200.200)  
IP-EIGRP Topology Table for AS(2)/ID(200.200.200.200)  
  
R1# show ip eigrp topology | include ID  
IP-EIGRP Topology Table for AS(1)/ID(200.200.200.200)  
R1#
```

```
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# eigrp router-id 120.100.3.1

R3# show ip eigrp topology | include ID
IP-EIGRP Topology Table for AS(1)/ID(120.100.3.1)

R3# show ip eigrp topology 150.101.2.0/24
IP-EIGRP (AS 1): Topology entry for 150.101.2.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
Routing Descriptor Blocks:
120.100.123.2, from Redistributed, Send flag is 0x0
    Composite metric is (156160/0), Route is External
    Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
    External data:
        Originating router is 120.100.3.1 (this system)
        AS number of route is 2
        External protocol is EIGRP, external metric is 156160
        Administrator tag is 0 (0x00000000)
IP-EIGRP (AS 2): Topology entry for 150.101.2.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 156160
Routing Descriptor Blocks:
120.100.123.2 (GigabitEthernet0/1), from 120.100.123.2, Send flag is 0x0
    Composite metric is (156160/128256), Route is Internal
    Vector metric:
        Minimum bandwidth is 100000 Kbit
        Total delay is 5100 microseconds
        Reliability is 255/255
        Load is 1/255
        Minimum MTU is 1500
        Hop count is 1
R1# show ip route 150.101.2.0
Routing entry for 150.101.2.0/24
Known via "eigrp 1", distance 170, metric 158720, type external
Redistributing via eigrp 1
Last update from 120.100.123.3 on GigabitEthernet0/0, 00:03:06 ago
Routing Descriptor Blocks:
```

```
* 120.100.123.3, from 120.100.123.3, 00:03:06 ago, via GigabitEthernet0/0
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# af-interface GigabitEthernet0/1
R3(config-router-af-interface)# no next-hop-self
```

```
R1# show ip route 150.101.2.0
Routing entry for 150.101.2.0/24
Known via "eigrp 1", distance 170, metric 158720, type external
Redistributing via eigrp 1
Last update from 120.100.123.2 on Gigabit0/0, 00:00:24 ago
Routing Descriptor Blocks:
* 120.100.123.2, from 120.100.123.3, 00:00:24 ago, via GigabitEthernet0/0
    Route metric is 158720, traffic share count is 1
    Total delay is 5200 microseconds, minimum bandwidth is 100000 Kbit
Reliability 255/255, minimum MTU 1500 bytes
Loading 1/255, Hops 1
```

```
R4(config-router)# router ospf 1
R4(config-router)# redistribute eigrp 1 subnets
R4(config-router)# default-metric 5000
R4(config-router)# router eigrp CCIE
R4(config-router)# address-family ipv4 unicast autonomous-system 1
R4(config-router-af)# topology base
R4(config-router-af-topology)# redistribute ospf 1
R4(config-router-af-topology)# default-metric 1544 20000 255 1 1500
```

```
R5(config-router)# router ospf 1
R5(config-router)# redistribute eigrp 1 subnets
R5(config-router)# default-metric 5000
R5(config-router)# router eigrp CCIE
R5(config-router)# address-family ipv4 unicast autonomous-system 1
R5(config-router-af)# topology base
R5(config-router-af-topology)# redistribute ospf 1
R5(config-router-af-topology)# default-metric 1544 20000 255 1 1500
R3# show ip route eigrp
    150.101.0.0/24 is subnetted, 2 subnets
D      150.101.1.0 [90/2297856] via 120.100.123.1, 00:05:05, Gigabit0/1
        120.0.0.0/8 is variably subnetted, 20 subnets, 3 masks
D EX    120.100.9.1/32
        [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
        [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.8.1/32
        [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
        [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.10.1/32
        [170/6780416] via 120.100.34.5, 00:00:22, GigabitEthernet0/0
        [170/6780416] via 120.100.34.4, 00:00:22, GigabitEthernet0/0
D EX    120.100.5.1/32
        [170/6780416] via 120.100.34.4, 00:01:51, GigabitEthernet0/0
D      120.100.4.0/24
        [90/156160] via 120.100.34.4, 00:07:17, GigabitEthernet0/0
D      120.100.5.0/24
        [90/156160] via 120.100.34.5, 00:07:17, GigabitEthernet0/0
D EX    120.100.4.1/32
        [170/6780416] via 120.100.34.5, 00:00:23, GigabitEthernet0/0
```

D EX 120.100.7.1/32
[170/6780416] via 120.100.34.5, 00:00:23, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:23, GigabitEthernet0/0

D EX 120.100.6.1/32
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0

D 120.100.0.0/16 [90/2172416] via 120.100.123.1, 00:05:07, GigabitEthernet0/1

D 120.100.1.0/24 [90/2297856] via 120.100.123.1, 00:05:07, GigabitEthernet0/1

D 120.100.2.0/24 [90/2297856] via 120.100.123.2, 00:05:07, GigabitEthernet0/1

D EX 120.100.63.0/24
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0

D EX 120.100.53.0/24
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0

D EX 120.100.46.0/24
[170/6780416] via 120.100.34.5, 00:00:24, GigabitEthernet0/0
[170/6780416] via 120.100.34.4, 00:00:24, GigabitEthernet0/0

D 120.100.100.0/24 [90/2172416] via 120.100.123.1, 00:05:07,
GigabitEthernet0/1

D 120.100.200.0/24 [90/2172416] via 120.100.123.2, 00:05:08,
GigabitEthernet0/1

```
R3# show ip route 120.100.63.0
Routing entry for 120.100.63.0/24
Known via "eigrp 1", distance 170, metric 6780416, type external
Redistributing via eigrp 1
Last update from 120.100.34.5 on GigabitEthernet0/0, 00:01:59 ago
Routing Descriptor Blocks:
  120.100.34.5, from 120.100.34.5, 00:01:59 ago, via GigabitEthernet0/0
    Route metric is 6780416, traffic share count is 1
    Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
  * 120.100.34.4, from 120.100.34.4, 00:01:59 ago, via GigabitEthernet0/0
    Route metric is 6780416, traffic share count is 1
    Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1
```

```
R3(config)# access-list 1 permit 120.100.34.4
R3(config)# access-list 2 permit 120.100.63.0
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv4 unicast autonomous-system 1
R3(config-router-af)# topology base
R3(config-router-af-topology)# distribute-list route-map PENALISE-VLAN63 in
  GigabitEthernet0/0
R3(config-router-af-topology)# exit
R3(config-router-af)# exit
R3(config-router)# exit
R3(config)# route-map PENALISE-VLAN63 permit 10
R3(config-route-map)# match ip address 2
R3(config-route-map)# match ip route-source 1
R3(config-route-map)# set metric +500000
R3(config-route-map)# route-map PENALISE-VLAN63 permit 20
```

```
R3# show ip route 120.100.63.0
Routing entry for 120.100.63.0/24
```

```
Known via "eigrp 1", distance 170, metric 6780416, type external
Redistributing via eigrp 1
Last update from 120.100.34.5 on GigabitEthernet0/0, 00:00:21 ago
Routing Descriptor Blocks:
* 120.100.34.5, from 120.100.34.5, 00:00:21 ago, via GigabitEthernet0/0
    Route metric is 6780416, traffic share count is 1
    Total delay is 200100 microseconds, minimum bandwidth is 1544 Kbit
    Reliability 255/255, minimum MTU 1500 bytes
    Loading 1/255, Hops 1

R3# show ip eigrp topology 120.100.63.0 255.255.255.0
IP-EIGRP (AS 1): Topology entry for 120.100.63.0/24
State is Passive, Query origin flag is 1, 1 Successor(s), FD is 6780416
Routing Descriptor Blocks:
120.100.34.5 (GigabitEthernet0/0), from 120.100.34.5, Send flag is 0x0
    Composite metric is (6780416/6777856), Route is External
    Vector metric:
        Minimum bandwidth is 1544 Kbit
Total delay is 200100 microseconds
    Reliability is 255/255
    Load is 1/255
    Minimum MTU is 1500
    Hop count is 1
    External data:
        Originating router is 120.100.5.1
        AS number of route is 1
        External protocol is OSPF, external metric is 2
        Administrator tag is 0 (0x00000000)
120.100.34.4 (GigabitEthernet0/0), from 120.100.34.4, Send flag is 0x0
    Composite metric is (128000000/6777856), Route is External
    Vector metric:
        Minimum bandwidth is 20 Kbit
        Total delay is 0 microseconds
        Reliability is 0/255
        Load is 0/255
        Minimum MTU is 0
        Hop count is 1
    External data:
        Originating router is 120.100.4.1
        AS number of route is 1
        External protocol is OSPF, external metric is 2
        Administrator tag is 0 (0x00000000)
```

```
R1(config)# router bgp 100
R1(config-router)# no auto-summary
R1(config-router)# no synchronization
R1(config-router)# neighbor 120.100.3.1 remote-as 100
R1(config-router)# neighbor 120.100.3.1 update-source Loopback0

R2(config)# router bgp 100
R2(config-router)# no auto-summary
R2(config-router)# no synchronization
R2(config-router)# neighbor 120.100.3.1 remote-as 100
R2(config-router)# neighbor 120.100.3.1 update-source Loopback0

R3(config)# router bgp 100
R3(config-router)# no auto-summary
R3(config-router)# no synchronization
R3(config-router)# neighbor AS100 peer-group
R3(config-router)# neighbor AS100 remote-as 100
R3(config-router)# neighbor AS100 update-source Loopback0
R3(config-router)# neighbor 120.100.1.1 peer-group AS100
R3(config-router)# neighbor 120.100.2.1 peer-group AS100
R3(config-router)# neighbor AS100 route-reflector-client
R3(config-router)# neighbor 120.100.34.4 remote-as 200
R3(config-router)# neighbor 120.100.34.5 remote-as 300

R4(config)# router bgp 200
R4(config-router)# router bgp 200
R4(config-router)# no auto-summary
R4(config-router)# no synchronization
R4(config-router)# neighbor AS200 peer-group
R4(config-router)# neighbor AS200 remote-as 200
R4(config-router)# neighbor AS200 update-source Loopback0
R4(config-router)# neighbor AS200 route-reflector-client
R4(config-router)# neighbor 120.100.6.1 peer-group AS200
R4(config-router)# neighbor 120.100.8.1 peer-group AS200
R4(config-router)# neighbor 120.100.34.3 remote-as 100

R5(config)# router bgp 300
R5(config-router)# no auto-summary
R5(config-router)# no synchronization
```

```
R5(config-router)# neighbor AS300 peer-group
R5(config-router)# neighbor AS300 remote-as 300
R5(config-router)# neighbor AS300 update-source Loopback0
R5(config-router)# neighbor AS300 route-reflector-client
R5(config-router)# neighbor 120.100.7.1 peer-group AS300
R5(config-router)# neighbor 120.100.9.1 peer-group AS300
R5(config-router)# neighbor 120.100.34.3 remote-as 100

R6(config)# router bgp 200
R6(config-router)# no auto-summary
R6(config-router)# no synchronization
R6(config-router)# neighbor 120.100.4.1 remote-as 200
R6(config-router)# neighbor 120.100.4.1 update-source Loopback0
R6(config-router)# neighbor 120.100.10.1 remote-as 400
R6(config-router)# neighbor 120.100.10.1 update-source Loopback0

SW1(config)# router bgp 300
SW1(config-router)# no auto-summary
SW1(config-router)# no synchronization
SW1(config-router)# neighbor 120.100.5.1 remote-as 300
SW1(config-router)# neighbor 120.100.5.1 update-source Loopback0

SW2(config)# router bgp 200
SW2(config-router)# no auto-summary
SW2(config-router)# no synchronization
SW2(config-router)# neighbor 120.100.4.1 remote-as 200
SW2(config-router)# neighbor 120.100.4.1 update-source Loopback0

SW3(config)# router bgp 300
SW3(config-router)# no auto-summary
SW3(config-router)# no synchronization
SW3(config-router)# neighbor 120.100.5.1 remote-as 300
SW3(config-router)# neighbor 120.100.5.1 update-source Loopback0
SW3(config-router)# neighbor 120.100.10.1 remote-as 400
SW3(config-router)# neighbor 120.100.10.1 update-source Loopback0
```

```
SW4(config)# router bgp 400
SW4(config-router)# no auto-summary
SW4(config-router)# no synchronization
SW4(config-router)# neighbor 120.100.6.1 remote-as 200
SW4(config-router)# neighbor 120.100.6.1 update-source Loopback0
SW4(config-router)# neighbor 120.100.9.1 remote-as 300
SW4(config-router)# neighbor 120.100.9.1 update-source Loopback0

SW4# sh ip bgp neigh 120.100.6.1 | include External
External BGP neighbor not directly connected.
SW4# show ip bgp neighbors 120.100.9.1 | include External
External BGP neighbor not directly connected.
SW4#

SW4# sh ip bgp neighbors 120.100.6.1 | include active
No active TCP connection
SW4# sh ip bgp neighbors 120.100.9.1 | include active
No active TCP connection

SW4(config-router)# neighbor 120.100.6.1 disable-connected-check
SW4(config-router)# neighbor 120.100.9.1 disable-connected-check

R6(config-router)# neighbor 120.100.10.1 disable-connected-check

SW3(config-router)# neighbor 120.100.10.1 disable-connected-check

SW4# show ip bgp neighbors 120.100.6.1 | include Established
BGP state = Established, up for 00:02:01
SW4# show ip bgp neighbors 120.100.9.1 | include Established
BGP state = Established, up for 00:02:05
```

```
R1# * 19:30:13.287: %BGP-3-NOTIFICATION: sent to neighbor 120.100.3.1 2/3 (BGP
identifier wrong) 4 bytes C8C8C8C8

R3# * 19:25:30.043: %BGP-3-NOTIFICATION: received from neighbor 120.100.1.1 2/
3 (BGP identifier wrong) 4 bytes C8C8C8C8

R1# show ip bgp summary | include identifier
BGP router identifier 200.200.200.200, local AS number 100

R3# show ip bgp summary | include identifier
BGP router identifier 200.200.200.200, local AS number 100

R1(config-router)# bgp router-id 120.100.1.1
*19:34:45.467: %BGP-5-ADJCHANGE: neighbor 120.100.3.1 Up
```

```
R1(config)# router bgp 100
R1(config-router)# neighbor 120.100.3.1 transport connection-mode passive
R1(config-router)# neighbor 120.100.3.1 activate

R2(config)# router bgp 100
R2(config-router)# neighbor 120.100.3.1 transport connection-mode passive
R2(config-router)# neighbor 120.100.3.1 activate

R3(config)# router bgp 100
R3(config-router)# neighbor AS100 transport connection-mode active
R3(config-router)# neighbor 120.100.1.1 activate
R3(config-router)# neighbor 120.100.2.1 activate
```

```
R3(config)# interface Loopback5
R3(config-if)# ip address 152.100.100.1 255.255.255.0
R3(config-if)# router bgp 100
R3(config-router)# network 152.100.100.0 mask 255.255.255.0
```

```
SW4(config)# interface Loopback5
SW4(config-if)# ip address 152.200.32.1 255.255.255.0
SW4(config-if)# interface Loopback6
SW4(config-if)# ip address 152.200.33.1 255.255.255.0
SW4(config-if)# interface Loopback7
SW4(config-if)# ip address 152.200.34.1 255.255.255.0
SW4(config-if)# interface Loopback8
SW4(config-if)# ip address 152.200.35.1 255.255.255.0
SW4(config-if)# router bgp 400
SW4(config-router)# network 152.200.32.0 mask 255.255.255.0
SW4(config-router)# network 152.200.33.0 mask 255.255.255.0
SW4(config-router)# network 152.200.34.0 mask 255.255.255.0
SW4(config-router)# network 152.200.35.0 mask 255.255.255.0
```

```
R3# show ip bgp
BGP table version is 10, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 152.100.100.0/24	0.0.0.0	0		32768	i
* 152.200.32.0/24	120.100.34.4		0	200	400 i
*>	120.100.34.5		0	300	400 i
* 152.200.33.0/24	120.100.34.4		0	200	400 i
*>	120.100.34.5		0	300	400 i
* 152.200.34.0/24	120.100.34.4		0	200	400 i
*>	120.100.34.5		0	300	400 i
* 152.200.35.0/24	120.100.34.4		0	200	400 i
*>	120.100.34.5		0	300	400 i

```

R3(config)# router bgp 100
R3(config-router)# neighbor 120.100.34.4 capability orf prefix-list send
R3(config-router)# neighbor 120.100.34.4 prefix-list FILTER in
R3(config)# ip prefix-list FILTER seq 5 deny 152.200.33.0/24
R3(config)# ip prefix-list FILTER seq 10 deny 152.200.34.0/24
R3(config)# ip prefix-list FILTER seq 15 deny 152.200.35.0/24
R3(config)# ip prefix-list FILTER seq 20 permit 0.0.0.0/0 le 32

R4(config)# router bgp 200
R4(config-router)# neighbor 120.100.34.3 capability orf prefix-list receive
R4(config-router)# exit
R4(config)# exit
R4# show ip bgp neighbors 120.100.34.3 advertised-routes
BGP table version is 17, local router ID is 120.100.4.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*->i152.200.32.0/24  120.100.10.1        0       100      0 400 i

Total number of prefixes 1

R3# clear ip bgp *
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*> 152.100.100.0/24  0.0.0.0            0       32768 i
*> 152.200.32.0/24   120.100.34.4        0       200 400 i
*               120.100.34.5        0       300 400 i
*> 152.200.33.0/24   120.100.34.5        0       300 400 i
*> 152.200.34.0/24   120.100.34.5        0       300 400 i
*> 152.200.35.0/24   120.100.34.5        0       300 400 i

```

```
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 152.100.100.0/24	0.0.0.0	0		32768	i
*> 152.200.32.0/24	120.100.34.4			0	200 400 i
*	120.100.34.5			0	300 400 i
*> 152.200.33.0/24	120.100.34.5			0	300 400 i
*> 152.200.34.0/24	120.100.34.5			0	300 400 i
*> 152.200.35.0/24	120.100.34.5			0	300 400 i

```
R5(config)# router bgp 300
R5(config-router)# neighbor 120.100.34.3 route-map PREPEND out
R5(config-router)# exit
R5(config)# access-list 1 permit 152.200.32.0
R5(config)# route-map PREPEND permit 10
R5(config-route-map)# match ip address 1
R5(config-route-map)# set as-path prepend 300
R5(config-route-map)# route-map PREPEND permit 20
R5(config-route-map)# match ip address 1
R5(config-route-map)# set as-path prepend 300
R5(config-route-map)# route-map PREPEND permit 30
```

```
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 152.100.100.0/24	0.0.0.0	0		32768	i
*> 152.200.32.0/24	120.100.34.4			0	200 400 i
*	120.100.34.5			0	300 300 400 i
*> 152.200.33.0/24	120.100.34.5			0	300 400 i
*> 152.200.34.0/24	120.100.34.5			0	300 400 i
*> 152.200.35.0/24	120.100.34.5			0	300 400 i

```
R5(config)# route-map PREPEND permit 10
R5(config-route-map)# continue 20

R3# clear ip bgp *
R3# show ip bgp
BGP table version is 6, local router ID is 200.200.200.200
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*> 152.100.100.0/24 0.0.0.0                  0        32768 i
*> 152.200.32.0/24 120.100.34.4            0 200 400 i
*          120.100.34.5            0 300 300 300 400 i
*> 152.200.33.0/24 120.100.34.5            0 300 400 i
*> 152.200.34.0/24 120.100.34.5            0 300 400 i
*> 152.200.35.0/24 120.100.34.5            0 300 400 i
```

```
R1(config)# ipv6 unicast-routing
R1(config)# interface GigabitEthernet0/1
R1(config-if)# ipv6 address 2007:C15:C0:10::1/64
R1(config-if)# interface tunnel0
R1(config-if)# ipv6 address 2007:C15:C0:11::1/64

R2(config)# ipv6 unicast-routing
R2(config)# interface fastethernet 0/1
R2(config-if)# ipv6 address 2007:C15:C0:13::2/64
R2(config-if)# interface tunnel0
R2(config-if)# ipv6 address 2007:C15:C0:12::2/64

R3(config)# ipv6 unicast-routing
R3(config)# int GigabitEthernet0/0
R3(config-if)# ipv6 address 2007:C15:C0:14::3/64
R3(config-if)# interface tunnel0
R3(config-if)# ipv6 address 2007:C15:C0:11::3/64
R3(config-if)# interface tunnell1
R3(config-if)# ipv6 address 2007:C15:C0:12::3/64

R4(config)# ipv6 unicast-routing
R4(config)# interface GigabitEthernet0/0
R4(config-if)# ipv6 address 2007:C15:C0:14::4/64
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ipv6 address 2007:C15:C0:15::4/64
R5(config)# ipv6 unicast-routing
R5(config)# interface GigabitEthernet0/0
R5(config-if)# ipv6 address 2007:C15:C0:14::5/64

R6(config)# ipv6 unicast-routing
R6(config)# interface GigabitEthernet0/0
R6(config-if)# ipv6 address 2007:C15:C0:15::6/64
```

```
R1(config-if)# interface Tunnel0
R1(config-if)# tunnel source Gigabit0/0
R1(config-if)# tunnel destination 120.100.123.3
R1(config-if)# tunnel mode ipv6ip
R1(config-if)# router eigrp CCIE
R1(config-router)# address-family ipv6 unicast autonomous-system 6
R1(config-router-af)# af-interface Tunnel0
R1(config-router-af-interface)# no shutdown
R1(config-router-af-interface)# af-interface Gigabit0/1
R1(config-router-af-interface)# no shutdown

R2(config-if)# interface Tunnel0
R2(config-if)# tunnel source fastethernet0/0
R2(config-if)# tunnel destination 120.100.123.3
R2(config-if)# tunnel mode ipv6ip
R2(config-if)# router eigrp CCIE
R2(config-router)# address-family ipv6 unicast autonomous-system 6
R2(config-router-af)# af-interface Tunnel0
R2(config-router-af-interface)# no shutdown
R2(config-router-af-interface)# af-interface fastethernet0/1
R2(config-router-af-interface)# no shutdown

R3(config-if)# tunnel source Gigabit0/1
R3(config-if)# tunnel destination 120.100.123.1
R3(config-if)# tunnel mode ipv6ip
R3(config-if)# interface Tunnel1

R3(config-if)# tunnel source Gigabit0/1
R3(config-if)# tunnel destination 120.100.123.2
R3(config-if)# tunnel mode ipv6ip
R3(config-if)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 6
R3(config-router-af)# af-interface GigabitEthernet0/0
R3(config-router-af-interface)# passive-interface
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface Tunnel0
R3(config-router-af-interface)# no shutdown
R3(config-router-af-interface)# af-interface Tunnel1
R3(config-router-af-interface)# no shutdown
```

```
R1# show ipv6 route eigrp
IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007:C15:C0:12::/64 [90/310044416]
    via FE80::7864:7B03, Tunnel0
D  2007:C15:C0:13::/64 [90/310070016]
    via FE80::7864:7B03, Tunnel0
```

```
R2# show ipv6 route eigrp
IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007:C15:C0:10::/64 [90/310070016]
    via FE80::7864:7C03, Tunnel0
D  2007:C15:C0:11::/64 [90/310044416]
    via FE80::7864:7C03, Tunnel0
```

```
R3# show ipv6 route eigrp
IPv6 Routing Table - 9 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007:C15:C0:10::/64 [90/297270016]
    via FE80::7864:7B01, Tunnel0
D  2007:C15:C0:13::/64 [90/297270016]
    via FE80::7864:7C02, Tunnel1
```

```
R3(config)# interface GigabitEthernet 0/0
R3(config-if)# ipv6 ospf 1 area 0

R4(config)# interface GigabitEthernet0/0
R4(config-if)# ipv6 ospf 1 area 0
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ipv6 ospf 1 area 1

R5(config)# interface GigabitEthernet0/0
R5(config-if)# ipv6 ospf 1 area 0

R6(config)# interface GigabitEthernet0/0
R6(config-if)# ipv6 ospf 1 area 1
R3# show ipv6 route ospf
IPv6 Routing Table - 11 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI  2007:C15:C0:15::/64 [110/2]
    via FE80::213:C3FF:FE7B:E4A0, GigabitEthernet0/0

R5# show ipv6 route ospf
IPv6 Routing Table - 5 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI  2007:C15:C0:15::/64 [110/2]
    via FE80::213:C3FF:FE7B:E4A0, GigabitEthernet0/0
```

```
R6# show ipv6 route ospf
IPv6 Routing Table - 5 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
OI  2007:C15:C0:14::/64 [110/2]
    via FE80::213:C3FF:FE7B:E4A1, GigabitEthernet0/0
```

```
R4(config)# ipv6 router ospf 1
R4(config-router)# area 1 authentication ipsec spi 500 md5
DECODECC1E0DBBA11B0BB0BBEDB00B00
```

```
R6(config)# ipv6 router ospf 1
R6(config-router)# area 1 authentication ipsec spi 500 md5
DECODECC1E0DBBA11B0BB0BBEDB00B00
```

```
R4(config)# ipv6 router ospf 1
R4(config-rtr)# area 0 range 2007::/16

R6# show ipv6 route ospf | include OI
OI 2007::/16 [110/2]
    via FE80::213:C3FF:FE7B:E4A1, GigabitEthernet0/0
```

```
R3(config)# ipv6 router ospf 1
R3(config-rtr)# redistribute eigrp 6 include-connected metric 5000

R4# show ipv6 route ospf
IPv6 Routing Table - 11 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external

O  2007::/16 [110/0]
  via ::, Null0
OE2 2007:C15:C0:10::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:11::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:12::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
OE2 2007:C15:C0:13::/64 [110/5000]
  via FE80::214:6AFF:FEFC:7390, GigabitEthernet0/0
```

```
R3(config)# router eigrp CCIE
R3(config-router)# address-family ipv6 unicast autonomous-system 6
R3(config-router-af)# af-interface Tunnel0
R3(config-router-af-interface)# summary-address 2007::/16
R3(config-router-af-interface)# af-interface Tunnel1
R3(config-router-af-interface)# summary-address 2007::/16
```

```
R1# show ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007::/16 [90/310044416]
    via FE80::7864:7B03, Tunnel0
```

```
R1# ping ipv6 2007:C15:C0:15::6
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:15::6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/7/8 ms
```

```
R2# show ipv6 route eigrp
IPv6 Routing Table - 6 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route, M - MIPv6
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
D  2007::/16 [90/310044416]
    via FE80::7864:7C03, Tunnel0
```

```
R2# ping ipv6 2007:C15:C0:15::6
```

```
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2007:C15:C0:15::6, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 4/7/8 ms
```

```
SW2(config)# interface range fastethernet 0/15-16
SW2(config-if-range)# switchport access vlan 200
SW2(config-if-range)# switchport host
SW2(config-if-range)# exit
SW2(config)# mls qos
SW2(config)# class-map VIDEO
SW2(config-cmap)# match access-group 100
SW2(config-cmap)# exit
SW2(config)# access-list 100 permit tcp any any range 3230 3231
SW2(config)# access-list 100 permit udp any any range 3230 3235
SW2(config)# policy-map VIDEO-MARK
SW2(config-pmap)# class VIDEO
SW2(config-pmap-c)# set dscp AF41
SW2(config-pmap-c)# exit
SW2(config)# interface range fastethernet 0/15-16
SW2(config-if-range)# service-policy input VIDEO-MARK
```

```
R2(config)# class-map match-all VIDEO
R2(config-cmap)# match dscp af41
R2(config-cmap)# policy-map VIDEO-QOS
R2(config-pmap)# class VIDEO
R2(config-pmap-c)# priority percent 40
R2(config-pmap-c)# compression header ip rtp

R2(config-pmap-c)# class class-default
R2(config-pmap-c)# bandwidth percent 60
R2(config-pmap-c)# random-detect
R2(config-pmap-c)# exit
R2(config)# interface fastethernet0/0
R2(config-if)# service-policy output VIDEO-QOS
```

```
R1(config)# ip multicast-routing
R1(config)# interface Loopback0
R1(config-if)# ip pim sparse-dense-mode
R1(config-if)# interface GigabitEthernet0/0
R1(config-if)# ip pim sparse-dense-mode
R1(config-if)# ip pim send-rp-announce Loopback0 scope 3 group-list GROUPS
R1(config)# ip access-list standard GROUPS
R1(config-std-nacl)# permit 225.225.0.1
R1(config-std-nacl)# permit 225.225.0.2
R1(config-std-nacl)# permit 225.225.0.3
R1(config-std-nacl)# permit 225.225.0.4
```

```
R2(config)# ip multicast-routing
R2(config)# interface Loopback0
R2(config-if)# ip pim sparse-dense-mode
R2(config-if)# interface fastethernet0/0
R2(config-if)# ip pim sparse-dense-mode
R2(config-if)# ip pim send-rp-announce Loopback0 scope 3 group-list GROUPS
R2(config)# ip access-list standard GROUPS
R2(config-std-nacl)# permit 225.225.0.1
R2(config-std-nacl)# permit 225.225.0.2
R2(config-std-nacl)# permit 225.225.0.3
R2(config-std-nacl)# permit 225.225.0.4
```

```
R3(config)# ip multicast-routing
R3(config)# interface Loopback0
R3(config-if)# ip pim sparse-dense-mode
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-dense-mode
R3(config-if)# interface GigabitEthernet0/1
R3(config-if)# ip pim sparse-dense-mode
R3(config-if)# exit
R3(config)# ip pim send-rp-discovery lo0 scope 2
```

```
R4(config-if)# ip multicast-routing
R4(config-if)# interface GigabitEthernet0/0
R4(config-if)# ip pim sparse-dense-mode
```

```
R4# show ip pim rp mapping
PIM Group-to-RP Mappings
```

```
Group(s) 225.225.0.1/32
RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:52
Group(s) 225.225.0.2/32
RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:56
Group(s) 225.225.0.3/32
RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:55
Group(s) 225.225.0.4/32
RP 120.100.2.1 (?), v2v1
    Info source: 120.100.34.3 (?), elected via Auto-RP
    Uptime: 00:00:03, expires: 00:02:55
```

```
R3(config)# ip pim rp-announce-filter rp-list R1 group-list R1-GROUPS
R3(config)# ip pim rp-announce-filter rp-list R2 group-list R2-GROUPS
R3(config)# ip access-list standard R1
R3(config-std-nacl)# permit 120.100.1.1
R3(config-std-nacl)# exit
R3(config)# ip access-list standard R2
R3(config-std-nacl)# permit 120.100.2.1
R3(config-std-nacl)# exit
R3(config# ip access-list standard R1-GROUPS
R3(config-std-nacl)# permit 225.225.0.1
R3(config-std-nacl)# permit 225.225.0.2
R3(config-std-nacl)# exit
R3(config)# ip access-list standard R2-GROUPS
R3(config-std-nacl)# permit 225.225.0.3
R3(config-std-nacl)# permit 225.225.0.4

R3# debug ip pim auto-rp
PIM Auto-RP debugging is on
Auto-RP(0): Received RP-announce, from 120.100.1.1, RP_cnt 1, ht 181
Auto-RP(0): Update (225.225.0.1/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Update (225.225.0.2/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Filtered 225.225.0.3/32 for RP 120.100.1.1
Auto-RP(0): Filtered 225.225.0.4/32 for RP 120.100.1.1
Auto-RP(0): Received RP-announce, from 120.100.1.1, RP_cnt 1, ht 181
Auto-RP(0): Update (225.225.0.1/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Update (225.225.0.2/32, RP:120.100.1.1), PIMv2 v1
Auto-RP(0): Filtered 225.225.0.3/32 for RP 120.100.1.1
Auto-RP(0): Filtered 225.225.0.4/32 for RP 120.100.1.1

R4# show ip pim rp mapping
PIM Group-to-RP Mappings
```

```
Group(s) 225.225.0.1/32
    RP 120.100.1.1 (?), v2v1
        Info source: 120.100.34.3 (?), elected via Auto-RP
        Uptime: 00:00:08, expires: 00:02:52
Group(s) 225.225.0.2/32
    RP 120.100.1.1 (?), v2v1
        Info source: 120.100.34.3 (?), elected via Auto-RP
        Uptime: 00:00:08, expires: 00:02:51
Group(s) 225.225.0.3/32
    RP 120.100.2.1 (?), v2v1
        Info source: 120.100.34.3 (?), elected via Auto-RP
        Uptime: 00:00:47, expires: 00:02:12
Group(s) 225.225.0.4/32
    RP 120.100.2.1 (?), v2v1
        Info source: 120.100.34.3 (?), elected via Auto-RP
        Uptime: 00:00:47, expires: 00:02:09
```

```
R1(config)# snmp-server host 120.100.100.100 traps public
R1(config)# snmp-server enable traps ipmulticast
R1(config)# ip multicast heartbeat 225.225.0.1 1 1 10

R1# debug snmp packets

R3# ping 225.225.0.1

R1# SNMP: Queuing packet to 120.100.100.100
SNMP: V1 Trap, ent ciscoExperiment.2.3.1, addr 120.100.100.1, gentrap 6,spectrap 1
ciscoIpMRouteHeartBeatEntry.2.225.225.0.1 = 120.100.123.3
ciscoIpMRouteHeartBeatEntry.3.225.225.0.1 = 10
ciscoIpMRouteHeartBeatEntry.4.225.225.0.1 = 1
ciscoIpMRouteHeartBeatEntry.5.225.225.0.1 = 0
```

```
R6(config)# ip tcp intercept list 100
R6(config)# access-list 100 permit tcp any 120.100.63.0 0.0.0.255
R6(config)# ip tcp intercept mode watch
R6(config)# ip tcp intercept drop-mode random
R6(config)# ip tcp intercept watch-timeout 20
```

```
R1(config-if)# ip access-list extended FILTER-IN
R1(config-ext-nacl)# permit icmp any any
R1(config-ext-nacl)# permit eigrp any any
R1(config-ext-nacl)# permit pim any any
R1(config-ext-nacl)# permit tcp host 120.100.3.1 host 120.100.1.1 eq bgp
R1(config-ext-nacl)# permit 41 host 120.100.123.3 host 120.100.123.1
R1(config-ext-nacl)# evaluate DYNAMIC-TCP
R1(config-ext-nacl)# ip access-list extended FILTER-OUT
R1(config-ext-nacl)# permit tcp any any reflect DYNAMIC-TCP
R1(config-ext-nacl)# exit
R1(config)# ip reflexive-list timeout 100
R1(config)# interface GigabitEthernet0/0
R1(config-if)# ip access-group FILTER-IN in
R1(config-if)# ip access-group FILTER-OUT out
```

```
SW1(config)# interface vlan 100
SW1(config-if)# ip add 120.100.100.100 255.255.255.0
SW1(config-if)# exit
SW1(config)# ip route 120.100.3.1 255.255.255.255 120.100.100.1
SW1(config)# exit
SW1# trace 120.100.3.1
```

```
Type escape sequence to abort.
Tracing the route to 120.100.3.1
 1 120.100.100.1 0 msec 4 msec 0 msec
 2 120.100.100.1 !A * !A
SW1# telnet 120.100.3.1
Trying 120.100.3.1 ... Open
```

User Access Verification

```
Password:
R3>enable
Password:
R3#
R1# show access-lists
Standard IP access list 1
 10 permit 120.100.1.0 (3 matches)
 20 permit 120.100.100.0 (3 matches)
```

```
Standard IP access list GROUPS
```

```
 10 permit 225.225.0.1  
 20 permit 225.225.0.2  
 30 permit 225.225.0.3  
 40 permit 225.225.0.4
```

```
Reflexive IP access list DYNAMIC-TCP
```

```
  permit tcp host 120.100.3.1 eq telnet host 120.100.100.100 eq 11034 (34  
    matches) (time left 90)
```

```
Extended IP access list FILTER-IN
```

```
 5 permit icmp any any (150 matches)  
 10 permit eigrp any any (1710 matches)  
 20 permit pim any any (92 matches)  
 25 permit tcp host 120.100.3.1 host 120.100.1.1 eq bgp (126 matches)  
 30 evaluate DYNAMIC-TCP
```

```
Extended IP access list FILTER-OUT
```

```
 10 permit tcp any any reflect DYNAMIC-TCP (18 matches)
```

```
R1(config)# ip domain-name toughtest.co.uk
R1(config)# crypto key generate rsa modulus 768
The name for the keys will be: R1.toughtest.co.uk

% The key modulus size is 768 bits
% Generating 768 bit RSA keys, keys will be non-exportable...[OK]

R1(config)# aaa new-model
R1(config)# aaa authentication login default local
R1(config)# aaa authorization exec default local
R1(config)# username cisco privilege 15 password 0 cisco
R1(config)# ip ssh time-out 120
R1(config)# ip ssh authentication-retries 2
R1(config)# ip scp server enable
R1(config)#
00:57:29.343: %SSH-5-ENABLED: SSH 1.99 has been enabled
```

```
SW1(config)# ipv6 unicast-routing
SW1(config)# ipv6 access-list RA
SW1(config-ipv6-acl)# deny icmp any any router-advertisement
SW1(config-ipv6-acl)# permit ipv6 any any
SW1(config-ipv6-acl)# exit
SW1(config)# int vlan 34
SW1(config-if)# ipv6 traffic-filter RA in
SW1(config-if)# ipv6 address 2007:C15:C0:15::10/64
```

```
SW1# show log
*Oct 4 17:58:23: %IPV6-6-ACCESSLOGDP: list RA/10 denied icmpv6
  FE80::219:AFF:FEBA:BE40 -> FF02::1 (134/0), 1 packet
*Oct 4 17:58:23: %IPV6-6-ACCESSLOGDP: list RA/10 denied icmpv6
  FE80::218:18FF:FEA2:3250 -> FF02::1 (134/0), 1 packet
```

```
Switch1# show run interface fastethernet 0/19
!
interface fastethernet0/19
switchport access vlan 200
switchport mode access

Switch1# show run interface fastethernet 0/1
!
interface fastethernet0/1
switchport trunk encapsulation dot1q
switchport trunk allowed vlan 10,50,200
switchport mode trunk

Switch2# show run interface fastethernet 0/19
!
interface fastethernet0/19
switchport access vlan 400
switchport mode access

Switch2# show run interface fastethernet 0/6
!
interface fastethernet0/6
switchport trunk encapsulation dot1q
switchport trunk allowed vlan 20,100,400
switchport mode trunk

R1# show run | begin interface GigabitEthernet0/0
!
interface GigabitEthernet0/0
no ip address
!
interface GigabitEthernet0/0.10
encapsulation dot1Q 10
!
interface GigabitEthernet0/0.50
encapsulation dot1Q 50
!
interface GigabitEthernet0/0.200
encapsulation dot1Q 200
```

```
R6# show run | begin interface GigabitEthernet0/1
!
interface GigabitEthernet0/1
no ip address
!
interface GigabitEthernet0/1.20
encapsulation dot1Q 20
!
interface GigabitEthernet0/0.100
encapsulation dot1Q 100
!
interface GigabitEthernet0/1.400
```

```
Switch3# show run interface fastethernet 0/19
!
interface fastethernet0/19
no switchport
ip address 1.1.1.1 255.255.255.0
```

```
Switch4# show run interface fastethernet 0/19
!
interface fastethernet0/19
no switchport
ip address 1.1.1.2 255.255.255.0
```

```
R1(config-if)# int lo0
R1(config-if)# ip ospf 1 area 0
R1(config-if)# int Gi0/1
R1(config-if)# ip ospf 1 area 0
```

```
R2(config-if)# int lo0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# int Fa0/0
R2(config-if)# ip ospf 1 area 0
R2(config-if)# int Fa0/1
R2(config-if)# ip ospf 1 area 0
```

```
R3(config-if)# int lo0
R3(config-if)# ip ospf 1 area 0
R3(config-if)# int Gi0/0
R3(config-if)# ip ospf 1 area 0
R3(config-if)# int Gi0/1
R3(config-if)# ip ospf 1 area 0
```

```
R4(config-if)# int lo0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# int gi0/0
R4(config-if)# ip ospf 1 area 0
R4(config-if)# int Gi0/1
R4(config-if)# ip ospf 1 area 0
```

```
R5(config-if)# int lo0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# int gi0/0
R5(config-if)# ip ospf 1 area 0
R5(config-if)# int Gi0/1
R5(config-if)# ip ospf 1 area 0
```

```
R6(config-if)# int lo0
R6(config-if)# ip ospf 1 area 0
R6(config-if)# int gi0/0
R6(config-if)# ip ospf 1 area 0
```

```
R1# show ip route ospf
    120.0.0.0/8 is variably subnetted, 11 subnets, 2 masks
O      120.100.2.1/32 [110/2] via 120.100.132.2, 00:05:00, GigabitEthernet0/1
O      120.100.3.1/32 [110/2] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.4.1/32 [110/12] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.5.1/32 [110/22] via 120.100.132.3, 00:02:36, GigabitEthernet0/1
O      120.100.6.1/32 [110/22] via 120.100.132.3, 00:01:19, GigabitEthernet0/1
O      120.100.25.0/24 [110/31] via 120.100.132.3, 00:02:26, GigabitEthernet0/1
O      120.100.34.0/24 [110/11] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
O      120.100.45.0/24 [110/21] via 120.100.132.3, 00:06:16, GigabitEthernet0/1
```

```
R1(config)# mpls label protocol tdp
R1(config)# interface Gi0/1
R1(config-if)# mpls label protocol ldp
R1(config-if)# mpls ip

R2(config)# mpls label protocol tdp
R2(config)# interface Fa0/0
R2(config-if)# mpls label protocol ldp
R2(config-if)# mpls ip
R2(config)# interface Fa0/1
R2(config-if)# mpls label protocol ldp
R2(config-if)# mpls ip

R3(config)# mpls label protocol tdp
R3(config)# interface Gi0/0
R3(config-if)# mpls label protocol ldp
R3(config-if)# mpls ip
R3(config-if)# interface Gi0/1
R3(config-if)# mpls label protocol ldp
R3(config-if)# mpls ip

R4(config)# mpls label protocol tdp
R4(config)# interface GigabitEthernet0/0
R4(config-if)# mpls label protocol ldp
R4(config-if)# mpls ip
R4(config-if)# interface Gi0/1
R4(config-if)# mpls label protocol ldp
R4(config-if)# mpls ip

R5(config)# mpls label protocol tdp
R5(config)# interface Gi0/0
R5(config-if)# mpls label protocol ldp
R5(config-if)# mpls ip
R5(config-if)# interface Gi0/1
R5(config-if)# mpls label protocol ldp
R5(config-if)# mpls ip

R6(config)# mpls label protocol tdp
R6(config)# interface Gi0/0
R6(config-if)# mpls label protocol ldp
R6(config-if)# mpls ip
```

```
R1# show mpls ldp neighbor
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.1.1:0
TCP connection: 120.100.2.1.40418 - 120.100.1.1.646
State: Oper; Msgs sent/rcvd: 69/71; Downstream
Up time: 00:47:20
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.123.2
Addresses bound to peer LDP Ident:
    120.100.123.2    120.100.25.2    120.100.2.1
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.1.1:0
TCP connection: 120.100.3.1.51369 - 120.100.1.1.646
State: Oper; Msgs sent/rcvd: 68/68; Downstream
Up time: 00:47:18
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.123.3
Addresses bound to peer LDP Ident:
    120.100.123.3    120.100.3.1    120.100.34.3
```

```
R2# show mpls ldp neighbor
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.3.1.16991 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 71/68; Downstream
Up time: 00:46:33
LDP discovery sources:
```

```
    fastethernet0/0, Src IP addr: 120.100.123.3
    fastethernet0/1, Src IP addr: 120.100.34.3
Addresses bound to peer LDP Ident:
    120.100.123.3    120.100.3.1    120.100.34.3
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.5.1.13826 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 73/76; Downstream
Up time: 00:46:24
LDP discovery sources:
    fastethernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
    120.100.25.5    120.100.5.1    5.5.5.5        120.100.45.5
    100.100.100.5
Peer LDP Ident: 120.100.1.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.1.1.646 - 120.100.2.1.40418
State: Oper; Msgs sent/rcvd: 69/68; Downstream
Up time: 00:46:07
LDP discovery sources:
    fastethernet0/0, Src IP addr: 120.100.123.1
Addresses bound to peer LDP Ident:
    120.100.123.1    120.100.1.1
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.2.1:0
TCP connection: 120.100.4.1.47401 - 120.100.2.1.646
State: Oper; Msgs sent/rcvd: 54/57; Downstream
Up time: 00:32:28
LDP discovery sources:
    fastethernet0/1, Src IP addr: 120.100.34.4
Addresses bound to peer LDP Ident:
    120.100.4.1    4.4.4.4        120.100.45.4    100.100.100.4
    120.100.34.4
```

```
R3# show mpls ldp neighbor
Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.2.1.646 - 120.100.3.1.16991
State: Oper; Msgs sent/rcvd: 69/72; Downstream
Up time: 00:47:11
LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.123.2
    GigabitEthernet0/1, Src IP addr: 120.100.25.2
Addresses bound to peer LDP Ident:
    120.100.123.2    120.100.25.2    120.100.2.1
Peer LDP Ident: 120.100.1.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.1.1.646 - 120.100.3.1.51369
```

```
State: Oper; Msgs sent/rcvd: 67/67; Downstream
Up time: 00:46:43
LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.123.1
Addresses bound to peer LDP Ident:
    120.100.123.1    120.100.1.1
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.5.1.53107 - 120.100.3.1.646
State: Oper; Msgs sent/rcvd: 67/74; Downstream
Up time: 00:45:22
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
    120.100.25.5    120.100.5.1    5.5.5.5          120.100.45.5
    100.100.100.5
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.3.1:0
TCP connection: 120.100.4.1.15940 - 120.100.3.1.646
State: Oper; Msgs sent/rcvd: 52/56; Downstream
Up time: 00:33:06
LDP discovery sources:
    GigabitEthernet0/1, Src IP addr: 120.100.34.4
Addresses bound to peer LDP Ident:
    120.100.4.1      4.4.4.4          120.100.45.4    100.100.100.4
    120.100.34.4
```

```
R4# show mpls ldp neighbor
Peer LDP Ident: 120.100.6.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.6.1.55234 - 120.100.4.1.646
State: Oper; Msgs sent/rcvd: 74/76; Downstream
Up time: 00:43:52
LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.6
Addresses bound to peer LDP Ident:
    120.100.6.1      6.6.6.6          100.100.100.6   120.100.45.6
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.5.1.57689 - 120.100.4.1.646
State: Oper; Msgs sent/rcvd: 72/74; Downstream
Up time: 00:43:48
LDP discovery sources:
    GigabitEthernet0/0, Src IP addr: 120.100.45.5
    GigabitEthernet0/1, Src IP addr: 120.100.25.5
Addresses bound to peer LDP Ident:
```

```
120.100.25.5    120.100.5.1      5.5.5.5          120.100.45.5
100.100.100.5

Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.2.1.646 - 120.100.4.1.47401
State: Oper; Msgs sent/rcvd: 55/52; Downstream
Up time: 00:30:52
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.25.2
Addresses bound to peer LDP Ident:
  120.100.123.2  120.100.25.2  120.100.2.1

Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.4.1:0
TCP connection: 120.100.3.1.646 - 120.100.4.1.15940
State: Oper; Msgs sent/rcvd: 54/50; Downstream
Up time: 00:30:52
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.34.3
Addresses bound to peer LDP Ident:
  120.100.123.3  120.100.3.1   120.100.34.3
```

```
R5# show mpls ldp neighbor

Peer LDP Ident: 120.100.2.1:0; Local LDP Ident 120.100.5.1:0
TCP connection: 120.100.2.1.646 - 120.100.5.1.13826
State: Oper; Msgs sent/rcvd: 80/77; Downstream
Up time: 00:49:55
LDP discovery sources:
  GigabitEthernet0/1, Src IP addr: 120.100.25.2
Addresses bound to peer LDP Ident:
  120.100.123.2  120.100.25.2  120.100.2.1

Peer LDP Ident: 120.100.6.1:0; Local LDP Ident 120.100.5.1:0
TCP connection: 120.100.6.1.18472 - 120.100.5.1.646
State: Oper; Msgs sent/rcvd: 81/81; Downstream
Up time: 00:48:58
LDP discovery sources:
  GigabitEthernet0/0, Src IP addr: 120.100.45.6
Addresses bound to peer LDP Ident:
  120.100.6.1      6.6.6.6        100.100.100.6  120.100.45.6

Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.5.1:0
TCP connection: 120.100.4.1.646 - 120.100.5.1.57689
State: Oper; Msgs sent/rcvd: 80/78; Downstream
Up time: 00:48:54
```

```
LDP discovery sources:  
    GigabitEthernet0/0, Src IP addr: 120.100.45.4  
    GigabitEthernet0/1, Src IP addr: 120.100.34.4  
Addresses bound to peer LDP Ident:  
    120.100.4.1      4.4.4.4          120.100.45.4      100.100.100.4  
    120.100.34.4  
Peer LDP Ident: 120.100.3.1:0; Local LDP Ident 120.100.5.1:0  
TCP connection: 120.100.3.1.646 - 120.100.5.1.53107  
State: Oper; Msgs sent/rcvd: 77/70; Downstream  
Up time: 00:48:17  
LDP discovery sources:  
    GigabitEthernet0/1, Src IP addr: 120.100.34.3  
Addresses bound to peer LDP Ident:  
    120.100.123.3    120.100.3.1      120.100.34.3
```

```
R6# show mpls ldp neighbor  
Peer LDP Ident: 120.100.5.1:0; Local LDP Ident 120.100.6.1:0  
TCP connection: 120.100.5.1.646 - 120.100.6.1.18472  
State: Oper; Msgs sent/rcvd: 82/82; Downstream  
Up time: 00:49:31  
LDP discovery sources:  
    GigabitEthernet0/0, Src IP addr: 120.100.45.5  
Addresses bound to peer LDP Ident:  
    120.100.25.5      120.100.5.1      5.5.5.5          120.100.45.5  
    100.100.100.5  
Peer LDP Ident: 120.100.4.1:0; Local LDP Ident 120.100.6.1:0  
TCP connection: 120.100.4.1.646 - 120.100.6.1.55234  
State: Oper; Msgs sent/rcvd: 82/80; Downstream  
Up time: 00:49:31  
LDP discovery sources:  
    GigabitEthernet0/0, Src IP addr: 120.100.45.4  
Addresses bound to peer LDP Ident:  
    120.100.4.1      4.4.4.4          120.100.45.4      100.100.100.4  
    120.100.34.4
```

```
R1(config)# ip vrf BLUE
R1(config-vrf)# rd 65001:100
R1(config-vrf)# route-target export 65001:100
R1(config-vrf)# route-target import 65001:100
R1(config-vrf)# !
R1(config-vrf)# ip vrf RED
R1(config-vrf)# rd 65001:200
R1(config-vrf)# route-target export 65001:200
R1(config-vrf)# route-target import 65001:200
R1(config-vrf)# exit
R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# ip vrf forwarding BLUE
R1(config-subif)# interface GigabitEthernet0/0.50
R1(config-subif)# ip vrf forwarding RED
```

```
R6(config)# ip vrf BLUE
R6(config-vrf)# rd 65001:100
R6(config-vrf)# route-target export 65001:100
R6(config-vrf)# route-target import 65001:100
R6(config-vrf)# !
R6(config-vrf)# ip vrf RED
R6(config-vrf)# rd 65001:200
R6(config-vrf)# route-target export 65001:200
R6(config-vrf)# route-target import 65001:200
R6(config-vrf)# exit
R6(config)# interface GigabitEthernet0/1.20
R6(config-subif)# ip vrf forwarding BLUE
R6(config)# interface GigabitEthernet0/1.100
R6(config-subif)# ip vrf forwarding RED
```

```
R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# ip add 10.10.10.1 255.255.255.252

Switch1(config)# vlan 10
Switch1(config-vlan)# exit
Switch1(config)# interface vlan 10
Switch1(config-if)# no shutdown
Switch1(config-if)# ip add 10.10.10.2 255.255.255.252

R1# ping vrf BLUE 10.10.10.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.10.10.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

```
R6(config)# interface GigabitEthernet0/1.20
R6(config-subif)# ip add 10.10.20.1 255.255.255.252

Switch2(config)# vlan 20
Switch2(config-vlan)# exit
Switch2(config)# interface vlan 20
Switch2(config-if)# no shutdown
Switch2(config-if)# ip add 10.10.20.2 255.255.255.252

R6# ping vrf BLUE 10.10.20.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.10.20.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

```
R1(config)# interface GigabitEthernet0/0.50
R1(config-subif)# ip add 130.50.50.1 255.255.255.252

Switch3(config)# interface vlan 50
Switch3(config-if)# no shutdown
Switch3(config-if)# ip add 130.50.50.2 255.255.255.252

R1# ping vrf RED 130.50.50.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 130.50.50.2, timeout is 2 seconds:
..!!!!

Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

```
R6(config)# interface GigabitEthernet0/1.100
R6(config-subif)# ip add 130.100.100.1 255.255.255.252

Switch4(config)# interface vlan 100
Switch4(config-if)# no shutdown
Switch4(config-if)# ip add 130.100.100.2 255.255.255.252

R6# ping vrf RED 130.100.100.2

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 130.100.100.2, timeout is 2 seconds:
..!!!
Success rate is 60 percent (3/5), round-trip min/avg/max = 1/1/1 ms
```

```
R1(config)# router bgp 65001
R1(config-router)# no synchronization
R1(config-router)# no auto-summary
R1(config-router)# neighbor 120.100.6.1 remote-as 65001
R1(config-router)# neighbor 120.100.6.1 update-source Loopback0
R1(config-router)# address-family vpng4
R1(config-router-af)# neighbor 120.100.6.1 activate
R1(config-router-af)# neighbor 120.100.6.1 next-hop-self
R1(config-router-af)# neighbor 120.100.6.1 send-community both
```

```
R6(config)# router bgp 65001
R6(config-router)# no sync
R6(config-router)# no auto-summary
R6(config-router)# neighbor 120.100.1.1 remote-as 65001
R6(config-router)# neighbor 120.100.1.1 update-source Loopback0

R6(config-router)# address-family vpng4
R6(config-router-af)# neighbor 120.100.1.1 activate
R6(config-router-af)# neighbor 120.100.1.1 next-hop-self
R6(config-router-af)# neighbor 120.100.1.1 send-community both
```

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R6(config-router-af)# network 10.10.20.0 0.0.0.3
```

```
Switch2(config)# ip routing
Switch2(config)# router eigrp 10
Switch2(config-router)# no auto-summary
Switch2(config-router)# network 10.10.20.0 0.0.0.3
Switch2(config-router)# network 10.2.2.0 0.0.0.255
Switch2(config-router)# network 10.2.3.0 0.0.0.255
Switch2(config-router)# network 10.2.4.0 0.0.0.255
```

```
R6# show ip eigrp vrf BLUE neighbors
IP-EIGRP neighbors for process 10
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
                                         (sec)          (ms)       Cnt Num
0   10.10.20.2        Gi0/1.20      11 00:04:18   1   200   0   1
R6#
R6# show ip route vrf BLUE eigrp
  10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks
D     10.2.2.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
D     10.2.3.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
D     10.2.4.0/24
      [90/156160] via 10.10.20.2, 00:04:36, GigabitEthernet0/1.20
```

```
R1(config)# router eigrp VPN
R1(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R1(config-router-af)# network 10.10.10.0 0.0.0.3
```

```
Switch1(config)# ip routing
Switch1(config)# router eigrp 10
Switch1(config-router)# no auto-summary
Switch1(config-router)# network 10.10.10.0 0.0.0.3
Switch1(config-router)# network 10.1.1.0 0.0.0.255
Switch1(config-router)# network 10.1.2.0 0.0.0.255
Switch1(config-router)# network 10.1.3.0 0.0.0.255
```

```
R1# show ip eigrp vrf BLUE neighbors
```

```
IP-EIGRP neighbors for process 10
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
                (sec)          (ms)        Cnt Num
0   10.10.10.2       Gi0/0.10     13 00:00:24    1    200   0   1
```

```
R1#
```

```
R1# show ip eigrp vrf BLUE neighbors
```

```
IP-EIGRP neighbors for process 10
H   Address           Interface      Hold Uptime    SRTT     RTO   Q   Seq
                (sec)          (ms)        Cnt Num
0   10.10.10.2       Gi0/0.10     13 00:00:24    1    200   0   1
```

```
R1# show ip route vrf BLUE eigrp
```

```
 10.0.0.0/8 is variably subnetted, 4 subnets, 2 masks
D       10.1.3.0/24
          [90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10
D       10.1.2.0/24
          [90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10
D       10.1.1.0/24
          [90/153856] via 10.10.10.2, 00:01:18, GigabitEthernet0/0.10
```

```
R1(config)# router eigrp VPN
R1(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R1(config-router)# topology base
R1(config-router-af-topology)# redistribute bgp 65001 metric 10000 100 255 1 1500
R1(config-router-af-topology)# router bgp 65001
R1(config-router)# address-family ipv4 vrf BLUE
R1(config-router-af)# redistribute eigrp 10 metric 50
```

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 vrf BLUE autonomous-system 10
R6(config-router)# topology base
R6(config-router-af-topology)# redistribute bgp 65001 metric 10000 100 255 1 1500
R6(config-router-af-topology)# router bgp 65001
R6(config-router)# address-family ipv4 vrf BLUE
R6(config-router-af)# redistribute eigrp 10 metric 50
```

```
SW1# show ip route eigrp
D      10.2.2.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.2.3.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.2.4.0/24 [90/156416] via 10.10.10.1, 00:32:05, Vlan10
D      10.10.20.0/30 [90/28416] via 10.10.10.1, 00:32:05, Vlan10
```

```
SW2# show ip route eigrp
D      10.1.3.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.1.2.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.1.1.0/24 [90/154112] via 10.10.20.1, 00:33:07, Vlan20
D      10.10.10.0/30 [90/26112] via 10.10.20.1, 00:33:07, Vlan20
```

```
R6# show ip bgp vpng4 vrf BLUE
BGP table version is 17, local router ID is 120.100.6.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*>i10.1.1.0/24	120.100.1.1	50	100	0	?
*>i10.1.2.0/24	120.100.1.1	50	100	0	?
*>i10.1.3.0/24	120.100.1.1	50	100	0	?
*> 10.2.2.0/24	10.10.20.2	50		32768	?
*> 10.2.3.0/24	10.10.20.2	50		32768	?
*> 10.2.4.0/24	10.10.20.2	50		32768	?
*>i10.10.10.0/30	120.100.1.1	0	100	0	?
*> 10.10.20.0/30	0.0.0.0		0		32768 ?

```
R1# show ip bgp vpng4 vrf BLUE
BGP table version is 17, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
              r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*> 10.1.1.0/24	10.10.10.2	50		32768	?
*> 10.1.2.0/24	10.10.10.2	50		32768	?
*> 10.1.3.0/24	10.10.10.2	50		32768	?
*>i10.2.2.0/24	120.100.6.1	50	100	0	?
*>i10.2.3.0/24	120.100.6.1	50	100	0	?
*>i10.2.4.0/24	120.100.6.1	50	100	0	?
*> 10.10.10.0/30	0.0.0.0		0		32768 ?
*>i10.10.20.0/30	120.100.6.1	0	100	0	?

```
R1# show mpls forwarding-table vrf BLUE
Local  Outgoing      Prefix          Bytes tag  Outgoing      Next Hop
tag    tag or VC    or Tunnel Id   switched   interface
26     Untagged     10.1.3.0/24[V]  0          Gi0/0.10   10.10.10.2
27     Untagged     10.1.2.0/24[V]  0          Gi0/0.10   10.10.10.2
28     Aggregate    10.10.10.0/30[V] 0
29     Untagged     10.1.1.0/24[V]  0          Gi0/0.10   10.10.10.2
```

```
R6# show mpls forwarding-table vrf BLUE
```

```
Local  Outgoing      Prefix          Bytes tag  Outgoing      Next Hop
tag    tag or VC    or Tunnel Id   switched   interface
26     Untagged     10.2.2.0/24[V]  0          Gi0/1.20   10.10.20.2
27     Untagged     10.2.3.0/24[V]  0          Gi0/1.20   10.10.20.2
28     Untagged     10.2.4.0/24[V]  0          Gi0/1.20   10.10.20.2
29     Aggregate    10.10.20.0/30[V] 0
```

```
SW3(config)# ip routing
SW3(config)# router ospf 3
SW3(config-router)# network 130.50.50.0 0.0.0.3 area 0
SW3(config-router)# network 10.33.33.0 0.0.0.255 area 1
SW3(config-router)# network 10.33.34.0 0.0.0.255 area 1
SW3(config-router)# network 10.33.35.0 0.0.0.255 area 1

SW4(config)# ip routing
SW4(config)# router ospf 2
SW4(config-router)# network 130.100.100.0 0.0.0.3 area 0
SW4(config-router)# network 10.44.44.0 0.0.0.255 area 2
SW4(config-router)# network 10.44.45.0 0.0.0.255 area 2
SW4(config-router)# network 10.44.46.0 0.0.0.255 area 2

R1(config)# router ospf 3 vrf RED
R1(config-router)# network 130.50.50.0 0.0.0.3 area 0
R1(config-router)# redistribute bgp 65001 subnets metric-type 1
R1(config-router)# router bgp 65001
R1(config-router)# address-family ipv4 vrf RED
R1(config-router-af)# redistribute ospf 3 match internal

R6(config)# router ospf 2 vrf RED
R6(config-router)# net 130.100.100.0 0.0.0.3 area 0
R6(config-router)# redistribute bgp 65001 subnets metric-type 1
R6(config-router)# router bgp 65001
R6(config-router)# address-family ipv4 vrf RED
R6(config-router-af)# redistribute ospf 2 match internal
```

```
R1# show ip route vrf RED ospf
```

Routing Table: RED

```
    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.33.34.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.35.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.33.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
```

```
R6# show ip route vrf RED ospf
```

Routing Table: RED

```
    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.44.46.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.45.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.44.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
```

```
SW3# show ip route ospf
```

```
    130.100.0.0/30 is subnetted, 1 subnets
O E1    130.100.100.0 [110/2] via 130.50.50.1, 00:06:08, Vlan50
    10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.44.46.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.45.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.44.1/32 [110/3] via 130.50.50.1, 00:02:55, Vlan50
```

```
SW4# show ip route ospf
```

```
    130.50.0.0/30 is subnetted, 1 subnets
O E1    130.50.50.0 [110/2] via 130.100.100.1, 00:03:37, Vlan100
    10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.33.34.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.35.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.33.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
```

```
R1# show ip route vrf RED ospf
```

Routing Table: RED

```
    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.33.34.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.35.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
O IA    10.33.33.1 [110/2] via 130.50.50.2, 00:04:48, GigabitEthernet0/0.50
```

```
R6# show ip route vrf RED ospf
```

Routing Table: RED

```
    10.0.0.0/32 is subnetted, 6 subnets
O IA    10.44.46.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.45.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
O IA    10.44.44.1 [110/2] via 130.100.100.2, 00:02:32, GigabitEthernet0/1.100
```

```
SW3# show ip route ospf
```

```
    130.100.0.0/30 is subnetted, 1 subnets
O E1    130.100.100.0 [110/2] via 130.50.50.1, 00:06:08, Vlan50
    10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.44.46.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.45.1/32 [110/3] via 130.50.50.1, 00:02:54, Vlan50
O E1    10.44.44.1/32 [110/3] via 130.50.50.1, 00:02:55, Vlan50
```

```
SW4# show ip route ospf
```

```
    130.50.0.0/30 is subnetted, 1 subnets
O E1    130.50.50.0 [110/2] via 130.100.100.1, 00:03:37, Vlan100
    10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks
O E1    10.33.34.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.35.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
O E1    10.33.33.1/32 [110/3] via 130.100.100.1, 00:03:37, Vlan100
```

```
R1# show ip ospf 3 | include Domain  
Domain ID type 0x0005, value 0.0.0.3
```

```
R6# show ip ospf 2 | include Domain  
Domain ID type 0x0005, value 0.0.0.2
```

```
R1(config)# router ospf 3 vrf RED  
R1(config-router)# domain-id 0.0.0.2
```

```
SW3# show ip route ospf  
 130.100.0.0/30 is subnetted, 1 subnets  
O IA    130.100.100.0 [110/2] via 130.50.50.1, 00:00:09, Vlan50  
      10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks  
O IA    10.44.46.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50  
O IA    10.44.45.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50  
O IA    10.44.44.1/32 [110/3] via 130.50.50.1, 00:00:09, Vlan50  
SW3#
```

```
SW4# show ip route ospf  
 130.50.0.0/30 is subnetted, 1 subnets  
O IA    130.50.50.0 [110/2] via 130.100.100.1, 00:00:07, Vlan100  
      10.0.0.0/8 is variably subnetted, 6 subnets, 2 masks  
O IA    10.33.34.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100  
O IA    10.33.35.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100  
O IA    10.33.33.1/32 [110/3] via 130.100.100.1, 00:00:07, Vlan100
```

```
Sw4(config)# interface Loopback0
Sw4(config-if)# ip ospf network point-to-point

R1(config)# ip vrf BLUE
R1(config-vrf)# export map SW1
R1(config-vrf)# access-list 10 permit 10.1.1.0 0.0.0.255
R1(config-vrf)# exit
R1(config)# route-map SW1 permit 10
R1(config-route-map)# match ip address 10
R1(config-route-map)# set extcommunity rt 65001:200 additive
```

```
R6(config)# ip vrf RED
R6(config-vrf)# export map SW4
R6(config-vrf)# access-list 10 permit 10.44.44.0 0.0.0.255
R6(config-vrf)# exit
R6(config)# route-map SW4 permit 10
R6(config-route-map)# match ip address 10
R6(config-route-map)# set extcommunity rt 65001:100 additive
```

! R1 is now sending 10.1.1.0 into VRF RED and R6 10.44.44.0 into VRF BLUE

```
R1# show ip bgp vpng4 vrf RED
BGP table version is 33, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:200 (default for vrf RED)					
*> 10.33.33.1/32	130.50.50.2	2		32768	?
*> 10.33.34.1/32	130.50.50.2	2		32768	?
*> 10.33.35.1/32	130.50.50.2	2		32768	?
*>i10.44.44.1/32	120.100.6.1	2	100	0	?
*>i10.44.45.1/32	120.100.6.1	2	100	0	?
*>i10.44.46.1/32	120.100.6.1	2	100	0	?
*> 130.50.50.0/30	0.0.0.0	0		32768	?
*>i130.100.100.0/30	120.100.6.1	0	100	0	?

! No sign of the 10.1.1.0 route, clear the BGP session to kick start the export map

```
R1# clear ip bgp *
R1# show ip bgp vpng4 vrf RED
BGP table version is 34, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:200 (default for vrf RED)					
*> 10.1.1.0/24	10.10.10.2	50	32768	?	
*> 10.33.33.1/32	130.50.50.2	2	32768	?	
*> 10.33.34.1/32	130.50.50.2	2	32768	?	
*> 10.33.35.1/32	130.50.50.2	2	32768	?	
*>i10.44.44.1/32	120.100.6.1	2	100	0	?
*>i10.44.45.1/32	120.100.6.1	2	100	0	?
*>i10.44.46.1/32	120.100.6.1	2	100	0	?
*> 130.50.50.0/30	0.0.0.0	0	32768	?	
*>i130.100.100.0/30	120.100.6.1	0	100	0	?

```
R6# show ip bgp vpng4 vrf BLUE
BGP table version is 35, local router ID is 120.100.6.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
Route Distinguisher: 65001:100 (default for vrf BLUE)					
*>i10.1.1.0/24	120.100.1.1	50	100	0	?
*>i10.1.2.0/24	120.100.1.1	50	100	0	?
*>i10.1.3.0/24	120.100.1.1	50	100	0	?
*> 10.2.2.0/24	10.10.20.2	50	32768	?	
*> 10.2.3.0/24	10.10.20.2	50	32768	?	
*> 10.2.4.0/24	10.10.20.2	50	32768	?	
*>i10.10.10.0/30	120.100.1.1	0	100	0	?
*> 10.10.20.0/30	0.0.0.0	0	32768	?	
*> 10.44.44.1/32	130.100.100.2	2	32768	?	

! Notice the 10.44.44.0 route is actually listed as a host route, change the loop-back interface on Sw4 to a point-to-point for OSPF to advertise it correctly

```
SW4(config)# interface lo0
SW4(config-if)# ip ospf network point-to-point

R6# show ip bgp vpngv4 vrf BLUE | include 10.44.44.0
*> 10.44.44.0/24      130.100.100.2          2          32768 ?

Switch1# show ip route | include 10.44.44.0
D EX    10.44.44.0/24 [170/281856] via 10.10.10.1, 00:00:51, Vlan10
Switch1# 

SW4# show ip route | include 10.1.1.0
O E1    10.1.1.0/24 [110/51] via 130.100.100.1, 00:02:45, Vlan10

! Now test with an extended ping to ensure that the loopback interface is used as
the source

SW1# ping
Protocol [ip]:
Target IP address: 10.44.44.1
Repeat count [5]:
Datagram size [100]:
Timeout in seconds [2]:
Extended commands [n]: y
Source address or interface: 10.1.1.1
Type of service [0]:
Set DF bit in IP header? [no]:
Validate reply data? [no]:
Data pattern [0xABCD]:
Loose, Strict, Record, Timestamp, Verbose[none]:
Sweep range of sizes [n]:
Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 10.44.44.1, timeout is 2 seconds:
Packet sent with a source address of 10.1.1.1
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/9/12 ms
```

```
R1# show mpls forwarding-table vrf BLUE
```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes	tag switched	Outgoing interface	Next Hop
34	Untagged	10.1.3.0/24[V]	0		Gi0/0.10	10.10.10.2
35	Untagged	10.1.2.0/24[V]	0		Gi0/0.10	10.10.10.2
36	Aggregate	10.10.10.0/30[V]	0			
37	Untagged	10.1.1.0/24[V]	590		Gi0/0.10	10.10.10.2

```
R1# show mpls forwarding-table vrf RED
```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes	tag switched	Outgoing interface	Next Hop
38	Aggregate	130.50.50.0/30[V]	0			
39	Untagged	10.33.34.1/32[V]	0		Gi0/0.50	130.50.50.2
40	Untagged	10.33.35.1/32[V]	0		Gi0/0.50	130.50.50.2
41	Untagged	10.33.33.1/32[V]	0		Gi0/0.50	130.50.50.2

! Note the Routes are not leaked within the MPLS forwarding-table

```
R6# show mpls forwarding-table vrf BLUE
```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes	tag switched	Outgoing interface	Next Hop
34	Untagged	10.2.2.0/24[V]	0		Gi0/1.20	10.10.20.2
35	Untagged	10.2.3.0/24[V]	0		Gi0/1.20	10.10.20.2
36	Untagged	10.2.4.0/24[V]	0		Gi0/1.20	10.10.20.2
37	Aggregate	10.10.20.0/30[V]	0			

```
R6# show mpls forwarding-table vrf RED
```

Local tag	Outgoing tag or VC	Prefix or Tunnel Id	Bytes	tag switched	Outgoing interface	Next Hop
38	Aggregate	130.100.100.0/30[V]	0			
39	Untagged	10.44.46.1/32[V]	0		Gi0/1.100	130.100.100.2
40	Untagged	10.44.45.1/32[V]	0		Gi0/1.100	130.100.100.2
42	Untagged	10.44.44.0/24[V]	1534		Gi0/1.100	130.100.100.2

! Note the Routes are not leaked within the MPLS forwarding-table

```
SW1# traceroute 10.2.2.1
```

Type escape sequence to abort.

Tracing the route to 10.2.2.1

```
1 10.10.10.1 0 msec 0 msec 0 msec  
2 120.100.123.2 12 msec 12 msec 16 msec  
3 120.100.25.5 8 msec 12 msec 8 msec  
4 10.10.20.1 8 msec 8 msec 8 msec  
5 10.10.20.2 8 msec * 4 msec
```

```
R1(config)# no mpls ip propagate-ttl
```

```
R6(config)# no mpls ip propagate-ttl
```

```
SW1# traceroute 10.2.2.1
```

Type escape sequence to abort.

Tracing the route to 10.2.2.1

```
1 10.10.10.1 4 msec 0 msec 0 msec  
2 10.10.20.1 12 msec 8 msec 12 msec  
3 10.10.20.2 4 msec * 4 msec
```

```
R1(config)# pseudowire-class PW-CLASS
R1(config-pw-class)# encapsulation l2tpv3
R1(config-pw-class)# protocol l2tpv3
R1(config-pw-class)# ip local interface Loopback0
R1(config-pw-class)# interface GigabitEthernet0/0.200
R1(config-subif)# xconnect 120.100.6.1 200 pw-class PW-CLASS
```

```
R6(config)# pseudowire-class PW-CLASS
R6(config-pw-class)# encapsulation l2tpv3
R6(config-pw-class)# protocol l2tpv3
R6(config-pw-class)# ip local interface Loopback0
R6(config-pw-class)# interface GigabitEthernet0/1.400
R6(config-subif)# xconnect 120.100.1.1 200 pw-class PW-CLASS
```

```
R1# show l2tp session
```

```
L2TP Session Information Total tunnels 1 sessions 1
```

LocID	RemID	Remote Name	State	Remote Address	Port	Sessions	L2TP Class/ VPDN Group
-------	-------	-------------	-------	----------------	------	----------	---------------------------

51446	36190	R6		est	120.100.6.1	0	1
-------	-------	----	--	-----	-------------	---	---

LocID	RemID	TunID	Username, Intf/ Vcid,	State	Last Chg	Uniq ID
			Circuit			

51003	9619	51446	200, Gi0/0.200:200	est	00:24:40	1
-------	------	-------	--------------------	-----	----------	---

```
R6# show l2t session
```

```
L2TP Tunnel and Session Information Total tunnels 1 sessions 1
```

LocID	RemID	Remote Name	State	Remote Address	Port	Sessions	L2TP Class/ VPDN Group
-------	-------	-------------	-------	----------------	------	----------	---------------------------

36190	51446	R1		est	120.100.1.1	0	1
-------	-------	----	--	-----	-------------	---	---

LocID	RemID	TunID	Username, Intf/ Vcid,	State	Last Chg	Uniq ID
			Circuit			

9619	51003	36190	200, Gi0/1.400:400	est	00:25:26	1
------	-------	-------	--------------------	-----	----------	---

```
SW3# ping 1.1.1.2
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:

.....

!Make sure you are logging on your CE devices

```
SW1(config)# logging console
```

```
SW1#
```

```
03:22:19: %SPANTREE-2-RECV_PVID_ERR: Received BPDU with inconsistent peer vlan id  
400 on fastethernet0/1 VLAN200.
```

```
03:22:19: %SPANTREE-2-BLOCK_PVID_LOCAL: Blocking fastethernet0/1 on VLAN0200.  
Inconsistent local vlan.
```

```
SW1# show spanning-tree blockedports
```

Name	Blocked Interfaces List
------	-------------------------

VLAN0200	Fa0/1
----------	-------

Number of blocked ports (segments) in the system : 1

```
SW2#03:22:21: %SPANTREE-2-RECV_PVID_ERR: Received BPDU with inconsistent peer vlan  
id 200 on fastethernet0/6 VLAN400.
```

```
03:22:21: %SPANTREE-2-BLOCK_PVID_PEER: Blocking fastethernet0/6 on VLAN0200.  
Inconsistent peer vlan.
```

```
SW2# show spanning-tree blockedports
```

Name	Blocked Interfaces List
------	-------------------------

VLAN0200	Fa0/6
----------	-------

VLAN0400	Fa0/6
----------	-------

Number of blocked ports (segments) in the system : 2

```
SW3# ping 1.1.1.2
```

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:

.....

Success rate is 0 percent (0/5)

```
SW1# show spanning-tree blockedports
```

Name	Blocked Interfaces List
------	-------------------------

VLAN0200

Fa0/1

Number of blocked ports (segments) in the system : 1

SW1(config)# int fast 0/1

SW1(config-if)# spanning-tree bpdufilter enable

SW1(config-if)#03:33:57: %SPANTREE-2-UNBLOCK_CONSIST_PORT: Unblocking fastethernet0/1 on VLAN0200. Port consistency restored.

SW3# ping 1.1.1.2

Type escape sequence to abort.

Sending 5, 100-byte ICMP Echos to 1.1.1.2, timeout is 2 seconds:

.!!!!

Success rate is 80 percent (4/5), round-trip min/avg/max = 8/12/17 ms

```
! Initial Multicast Setup for the MPLS Core Routers
```

```
R1(config)# ip multicast-routing
R1(config-vrf)# interface Loopback0
R1(config-if)# ip pim sparse-mode
R1(config-if)# interface GigabitEthernet0/1
R1(config-if)# ip pim sparse-mode
```

```
R2(config)# ip multicast-routing
R2(config)# interface fastethernet0/0
R2(config-if)# ip pim sparse-mode
R2(config-if)# interface fastethernet0/1
R2(config-if)# ip pim sparse-mode
```

```
R3(config)# ip multicast-routing
R3(config)# interface GigabitEthernet0/0
R3(config-if)# ip pim sparse-mode
R3(config-if)# interface GigabitEthernet0/1
R3(config-if)# ip pim sparse-mode
```

```
R4(config)# ip multicast-routing
R4(config)# interface GigabitEthernet0/0
R4(config-if)# ip pim sparse-mode
R4(config-if)# interface GigabitEthernet0/1
R4(config-if)# ip pim sparse-mode
```

```
R5(config)# ip multicast-routing
R5(config)# interface GigabitEthernet0/0
R5(config-if)# ip pim sparse-mode
R5(config-if)# interface GigabitEthernet0/1
R5(config-if)# ip pim sparse-mode
```

```
R6(config)# ip multicast-routing
R6(config)# interface Loopback0
R6(config-if)# ip pim sparse-mode
R6(config)# interface GigabitEthernet0/0
R6(config-if)# ip pim sparse-mode
```

```
! PE Specific mVRF and MDT Configuration
```

```
R1(config)# ip multicast-routing vrf RED
R1(config)# ip vrf RED
R1(config-vrf)# mdt default 232.0.0.11
R1(config-vrf)# interface GigabitEthernet0/0.50
R1(config-subif)# ip pim sparse-mode
R1(config-subif)# exit
R1(config)# ip pim vrf RED rp-address 130.100.100.1
R1(config)# ip pim ssm default
```

```
R6(config)# ip vrf RED
R6(config-vrf)# mdt default 232.0.0.11
R6(config-vrf)# interface GigabitEthernet0/1.100
R6(config-subif)# ip pim sparse-mode
R6(config-subif)# exit
R6(config)# ip pim vrf RED rp-address 130.100.100.1
R6(config)# ip pim ssm default
```

```
! CE Specific Configuration
```

```
SW3(config)# ip multicast-routing distributed
SW3(config)# int vlan 50
SW3(config-if)# ip pim sparse-mode
SW3(config-if)# exit
SW3(config)# ip pim rp-address 130.100.100.1
```

```
SW4(config)# ip multicast-routing distributed
SW4(config)# interface vlan 100
SW4(config-if)# ip pim sparse-mode
SW4(config-if)# ip igmp join-group 226.2.2.2
SW4(config-if)# exit
SW4(config)# ip pim rp-address 130.100.100.1
```

```
R6# show ip pim vrf RED neigh
PIM Neighbor Table
Mode: B - Bidir Capable, DR - Designated Router, N - Default DR Priority,
      S - State Refresh Capable
Neighbor           Interface          Uptime/Expires   Ver   DR
Address
Prio/Mode
130.100.100.2    GigabitEthernet0/1.100 00:02:08/00:01:34 v2    1 / DR S
120.100.1.1       Tunnel1           00:00:05/00:01:39 v2    1 / S
```

```
R1# ping vrf RED 226.2.2.2
```

Type escape sequence to abort.

Sending 1, 100-byte ICMP Echos to 226.2.2.2, timeout is 2 seconds:

Reply to request 0 from 130.100.100.2, 12 ms

```
SW3# ping 226.2.2.2
```

Type escape sequence to abort.

Sending 1, 100-byte ICMP Echos to 226.2.2.2, timeout is 2 seconds:

Reply to request 0 from 130.100.100.2, 9 ms

```
SW3# show ip pim rp
```

Group: 226.2.2.2, RP: 130.100.100.1, v2, uptime 00:00:37, expires never

Group: 224.0.1.40, RP: 130.100.100.1, v2, uptime 01:01:24, expires never

```
R1# show ip pim mdt bgp
```

Peer (Route Distinguisher + IPv4)	Next Hop
MDT group 232.0.0.11	
2:65001:200:120.100.6.1	120.100.6.1

```
R6# show ip pim mdt bgp
```

Peer (Route Distinguisher + IPv4)	Next Hop
MDT group 232.0.0.11	
2:65001:200:120.100.1.1	120.100.1.1

```
R1(config)# ipv6 unicast-routing
R1(config)# ipv6 cef
R1(config)# mpls ipv6 source-interface Loopback0
R1(config)# interface loopback0
R1(config-if)# ipv6 add 2010:C15:C0:1::1/64
R1(config-if)# interface GigabitEthernet0/0.10
R1(config-subif)# ipv6 address 2010:C15:C0:11::1/64
R1(config-subif)# router bgp 65001
R1(config-router)# no bgp default ipv4-unicast
R1(config-router)# address-family ipv6
R1(config-router-af)# neighbor 120.100.6.1 activate
R1(config-router-af)# neighbor 120.100.6.1 send-lsa
R1(config-router-af)# network 2010:C15:C0:11::0/64
R1(config-router-af)# network 2010:C15:C0:1::/64
R1(config-router-af)# exit-address-family
```

```
R6(config)# ipv6 unicast-routing
R6(config)# ipv6 cef
R6(config)# mpls ipv6 source-interface Loopback0
R6(config)# interface loopback0
R6(config-if)# ipv6 add 2010:C15:C0:6::1/64
R6(config-if)# interface GigabitEthernet0/1.20
R6(config-subif)# ipv6 address 2010:C15:C0:62::1/64
R6(config-subif)# router bgp 65001
R6(config-router)# no bgp default ipv4-unicast
R6(config-router)# address-family ipv6
R6(config-router-af)# neighbor 120.100.1.1 activate
R6(config-router-af)# neighbor 120.100.1.1 send-labe
R6(config-router-af)# network 2010:C15:C0:62::/64
R6(config-router-af)# network 2010:C15:C0:6::/64
R6(config-router-af)# exit-address-family
```

```
R1# show ip bgp ipv6 unicast
BGP table version is 5, local router ID is 120.100.1.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete
```

Network	Next Hop	Metric	LocPrf	Weight	Path
*> 2010:C15:C0:1::/64	...	0		32768	i

```
*>i2010:C15:C0:6::/64
    ::FFFF:120.100.6.1
        0      100      0 i
*> 2010:C15:C0:11::/64
    :::
        0          32768 i
*>i2010:C15:C0:62::/64
    ::FFFF:120.100.6.1
        0      100      0 i

R6# show ip bgp ipv6 unicast
BGP table version is 5, local router ID is 120.100.6.1
Status codes: s suppressed, d damped, h history, * valid, > best, i - internal,
               r RIB-failure, S Stale
Origin codes: i - IGP, e - EGP, ? - incomplete

      Network          Next Hop           Metric LocPrf Weight Path
*>i2010:C15:C0:1::/64
    ::FFFF:120.100.1.1
        0      100      0 i
*> 2010:C15:C0:6::/64
    :::
        0          32768 i
*>i2010:C15:C0:11::/64
    ::FFFF:120.100.1.1
        0      100      0 i
*> 2010:C15:C0:62::/64
    :::
        0          32768 i

R1# ping ipv6 2010:C15:C0:62::1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2010:C15:C0:62::1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms
R1# ping ipv6 2010:C15:C0:6::1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2010:C15:C0:6::1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms
```

```
R6# ping ipv6 2010:C15:C0:11::1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2010:C15:C0:11::1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/8/12 ms

R6# ping ipv6 2010:C15:C0:1::1

Type escape sequence to abort.
Sending 5, 100-byte ICMP Echos to 2010:C15:C0:1::1, timeout is 2 seconds:
!!!!!
Success rate is 100 percent (5/5), round-trip min/avg/max = 8/9/12 ms

R1# show ipv6 route

IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external

C    2010:C15:C0:1::/64 [0/0]
    via ::, Loopback0
L    2010:C15:C0:1::1/128 [0/0]
    via ::, Loopback0
B    2010:C15:C0:6::/64 [200/0]
    via ::FFFF:120.100.6.1, IPv6-mpls
C    2010:C15:C0:11::/64 [0/0]
    via ::, GigabitEthernet0/0.10
L    2010:C15:C0:11::1/128 [0/0]
    via ::, GigabitEthernet0/0.10
B    2010:C15:C0:62::/64 [200/0]
    via ::FFFF:120.100.6.1, IPv6-mpls
L    FE80::/10 [0/0]
    via ::, Null0
L    FF00::/8 [0/0]
    via ::, Null0
```

```
R6# show ipv6 route
IPv6 Routing Table - 8 entries
Codes: C - Connected, L - Local, S - Static, R - RIP, B - BGP
      U - Per-user Static route
      I1 - ISIS L1, I2 - ISIS L2, IA - ISIS interarea, IS - ISIS summary
      O - OSPF intra, OI - OSPF inter, OE1 - OSPF ext 1, OE2 - OSPF ext 2
      ON1 - OSPF NSSA ext 1, ON2 - OSPF NSSA ext 2
      D - EIGRP, EX - EIGRP external
B  2010:C15:C0:1::/64 [200/0]
    via ::FFFF:120.100.1.1, IPv6-mpls
C  2010:C15:C0:6::/64 [0/0]
    via ::, Loopback0
L  2010:C15:C0:6::1/128 [0/0]
    via ::, Loopback0
B  2010:C15:C0:11::/64 [200/0]
    via ::FFFF:120.100.1.1, IPv6-mpls
C  2010:C15:C0:62::/64 [0/0]
    via ::, GigabitEthernet0/1.10
L  2010:C15:C0:62::1/128 [0/0]
    via ::, GigabitEthernet0/1.20
L  FE80::/10 [0/0]
    via ::, Null0
L  FF00::/8 [0/0]
    via ::, Null0
```

```
R1(config)# class-map match-any VOICE
R1(config-cmap)# match ip dscp ef
R1(config-cmap)# match ip dscp cs5
R1(config-cmap)# class-map match-any MISSION-CRITICAL
R1(config-cmap)# match ip dscp cs6
R1(config-cmap)# match ip dscp af31
R1(config-cmap)# match ip dscp af32
R1(config-cmap)# match ip dscp cs3
R1(config-cmap)# policy-map PE-CE-CHILD
R1(config-pmap)# class VOICE
R1(config-pmap-c)# priority percent 35
R1(config-pmap-c)# class MISSION-CRITICAL
R1(config-pmap-c)# bandwidth percent 40
R1(config-pmap-c)# random-detect dscp-based
R1(config-pmap-c)# class class-default
R1(config-pmap-c)# bandwidth percent 25
R1(config-pmap-c)# random-detect
R1(config-pmap-c)# exit
R1(config-cmap)# policy-map PE-CE-PARENT
R1(config-pmap-c)# class class-default
R1(config-pmap-c)# shape average 1000000
R1(config-pmap-c)# service-policy PE-CE-CHILD
R1(config-pmap-c)# exit
R1(config-pmap)# exit

R1(config)# interface GigabitEthernet0/0.10
R1(config-subif)# service-policy output PE-CE-PARENT
```

```
R1(config)# policy-map CE-PE-SHAPE
R1(config-pmap)# class VOICE
R1(config-pmap-c)# police cir 350000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 5
R1(config-pmap-c-police)# exceed-action drop
R1(config-pmap-c-police)# class MISSION-CRITICAL
R1(config-pmap-c)# police cir 400000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 3
R1(config-pmap-c-police)# exceed-action set-mpls-exp-topmost-transmit 7
R1(config-pmap-c-police)# class class-default
R1(config-pmap-c)# police cir 250000
R1(config-pmap-c-police)# conform-action set-mpls-exp-topmost-transmit 0
R1(config-pmap-c-police)# exceed-action set-mpls-exp-topmost-transmit 4
R1(config-pmap-c-police)# interface GigabitEthernet0/0.10
R1(config-subif)# service-policy input CE-PE-SHAPE
```

```
R4(config)# interface loopback1
R4(config-if)# ip add 4.4.4.4 255.255.255.0
R4(config-if)# router eigrp VPN
R4(config-router)# address-family ipv4 autonomous-system 1
R4(config-router-af)# network 100.100.100.0 0.0.0.255
R4(config-router-af)# network 4.4.4.0 0.0.0.255

R5(config)# interface loopback1
R5(config-if)# ip address 5.5.5.5 255.255.255.0
R5(config-if)# router eigrp VPN
R5(config-router)# address-family ipv4 autonomous-system 1
R5(config-router-af)# network 100.100.100.0 0.0.0.255
R5(config-router-af)# network 5.5.5.0 0.0.0.255

R6(config)# interface loopback1
R6(config-if)# ip address 6.6.6.6 255.255.255.0
R6(config-if)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
R6(config-router-af)# network 100.100.100.0 0.0.0.255
R6(config-router-af)# network 6.6.6.0 0.0.0.255

R6(config)# crypto isakmp policy 1
R6(config-isakmp)# authentication pre-share
R6(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R6(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R6(crypto-profile-trans)# crypto ipsec profile IPSEC
R6(ipsec-profile)# set transform-set DMVPN
R6(ipsec-profile)# interface Tunnell1
R6(config-if)# ip address 100.100.100.6 255.255.255.0
R6(config-if)# ip mtu 1416
R6(config-if)# ip nhrp authentication SECRET
R6(config-if)# ip nhrp map multicast dynamic
R6(config-if)# ip nhrp network-id 10
R6(config-if)# ip nhrp holdtime 100
R6(config-if)# delay 2000
R6(config-if)# tunnel source gig 0/0
R6(config-if)# tunnel mode gre multipoint
R6(config-if)# tunnel key 1
R6(config-if)# tunnel protection ipsec profile IPSEC
```

```
R4(config)# crypto isakmp policy 1
R4(config-isakmp)# authentication pre-share
R4(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R4(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R4(cfg-crypto-trans)# crypto ipsec profile IPSEC
R4(ipsec-profile)# set transform-set DMVPN
R4(ipsec-profile)# interface Tunnel0
R4(config-if)# ip address 100.100.100.4 255.255.255.0
R4(config-if)# ip mtu 1416
R4(config-if)# ip nhrp authentication SECRET
R4(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R4(config-if)# ip nhrp map multicast 120.100.45.6
R4(config-if)# ip nhrp network-id 10
R4(config-if)# ip nhrp holdtime 100
R4(config-if)# ip nhrp nhs 100.100.100.6
R4(config-if)# delay 2000
R4(config-if)# tunnel source gig 0/0
R4(config-if)# tunnel mode gre multipoint
R4(config-if)# tunnel key 1
R4(config-if)# tunnel protection ipsec profile IPSEC
```

```
R5(config)# crypto isakmp policy 1
R5(config-isakmp)# authentication pre-share
R5(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R5(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R5(cfg-crypto-trans)# crypto ipsec profile IPSEC
R5(ipsec-profile)# set transform-set DMVPN
R5(ipsec-profile)# interface Tunnel0
R5(config-if)# ip address 100.100.100.5 255.255.255.0
R5(config-if)# ip mtu 1416
R5(config-if)# ip nhrp authentication SECRET
R5(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R5(config-if)# ip nhrp map multicast 120.100.45.6
R5(config-if)# ip nhrp network-id 10
R5(config-if)# ip nhrp holdtime 100
R5(config-if)# ip nhrp nhs 100.100.100.6
R5(config-if)# delay 2000
R5(config-if)# tunnel source gig 0/0
R5(config-if)# tunnel mode gre multipoint
R5(config-if)# tunnel key 1
R5(config-if)# tunnel protection ipsec profile IPSEC
```

```
R4# show ip route eigrp
  6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:02:42, Tunnel0
```

```
R5# show ip route eigrp
  6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:50, Tunnel0
```

```
R6# show ip route eigrp
  4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/285084416] via 100.100.100.4, 00:03:06, Tunnel0
  5.0.0.0/24 is subnetted, 1 subnets
D      5.5.5.0 [90/285084416] via 100.100.100.5, 00:01:02, Tunnel0
```

!R6 has both spoke routes yet each spoke (R4 and R5) only have the hub network route, !a classic split horizon issue.

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
```

```
R6(config-router-af)# af-interface Tunnel0
R6(config-router-af-interface)# no split-horizon
```

```
R4# show ip route eigrp
  5.0.0.0/24 is subnetted, 1 subnets
D      5.5.5.0 [90/285596416] via 100.100.100.6, 00:00:22, Tunnel0
  6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:04:14, Tunnel0
```

```
R5# show ip route eigrp
  4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/285596416] via 100.100.100.6, 00:00:33, Tunnel0
  6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:02:20, Tunnel0
R5#
```

! The next-hop for spoke to spoke routes shows as the hub router (100.100.100.6) yet
! the question states traffic must flow directly between spokes so the next-hop must
! be modified

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
R6(config-router-af)# af-interface Tunnel1
R6(config-router-af-interface)# no next-hop-self

R4# show ip route eigrp
 5.0.0.0/24 is subnetted, 1 subnets
D      5.5.5.0 [90/285596416] via 100.100.100.5, 00:00:28, Tunnel0
 6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:29, Tunnel0

R5# show ip route eigrp
 4.0.0.0/24 is subnetted, 1 subnets
D      4.4.4.0 [90/285596416] via 100.100.100.4, 00:00:39, Tunnel0
 6.0.0.0/24 is subnetted, 1 subnets
D      6.6.6.0 [90/285084416] via 100.100.100.6, 00:00:39, Tunnel0
```

```
R5# show crypto map
Crypto Map "Tunnel0-head-0" 65536 ipsec-isakmp
    Profile name: IPSEC
    Security association lifetime: 4608000 kilobytes/3600 seconds
    PFS (Y/N) : N
    Transform sets={
        DMVPN,
    }
```

```
Crypto Map "Tunnel0-head-0" 65537 ipsec-isakmp
    Map is a PROFILE INSTANCE.
    Peer = 120.100.45.6
    Extended IP access list
        access-list permit gre host 120.100.45.5 host 120.100.45.6
    Current peer: 120.100.45.6
    Security association lifetime: 4608000 kilobytes/3600 seconds
    PFS (Y/N) : N
    Transform sets={
        DMVPN,
    }
    Interfaces using crypto map Tunnel0-head-0:
        Tunnel0
```

```
R5# show crypto isakmp sa
IPv4 Crypto ISAKMP SA
dst          src          state      conn-id slot status
120.100.45.6 120.100.45.5 QM_IDLE      4001     0 ACTIVE
```

```
IPv6 Crypto ISAKMP SA
```

R5 spoke router only has a connection to the hub router. An extended ping sourced from the loopback interface of one spoke to another is required to bring up the dynamic spoke to spoke connection.

```
R5# ping
Protocol [ip]:
Target IP address: 4.4.4.4
Repeat count [5]:
Datagram size [100]:
```

```
Timeout in seconds [2]:  
Extended commands [n]: y  
Source address or interface: 5.5.5.5  
Type of service [0]:  
Set DF bit in IP header? [no]:  
Validate reply data? [no]:  
Data pattern [0xABCD]:  
Loose, Strict, Record, Timestamp, Verbose [none]:  
Sweep range of sizes [n]:  
Type escape sequence to abort.  
Sending 5, 100-byte ICMP Echos to 4.4.4.4, timeout is 2 seconds:  
Packet sent with a source address of 5.5.5.5  
!!!!!  
Success rate is 100 percent (5/5), round-trip min/avg/max = 1/2/4 ms
```

```
R5# show crypto isakmp sa
```

```
IPv4 Crypto ISAKMP SA
```

dst	src	state	conn-id	slot	status
120.100.45.5	120.100.45.4	QM_IDLE	4002	0	ACTIVE
120.100.45.6	120.100.45.5	QM_IDLE	4001	0	ACTIVE

```
IPv6 Crypto ISAKMP SA
```

```
R5# show crypto isakmp sa
```

```
IPv4 Crypto ISAKMP SA
```

dst	src	state	conn-id	slot	status
120.100.45.5	120.100.45.4	QM_IDLE	4002	0	ACTIVE
120.100.45.6	120.100.45.5	QM_IDLE	4001	0	ACTIVE

```
IPv6 Crypto ISAKMP SA
```

```
R2(config-if)# router eigrp VPN
R2(config-router)# address-family ipv4 autonomous-system 1
R2(config-router-af)# network 100.100.100.0 0.0.0.255
R2(config-router-af)# exit-address-family
R2(config-router)# crypto isakmp policy 1
R2(config-isakmp)# authentication pre-share
R2(config-isakmp)# crypto isakmp key CCIE address 0.0.0.0
R2(config-isakmp)# crypto ipsec transform-set DMVPN esp-des esp-md5-hmac
R2(crypto-trans)# crypto ipsec profile IPSEC
R2(ipsec-profile)# set transform-set DMVPN
R2(ipsec-profile)# interface Tunnel0
R2(config-if)# ip address 100.100.100.2 255.255.255.0
R2(config-if)# ip mtu 1416
R2(config-if)# ip nhrp authentication SECRET
R2(config-if)# ip nhrp map 100.100.100.6 120.100.45.6
R2(config-if)# ip nhrp map multicast 120.100.45.6
R2(config-if)# ip nhrp network-id 10
R2(config-if)# ip nhrp holdtime 100
R2(config-if)# ip nhrp nhs 100.100.100.6
R2(config-if)# delay 2000
R2(config-if)# tunnel source fastethernet0/1
R2(config-if)# tunnel mode gre multipoint
R2(config-if)# tunnel key 1
R2(config-if)# tunnel protection ipsec profile IPSEC
R4(config)# interface loopback2
R4(config-if)# ip add 45.45.45.45 255.255.255.0
R4(config-if)# router eigrp VPN
R4(config-router)# address-family ipv4 autonomous-system 1
R4(config-router-af)# network 45.45.45.0 0.0.0.255

R5(config)# interface loopback2
R5(config-if)# ip add 45.45.45.45 255.255.255.0
R5(config-if)# router eigrp VPN
R5(config-router)# address-family ipv4 autonomous-system 1
R5(config-router-af)# network 45.45.45.0 0.0.0.255
```

```
R6# show ip route eigrp
    4.0.0.0/24 is subnetted, 1 subnets
D        4.4.4.0 [90/61440640] via 100.100.100.4, 00:00:16, Tunnel0
    5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        5.5.5.0/24 [90/61440640] via 100.100.100.5, 00:00:16, Tunnel0
    45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        45.45.45.0/24 [90/61440640] via 100.100.100.5, 00:01:10, Tunnel0
                                         [90/61440640] via 100.100.100.4, 00:01:10, Tunnel0
```

```
R2# show ip route eigrp
    4.0.0.0/24 is subnetted, 1 subnets
D        4.4.4.0 [90/71680640] via 100.100.100.4, 00:01:40, Tunnel0
    5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        5.5.5.0/24 [90/71680640] via 100.100.100.5, 00:01:40, Tunnel0
    6.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        6.6.6.0/24 [90/61440640] via 100.100.100.6, 00:07:05, Tunnel0
    45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        45.45.45.0/24 [90/71680640] via 100.100.100.4, 00:03:39, Tunnel0
```

```
R6(config)# router eigrp VPN
R6(config-router)# address-family ipv4 autonomous-system 1
R6(config-router-af)# af-interface Tunnel0
R6(config-router-af-interface)# add-paths 2

R2# show ip route eigrp
    4.0.0.0/24 is subnetted, 1 subnets
D        4.4.4.0 [90/71680640] via 100.100.100.4, 00:01:16, Tunnel0
    5.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        5.5.5.0/24 [90/71680640] via 100.100.100.5, 00:01:14, Tunnel0
    6.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        6.6.6.0/24 [90/61440640] via 100.100.100.6, 00:01:22, Tunnel0
    45.0.0.0/8 is variably subnetted, 2 subnets, 2 masks
D        45.45.45.0/24 [90/71680640] via 100.100.100.4, 00:01:16, Tunnel0
                                         [90/61440640] via 100.100.100.5, 00:01:16, Tunnel0
```

```
R1(config)# line con 0
R1(config-line)# activation-character ?
CHAR or <0-127> Activation character or its decimal equivalent
R1(config-line)# activation-character 49
```